

Operating Systems & C

Lecture 2: Processors

Willard Rafnsson

IT University of Copenhagen

Digital Logic

Overview of Logic Design

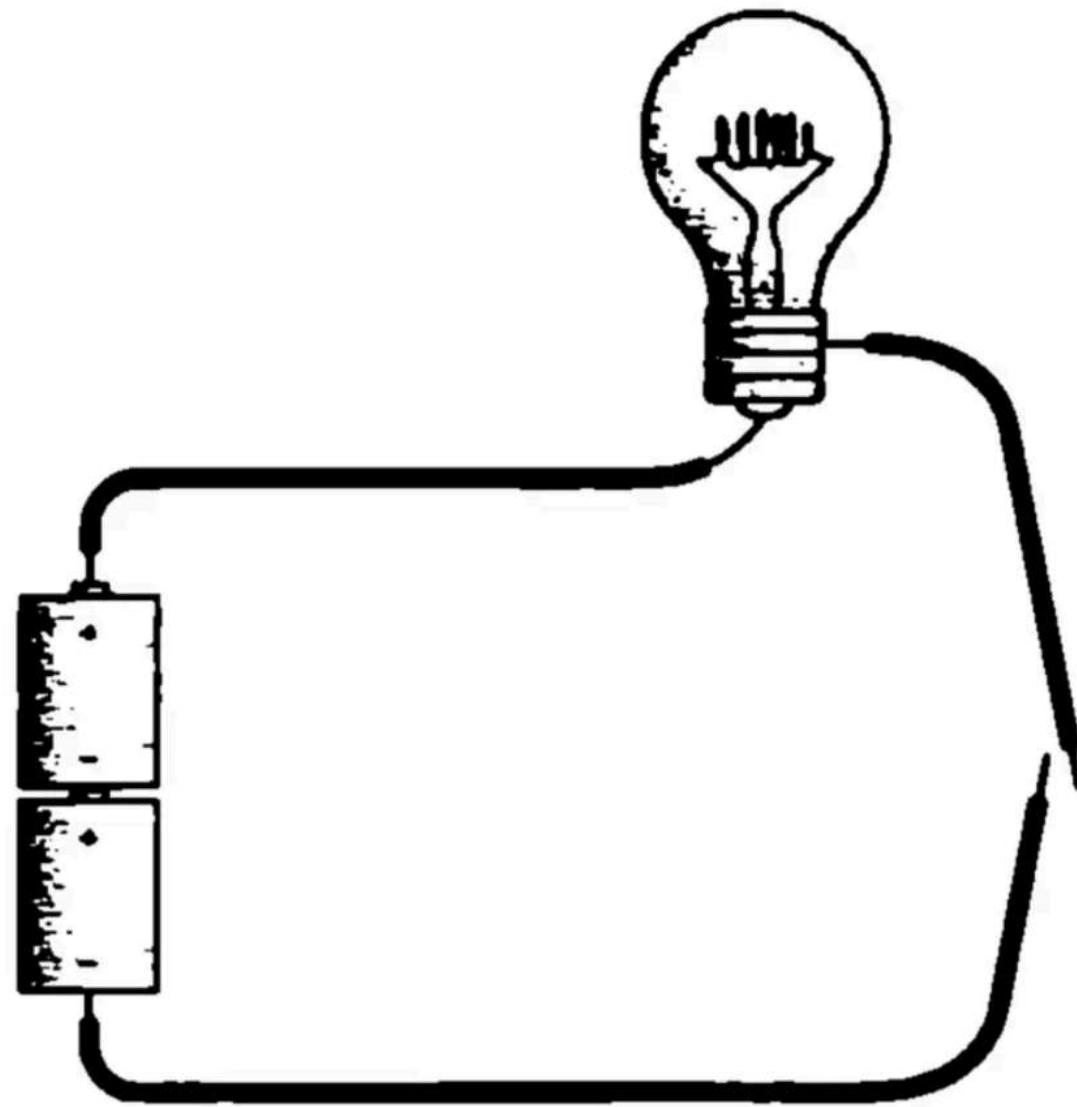
Fundamental Hardware Requirements

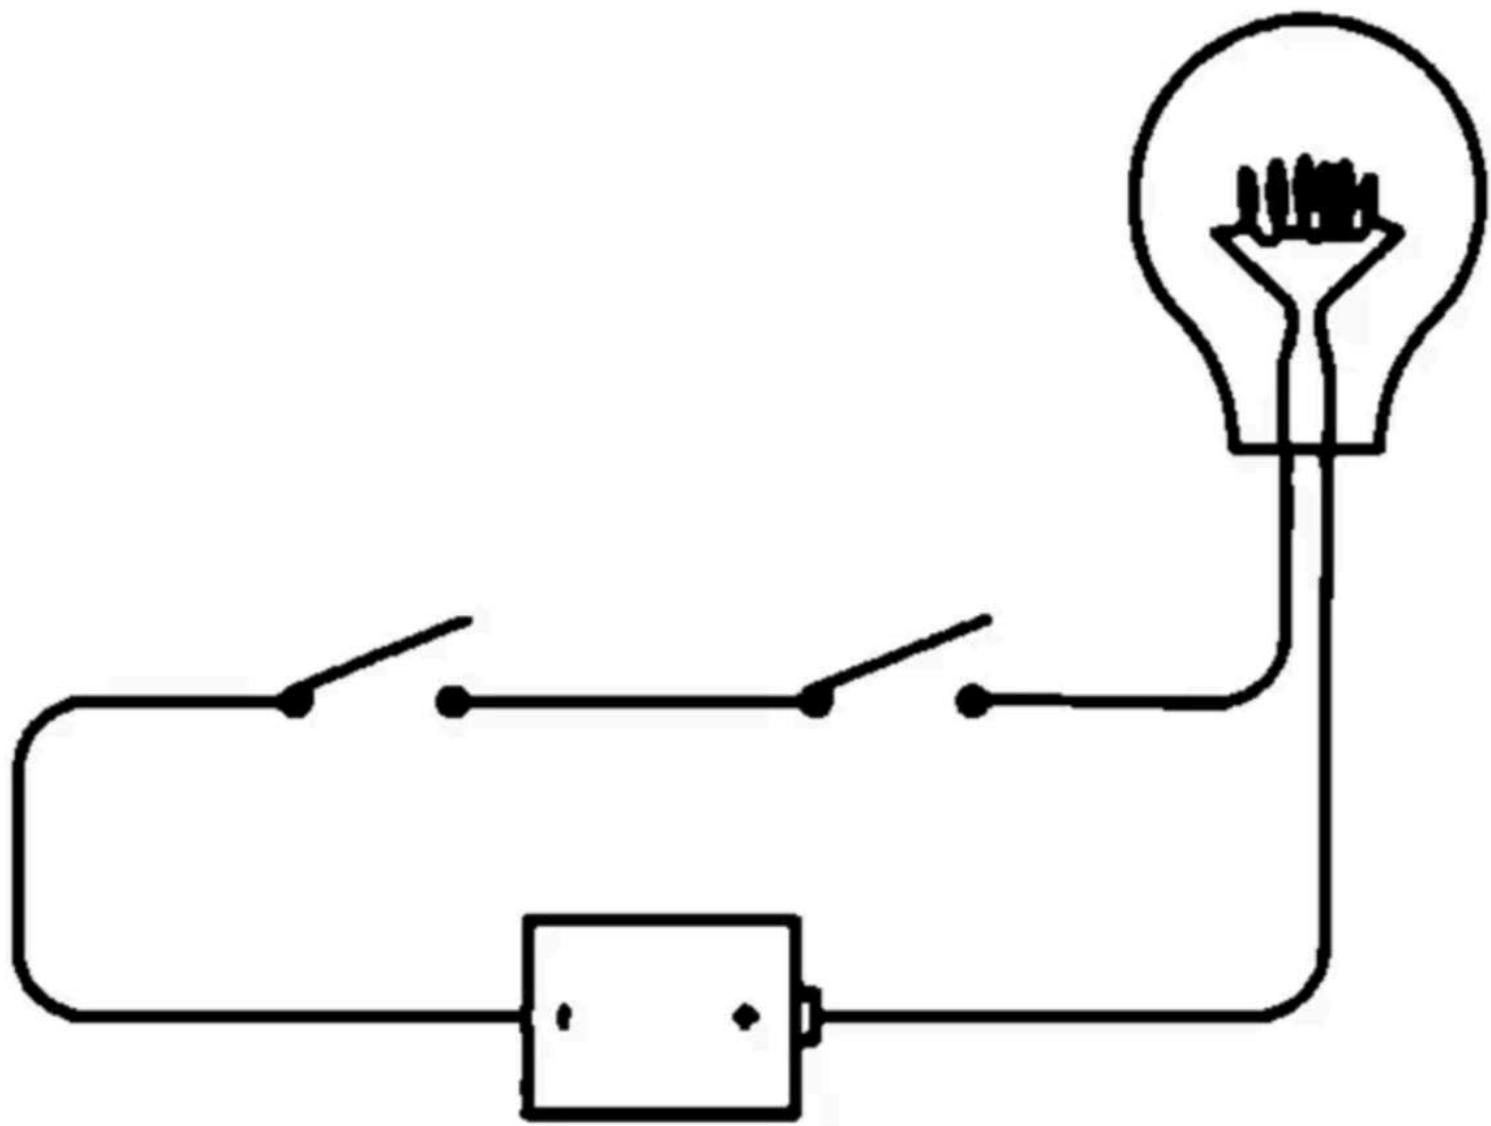
- Communication
 - How to get values from one place to another
- Computation
- Storage

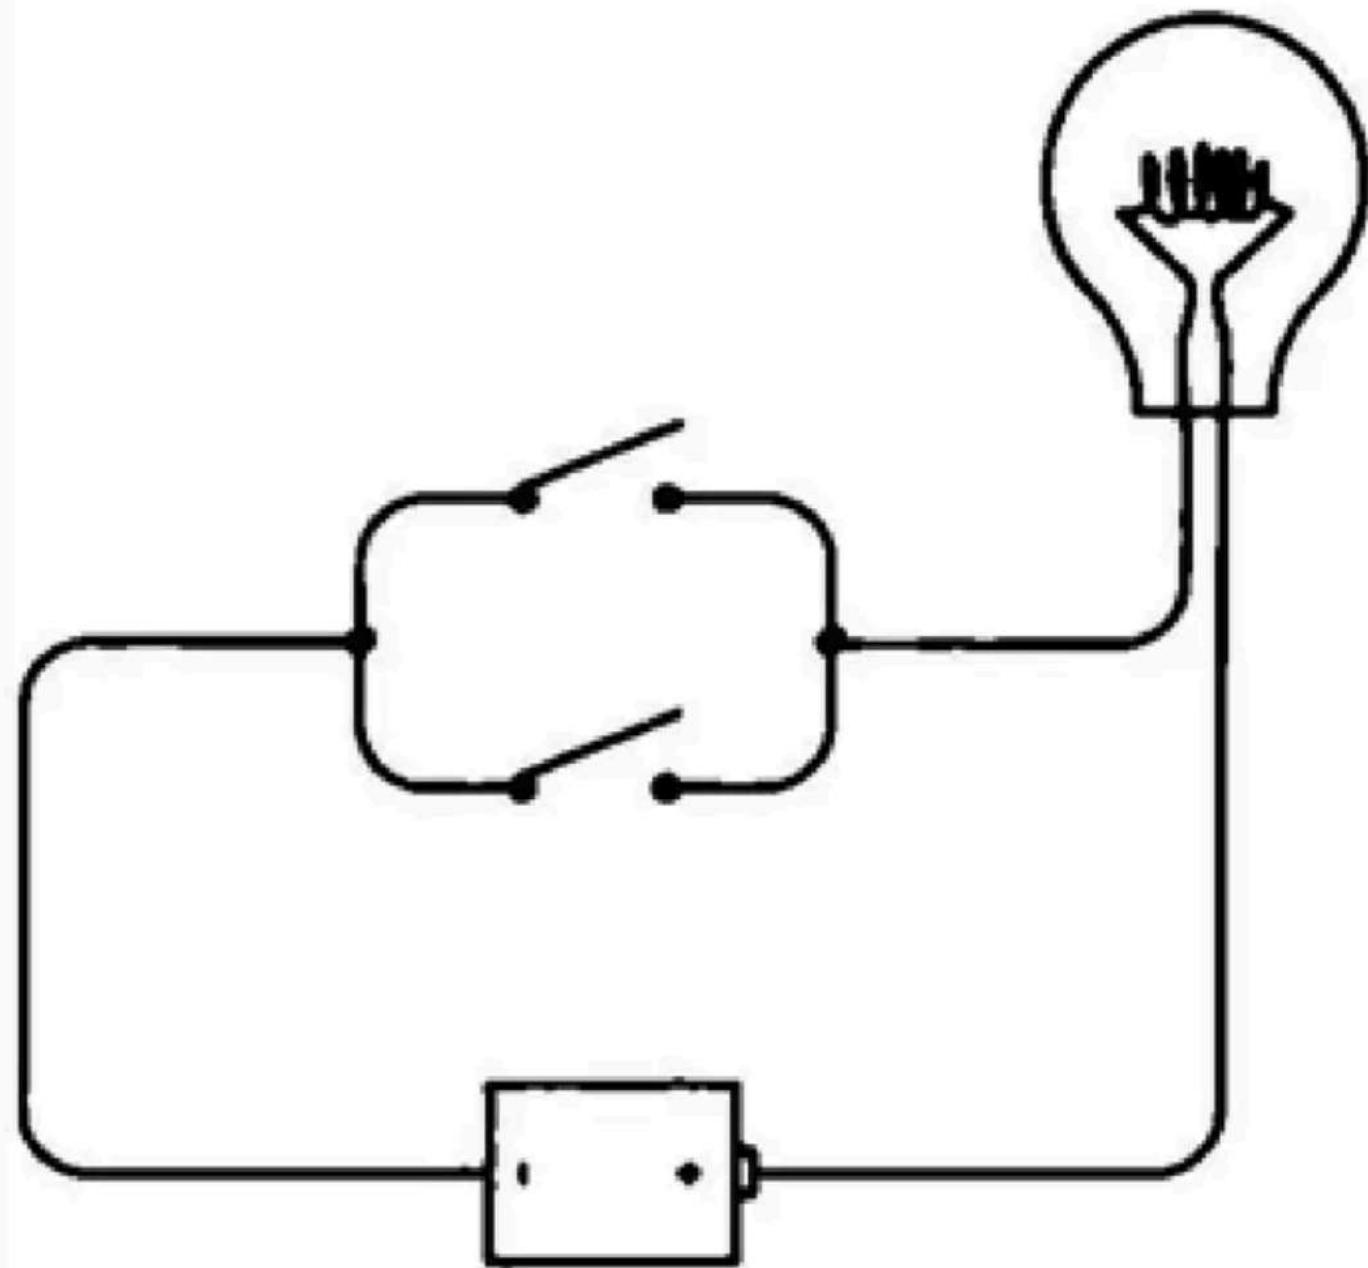
Bits are Our Friends

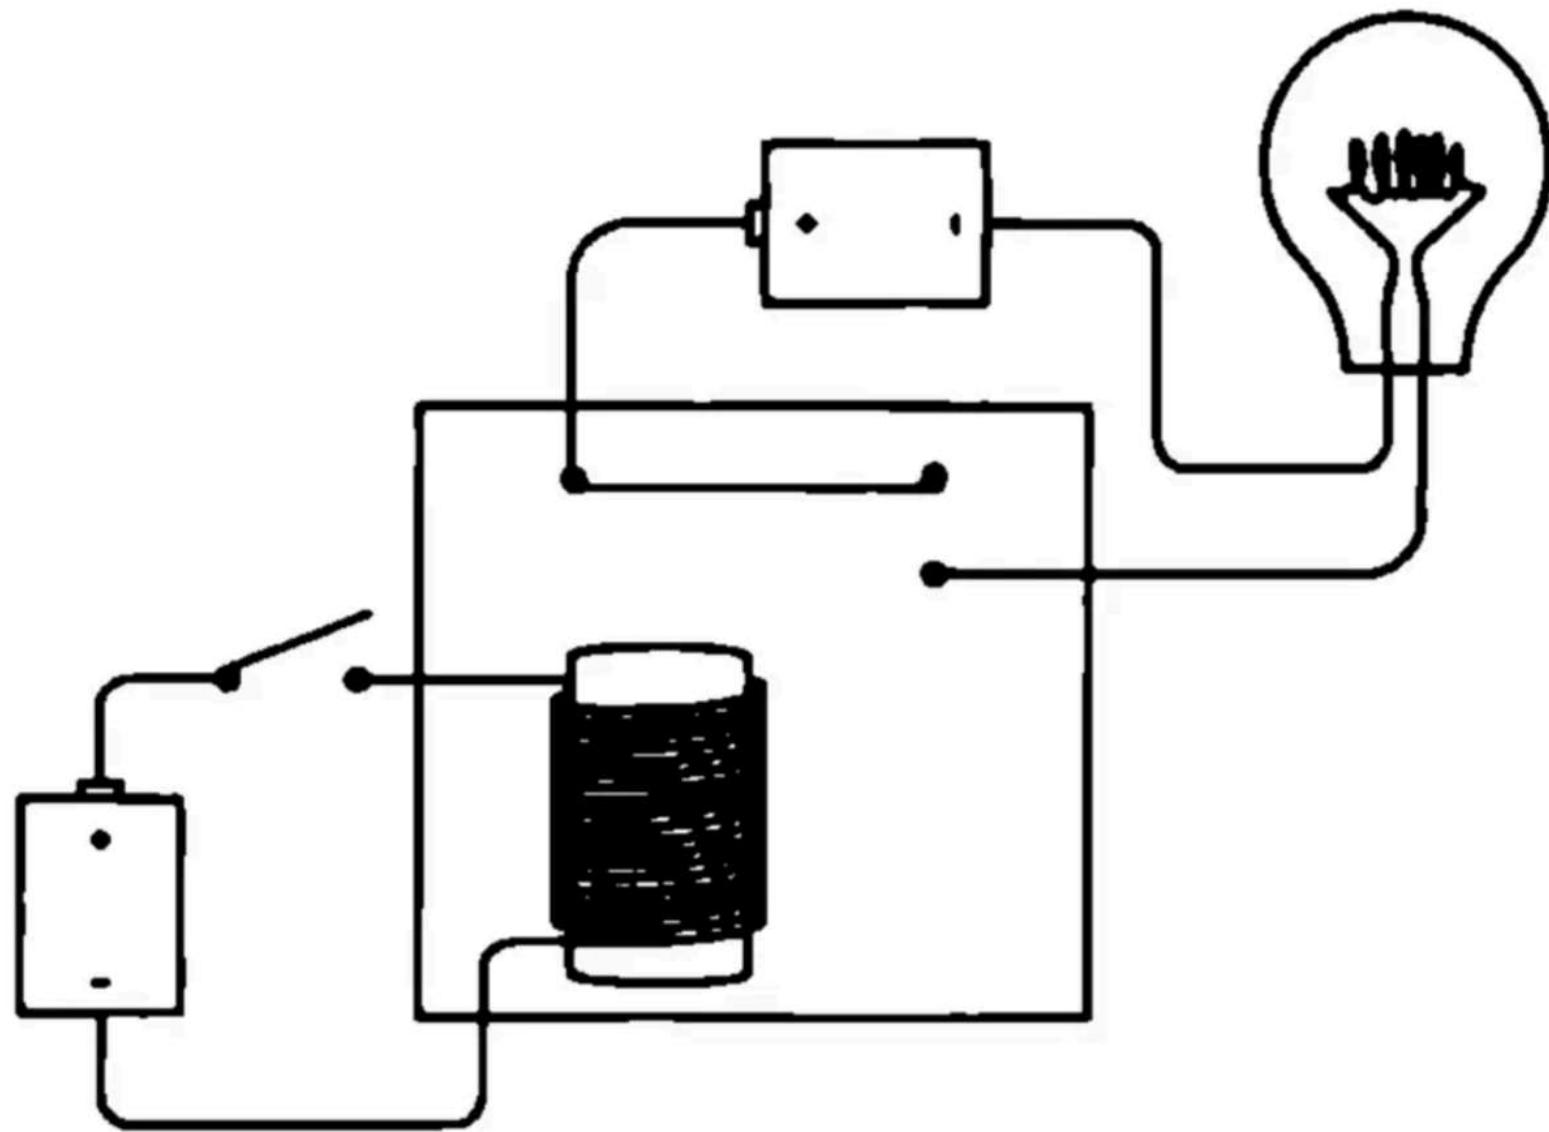
- Everything expressed in terms of values 0 and 1
- Communication
 - Low or high voltage on wire
- Computation
 - Compute Boolean functions
- Storage
 - Store bits of information

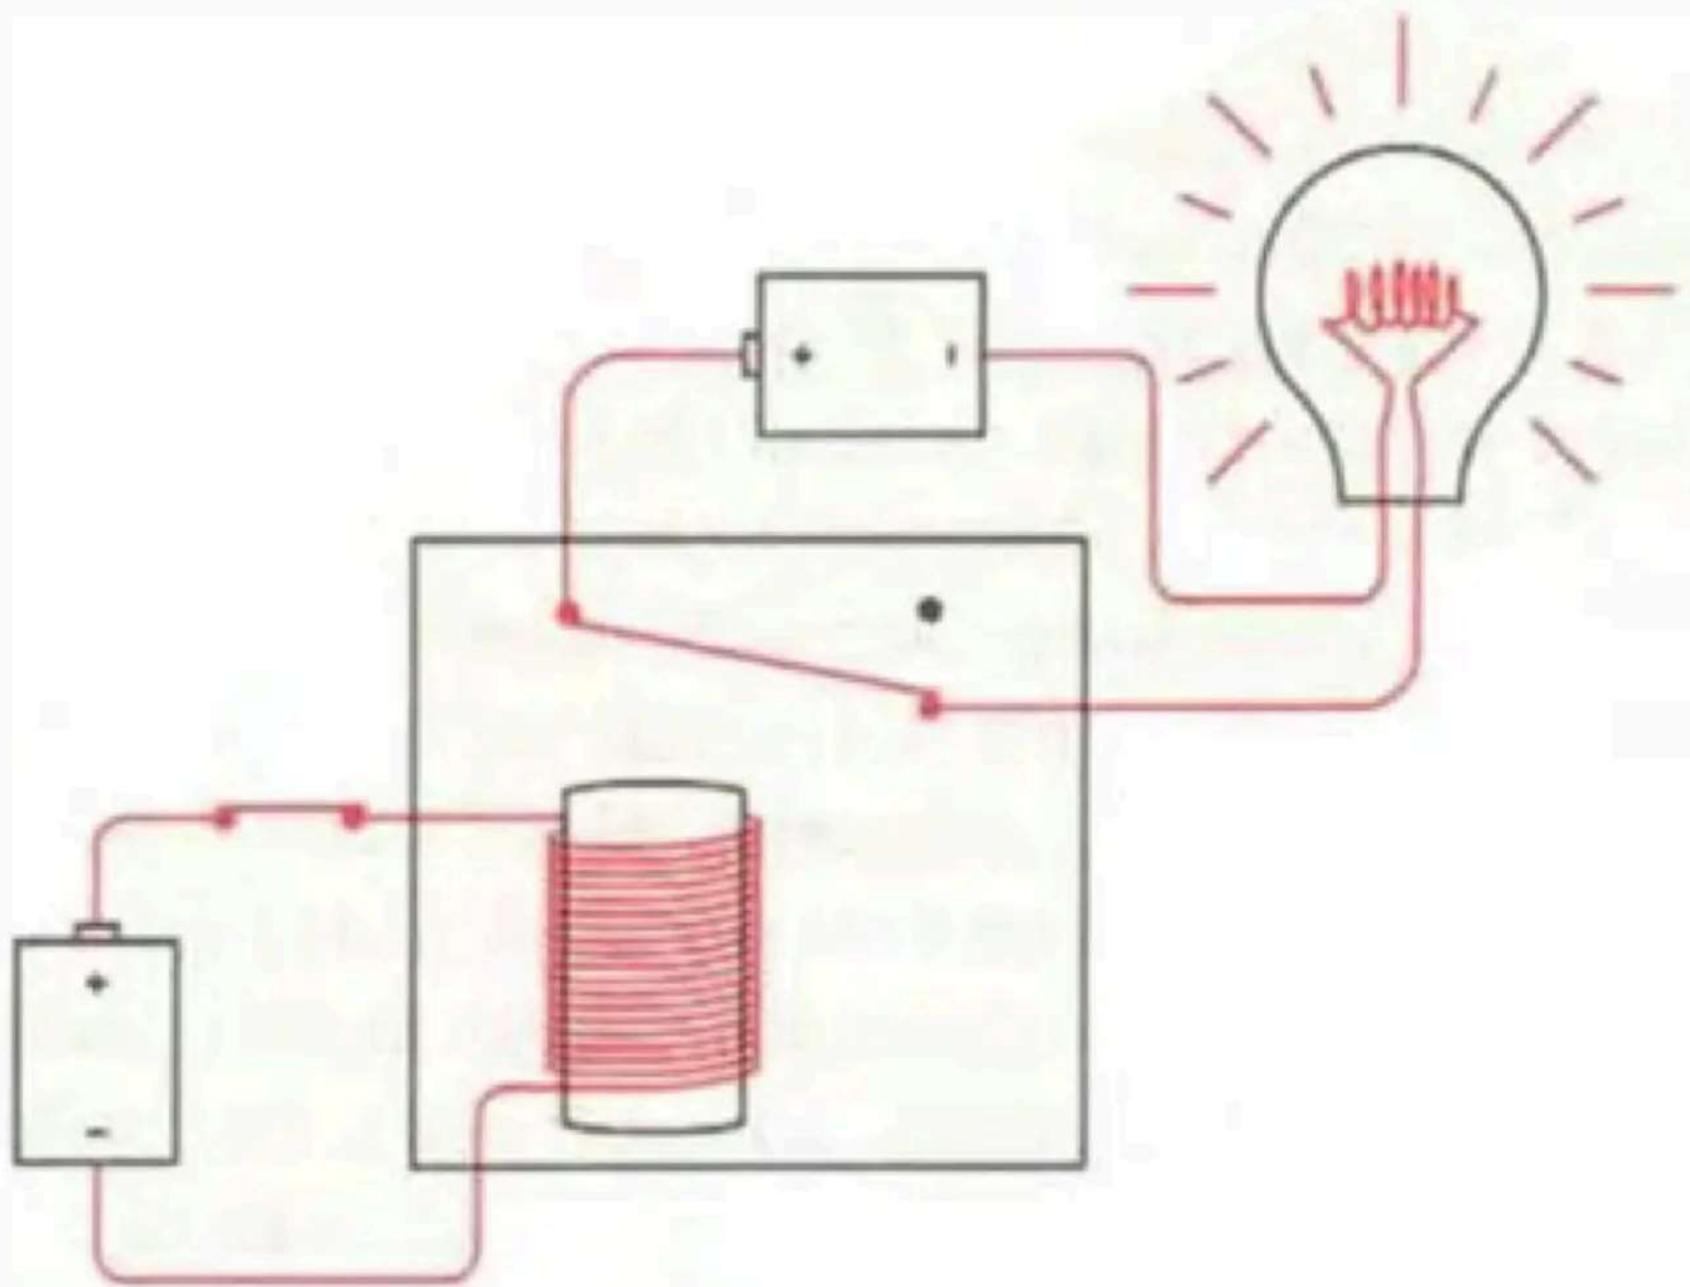
Anatomy of a Flashlight

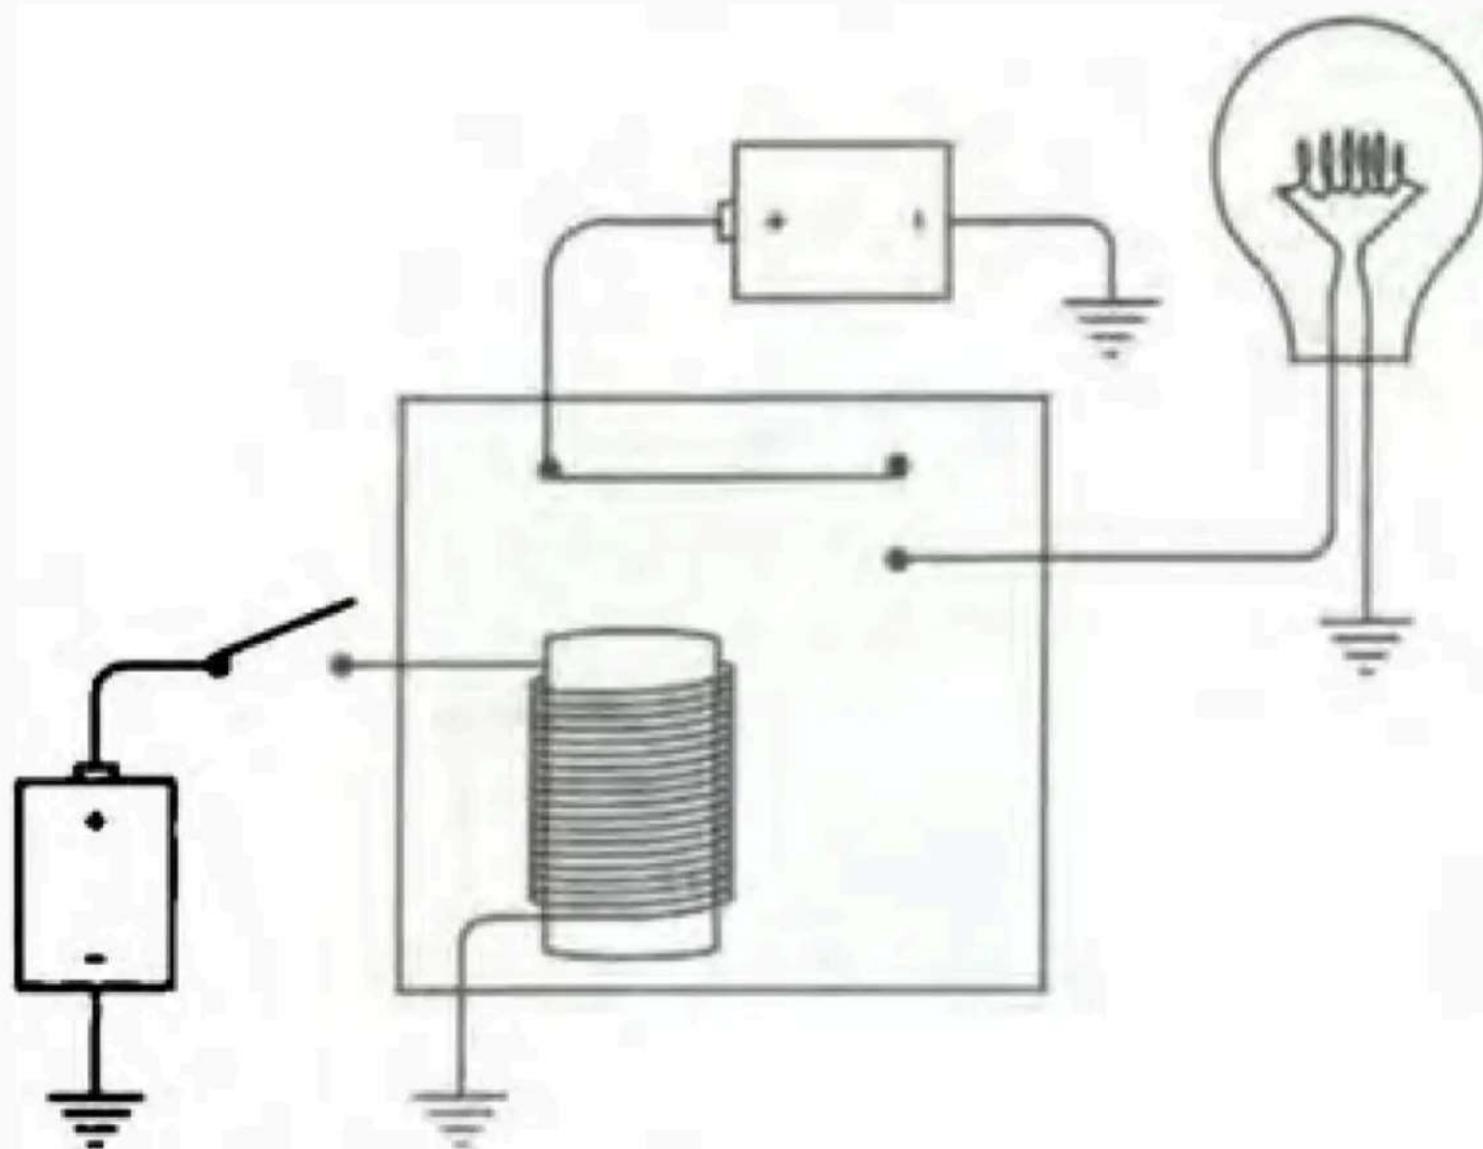


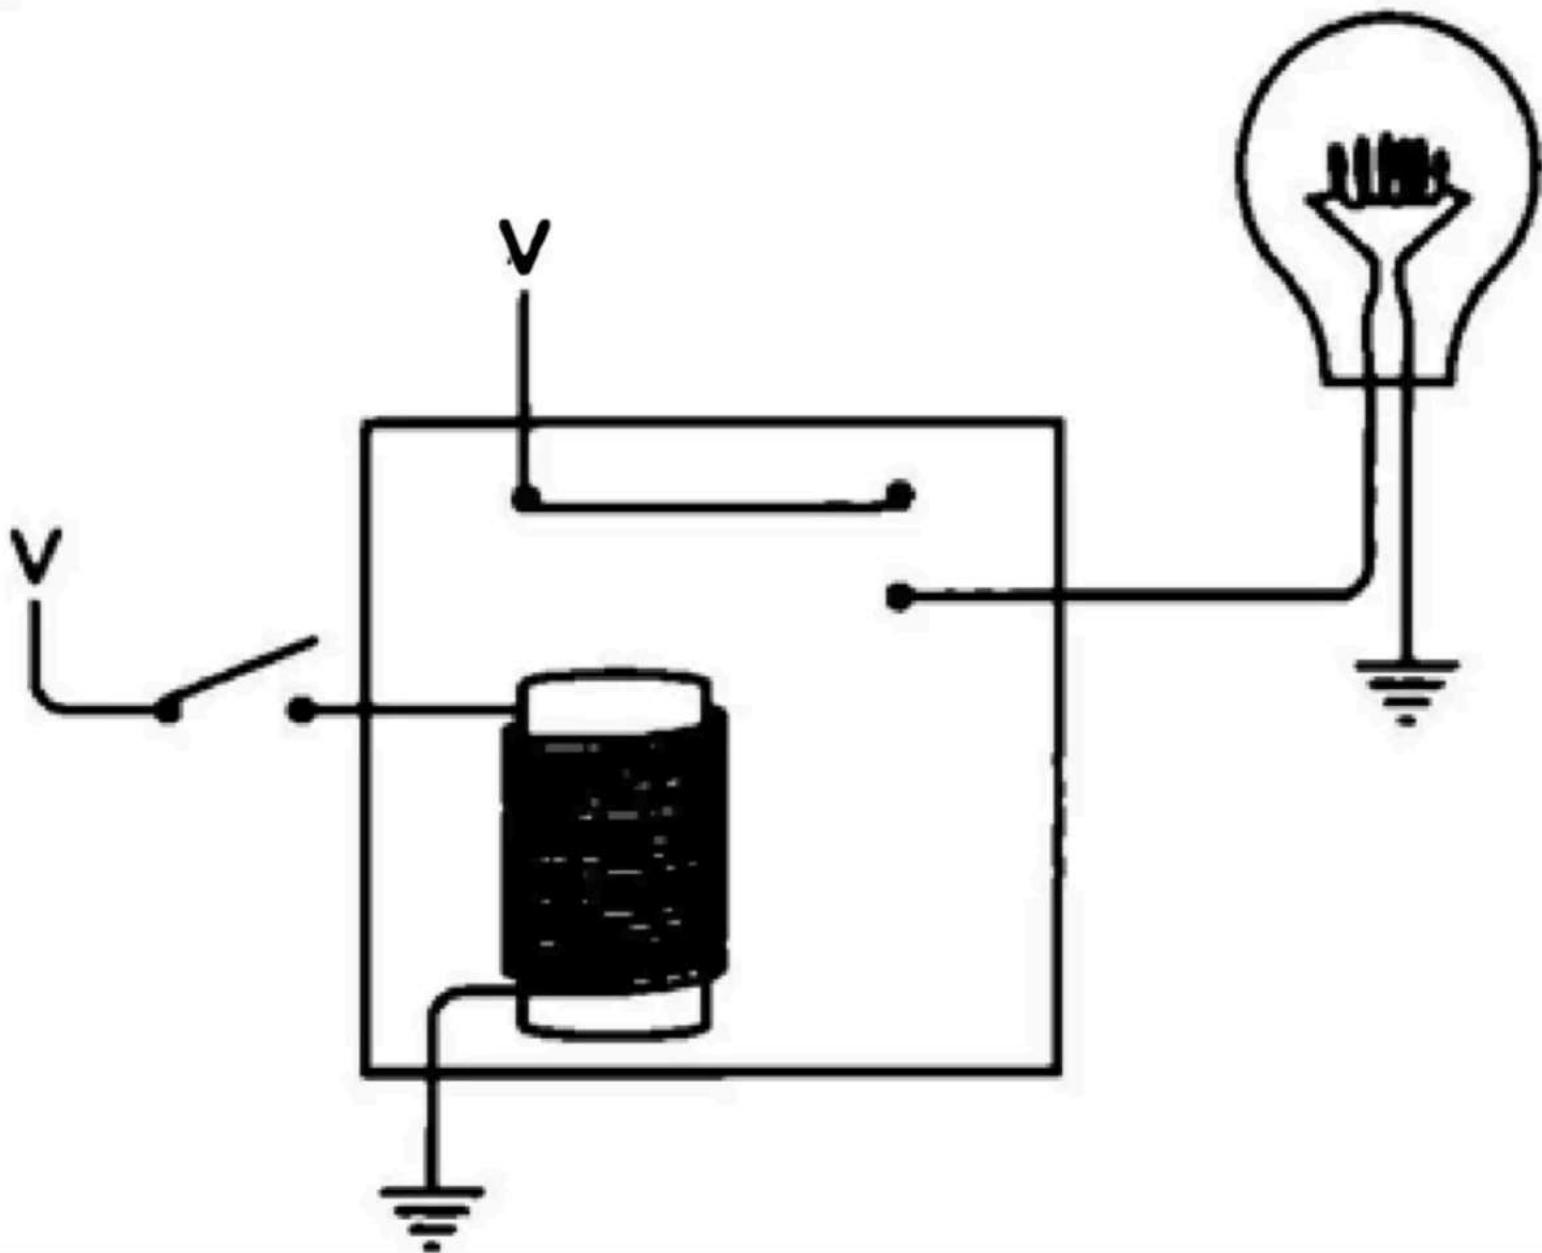












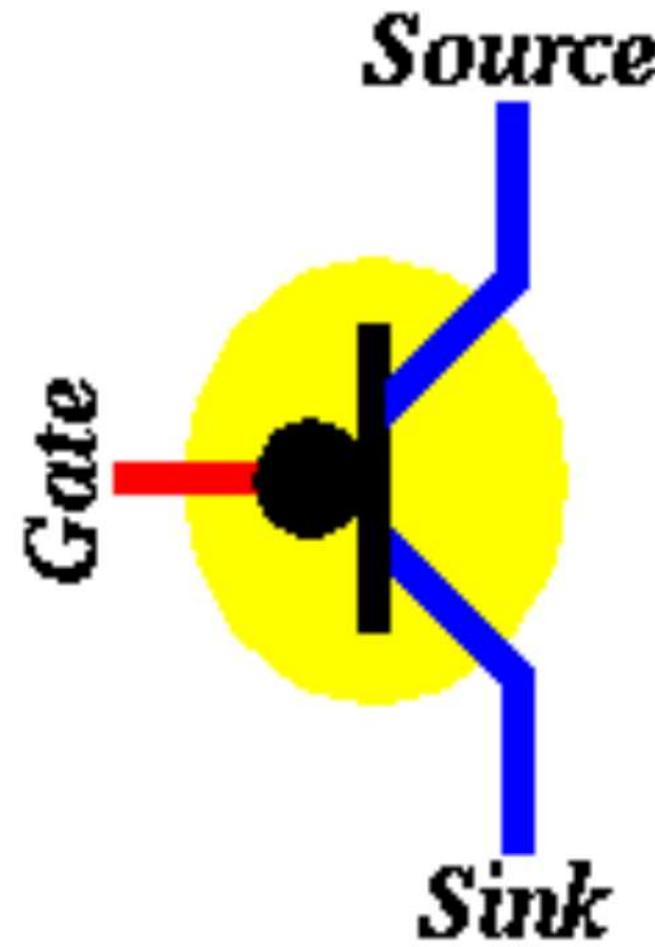


figure 4: Symbol used to depict a complementary transistor

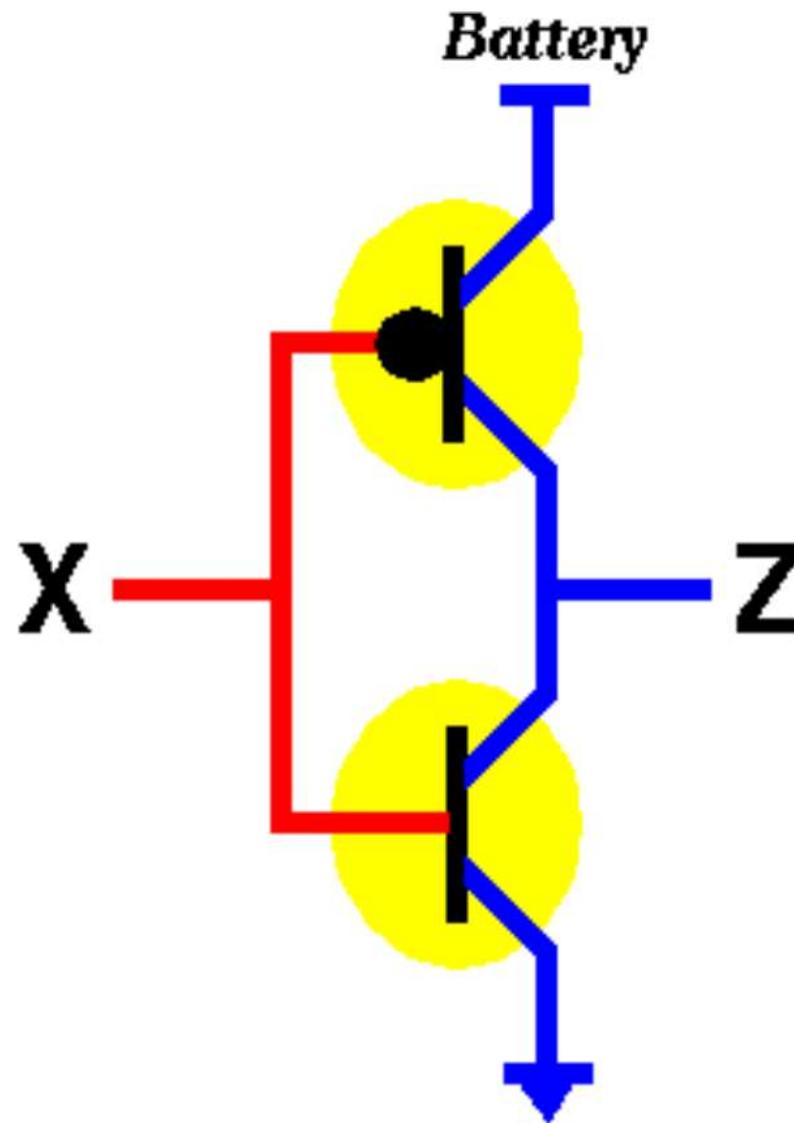


Figure 5: An inverter circuit (known as the NOT gate).

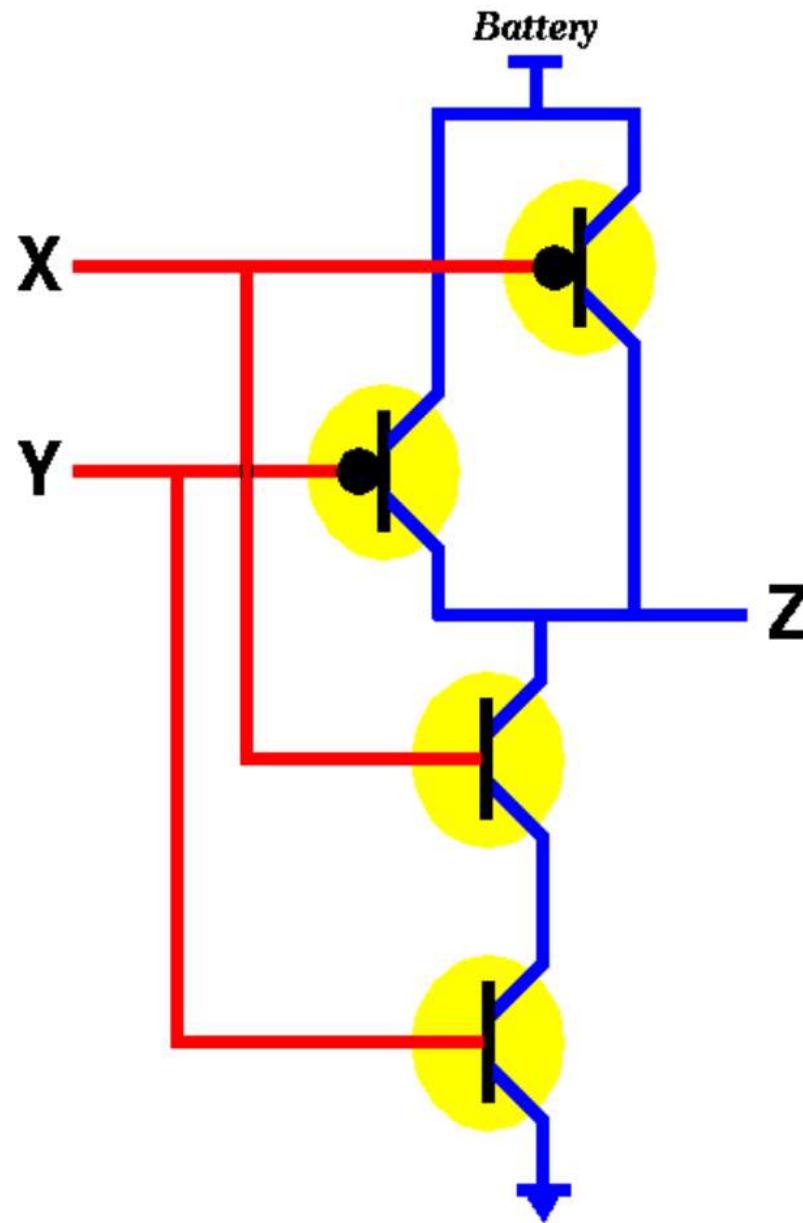
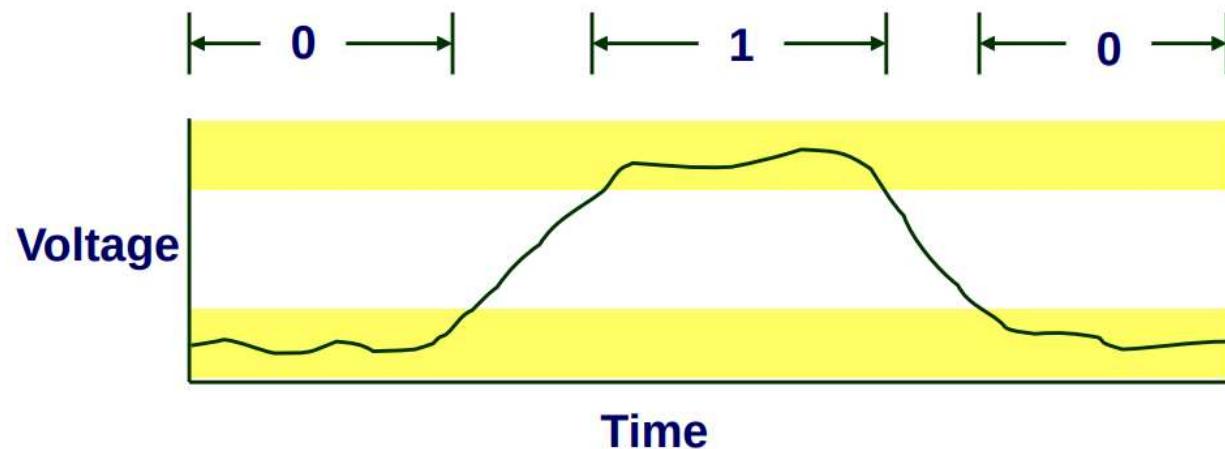


Figure 7: The "not both X and Y" circuit (the NAND function).

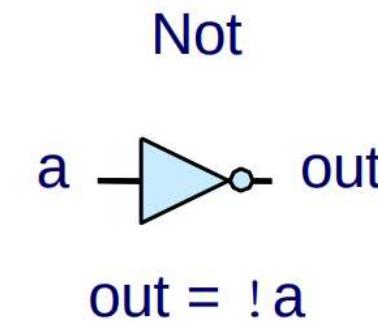
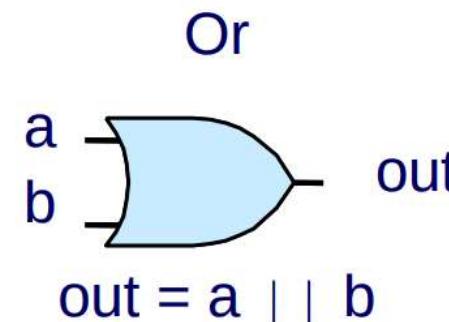
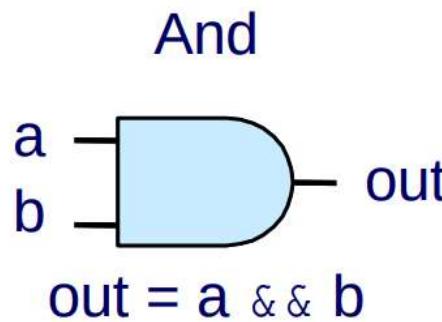
ITU CPH

Digital Signals

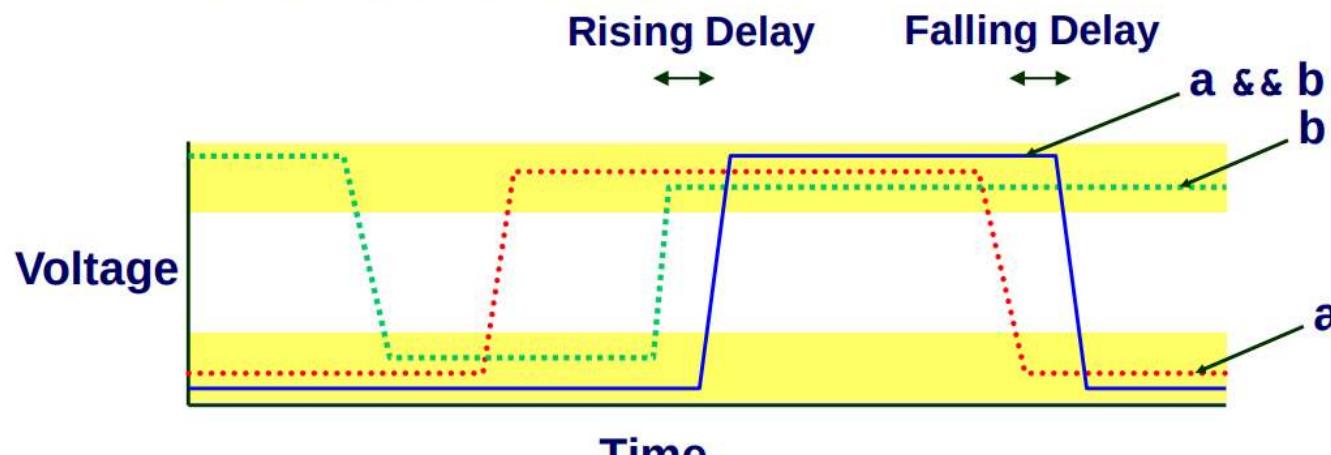


- Use voltage thresholds to extract discrete values from continuous signal
- Simplest version: 1-bit signal
 - Either high range (1) or low range (0)
 - With guard range between them
- Not strongly affected by noise or low quality circuit elements
 - Can make circuits simple, small, and fast

Computing with Logic Gates



- Outputs are Boolean functions of inputs
- Respond continuously to changes in inputs
 - With some, small delay

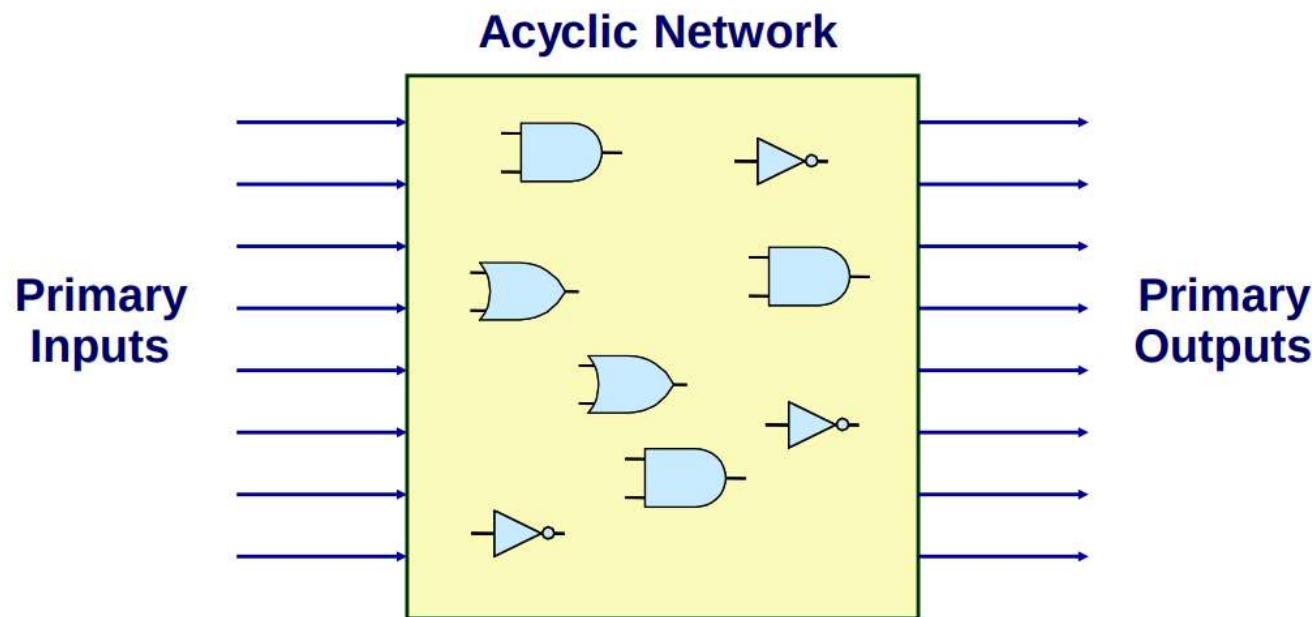


Today

- Digital Logic
 - Combinatorial Logic Circuits (functions)
 - Sequential Logic Circuits (state)
- Processor: Y86-64
 - Design (ISA), Semantics, Implementation (briefly)
- Stack
 - Procedure calls
 - Buffer overflow

Combinatorial Logic Circuits

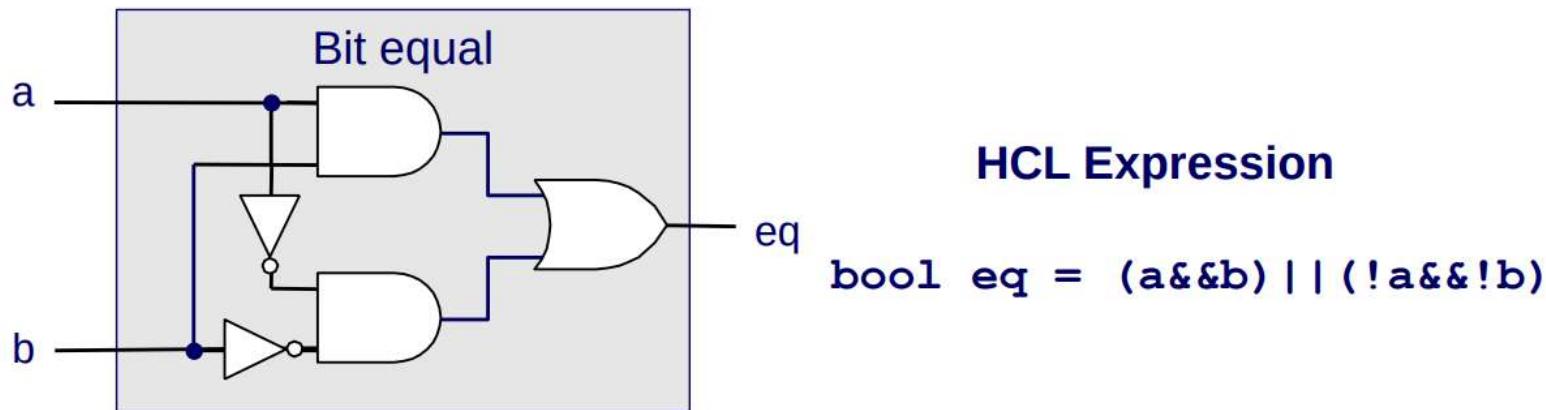
Combinational Circuits



Acyclic Network of Logic Gates

- Continuously responds to changes on primary inputs
- Primary outputs become (after some delay) Boolean functions of primary inputs

Bit Equality

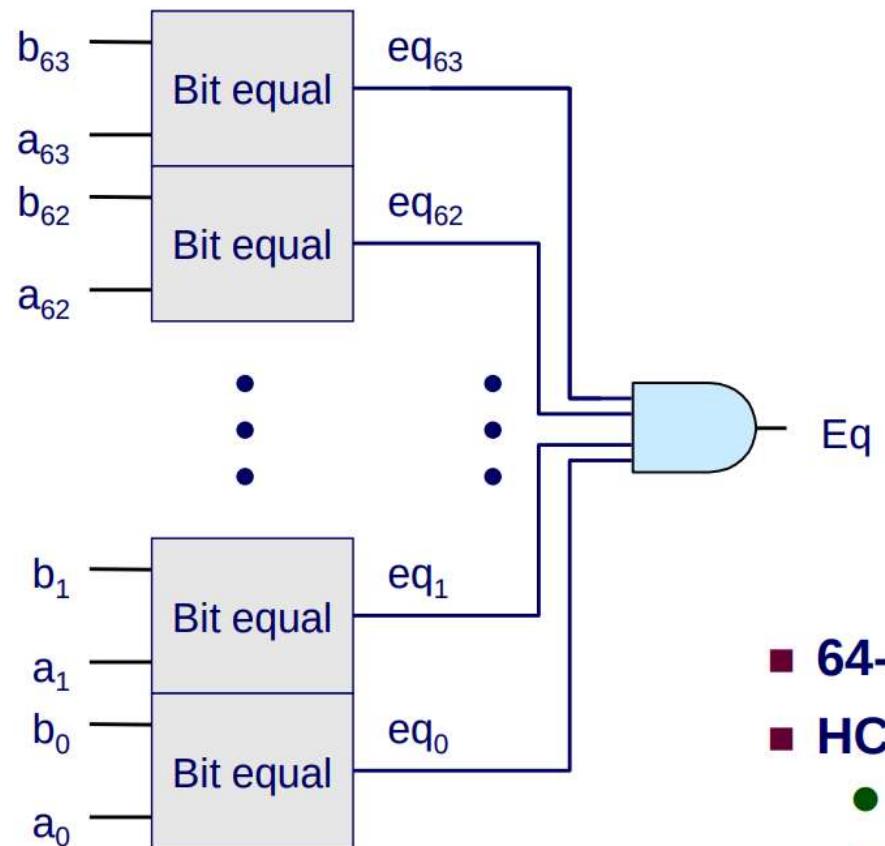


- Generate 1 if *a* and *b* are equal

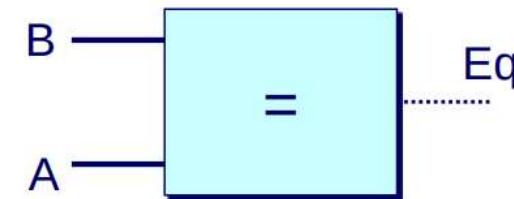
Hardware Control Language (HCL)

- Very simple hardware description language
 - Boolean operations have syntax similar to C logical operations
- We'll use it to describe control logic for processors

Word Equality



Word-Level Representation

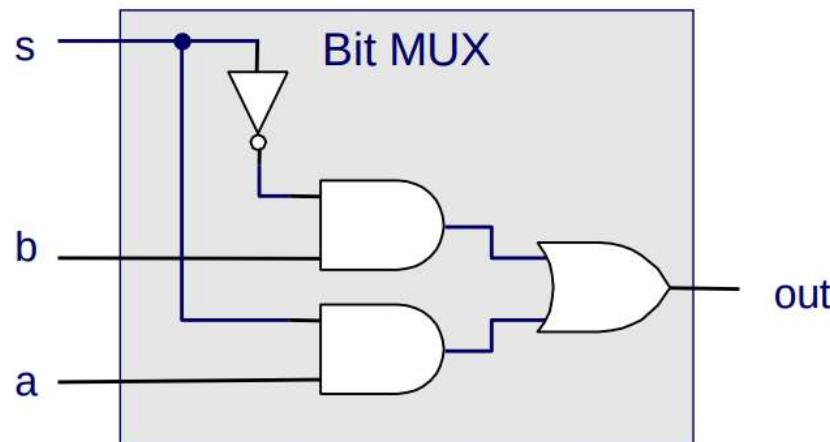


HCL Representation

`bool Eq = (A == B)`

- **64-bit word size**
- **HCL representation**
 - Equality operation
 - Generates Boolean value

Bit-Level Multiplexor

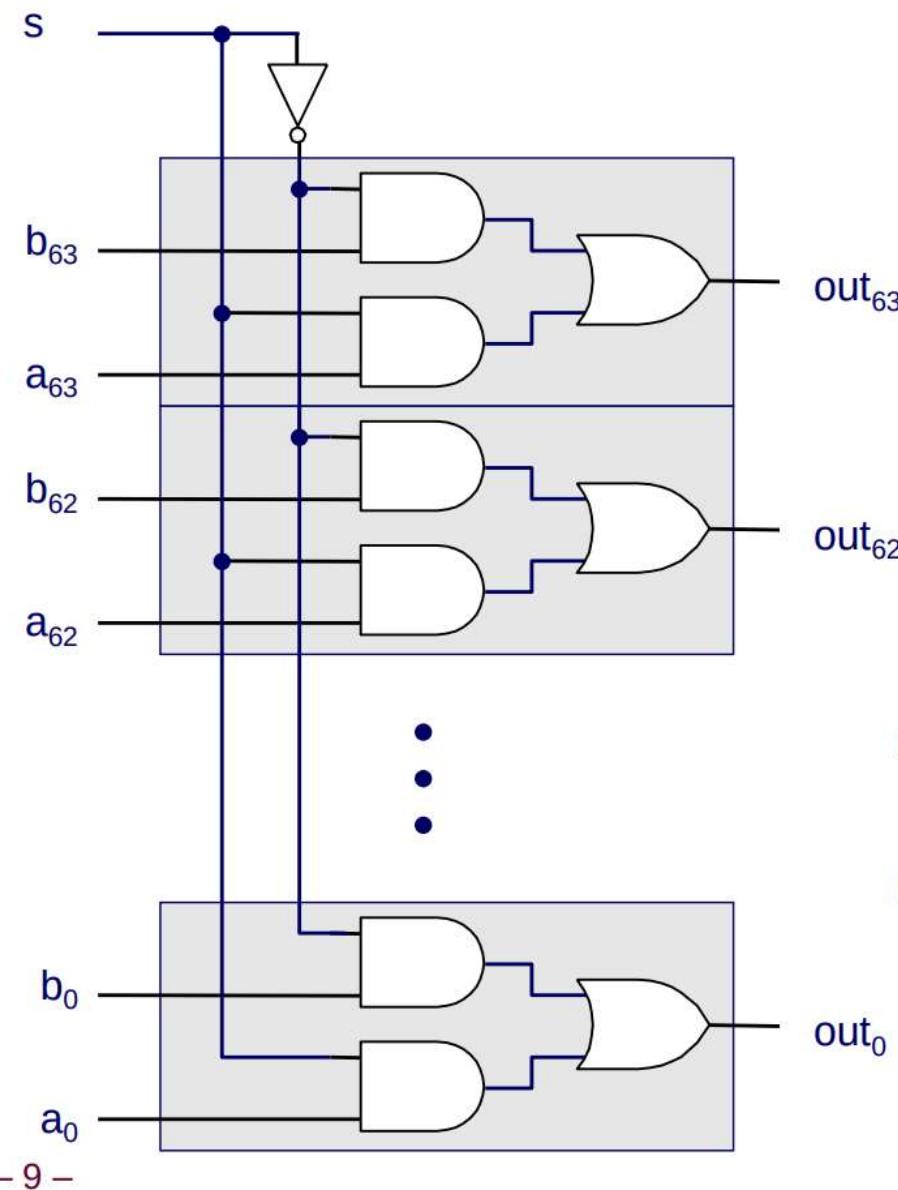


HCL Expression

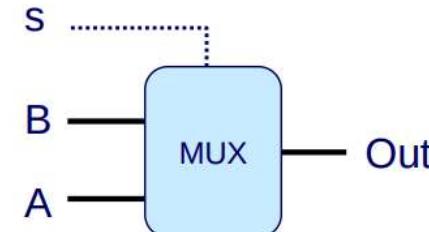
```
bool out = (s&&a) || (!s&&b)
```

- Control signal s
- Data signals a and b
- Output a when $s=1$, b when $s=0$

Word Multiplexor



Word-Level Representation



HCL Representation

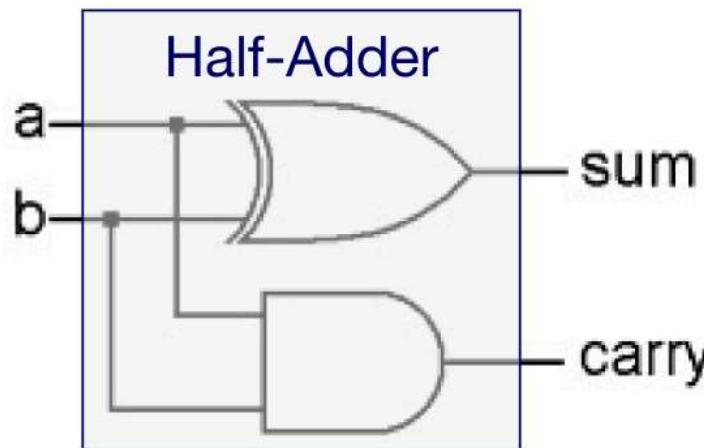
```
int Out = [
    s : A;
    1 : B;
];
```

- Select input word A or B depending on control signal s
- HCL representation
 - Case expression
 - Series of test : value pairs
 - Output value for first successful test

CS:APP3e

ITU CPH

Adder, Half-

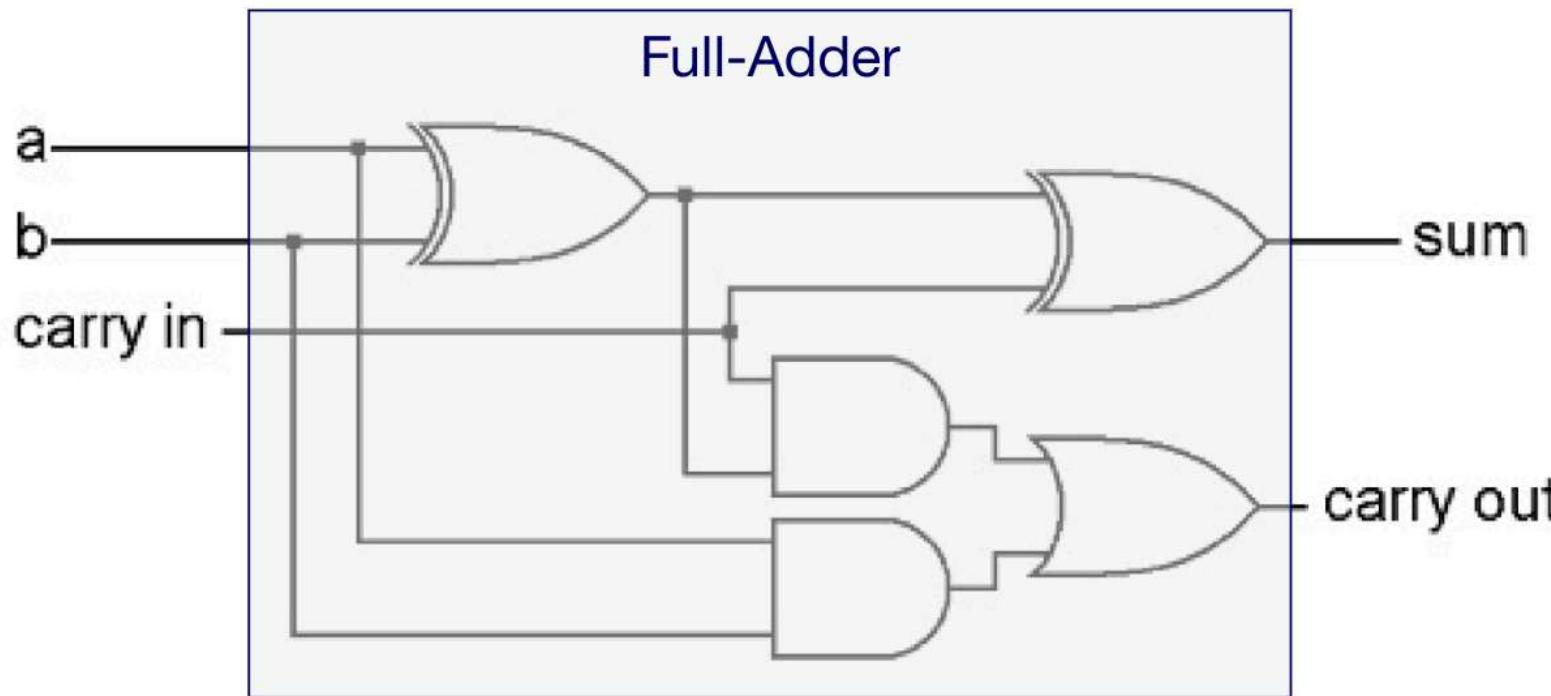


HCL Expression

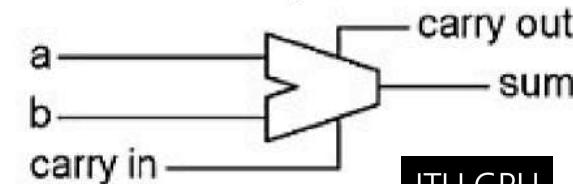
```
bool sum = (a^^b)  
bool carry = a&&b
```

- Output binary-sum to **sum**, carry-bit to **carry**.

Adder, Full-

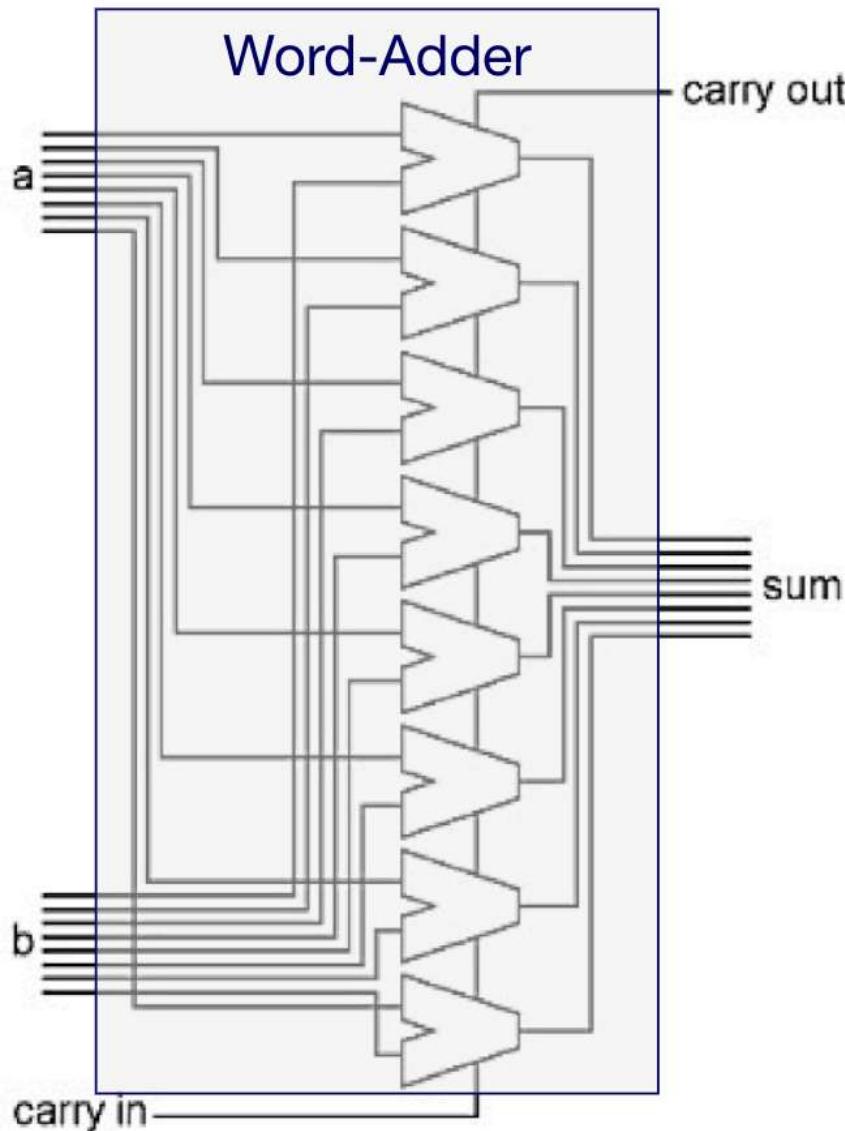


- A Half-adder with a carry-in bit.



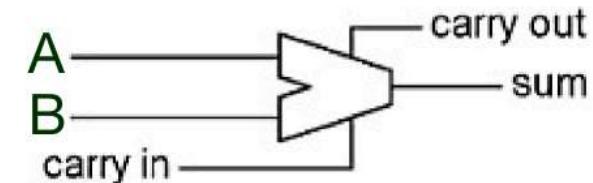
ITU CPH

Adder, Word-

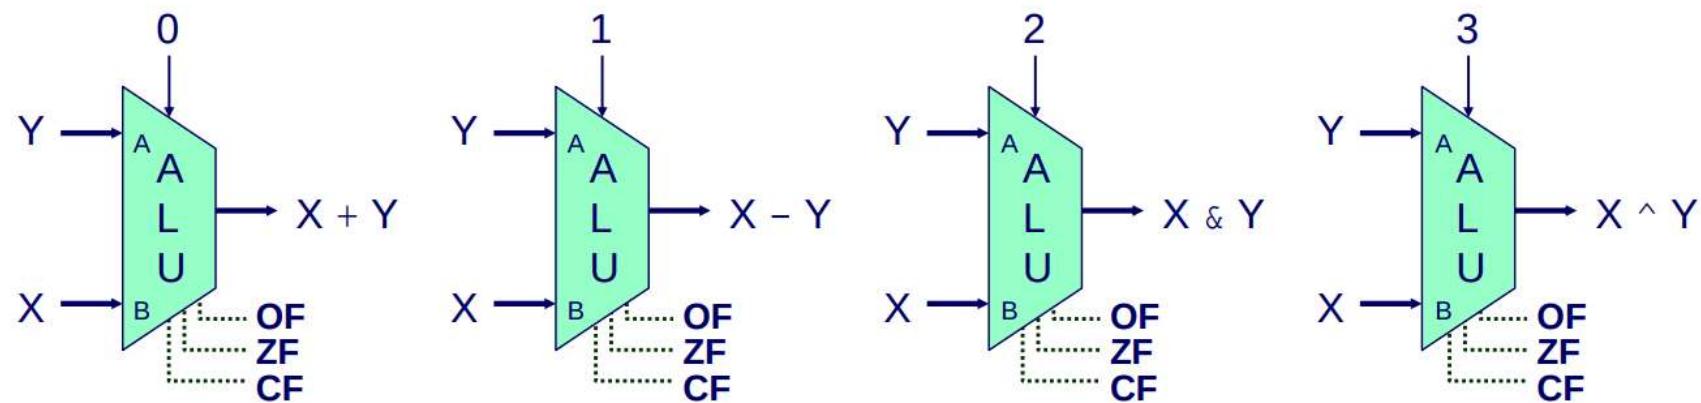


- Chain the 1-bit adders, carry-out to carry-in, to get a word-adder.

**Word-Level
Representation**



Arithmetic Logic Unit

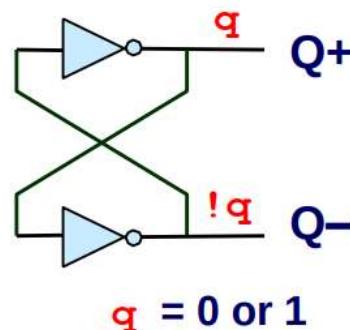


- Combinational logic
 - Continuously responding to inputs
- Control signal selects function computed
 - Corresponding to 4 arithmetic/logical operations in Y86-64
- Also computes values for condition codes

Sequential Logic Circuits

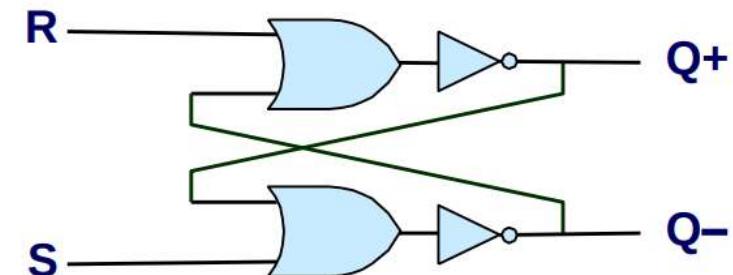
Storing and Accessing 1 Bit

Bistable Element

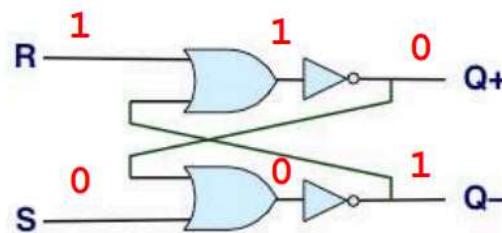


$q = 0 \text{ or } 1$

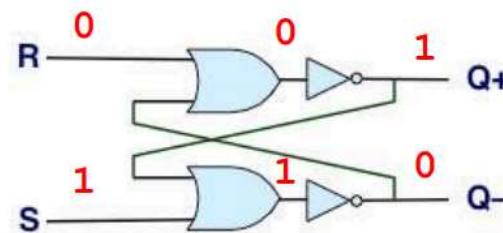
R-S Latch



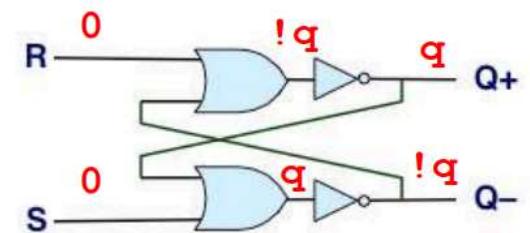
Resetting



Setting

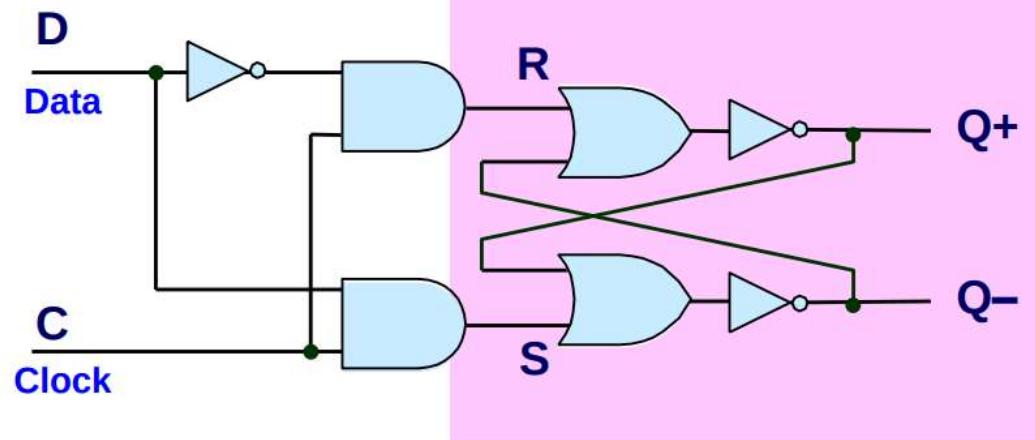


Storing

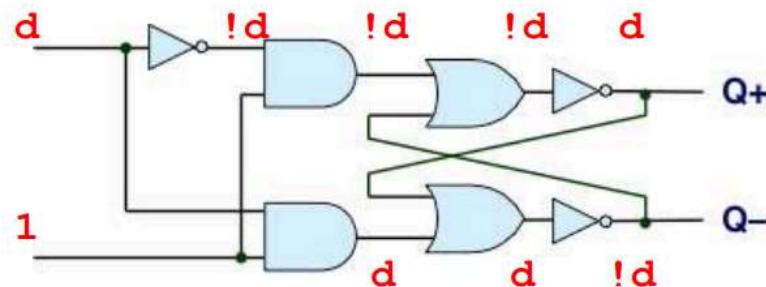


1-Bit Latch

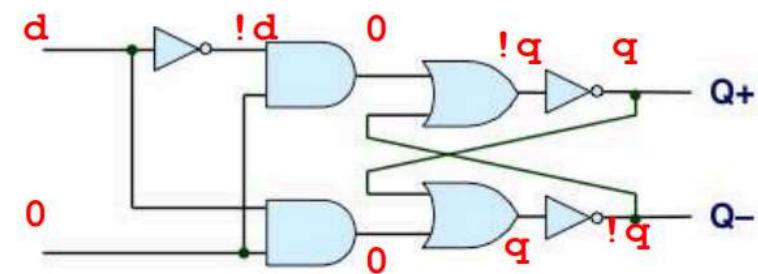
D Latch



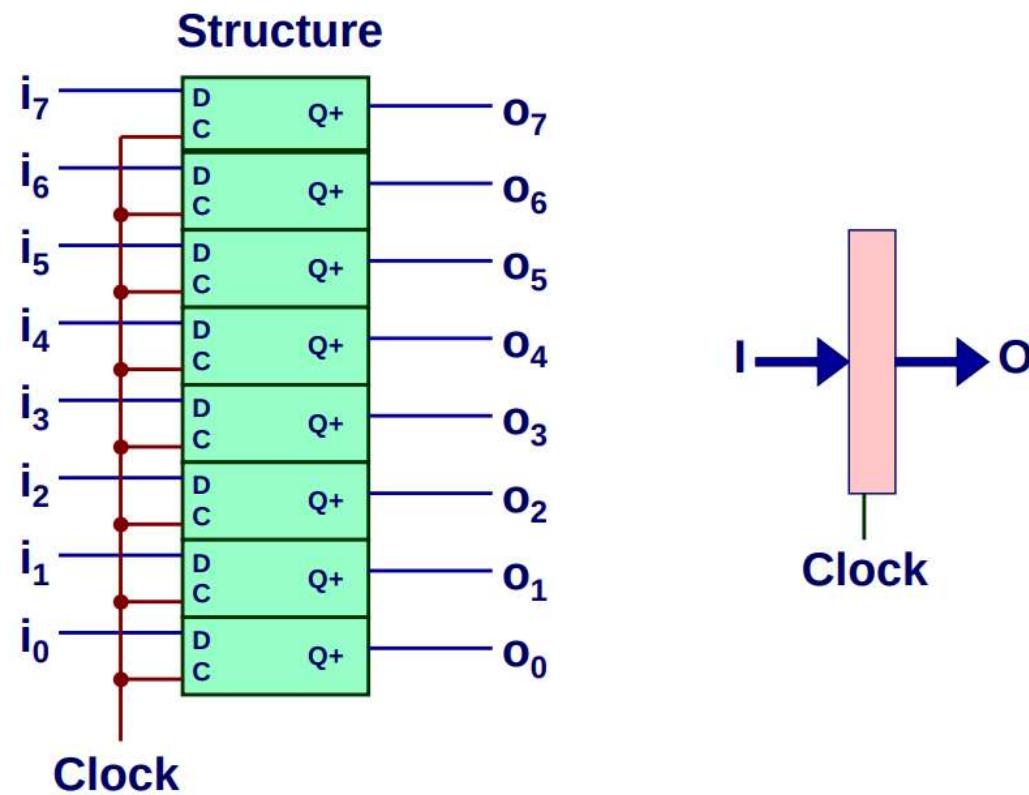
Latching



Storing



Registers



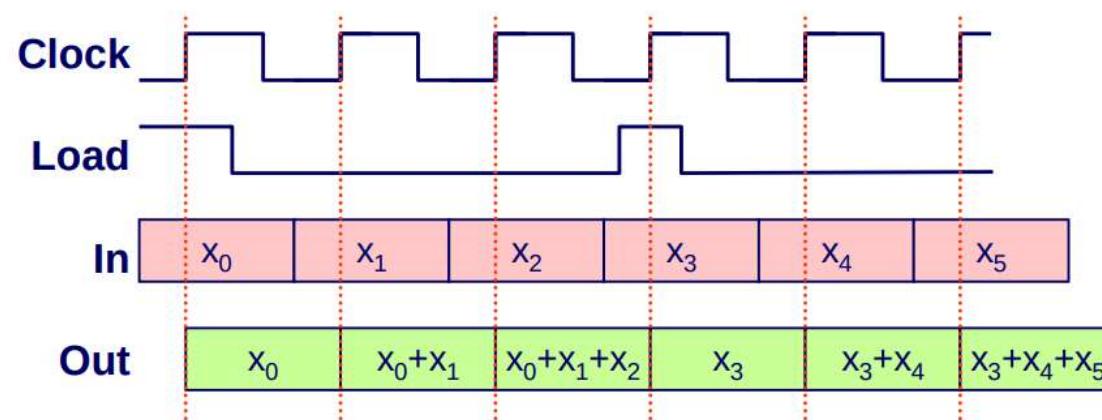
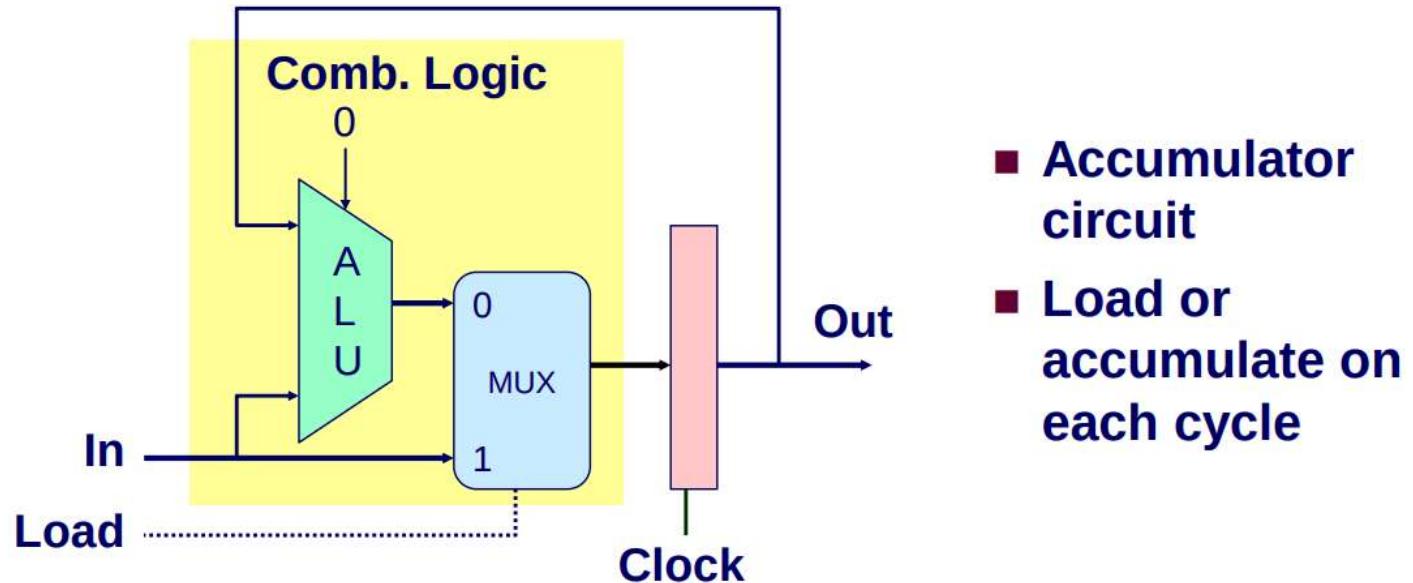
- Stores word of data
 - Different from *program registers* seen in assembly code
- Collection of edge-triggered latches
- Loads input on rising edge of clock

Register Operation

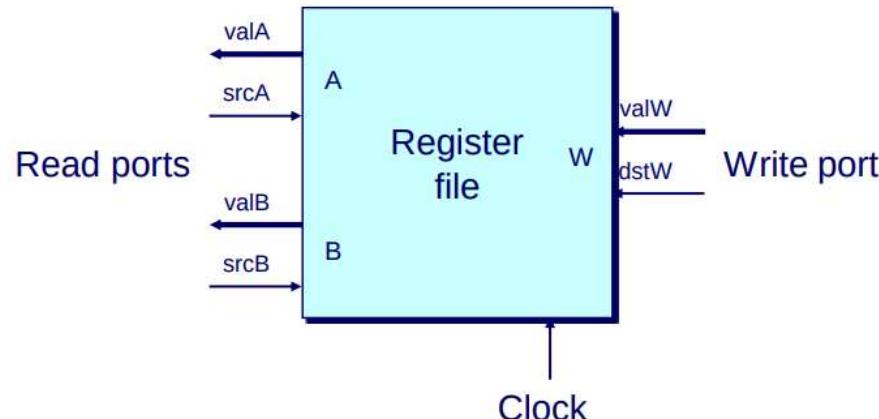


- Stores data bits
- For most of time acts as barrier between input and output
- As clock rises, loads input

State Machine Example



Random-Access Memory



- Stores multiple words of memory
 - Address input specifies which word to read or write
- Register file
 - Holds values of program registers
 - %rax, %rsp, etc.
 - Register identifier serves as address
 - » ID 15 (0xF) implies no read or write performed
- Multiple Ports
 - Can read and/or write multiple words in one cycle
 - » Each has separate address and data input/output

Summary

Computation

- Performed by combinational logic
- Computes Boolean functions
- Continuously reacts to input changes

Storage

- Registers
 - Hold single words
 - Loaded as clock rises
- Random-access memories
 - Hold multiple words
 - Possible multiple read or write ports
 - Read word when address input changes
 - Write word as clock rises

A Processor: Y86-64 Design (ISA)

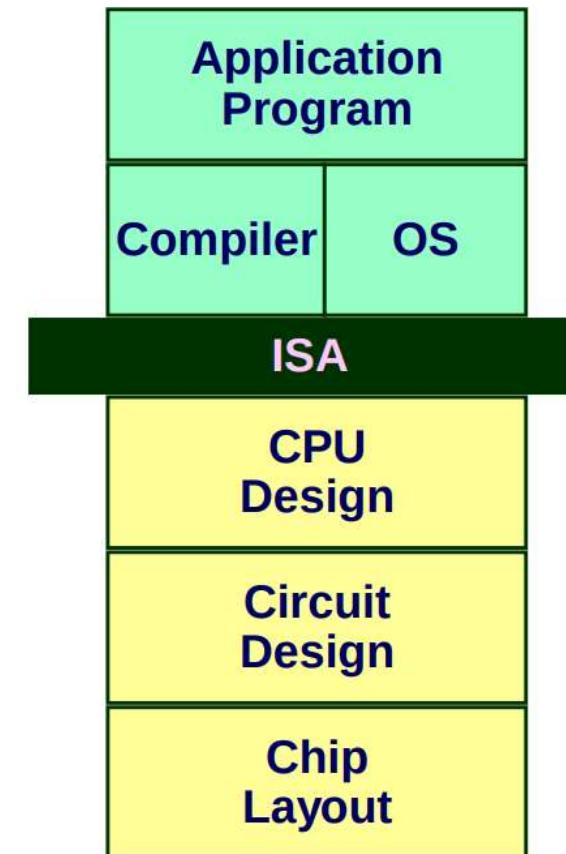
Instruction Set Architecture

Assembly Language View

- Processor state
 - Registers, memory, ...
- Instructions
 - `addq, pushq, ret, ...`
 - How instructions are encoded as bytes

Layer of Abstraction

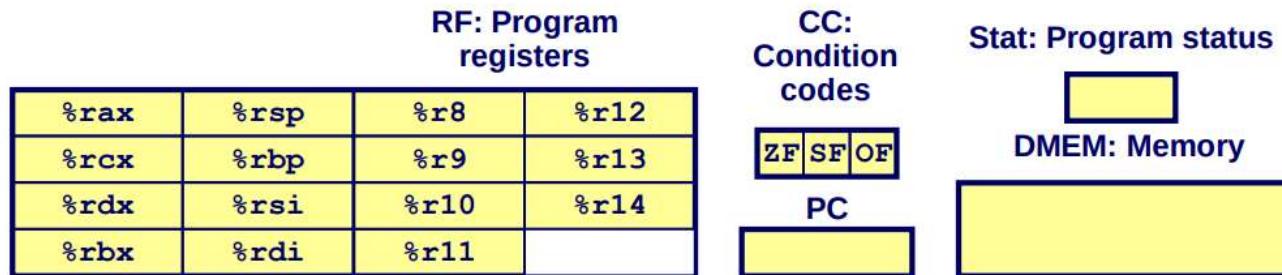
- Above: how to program machine
 - Processor executes instructions in a sequence
- Below: what needs to be built
 - Use variety of tricks to make it run fast
 - E.g., execute multiple instructions simultaneously



- 2 -

CS:APP3e

Y86-64 Processor State



■ Program Registers

- 15 registers (omit %r15). Each 64 bits

■ Condition Codes

- Single-bit flags set by arithmetic or logical instructions

» ZF: Zero

SF:Negative

OF: Overflow

■ Program Counter

- Indicates address of next instruction

■ Program Status

- Indicates either normal operation or some error condition

■ Memory

- Byte-addressable storage array
- Words stored in little-endian byte order

Y86-64 Instructions

Format

- 1–10 bytes of information read from memory
 - Can determine instruction length from first byte
 - Not as many instruction types, and simpler encoding than with x86-64
- Each accesses and modifies some part(s) of the program state

Y86-64 Instruction Set #1

Byte	0	1	2	3	4	5	6	7	8	9
halt	0	0								
nop	1	0								
cmoveXX rA, rB	2	fn	rA	rB						
irmovq V, rB	3	0	F	rB					V	
rmmovq rA, D(rB)	4	0	rA	rB					D	
mrmovq D(rB), rA	5	0	rA	rB					D	
OPq rA, rB	6	fn	rA	rB						
jXX Dest	7	fn							Dest	
call Dest	8	0							Dest	
ret	9	0								
pushq rA	A	0	rA	F						
popq rA	B	0	rA	F						

Y86-64 Instruction Set #2

Byte	0	1	2	3	4	5	6	
halt	0	0						
nop	1	0						
cmoveXX rA, rB	2	fn	rA	rB				
irmovq V, rB	3	0	F	rB	V			
rmmovq rA, D(rB)	4	0	rA	rB	D			
mrmovq D(rB), rA	5	0	rA	rB	D			
OPq rA, rB	6	fn	rA	rB				
jXX Dest	7	fn			Dest			
call Dest	8	0			Dest			
ret	9	0						
pushq rA	A	0	rA	F				
popq rA	B	0	rA	F				

Y86-64 Instruction Set #3

Byte	0	1	2	3	4	5	6	7	8	9
halt	0	0								
nop	1	0								
cmoveXX rA, rB	2	fn	rA	rB						
irmovq V, rB	3	0	F	rB	V					
rmmovq rA, D(rB)	4	0	rA	rB	D					
mrmovq D(rB), rA	5	0	rA	rB	D					
OPq rA, rB	6	fn	rA	rB						
jXX Dest	7	fn			Dest					
call Dest	8	0			Dest					
ret	9	0								
pushq rA	A	0	rA	F						
popq rA	B	0	rA	F						

{ addq 6 0
 subq 6 1
 andq 6 2
 xorq 6 3}

Y86-64 Instruction Set #4

Byte	0	1	2	3	4	5	6	7	
halt	0	0							
nop	1	0							
cmoveXX rA, rB	2	fn	rA	rB					
irmovq V, rB	3	0	F	rB	V				jle
rmmovq rA, D(rB)	4	0	rA	rB	D				je
mrmovq D(rB), rA	5	0	rA	rB	D				jne
OPq rA, rB	6	fn	rA	rB					jge
jXX Dest	7	fn			Dest				jg
call Dest	8	0			Dest				
ret	9	0							
pushq rA	A	0	rA	F					
popq rA	B	0	rA	F					

Encoding Registers

Each register has 4-bit ID

%rax	0
%rcx	1
%rdx	2
%rbx	3
%rsp	4
%rbp	5
%rsi	6
%rdi	7
%r8	8
%r9	9
%r10	A
%r11	B
%r12	C
%r13	D
%r14	E
No Register	F

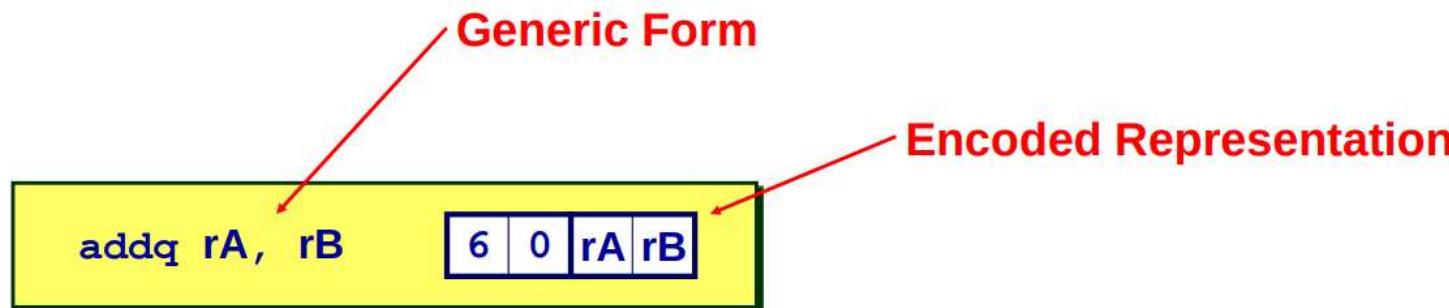
- Same encoding as in x86-64

Register ID 15 (0xF) indicates “no register”

- Will use this in our hardware design in multiple places

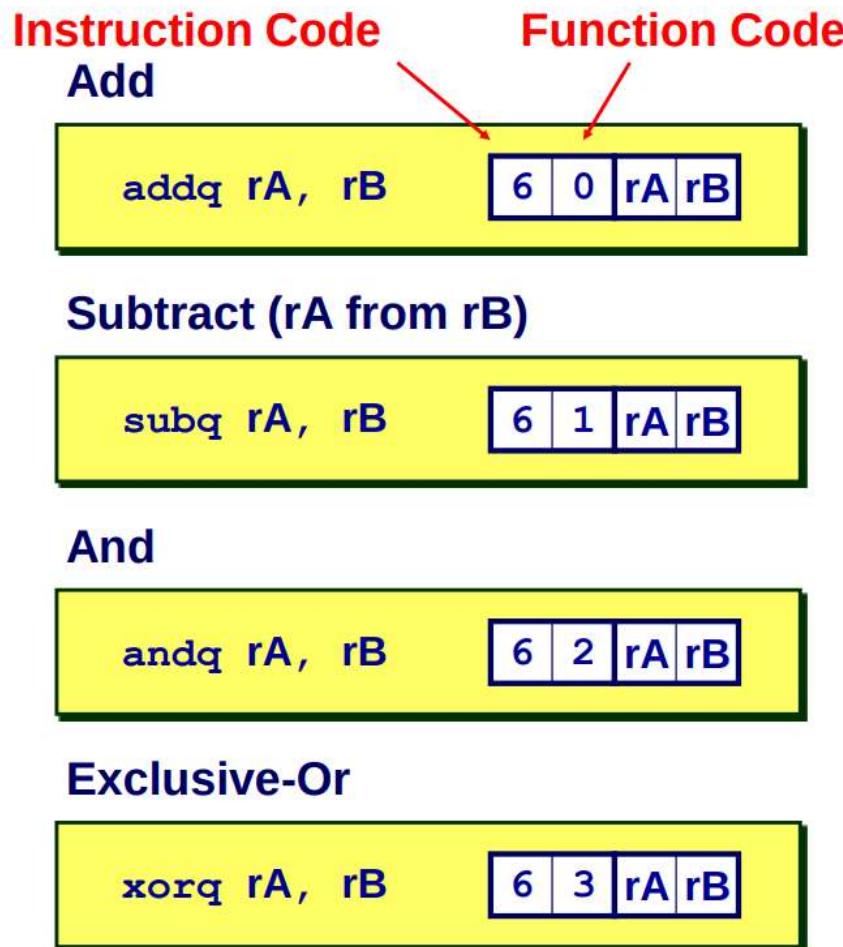
Instruction Example

Addition Instruction



- Add value in register rA to that in register rB
 - Store result in register rB
 - Note that Y86-64 only allows addition to be applied to register data
- Set condition codes based on result
- e.g., `addq %rax,%rsi` Encoding: 60 06
- Two-byte encoding
 - First indicates instruction type
 - Second gives source and destination registers

Arithmetic and Logical Operations



- Refer to generically as “OPq”
- Encodings differ only by “function code”
 - Low-order 4 bytes in first instruction word
- Set condition codes as side effect

Move Operations

Register → Register

`rrmovq rA, rB`

2	0
---	---

Immediate → Register

`irmovq V, rB`

3	0	F	rB	V
---	---	---	----	---

Register → Memory

`rmmovq rA, D(rB)`

4	0	rA	rB	D
---	---	----	----	---

Memory → Register

`mrmovq D(rB), rA`

5	0	rA	rB	D
---	---	----	----	---

- Like the x86-64 `movq` instruction
- Simpler format for memory addresses
- Give different names to keep them distinct

Conditional Move Instructions

Move Unconditionally

`rrmovq rA, rB`

2	0	rA	rB
---	---	----	----

Move When Less or Equal

`cmovele rA, rB`

2	1	rA	rB
---	---	----	----

Move When Less

`cmovl rA, rB`

2	2	rA	rB
---	---	----	----

Move When Equal

`cmove rA, rB`

2	3	rA	rB
---	---	----	----

Move When Not Equal

`cmovne rA, rB`

2	4	rA	rB
---	---	----	----

Move When Greater or Equal

`cmovge rA, rB`

2	5	rA	rB
---	---	----	----

Move When Greater

`cmovg rA, rB`

2	6	rA	rB
---	---	----	----

- Refer to generically as “`cmovXX`”
- Encodings differ only by “function code”
- Based on values of condition codes
- Variants of `rrmovq` instruction
 - (Conditionally) copy value from source to destination register

Jump Instructions

Jump (Conditionally)

jxx Dest	7	fn	Dest
----------	---	----	------

- Refer to generically as “jxx”
- Encodings differ only by “function code” fn
- Based on values of condition codes
- Same as x86-64 counterparts
- Encode full destination address
 - Unlike PC-relative addressing seen in x86-64

Jump Instructions

Jump Unconditionally

jmp Dest

7	0
---	---

 Dest

Jump When Less or Equal

jle Dest

7	1
---	---

 Dest

Jump When Less

jl Dest

7	2
---	---

 Dest

Jump When Equal

je Dest

7	3
---	---

 Dest

Jump When Not Equal

jne Dest

7	4
---	---

 Dest

Jump When Greater or Equal

jge Dest

7	5
---	---

 Dest

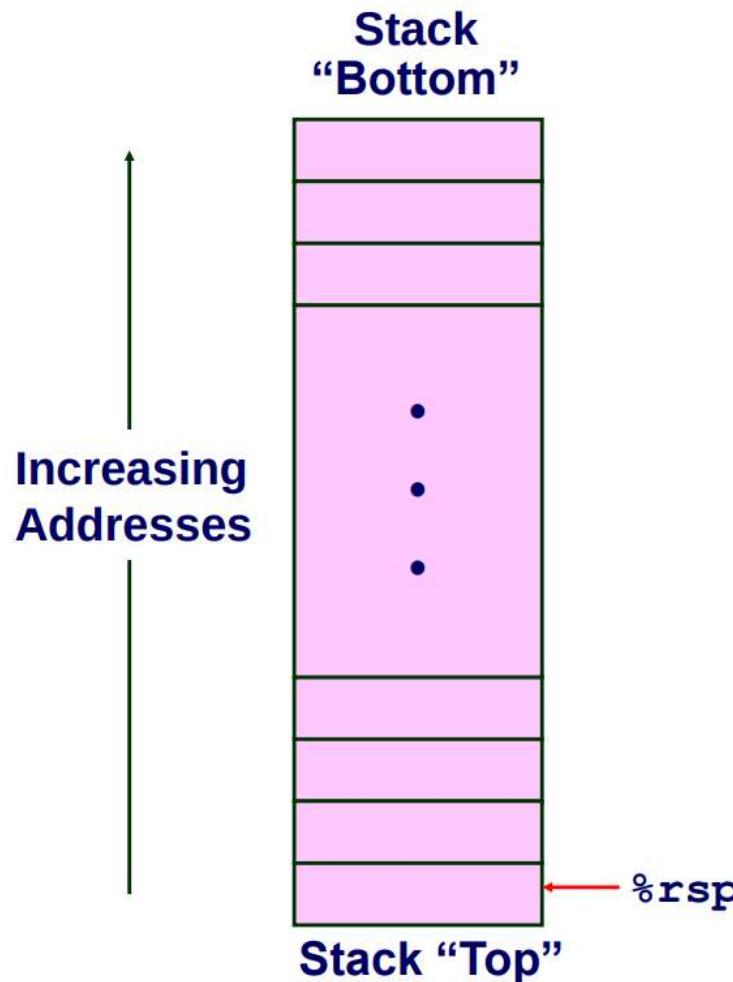
Jump When Greater

jg Dest

7	6
---	---

 Dest

Y86-64 Program Stack



- Region of memory holding program data
- Used in Y86-64 (and x86-64) for supporting procedure calls
- Stack top indicated by `%rsp`
 - Address of top stack element
- Stack grows toward lower addresses
 - Top element is at highest address in the stack
 - When pushing, must first decrement stack pointer
 - After popping, increment stack pointer

Stack Operations

pushq rA



- Decrement %rsp by 8
- Store word from rA to memory at %rsp
- Like x86-64

popq rA



- Read word from memory at %rsp
- Save in rA
- Increment %rsp by 8
- Like x86-64

Subroutine Call and Return

call Dest

8	0
---	---

Dest

- Push address of next instruction onto stack
- Start executing instructions at Dest
- Like x86-64

ret

9	0
---	---

- Pop value from stack
- Use as address for next instruction
- Like x86-64

Miscellaneous Instructions



- Don't do anything



- Stop executing instructions
- x86-64 has comparable instruction, but can't execute it in user mode
- We will use it to stop the simulator
- Encoding ensures that program hitting memory initialized to zero will halt

Status Conditions

Mnemonic	Code
AOK	1

- Normal operation

Mnemonic	Code
HLT	2

- Halt instruction encountered

Mnemonic	Code
ADR	3

- Bad address (either instruction or data) encountered

Mnemonic	Code
INS	4

- Invalid instruction encountered

Desired Behavior

- If AOK, keep going
- Otherwise, stop program execution

CISC Instruction Sets

- Complex Instruction Set Computer
- IA32 is example

Stack-oriented instruction set

- Use stack to pass arguments, save program counter
- Explicit push and pop instructions

Arithmetic instructions can access memory

- `addq %rax, 12(%rbx,%rcx,8)`
 - requires memory read and write
 - Complex address calculation

Condition codes

- Set as side effect of arithmetic and logical instructions

Philosophy

- Add instructions to perform “typical” programming tasks

RISC Instruction Sets

- Reduced Instruction Set Computer
- Internal project at IBM, later popularized by Hennessy (Stanford) and Patterson (Berkeley)

Fewer, simpler instructions

- Might take more to get given task done
- Can execute them with small and fast hardware

Register-oriented instruction set

- Many more (typically 32) registers
- Use for arguments, return pointer, temporaries

Only load and store instructions can access memory

- Similar to Y86-64 `mrmovq` and `rmmovq`

No Condition codes

- Test instructions return 0/1 in register

CISC vs. RISC

Original Debate

- Strong opinions!
- CISC proponents---easy for compiler, fewer code bytes
- RISC proponents---better for optimizing compilers, can make run fast with simple chip design

Current Status

- For desktop processors, choice of ISA not a technical issue
 - With enough hardware, can make anything run fast
 - Code compatibility more important
- x86-64 adopted many RISC features
 - More registers; use them for argument passing
- For embedded processors, RISC makes sense
 - Smaller, cheaper, less power
 - Most cell phones use ARM processor

Summary

Y86-64 Instruction Set Architecture

- Similar state and instructions as x86-64
- Simpler encodings
- Somewhere between CISC and RISC

How Important is ISA Design?

- Less now than before
 - With enough hardware, can make almost anything go fast

A Processor: Y86-64 Semantics (SEQ)

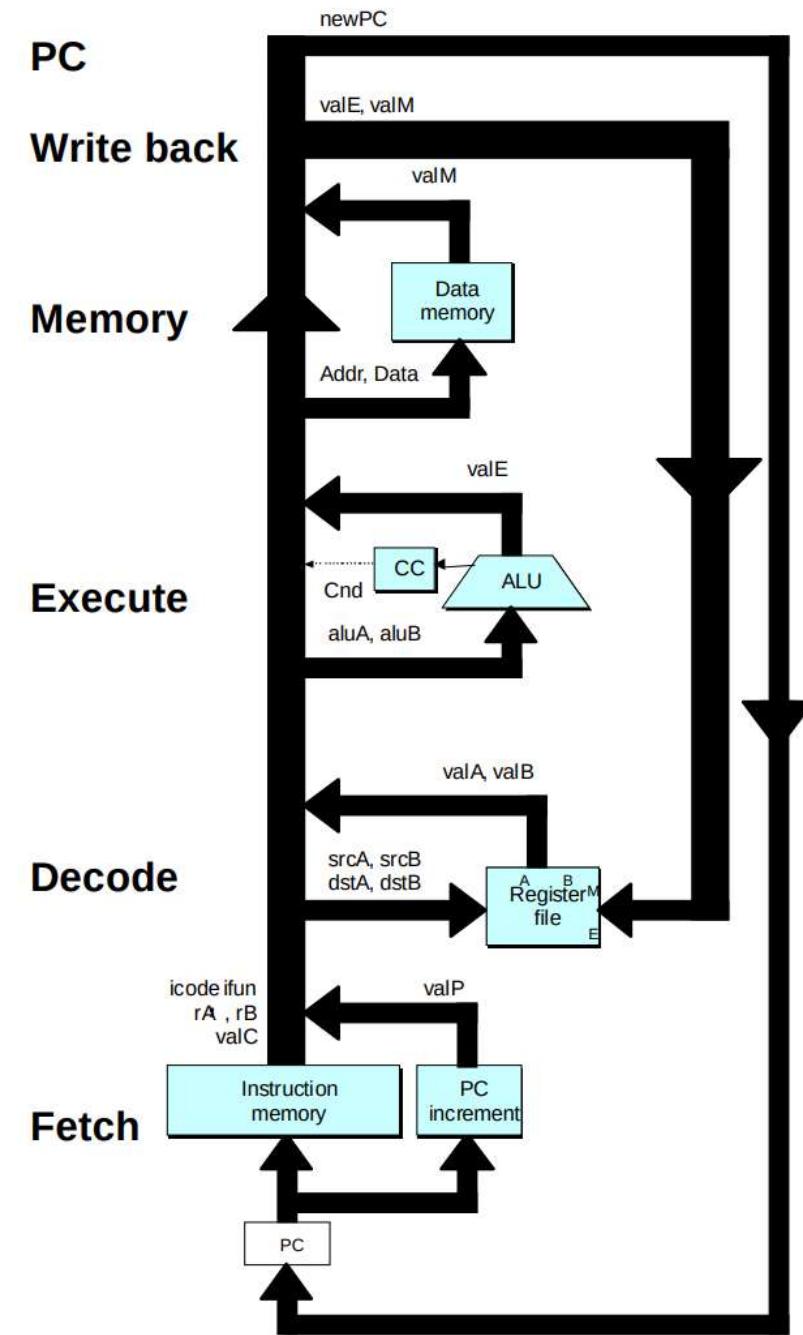
SEQ State & Flow

State

- Program counter register (PC)
- Condition code register (CC)
- Register File
- Memories
 - Access same memory space
 - Data: for reading/writing program data
 - Instruction: for reading instructions

Instruction Flow

- Read instruction at address specified by PC
- Process through stages
- Update program counter



– 10 –

SEQ Stages

Fetch

- Read instruction from instruction memory

Decode

- Read program registers

Execute

- Compute value or address

Memory

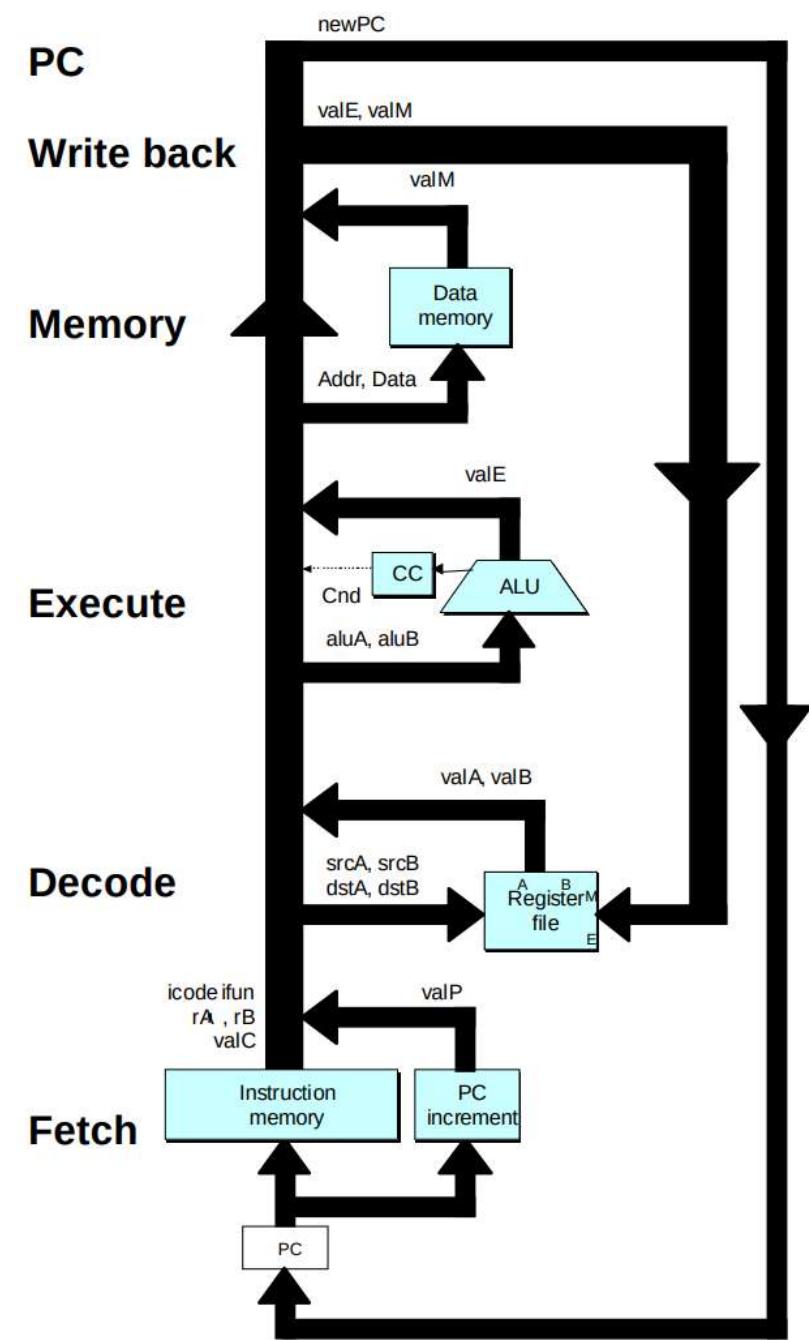
- Read or write data

Write Back

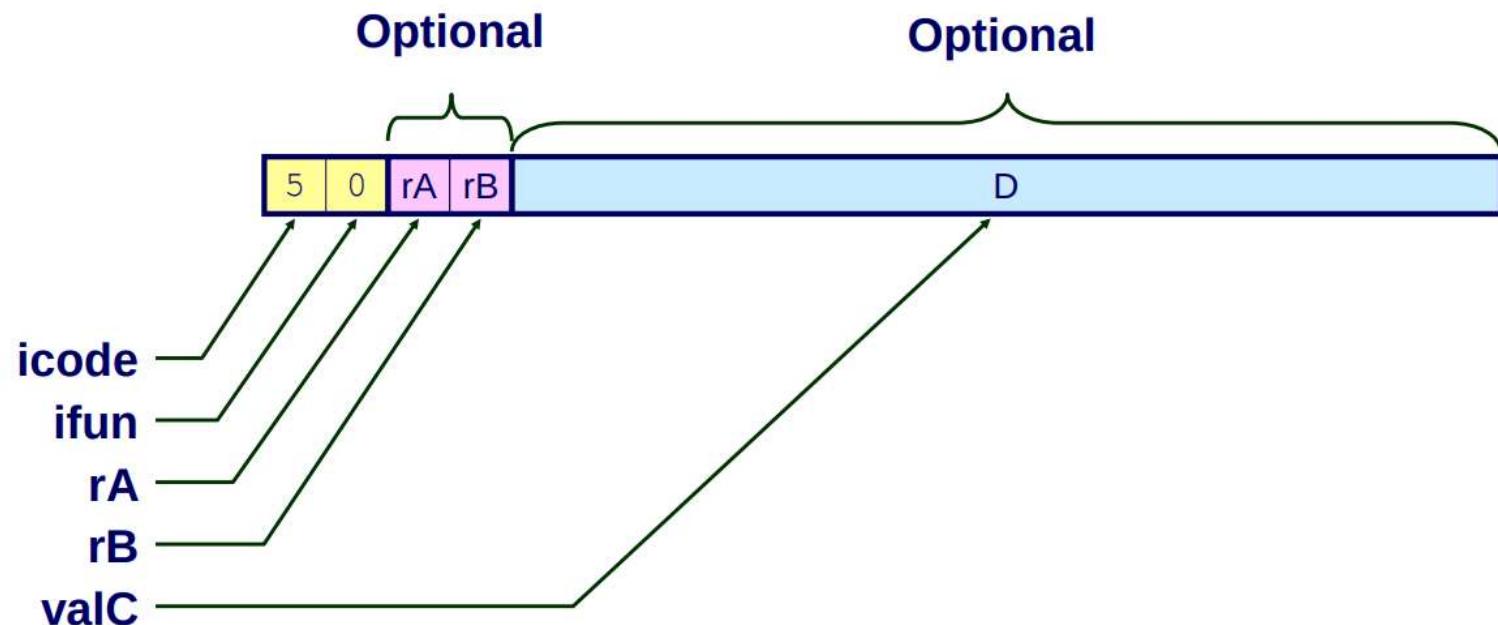
- Write program registers

PC

- Update program counter



Instruction Decoding



Instruction Format

- **Instruction byte** **icode:ifun**
- **Optional register byte** **rA:rB**
- **Optional constant word** **valC**

Executing Arith./Logical Operation



Fetch

- Read 2 bytes

Decode

- Read operand registers

Execute

- Perform operation
- Set condition codes

Memory

- Do nothing

Write back

- Update register

PC Update

- Increment PC by 2

Stage Computation: Arith/Log. Ops

	$OPq\ rA, rB$	
Fetch	$icode:ifun \leftarrow M_1[PC]$ $rA:rB \leftarrow M_1[PC+1]$ $valP \leftarrow PC+2$	Read instruction byte Read register byte Compute next PC
Decode	$valA \leftarrow R[rA]$ $valB \leftarrow R[rB]$	Read operand A Read operand B
Execute	$valE \leftarrow valB\ OP\ valA$ Set CC	Perform ALU operation Set condition code register
Memory		
Write back	$R[rB] \leftarrow valE$	Write back result
PC update	$PC \leftarrow valP$	Update PC

- Formulate instruction execution as sequence of simple steps
- Use same general form for all instructions

Executing `rmmovq`

`rmmovq rA, D(rB)`

4	0	rA	rB
---	---	----	----

D

Fetch

- Read 10 bytes

Decode

- Read operand registers

Execute

- Compute effective address

Memory

- Write to memory

Write back

- Do nothing

PC Update

- Increment PC by 10

Stage Computation: `rmmovq`

<code>rmmovq rA, D(rB)</code>	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{rA:rB} \leftarrow M_1[\text{PC+1}]$ $\text{valC} \leftarrow M_8[\text{PC+2}]$ $\text{valP} \leftarrow \text{PC+10}$
Decode	$\text{valA} \leftarrow R[\text{rA}]$ $\text{valB} \leftarrow R[\text{rB}]$
Execute	$\text{valE} \leftarrow \text{valB} + \text{valC}$
Memory	$M_8[\text{valE}] \leftarrow \text{valA}$
Write back	
PC update	$\text{PC} \leftarrow \text{valP}$

- Use ALU for address computation

Executing popq

popq rA

b	0	rA	8
---	---	----	---

Fetch

- Read 2 bytes

Decode

- Read stack pointer

Execute

- Increment stack pointer by 8

Memory

- Read from old stack pointer

Write back

- Update stack pointer
- Write result to register

PC Update

- Increment PC by 2

Stage Computation: popq

	<code>popq rA</code>	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{rA:rB} \leftarrow M_1[\text{PC+1}]$ $\text{valP} \leftarrow \text{PC+2}$	Read instruction byte Read register byte Compute next PC
Decode	$\text{valA} \leftarrow R[\%rsp]$ $\text{valB} \leftarrow R[\%rsp]$	Read stack pointer Read stack pointer
Execute	$\text{valE} \leftarrow \text{valB} + 8$	Increment stack pointer
Memory	$\text{valM} \leftarrow M_8[\text{valA}]$	Read from stack
Write back	$R[\%rsp] \leftarrow \text{valE}$ $R[rA] \leftarrow \text{valM}$	Update stack pointer Write back result
PC update	$\text{PC} \leftarrow \text{valP}$	Update PC

- Use ALU to increment stack pointer
- Must update two registers
 - Popped value
 - New stack pointer

Executing Conditional Moves

cmovxx rA, rB 2 fn rA rB

Fetch

- Read 2 bytes

Decode

- Read operand registers

Execute

- If !cnd, then set destination register to 0xF

Memory

- Do nothing

Write back

- Update register (or not)

PC Update

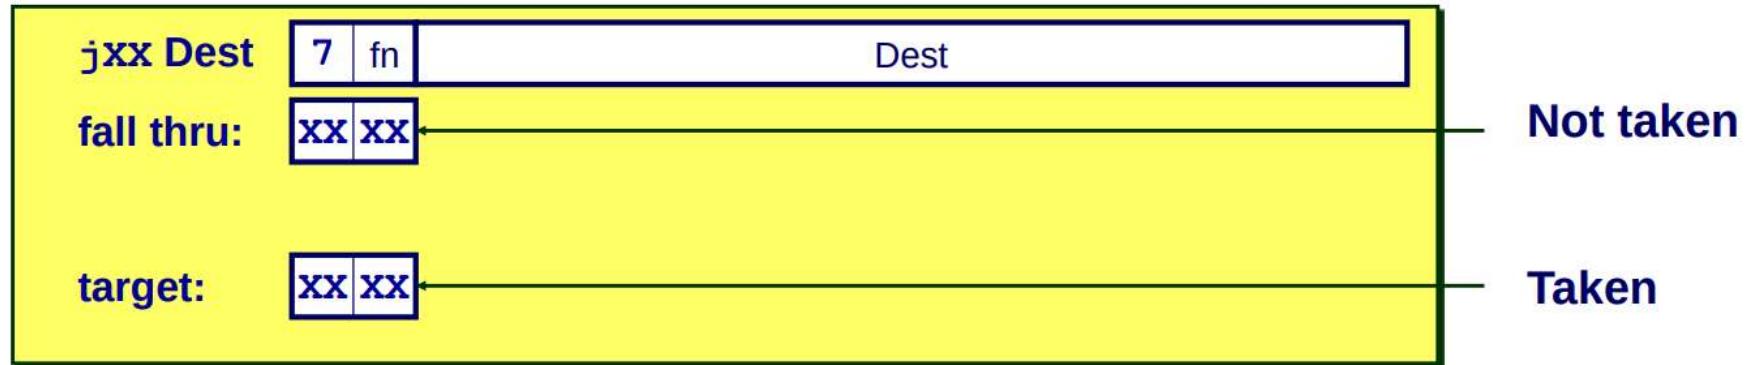
- Increment PC by 2

Stage Computation: Cond. Move

	<code>cmoveXX rA, rB</code>	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $rA:rB \leftarrow M_1[\text{PC+1}]$ $\text{valP} \leftarrow \text{PC+2}$	Read instruction byte Read register byte Compute next PC
Decode	$\text{valA} \leftarrow R[rA]$ $\text{valB} \leftarrow 0$	Read operand A
Execute	$\text{valE} \leftarrow \text{valB} + \text{valA}$ $\text{If } ! \text{Cond(CC,ifun)} \text{ } rB \leftarrow 0xF$	Pass valA through ALU (Disable register update)
Memory		
Write back	$R[rB] \leftarrow \text{valE}$	Write back result
PC update	$\text{PC} \leftarrow \text{valP}$	Update PC

- Read register rA and pass through ALU
- Cancel move by setting destination register to 0xF
 - If condition codes & move condition indicate no move

Executing Jumps



Fetch

- Read 9 bytes
- Increment PC by 9

Decode

- Do nothing

Execute

- Determine whether to take branch based on jump condition and condition codes

Memory

- Do nothing

Write back

- Do nothing

PC Update

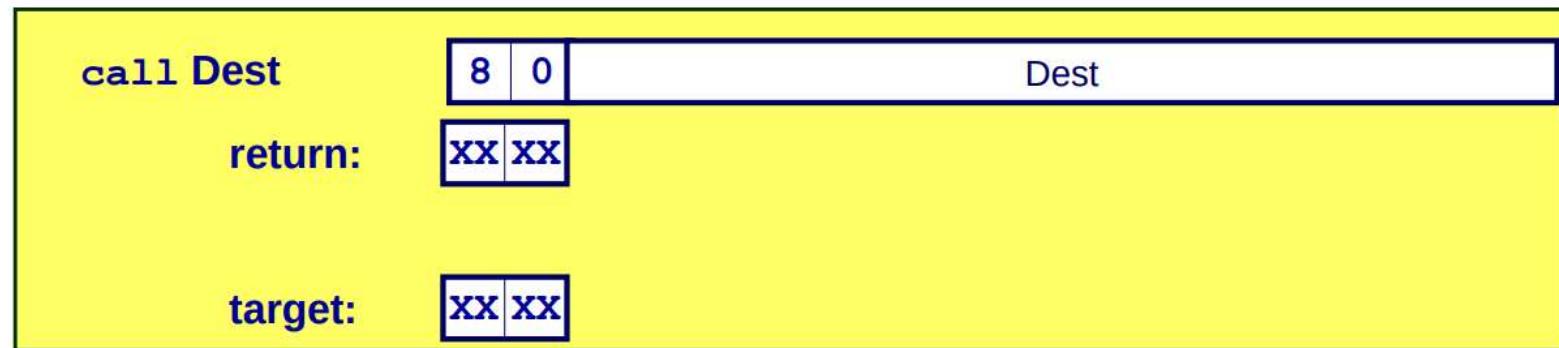
- Set PC to Dest if branch taken or to incremented PC if not branch

Stage Computation: Jumps

	jXX Dest	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{valC} \leftarrow M_8[\text{PC}+1]$ $\text{valP} \leftarrow \text{PC}+9$	Read instruction byte Read destination address Fall through address
Decode		
Execute	$\text{Cnd} \leftarrow \text{Cond}(\text{CC}, \text{ifun})$	Take branch?
Memory		
Write back		
PC update	$\text{PC} \leftarrow \text{Cnd} ? \text{valC} : \text{valP}$	Update PC

- Compute both addresses
- Choose based on setting of condition codes and branch condition

Executing call



Fetch

- Read 9 bytes
- Increment PC by 9

Decode

- Read stack pointer

Execute

- Decrement stack pointer by 8

Memory

- Write incremented PC to new value of stack pointer

Write back

- Update stack pointer

PC Update

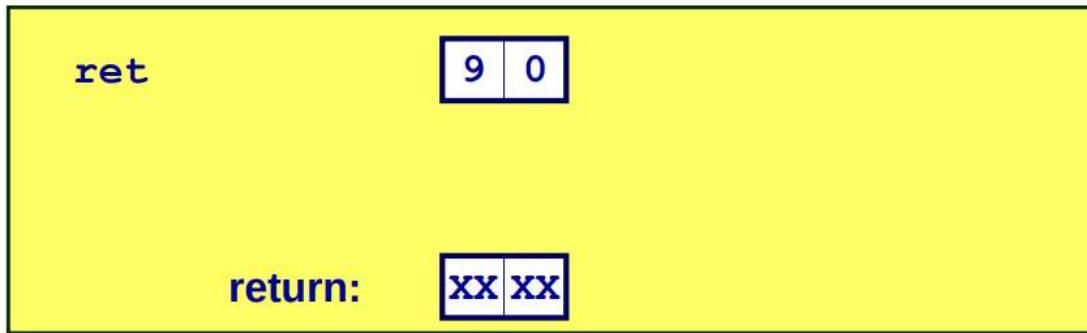
- Set PC to Dest

Stage Computation: call

	call Dest	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$ $\text{valC} \leftarrow M_8[\text{PC}+1]$ $\text{valP} \leftarrow \text{PC}+9$	Read instruction byte Read destination address Compute return point
Decode	$\text{valB} \leftarrow R[\%rsp]$	Read stack pointer
Execute	$\text{valE} \leftarrow \text{valB} + -8$	Decrement stack pointer
Memory	$M_8[\text{valE}] \leftarrow \text{valP}$	Write return value on stack
Write back	$R[\%rsp] \leftarrow \text{valE}$	Update stack pointer
PC update	$\text{PC} \leftarrow \text{valC}$	Set PC to destination

- Use ALU to decrement stack pointer
- Store incremented PC

Executing ret



Fetch

- Read 1 byte

Decode

- Read stack pointer

Execute

- Increment stack pointer by 8

Memory

- Read return address from old stack pointer

Write back

- Update stack pointer

PC Update

- Set PC to return address

Stage Computation: `ret`

<code>ret</code>	
Fetch	$\text{icode:ifun} \leftarrow M_1[\text{PC}]$
Decode	$\text{valA} \leftarrow R[\%rsp]$ $\text{valB} \leftarrow R[\%rsp]$
Execute	$\text{valE} \leftarrow \text{valB} + 8$
Memory	$\text{valM} \leftarrow M_8[\text{valA}]$
Write back	$R[\%rsp] \leftarrow \text{valE}$
PC update	$\text{PC} \leftarrow \text{valM}$

Read instruction byte
Read operand stack pointer
Read operand stack pointer
Increment stack pointer
Read return address
Update stack pointer
Set PC to return address

- Use ALU to increment stack pointer
- Read return address from memory

Computation Steps

		OPq rA, rB	
Fetch	icode,ifun	icode:ifun $\leftarrow M_1[PC]$	Read instruction byte
	rA,rB	rA:rB $\leftarrow M_1[PC+1]$	Read register byte
	valC		[Read constant word]
	valP	valP $\leftarrow PC+2$	Compute next PC
Decode	valA, srcA	valA $\leftarrow R[rA]$	Read operand A
	valB, srcB	valB $\leftarrow R[rB]$	Read operand B
Execute	valE	valE $\leftarrow valB \text{ OP } valA$	Perform ALU operation
	Cond code	Set CC	Set/use cond. code reg
Memory	valM		[Memory read/write]
Write back	dstE	R[rB] $\leftarrow valE$	Write back ALU result
	dstM		[Write back memory result]
PC update	PC	PC $\leftarrow valP$	Update PC

- All instructions follow same general pattern
- Differ in what gets computed on each step

Computation Steps

		call Dest	
Fetch	icode,ifun	$icode:ifun \leftarrow M_1[PC]$	Read instruction byte [Read register byte]
	rA,rB		Read constant word
	valC	$valC \leftarrow M_8[PC+1]$	Compute next PC
	valP	$valP \leftarrow PC+9$	
Decode	valA, srcA		[Read operand A]
	valB, srcB	$valB \leftarrow R[%rsp]$	Read operand B
Execute	valE	$valE \leftarrow valB + -8$	Perform ALU operation
	Cond code		[Set /use cond. code reg]
Memory	valM	$M_8[valE] \leftarrow valP$	Memory read/write
Write back	dstE	$R[%rsp] \leftarrow valE$	Write back ALU result
	dstM		[Write back memory result]
PC update	PC	$PC \leftarrow valC$	Update PC

- All instructions follow same general pattern
- Differ in what gets computed on each step

Computed Values

Fetch

icode	Instruction code
ifun	Instruction function
rA	Instr. Register A
rB	Instr. Register B
valC	Instruction constant
valP	Incremented PC

Execute

- **valE** ALU result
- **Cnd** Branch/move flag

Memory

- **valM** Value from memory

Decode

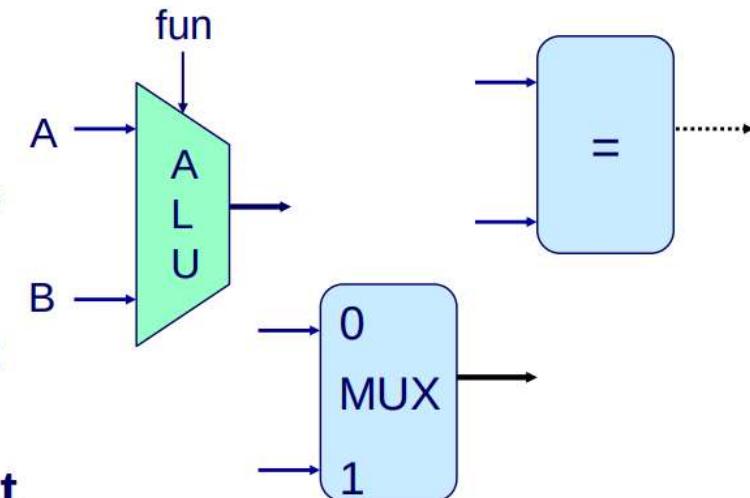
srcA	Register ID A
srcB	Register ID B
dstE	Destination Register E
dstM	Destination Register M
valA	Register value A
valB	Register value B

A Processor: Y86-64 Implementation (SEQ)

Building Blocks

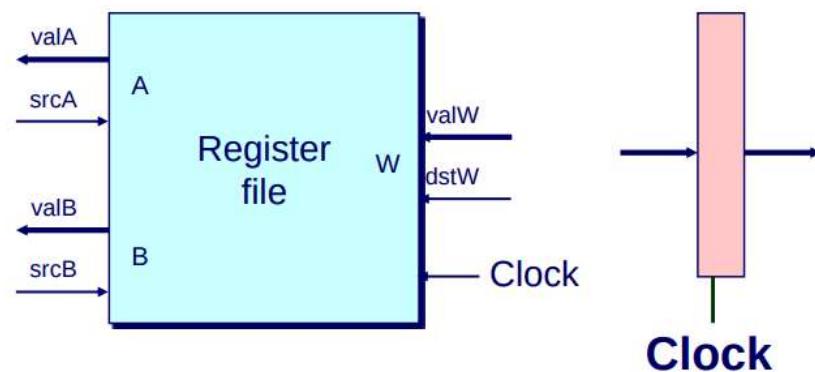
Combinational Logic

- Compute Boolean functions of inputs
- Continuously respond to input changes
- Operate on data and implement control



Storage Elements

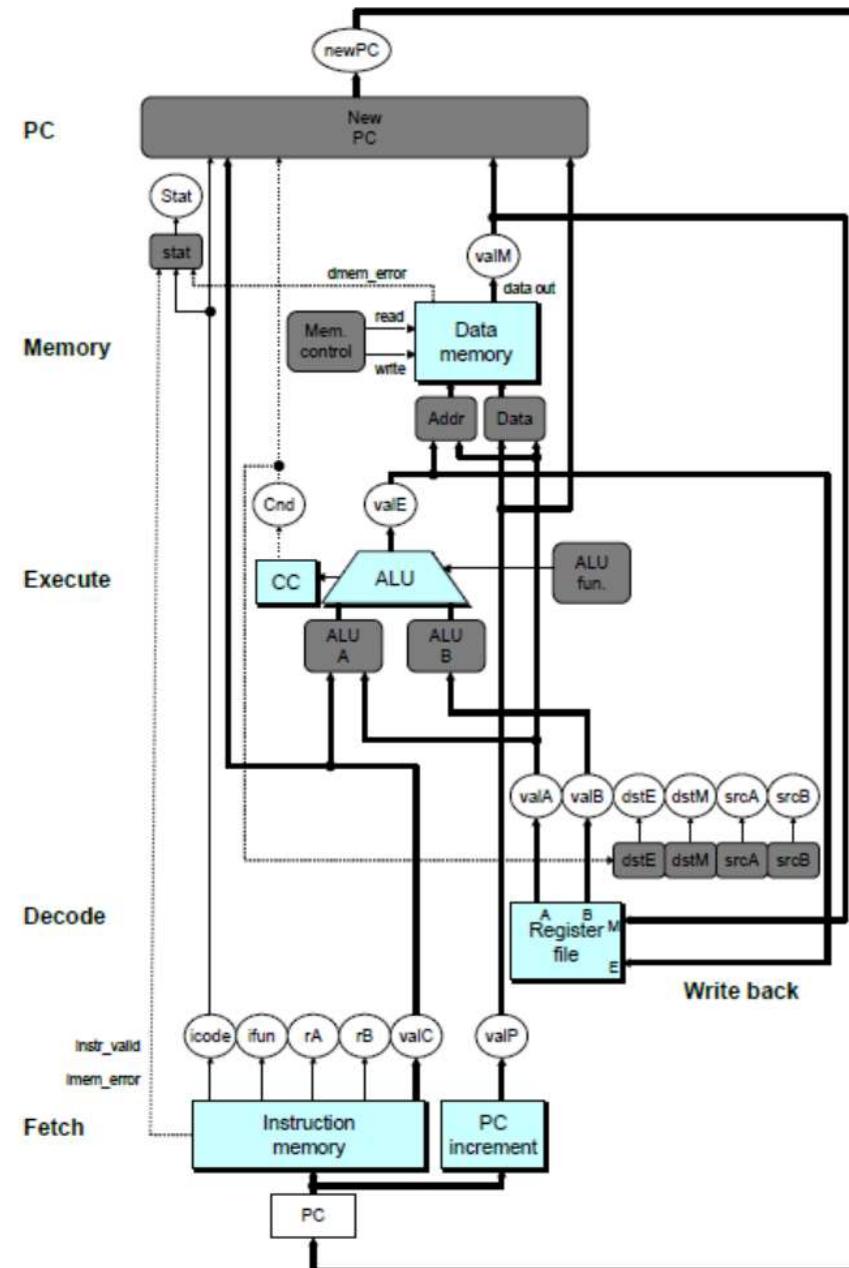
- Store bits
- Addressable memories
- Non-addressable registers
- Loaded only as clock rises



SEQ Hardware

Key

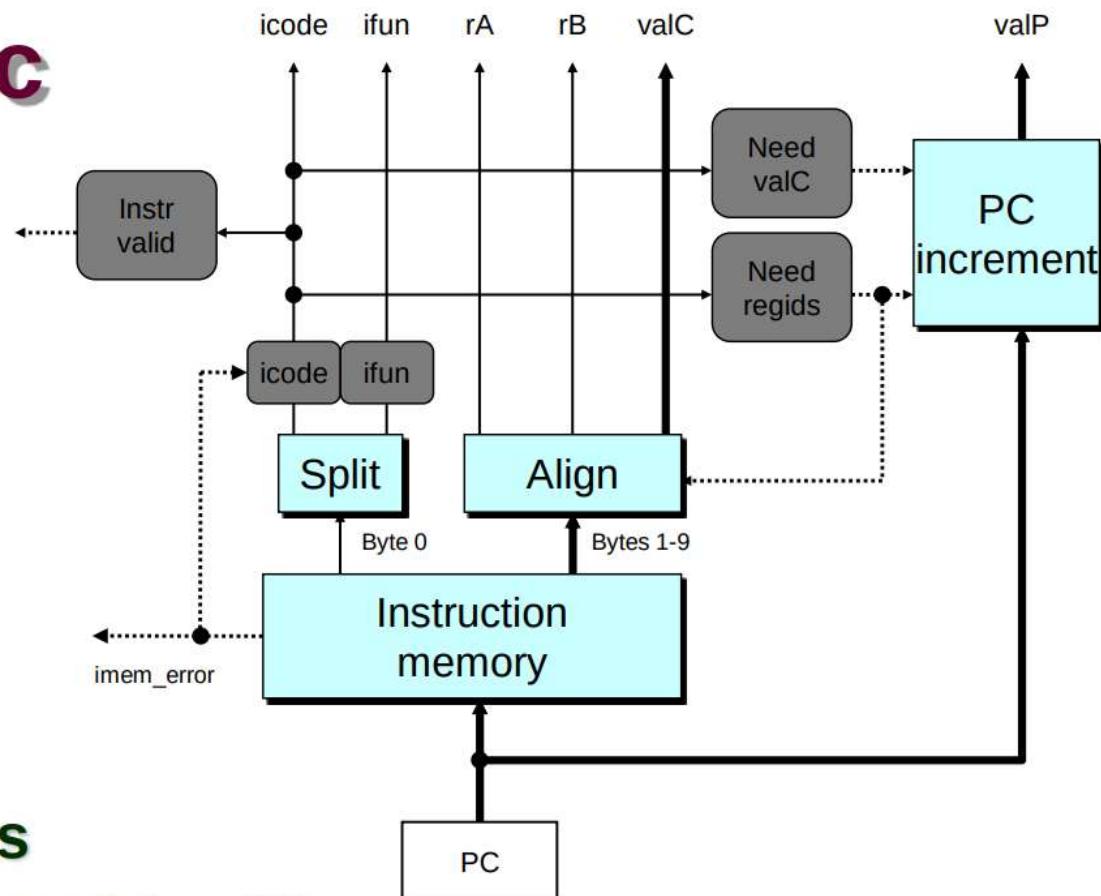
- Blue boxes: predefined hardware blocks
 - E.g., memories, ALU
- Gray boxes: control logic
 - Describe in HCL
- White ovals: labels for signals
- Thick lines: 64-bit word values
- Thin lines: 4-8 bit values
- Dotted lines: 1-bit values



– 29 –

CS:APP3e

Fetch Logic



Predefined Blocks

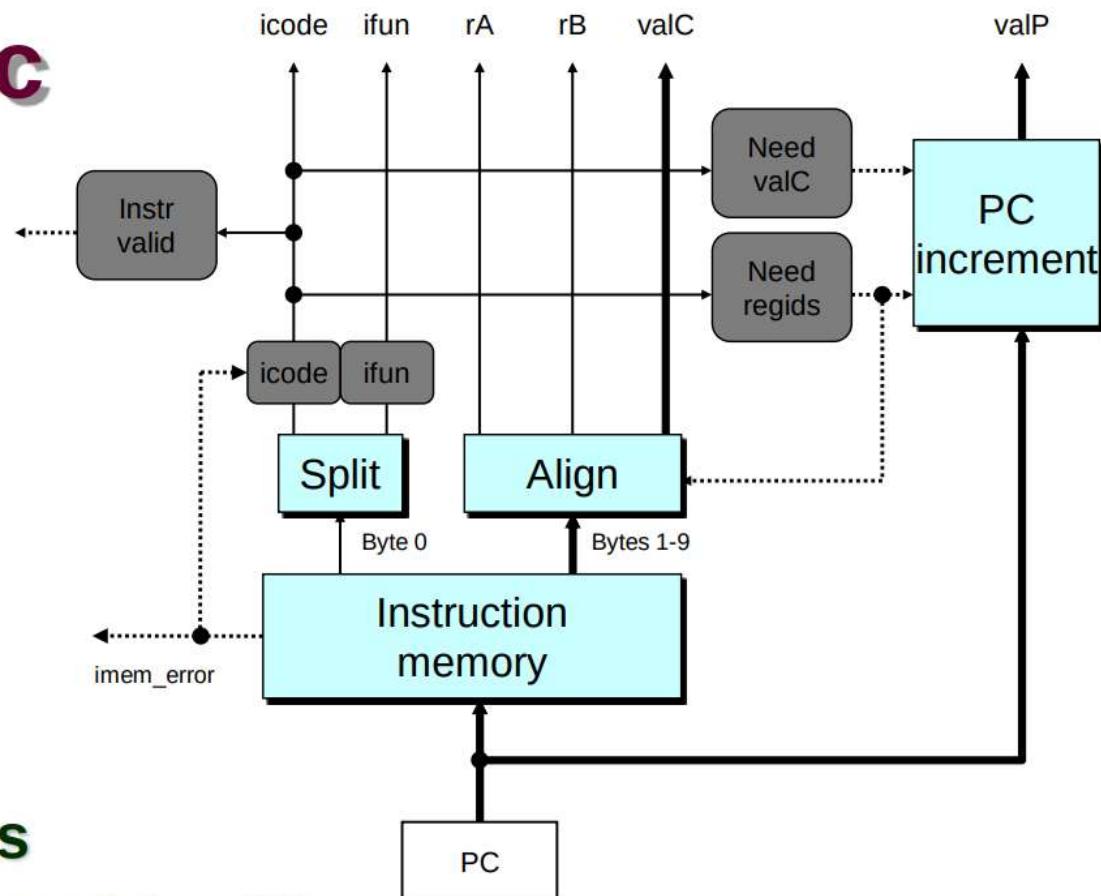
- **PC:** Register containing PC
- **Instruction memory:** Read 10 bytes (PC to PC+9)
 - Signal invalid address
- **Split:** Divide instruction byte into icode and ifun
- **Align:** Get fields for rA, rB, and valC

– 30 –

CS:APP3e

ITU CPH

Fetch Logic



Predefined Blocks

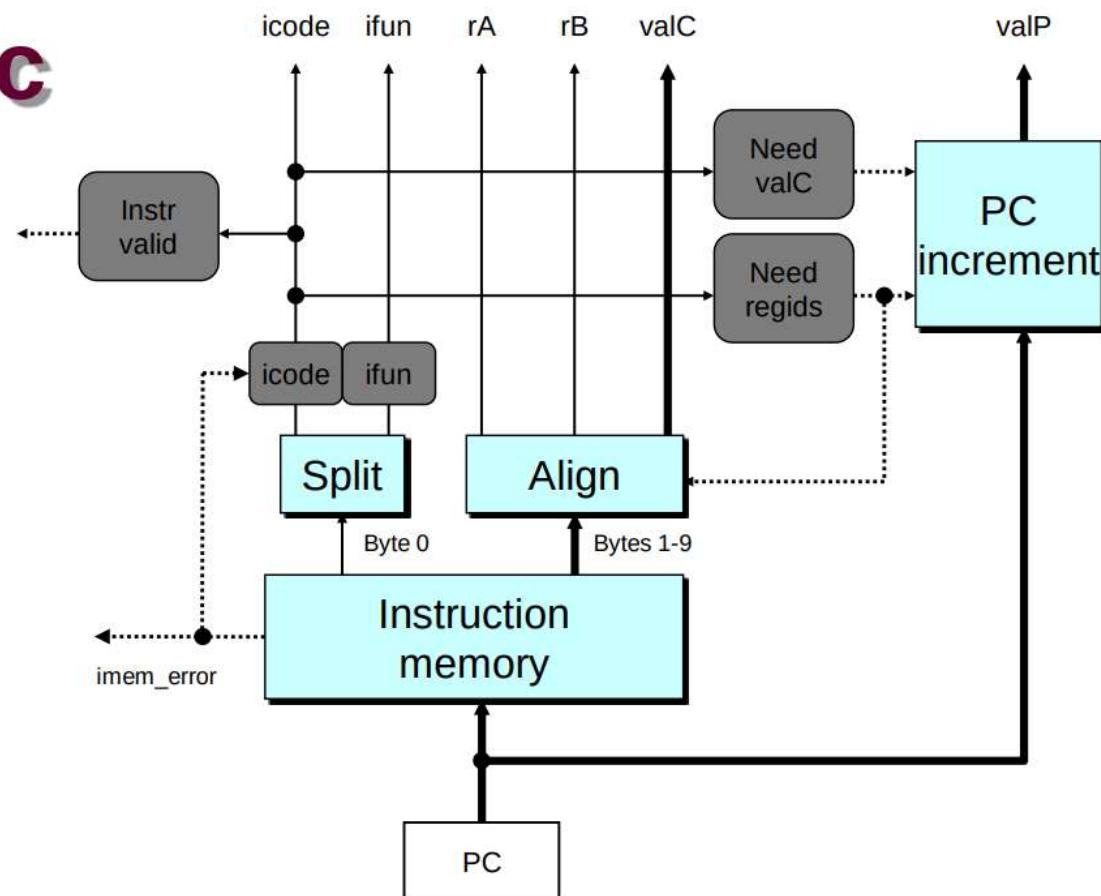
- **PC:** Register containing PC
- **Instruction memory:** Read 10 bytes (PC to PC+9)
 - Signal invalid address
- **Split:** Divide instruction byte into icode and ifun
- **Align:** Get fields for rA, rB, and valC

– 30 –

CS:APP3e

ITU CPH

Fetch Logic



Control Logic

- Instr. Valid: Is this instruction valid?
- icode, ifun: Generate no-op if invalid address
- Need regids: Does this instruction have a register byte?
- Need valC: Does this instruction have a constant word?

Decode Logic

Register File

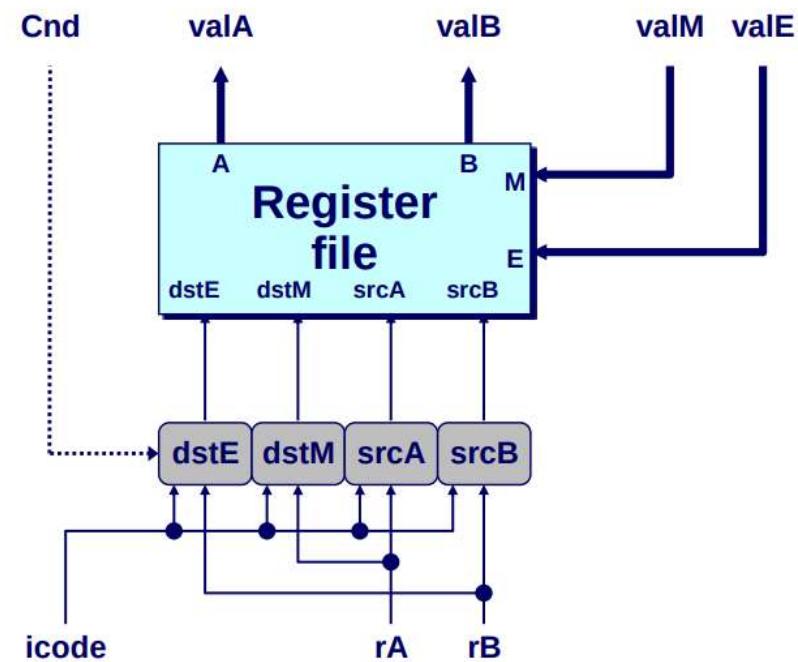
- Read ports A, B
- Write ports E, M
- Addresses are register IDs or 15 (0xF) (no access)

Control Logic

- srcA, srcB: read port addresses
- dstE, dstM: write port addresses

Signals

- Cnd: Indicate whether or not to perform conditional move
 - Computed in Execute stage



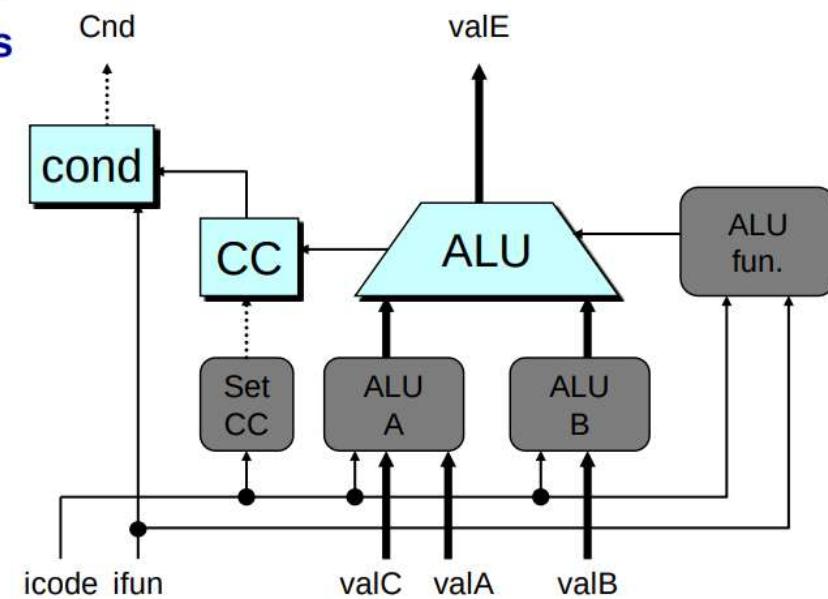
Execute Logic

Units

- ALU
 - Implements 4 required functions
 - Generates condition code values
- CC
 - Register with 3 condition code bits
- cond
 - Computes conditional jump/move flag

Control Logic

- Set CC: Should condition code register be loaded?
- ALU A: Input A to ALU
- ALU B: Input B to ALU
- ALU fun: What function should ALU compute?



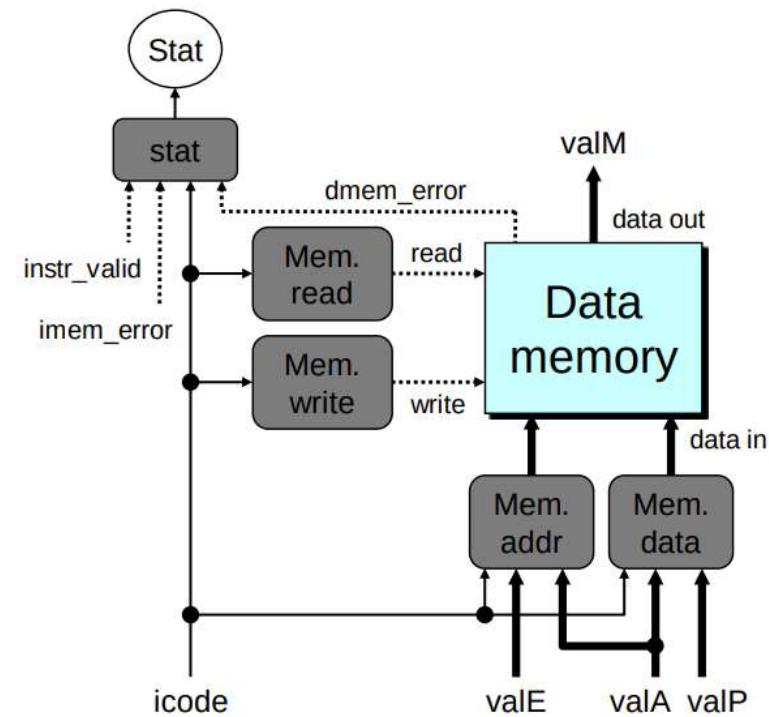
Memory Logic

Memory

- Reads or writes memory word

Control Logic

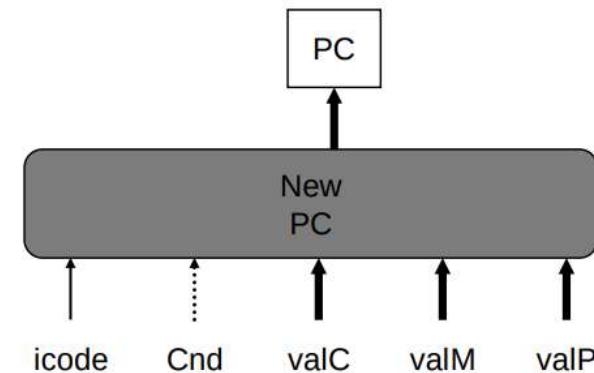
- stat: What is instruction status?
- Mem. read: should word be read?
- Mem. write: should word be written?
- Mem. addr.: Select address
- Mem. data.: Select data



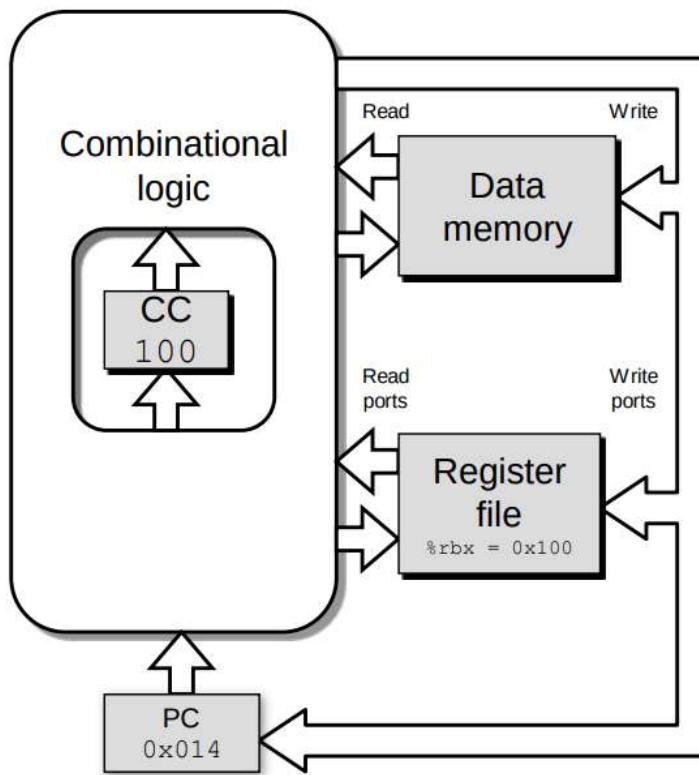
PC Update Logic

New PC

- Select next value of PC



SEQ Operation



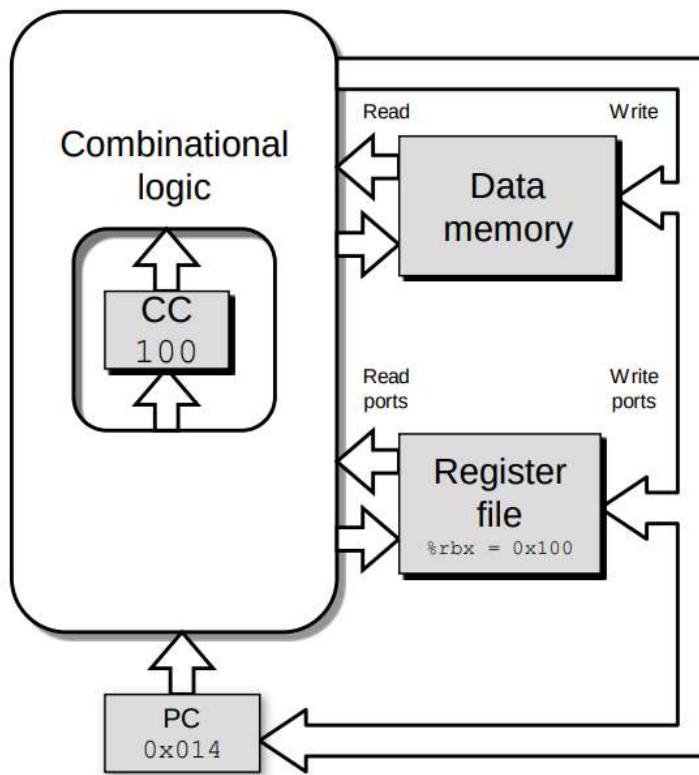
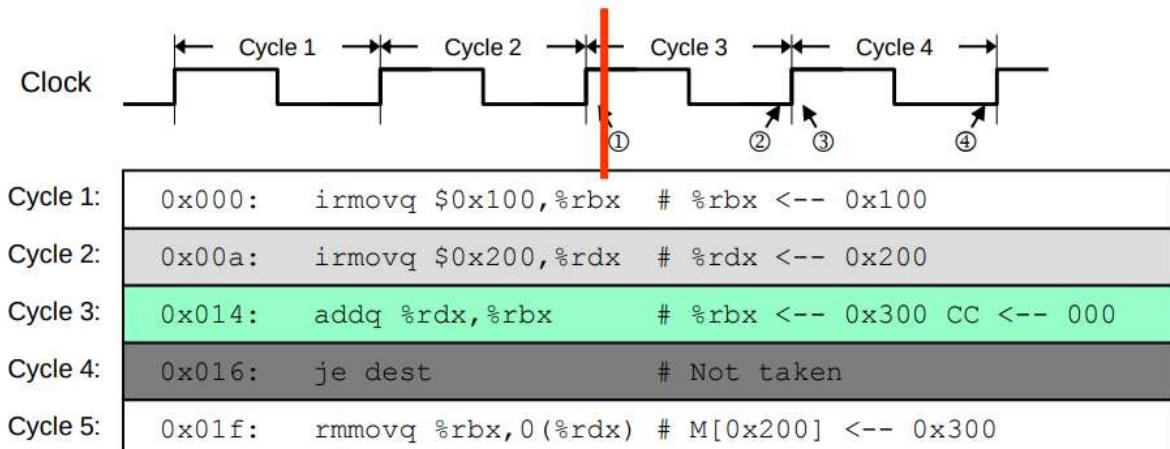
State

- **PC register**
 - **Cond. Code register**
 - **Data memory**
 - **Register file**
- All updated as clock rises*

Combinational Logic

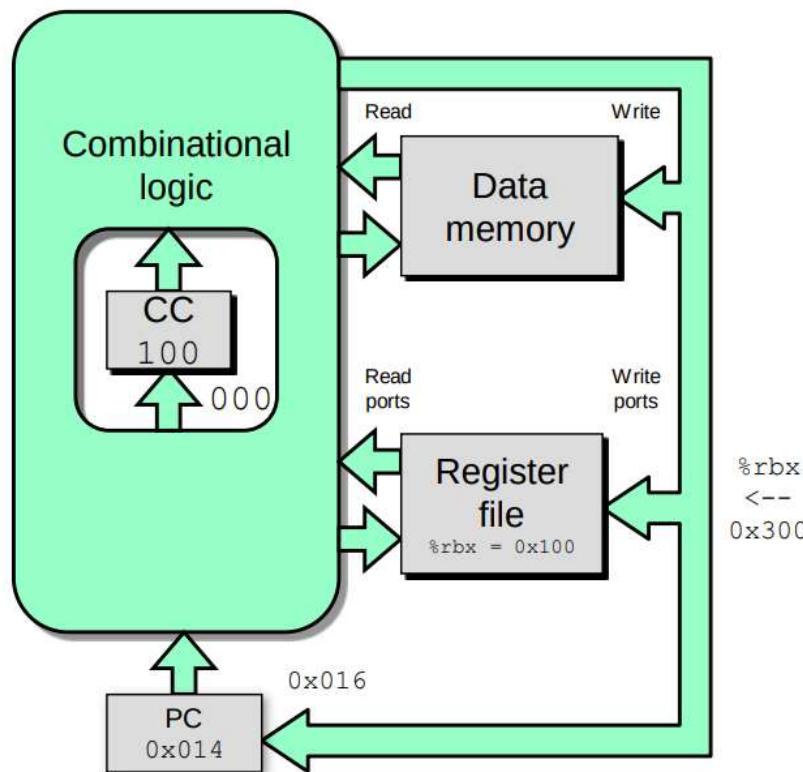
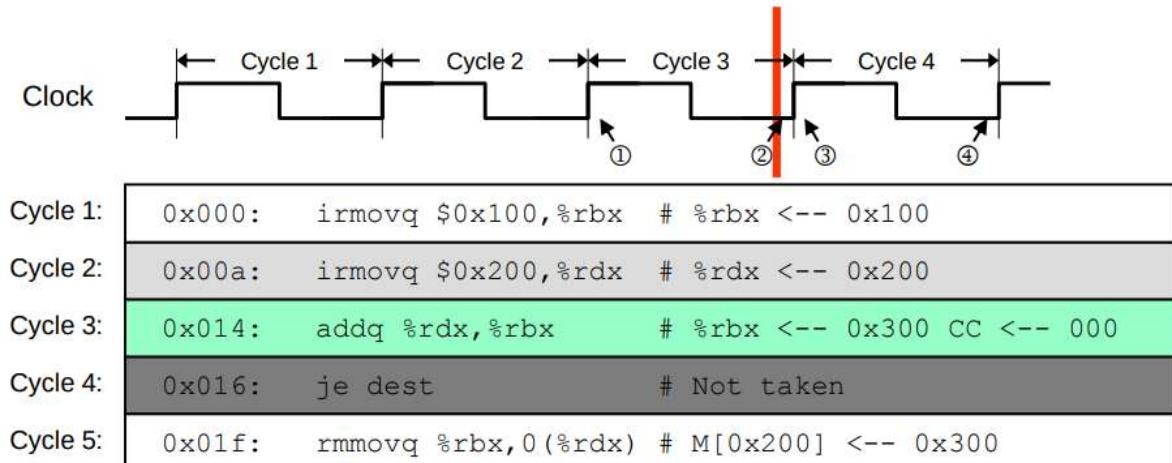
- **ALU**
- **Control logic**
- **Memory reads**
 - **Instruction memory**
 - **Register file**
 - **Data memory**

SEQ Operation #2



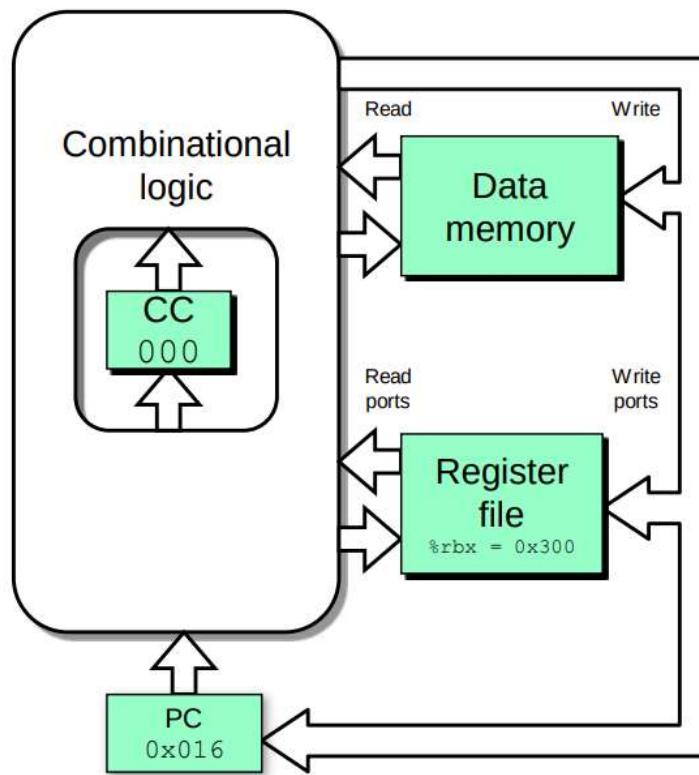
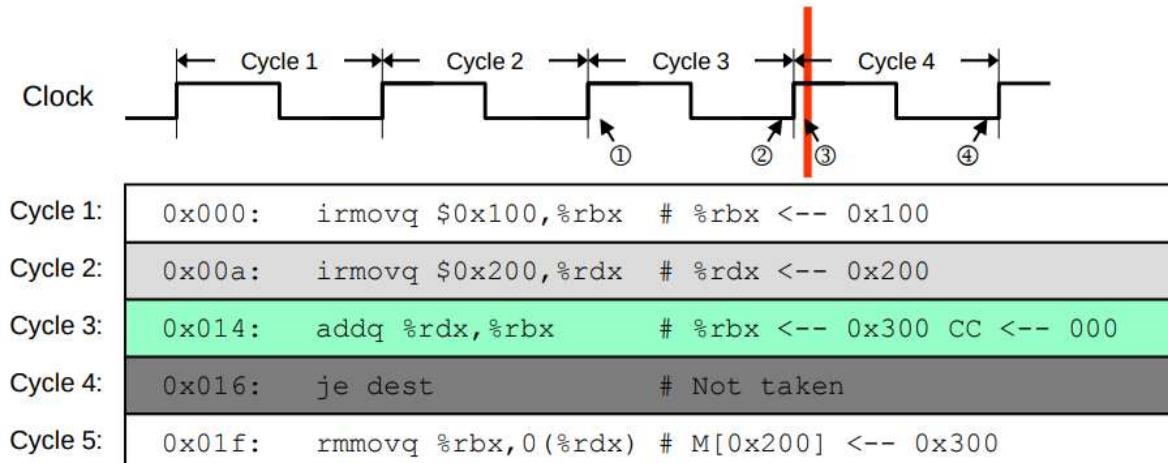
- state set according to second `irmovq` instruction
- combinational logic starting to react to state changes

SEQ Operation #3



- state set according to second `irmovq` instruction
- combinational logic generates results for `addq` instruction

SEQ Operation #4

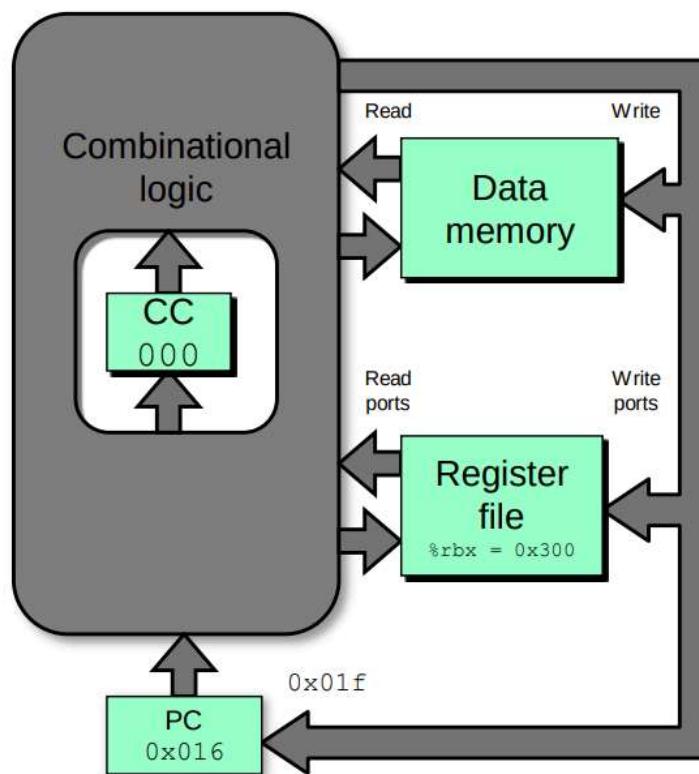


- state set according to addq instruction
- combinational logic starting to react to state changes

SEQ Operation #5

Clock

Cycle 1:	0x000: irmovq \$0x100,%rbx # %rbx <-- 0x100
Cycle 2:	0x00a: irmovq \$0x200,%rdx # %rdx <-- 0x200
Cycle 3:	0x014: addq %rdx,%rbx # %rbx <-- 0x300 CC <-- 000
Cycle 4:	0x016: je dest # Not taken
Cycle 5:	0x01f: rmmovq %rbx,0(%rdx) # M[0x200] <-- 0x300



- state set according to addq instruction
- combinational logic generates results for je instruction

SEQ Summary

Implementation

- Express every instruction as series of simple steps
- Follow same general flow for each instruction type
- Assemble registers, memories, predesigned combinational blocks
- Connect with control logic

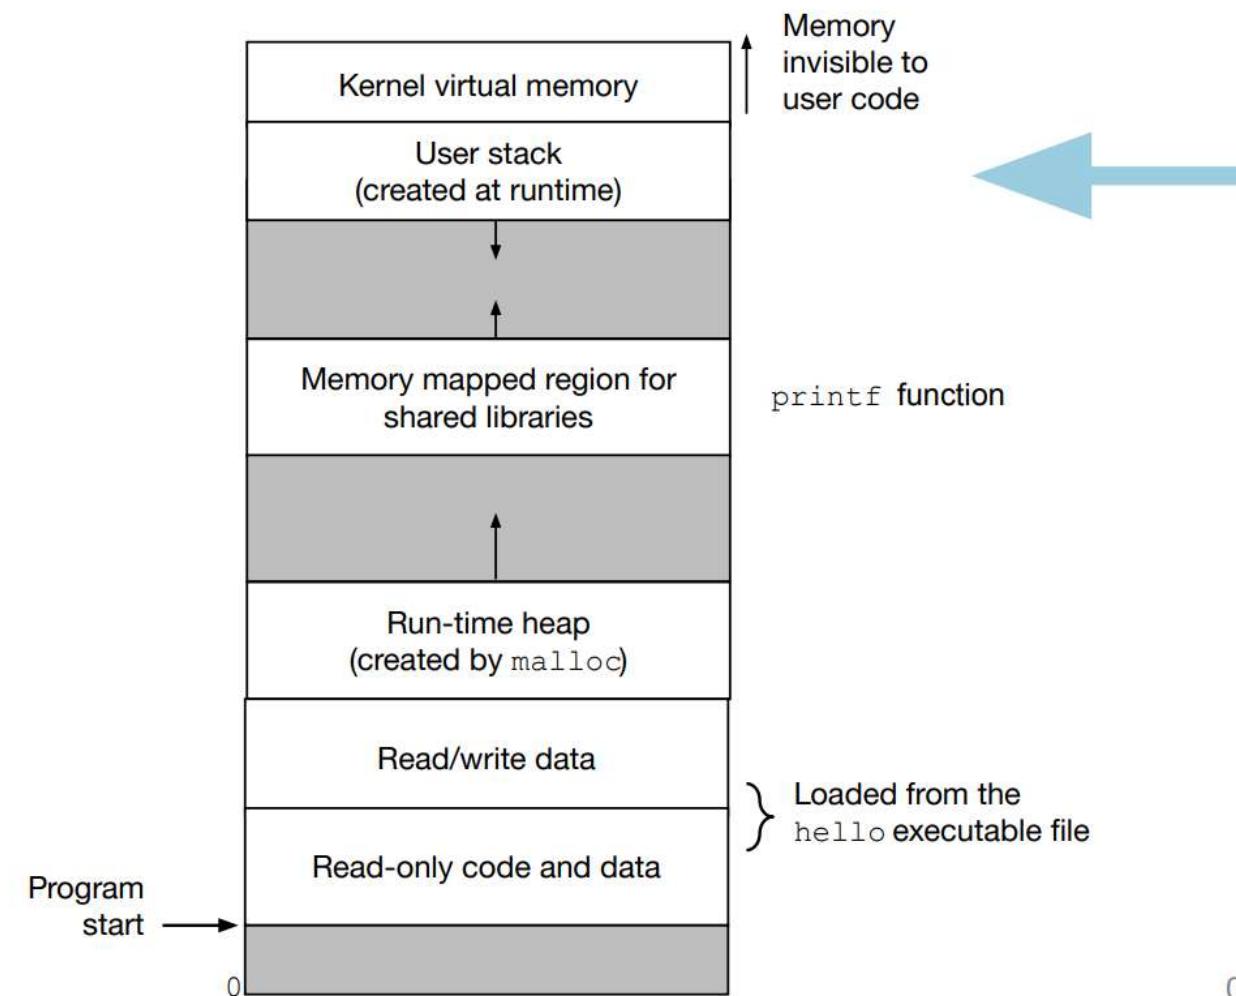
Limitations

- Too slow to be practical
- In one cycle, must propagate through instruction memory, register file, ALU, and data memory
- Would need to run clock very slowly
- Hardware units only active for fraction of clock cycle

Stack, Procedure Calls

Virtual Memory

What is the stack?

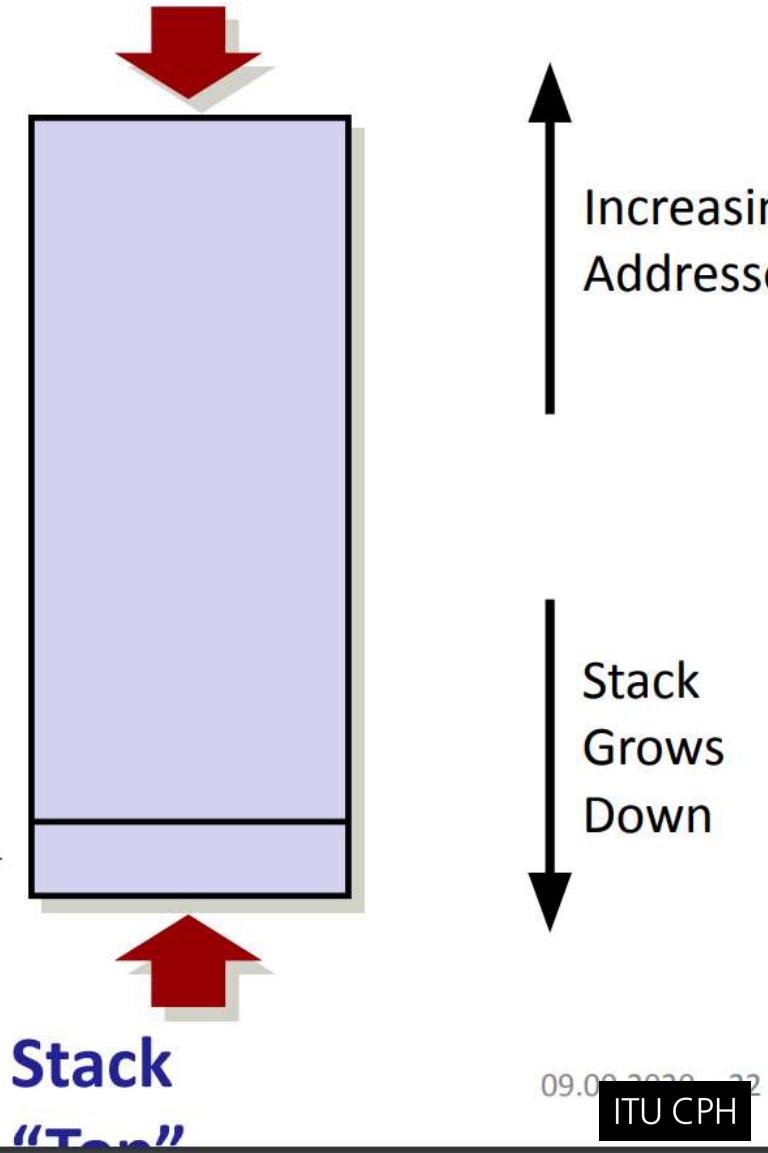


The Stack

- Region of memory managed with stack discipline
- Grows towards lower addresses
- Register %rbp points to the bottom of the stack
- Register %rsp points to the top of the stack, it is the **stack pointer**

Stack Pointer: %rsp

Stack “Bottom”



Stack - Push

Operator push

pushq %r10

1. Decrement %rsp by 8
2. Write contents of %r10 at address given by %rsp

Stack Pointer: %rsp 

Stack “Bottom”



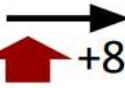
**Stack
“Top”**

Stack - Pop

Operator pop

popq %r10

1. Copy contents at address %rsp to %r10
2. Increment %rsp by 8

Stack Pointer: %rsp 

Stack “Bottom”



Stack
“Top”

Procedure Call

Instructions:

- **call:** push **return address** on stack; jump to label/address
 - Return address is address of instruction right after call instruction
- **ret:** pop address from stack; jump to address

Example

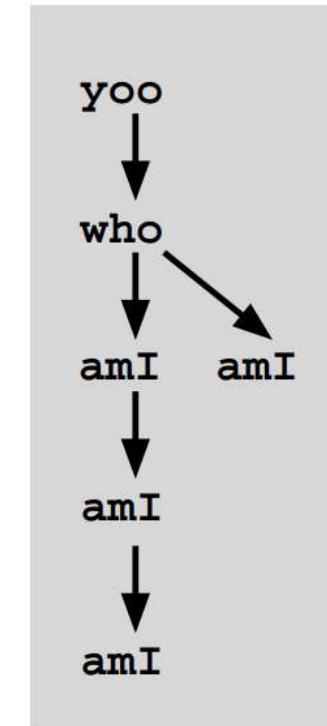
why/how is the stack relevant when talking about procedure calls?

```
yoo (...)  
{  
    •  
    •  
    who () ;  
    •  
    •  
}
```

```
who (...)  
{  
    • • •  
    amI () ;  
    • • •  
    amI () ;  
    • • •  
}
```

```
amI (...)  
{  
    •  
    •  
    amI () ;  
    •  
    •  
}
```

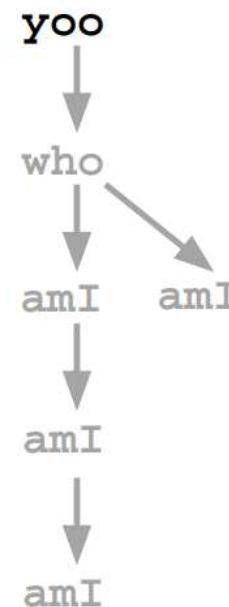
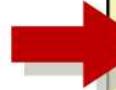
Example Call Chain



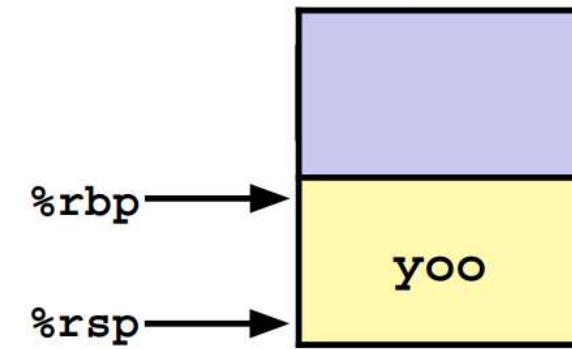
Procedure amI () is recursive

Example

```
yoo (...) {  
    .  
    .  
    who () ;  
    .  
    .  
}
```



Stack



stack frame

- caller's saved registers⁽¹⁾
- local vars
- args
- return address⁽²⁾

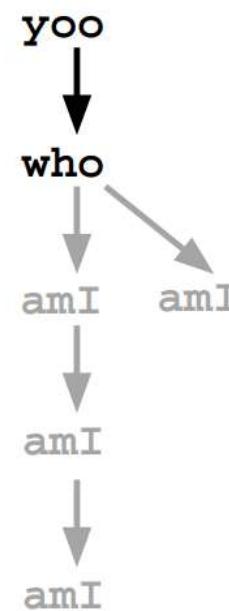
⁽¹⁾: to restore caller state on return
(bottom frame has no caller).

⁽²⁾: pushes it when it calls
(top frame does not need this)

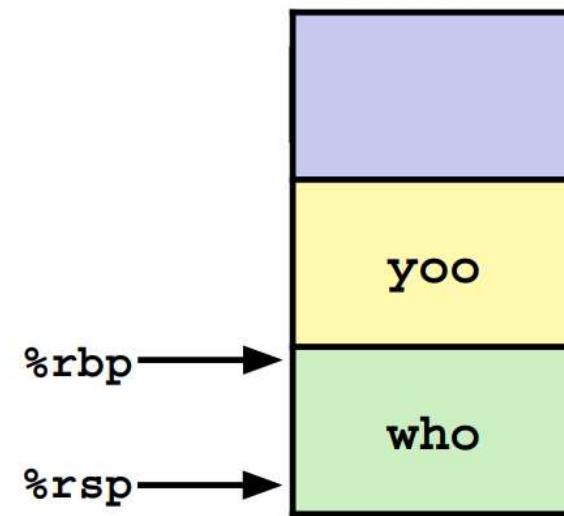
Example

```
yoo()
{
    who (...)

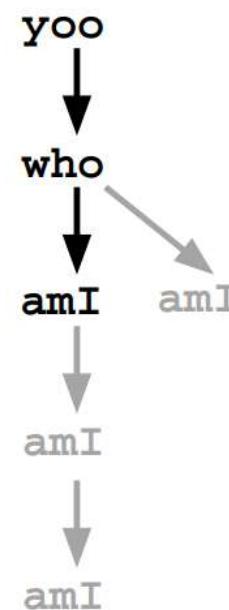
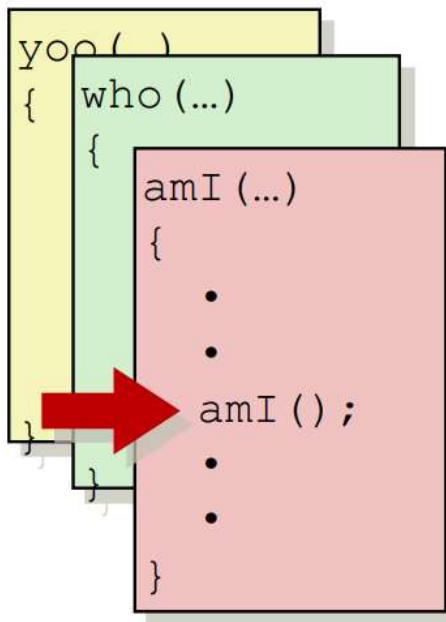
    {
        . . .
        amI ();
        . . .
        amI ();
        . . .
    }
}
```



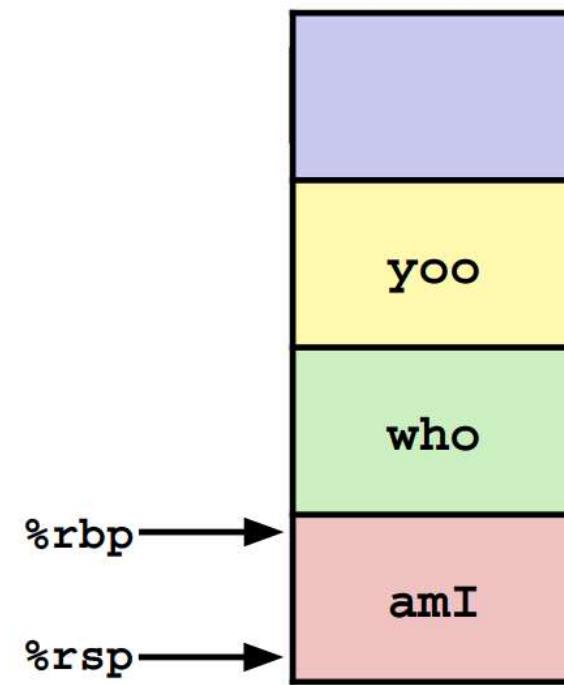
Stack



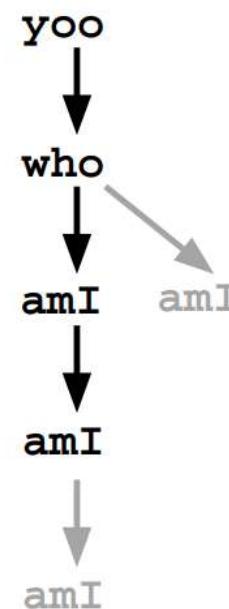
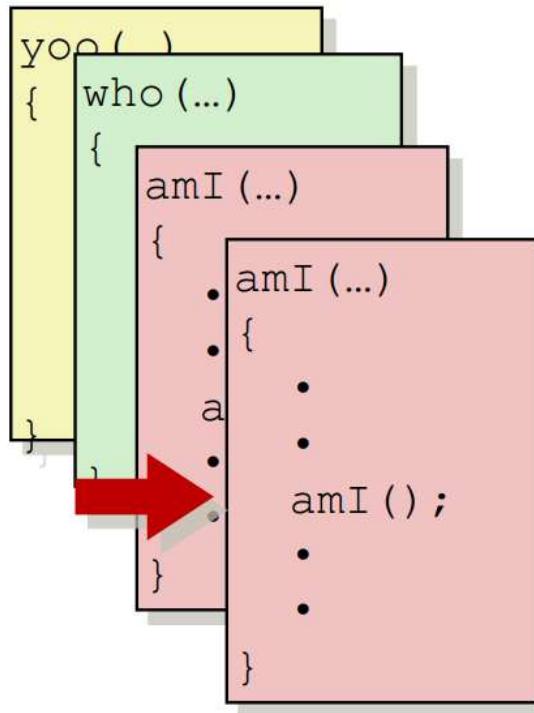
Example



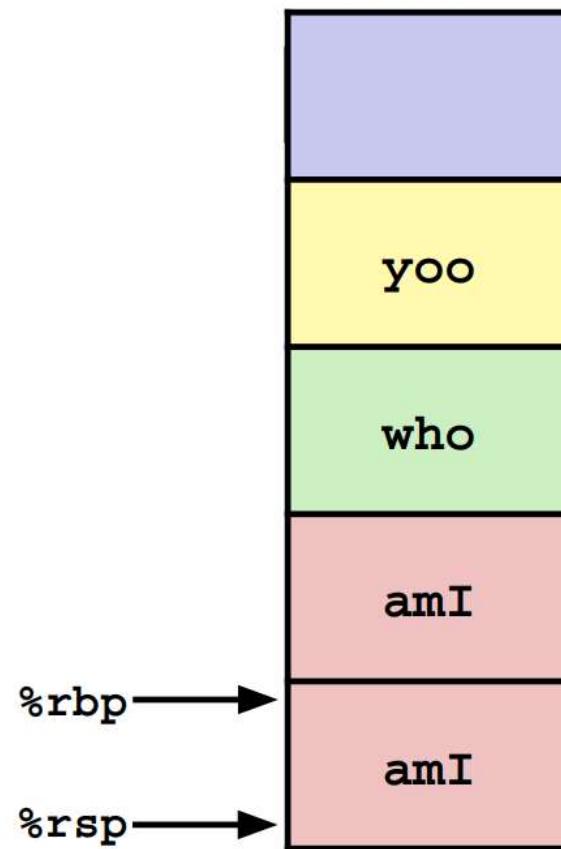
Stack



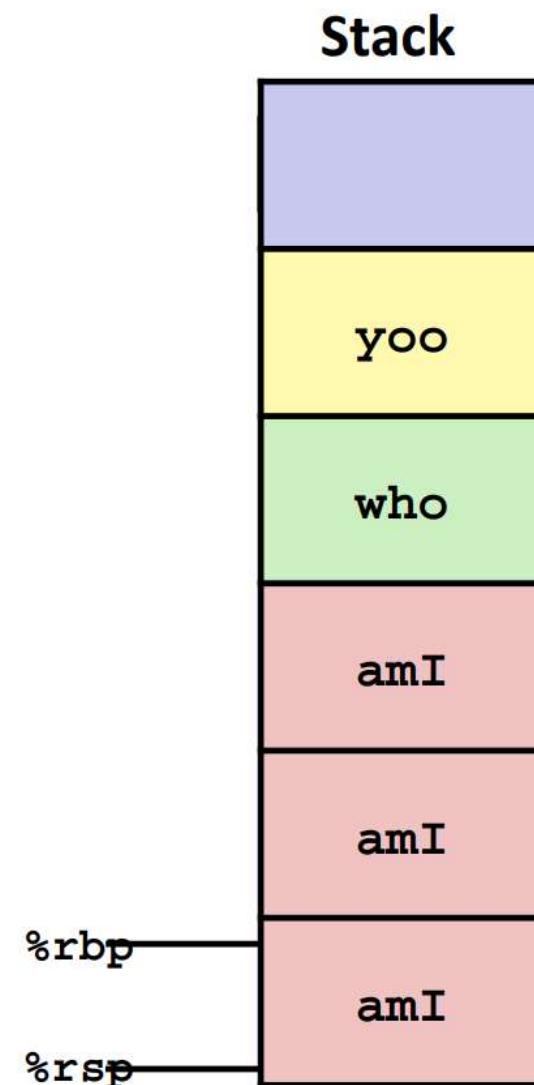
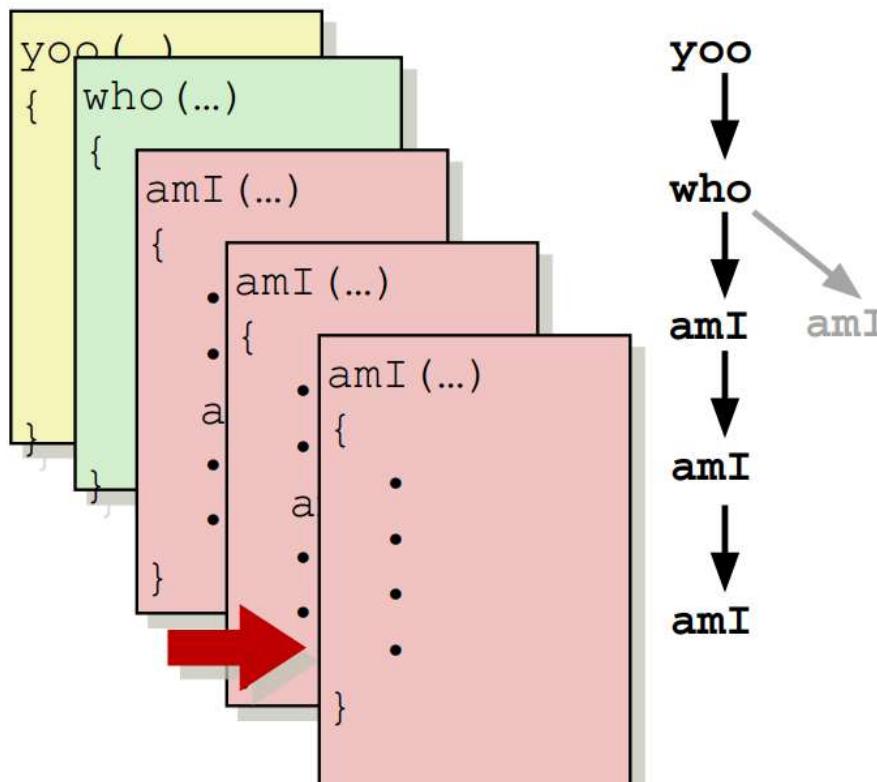
Example



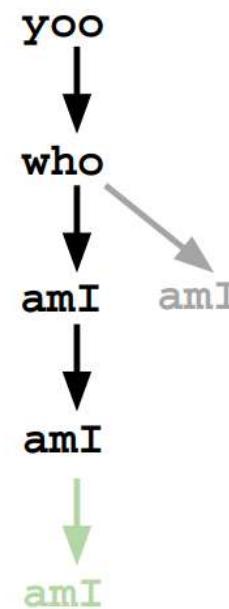
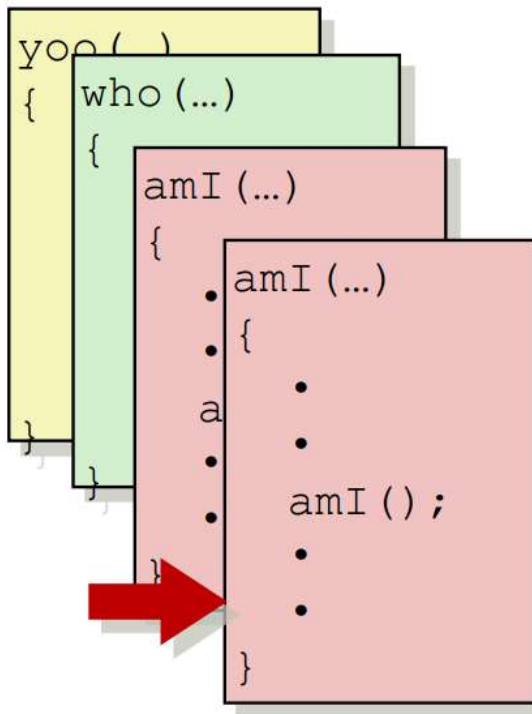
Stack



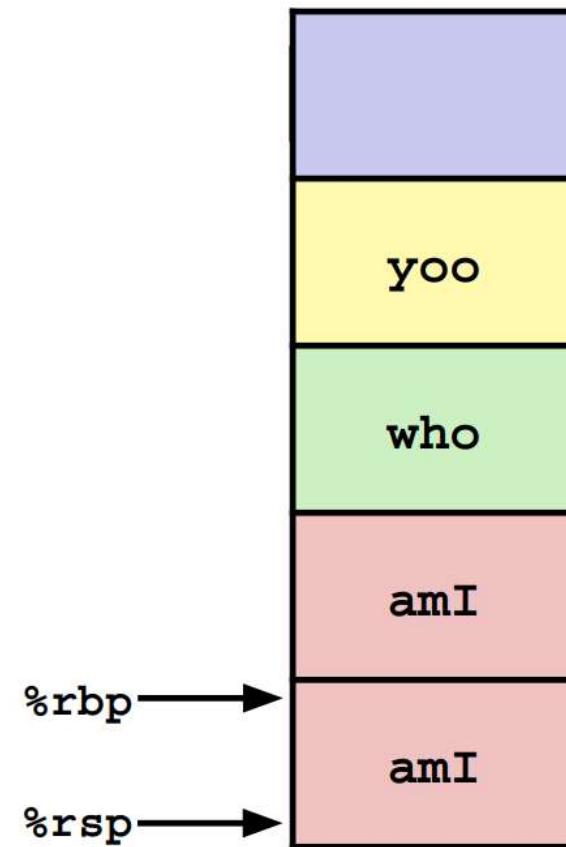
Example



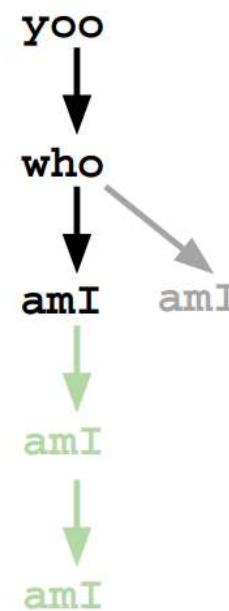
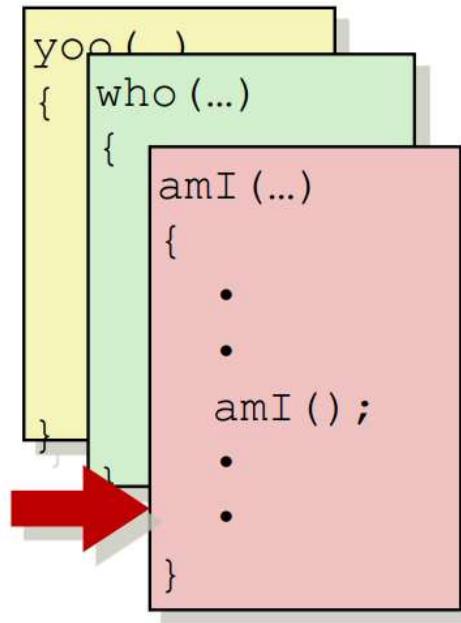
Example



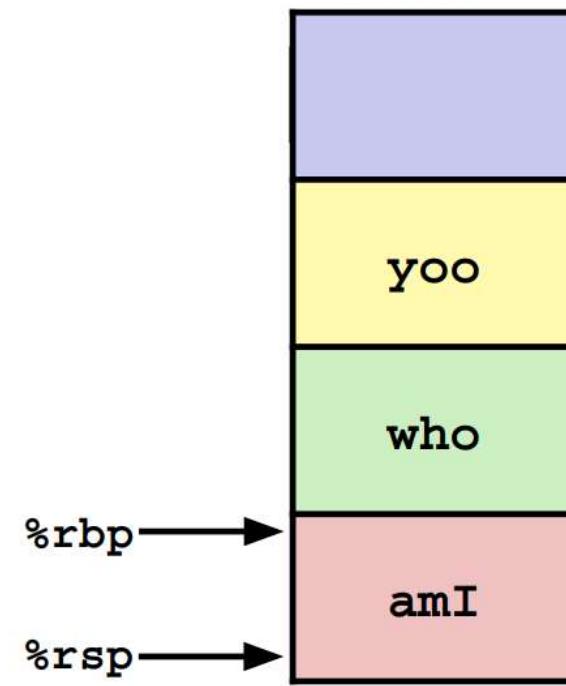
Stack



Example



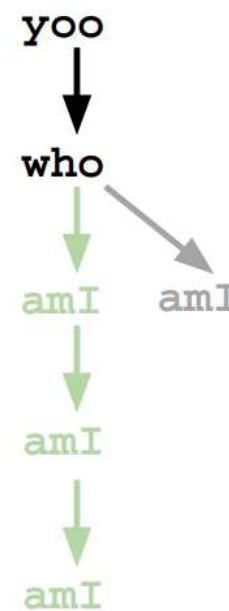
Stack



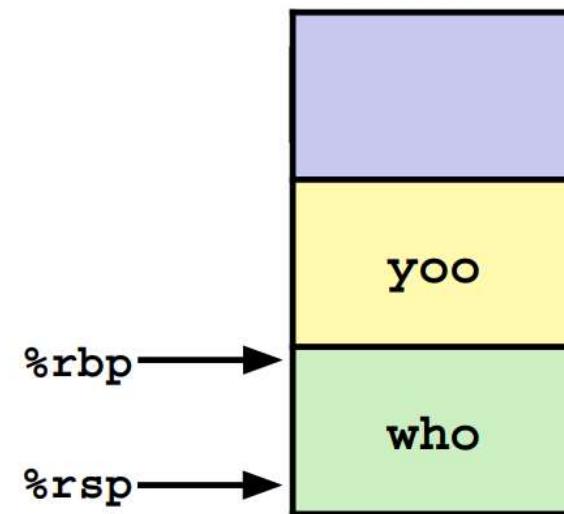
Example

```
yoo()
{
    who (...)

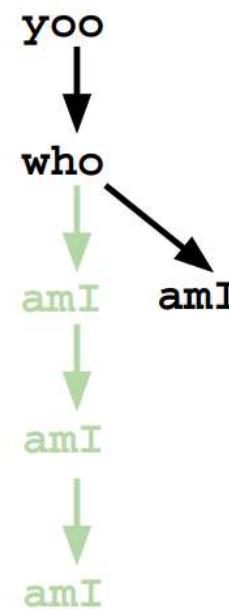
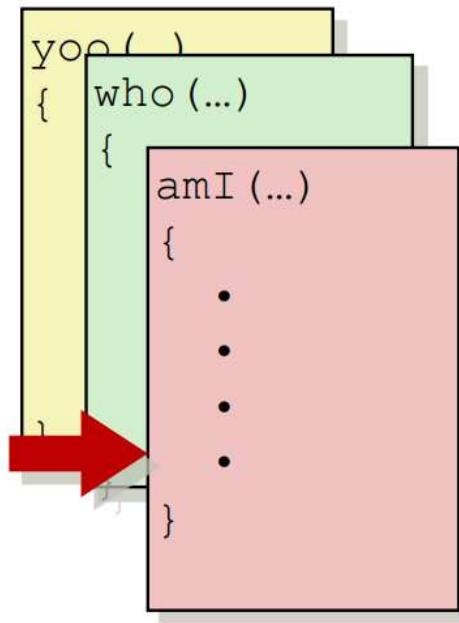
    {
        . . .
        amI ();
        . . .
        amI ();
        . . .
    }
}
```



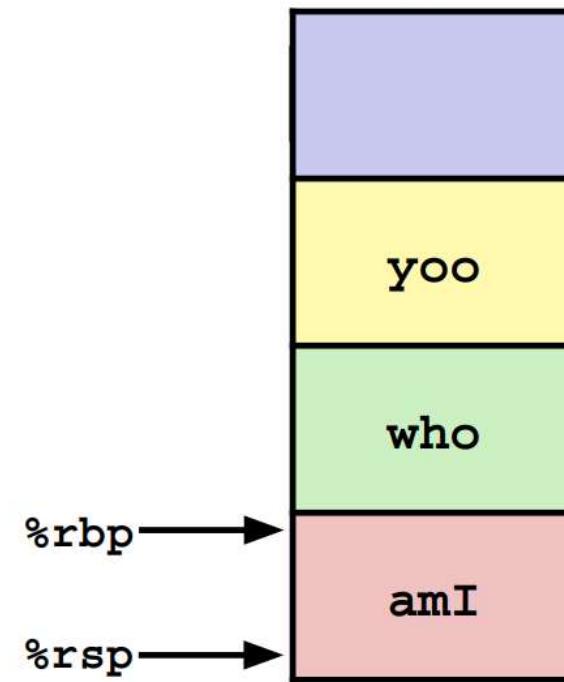
Stack



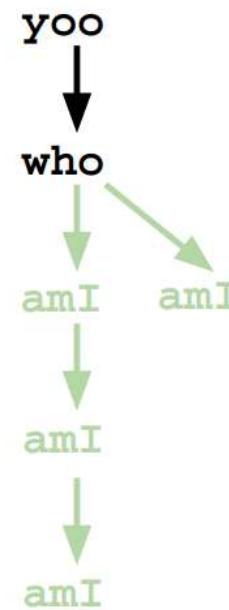
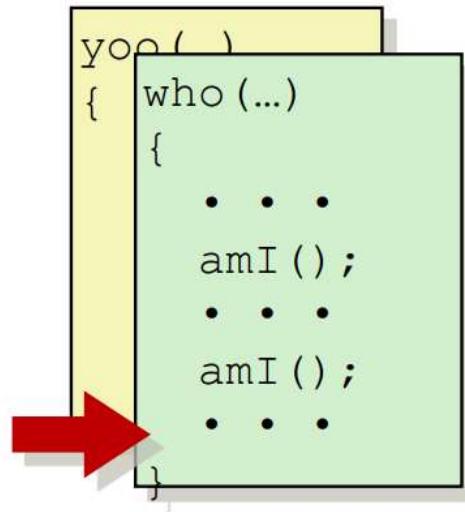
Example



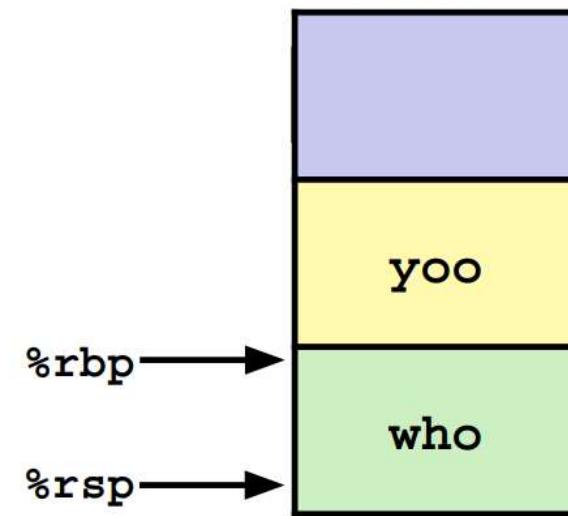
Stack



Example

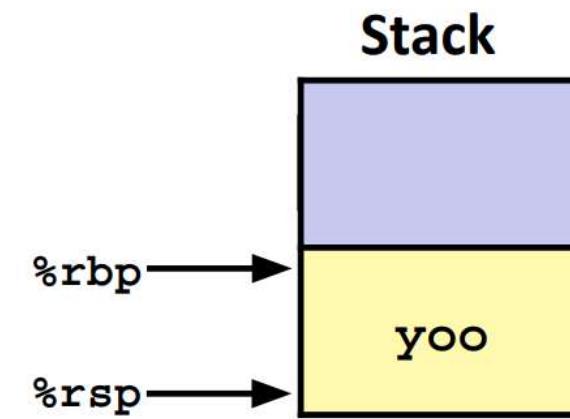
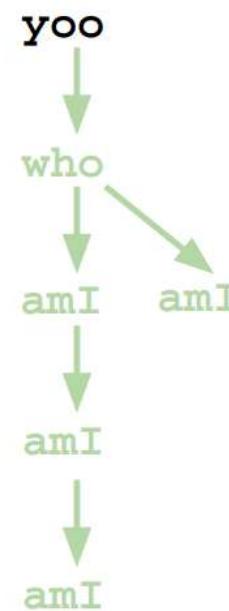


Stack



Example

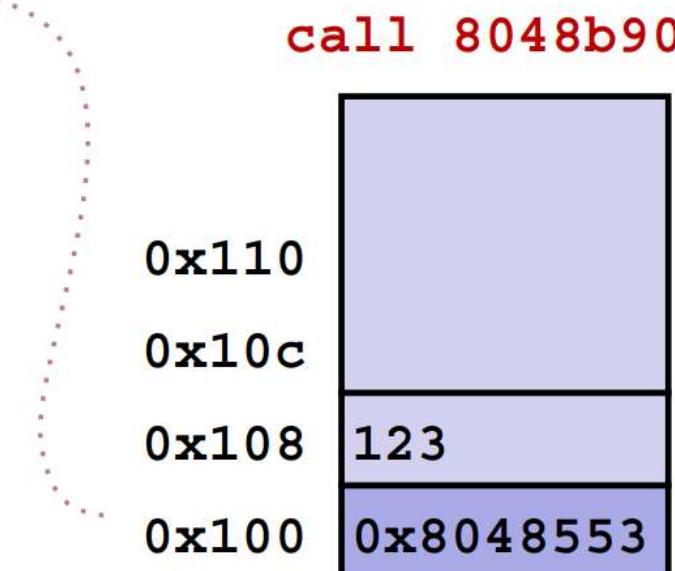
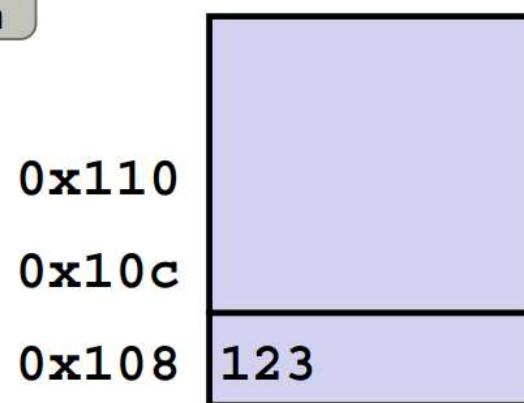
```
yoo (...)  
{  
    .  
    .  
    who ();  
    .  
    .  
}  
  
A red arrow points to the closing brace '}'.
```



Procedure Call Example

```
804854e: e8 3d 06 00 00    call    8048b90 <main>
8048553: 50                pushl   %eax
```

return address
= address of next instruction



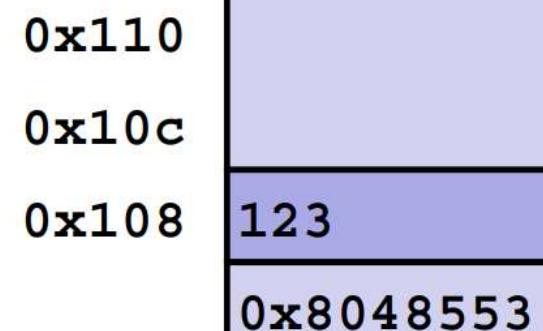
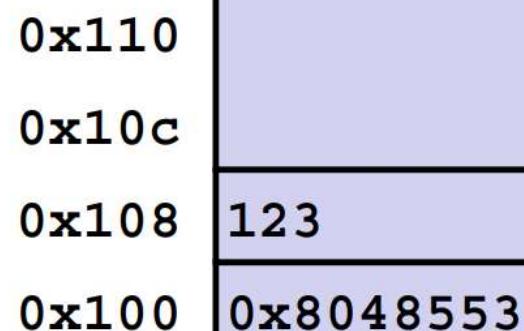
instruction pointer
updated to callee



Procedure Call Example

8048591: c3

ret



(garbage)

%rsp 0x100
%rip 0x8048591

%rsp 0x108
%rip 0x8048553

just a pop

Q: what happens if user input overwrites stack pointer, or return address?

Procedure Call Example

8048591: c3

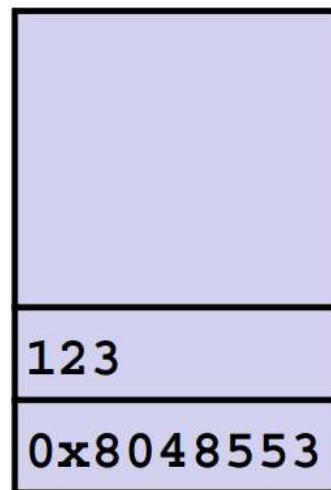
ret

0x110

0x10c

0x108

0x100



Q: what happens if user input overwrites stack pointer, or return address?

A: user has full control over control flow.
(exploit)

%rsp 0x100

%rip 0x8048591

0x110

0x10c

0x108

0x8048553

%rsp 0x108

%rip 0x8048553

ret

(garbage)

just a pop

Calling a function (x86-64)

To call a function, a program:

1. **Places the first six integer or pointer parameters in %rdi, %rsi, %rdx, %rcx, %r8 and %r9**
2. Pushes onto the stack subsequent parameters and parameters larger than 8B (in order).
3. Executes the call instruction, which:
 - Pushes the return address onto the stack
 - Jumps to the start of the specified function

GP Register Usage during function calls

First six arguments of a function stored in
di, si, dx, cx, r8, r9

Remaining arguments are on the
stack (more later)

Return value is in rax

Calling a function preserves rbp, rbx, r12-15. The other
registers **might be overwritten**.

What is the stack?
How is the stack used?

Executing a function

The C run-time system introduces instruction to set-up and clean-up the stack in each procedure.

Set-up consists in allocation and initialization of a stack-frame. Clean-up: deallocating a stack frame.

A stack-frame is the space needed on the stack by a procedure for storing:

- The return address
- (some) parameters
- Local variables

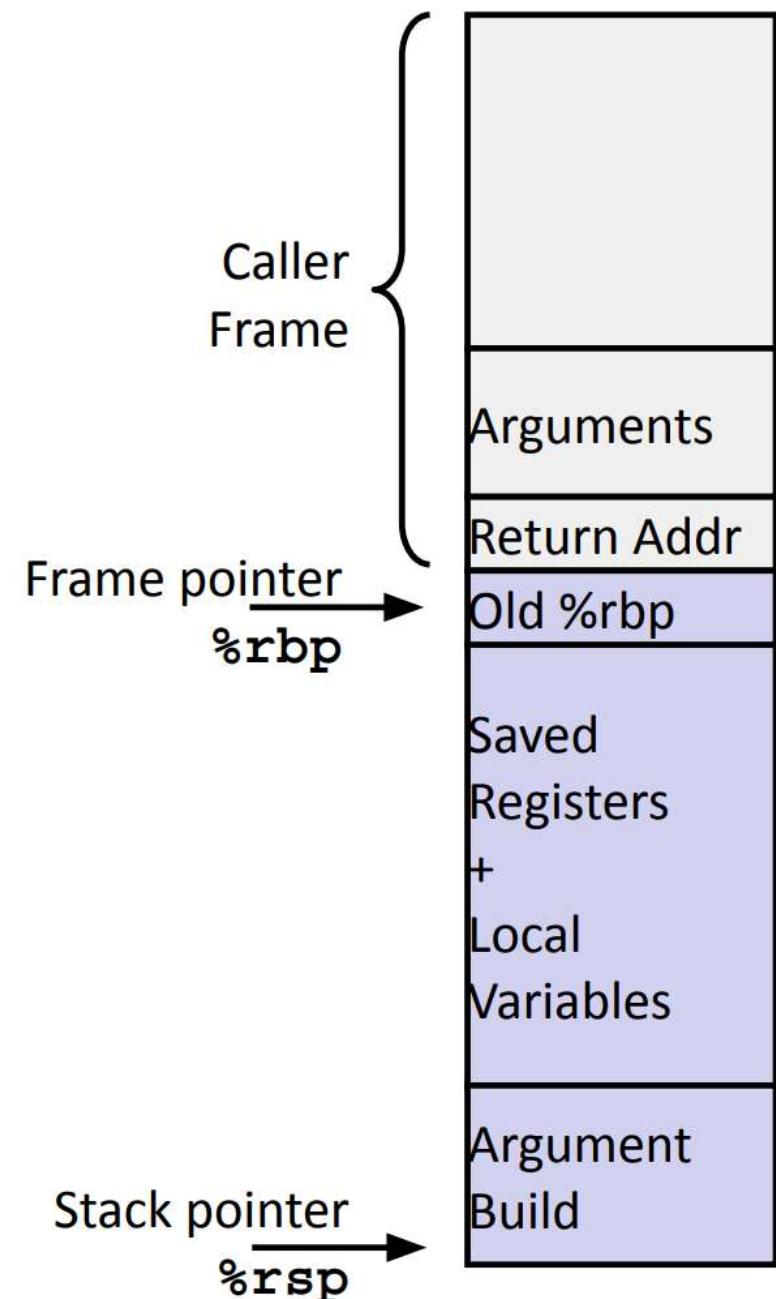
Stack Frame

Caller:

- Arguments
 - pushed by program (if needed)
- Return address
 - pushed by call

Callee:

- Previous frame pointer (%rbp)
- Other callee-save registers (%rbx, %r12-15)
- Space for local variables
- Arguments for next function (when about to call another function)



Example

```
arith.s ➤ arith.c ➤
1 #include <stdio.h>
2 #include <stdlib.h>
3
4 int logical(int x, int y)
5 {
6     int t1 = x^y;
7     int t2 = t1 >> 17;
8     int mask = (1<<13) - 7;
9     int rval = t2 & mask;
10    return rval;
11 }
12
13 int main(int argc, char* argv[])
14 {
15     if (argc != 3) {
16         printf("Usage: arith x y\n");
17         return 1;
18     }
19
20     int x = atoi(argv[1]);
21     int y = atoi(argv[2]);
22     printf("Arguments x: %d, y: %d\n", x, y);
23     printf("Logical returns: %d\n", logical(x,y));
24     printf("\n");
25
26     return 0;
27 }
```

Example

```

5 logical:
6 .LFB2:
7     .cfi_startproc
8     pushq %rbp
9     .cfi_def_cfa_offset 16
0     .cfi_offset %rbp, -16
1     movq %rsp, %rbp
2     .cfi_def_cfa_register 6
3     movl %edi, -20(%rbp)
4     movl %esi, -24(%rbp)
5     movl -20(%rbp), %eax
6     xorl -24(%rbp), %eax
7     movl %eax, -16(%rbp)
8     movl -16(%rbp), %eax
9     sarl $17, %eax
0     movl %eax, -12(%rbp)
1     movl $8185, -8(%rbp)
2     movl -12(%rbp), %eax
3     andl -8(%rbp), %eax
4     movl %eax, -4(%rbp)
5     movl -4(%rbp), %eax
6     popq %rbp
7     .cfi_def_cfa 7, 8
8     ret
9     .cfi_endproc

```

Set-up:

- Previous stack frame base %rbp pushed on stack
- %rbp is the only callee save register
- Frame pointer re-initialised

Function:

- 4 local variables at positions relative to stack frame base %rbp
 - t1: -16(%rbp)
 - t2: -12(%rbp)
 - mask: -8(%rbp)
 - rval: -4(%rbp)
- %eax holds intermediate results
- %eax holds return value at the end of the function

remember:
stack
grows
down

Clean-up:

- Previous stack frame base restored
- ret manipulates %rsp and %rip to return control to return address

Buffer Overflows

Vulnerable C code

example of vulnerable C code

```
echo.c
1 #include <stdio.h>
2
3 void echo()
4 {
5     char buf[4];
6     gets(buf);
7     puts(buf);
8
9 }
10
11 int main()
12 {
13     echo();
14
15     return 0;
16 }
```

disable stack protection
(more on that later today)
(so it's not really a problem
today; compilers prevent
problem by default)

```
% gcc -fno-stack-protector echo.c -o echonp
```

```
% ./echonp
12345678
12345678
% ./echonp
123456789
123456789
% ./echonp
12345678901234567890123
12345678901234567890123
% ./echonp
123456789012345678901234567890123
123456789012345678901234567890123
[1] 19108 segmentation fault (core dumped) ./echonp
```

process tried I/O on
kernel virtual memory
(kernel-space).
what's happening?
let's investigate.

Vulnerable C code

objdump -D echonp >> echonp.d

The stack is 16B aligned

=>

rsp is decreased with 16B sub \$0x10, %rsp

tip: for assignment, always use objdump. with it, you get assembly, and the object, and the address where it is

```

322
323 0000000000400566 <echo>:
324 400566:    55
325 400567:    48 89 e5
326 40056a:    48 83 ec 10
327 40056e:    48 8d 45 f0
328 400572:    48 89 c7
329 400575:    b8 00 00 00 00
330 40057a:    e8 d1 fe ff ff
331 40057f:    48 8d 45 f0
332 400583:    48 89 c7
333 400586:    e8 a5 fe ff ff
334 40058b:    90
335 40058c:    c9
336 40058d:    c3
337
338 000000000040058e <main>:
339 40058e:    55
340 40058f:    48 89 e5
341 400592:    b8 00 00 00 00
342 400597:    e8 ca ff ff ff
343 40059c:    b8 00 00 00 00
344 4005a1:    5d
345 4005a2:    c3
346 4005a3:    66 2e 0f 1f 84 00 00
347 4005aa:    00 00 00
348 4005ad:    0f 1f 00
349

```

enough space for buf

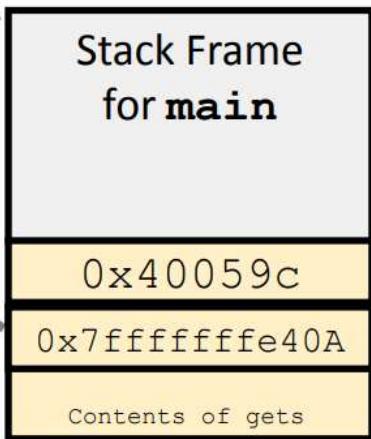
"load effective address" (arith on addr, store in specific register)

push	%rbp	push old base pointer on stack
mov	%rsp,%rbp	new base pointer := old stack pointer
sub	\$0x10,%rsp	new stack pointer := old stack pointer - 16
lea	-0x10(%rbp),%rax	= %rsp
mov	%rax,%rdi	address of buff into rax
mov	\$0x0,%eax	then into 1st arg
callq	400450 <gets@plt>	initialize return
lea	-0x10(%rbp),%rax	
mov	%rax,%rdi	
callq	400430 <puts@plt>	
nop		evict stack frame
leaveq		
retq		

Vulnerable code

so, why vulnerable?

Before call to gets



0x7fffffe40A
0x40059c
0x7fffffe40A
0x7fffffe480

```
% gdb ./echonp
GNU gdb (Ubuntu 7.11.1-0ubuntu1~16.5) 7.11.1
Copyright (C) 2016 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from ./echonp...done.
(gdb) break echo
Breakpoint 1 at 0x40056e: file echo.c, line 6.
(gdb) r
Starting program: /home/phbo/Class/C/TMP/echonp

Breakpoint 1, echo () at echo.c:6
6           gets(buf);
(gdb) p/x $rbp
$1 = 0x7fffffe490
(gdb) p/x *((unsigned long*)$rbp)
$2 = 0x7fffffe4a0
(gdb) p/x *((unsigned long*)$rbp+1)
$3 = 0x40059c
```

`%rbp+1` is return instruction

```
$3 = 0x40059c
(gdb) p/x $rsp
$4 = 0x7fffffe480
```

342	400597:	e8 ca ff ff ff	callq 400566 <echo>
343	40059c:	b8 00 00 00 00	mov \$0x0,%eax

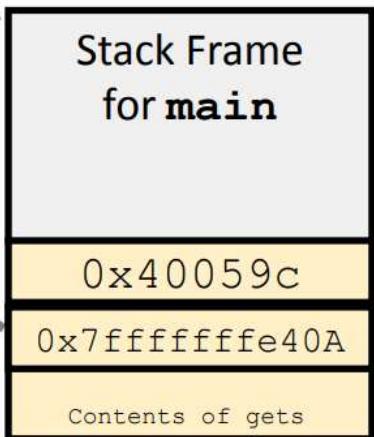
if I write more than 16 bytes,
then I overwrite return address.
I can take over the process
(or provoke segmentation fault)

UNIVERSITY OF COPENHAGEN

ITU CPH

Vulnerable code

Before call to gets



urn
struction

s buffer
aligned

so, why vulnerable?
Press Esc to exit full screen % gdb ./echonp

```
GNU gdb (Ubuntu 7.11.1-0ubuntu1~16.5) 7.11.1
Copyright (C) 2016 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from ./echonp...done.
(gdb) break echo
Breakpoint 1 at 0x40056e: file echo.c, line 6.
(gdb) r
Starting program: /home/phbo/Class/C/TMP/echonp

Breakpoint 1, echo () at echo.c:6
6          gets(buf);
(gdb) p/x $rbp
$1 = 0x7fffffe490
(gdb) p/x *((unsigned long*)$rbp)
$2 = 0x7fffffe4a0
(gdb) p/x *((unsigned long*)$rbp+1)
$3 = 0x40059c
```

%rbp+1 is return instruction

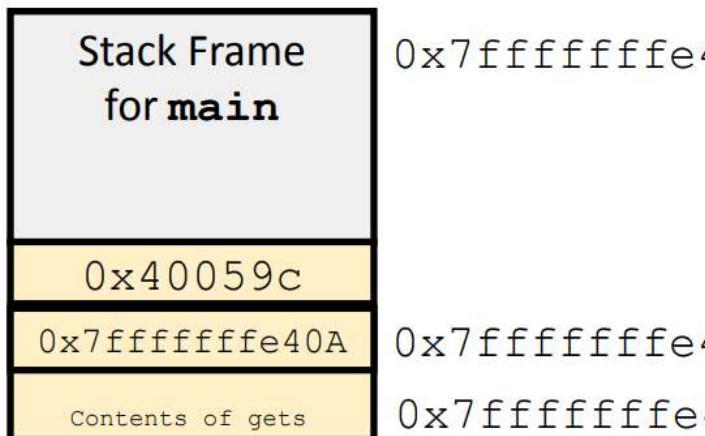
```
$5 = 0x40059c
(gdb) p/x $rsp
$4 = 0x7fffffe480
```

if I write more than 16 bytes,
then I overwrite return address.
I can take over the process
(or provoke segmentation fault)

342	400597:	e8 ca ff ff ff	callq 400566 <echo>
343	40059c:	b8 00 00 00 00	mov \$0x0,%eax

Buffer Overflow Stack Example

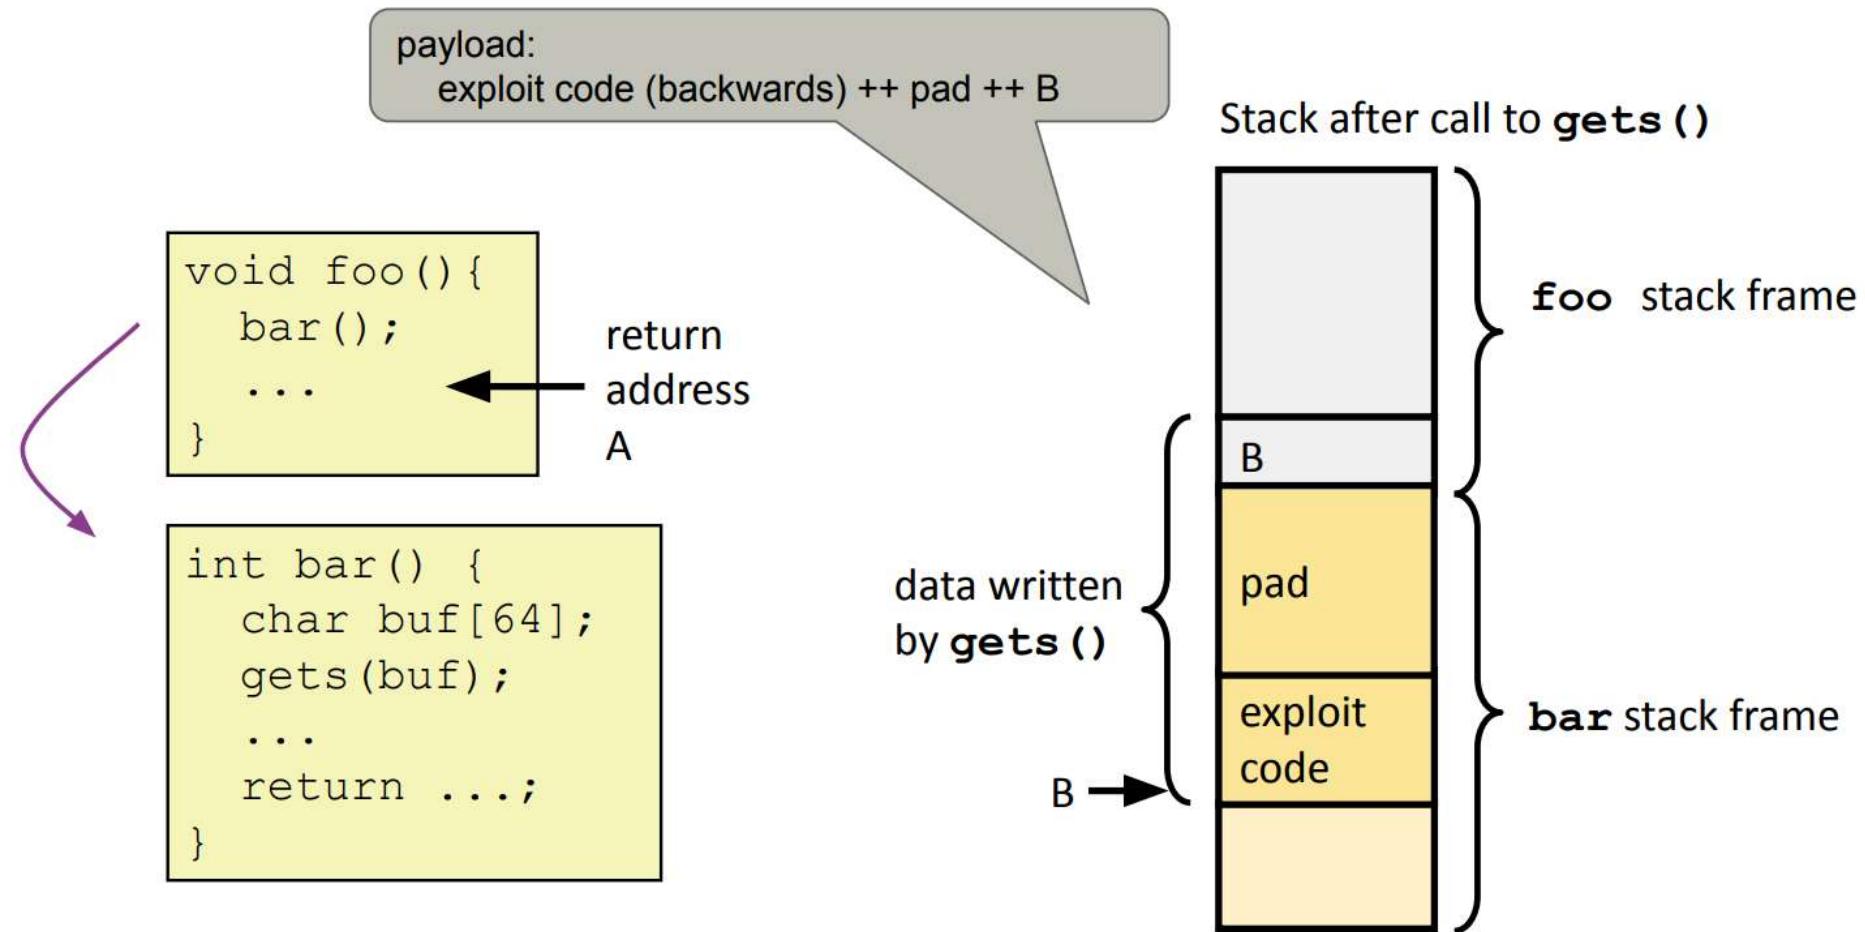
Before call to gets



depending on input,
different termination behavior.

```
quite anyway. (y or .., ,  
% ./echonp  
1234567890123456789012345  
1234567890123456789012345  
[1] 23189 segmentation fault (core dumped) ./echonp  
% ./echonp  
123456789012345678901234  
123456789012345678901234  
[1] 23255 illegal hardware instruction (core dumped) ./echonp  
% ./echonp  
12345678901234567890123  
12345678901234567890123
```

Malicious Use of Buffer Overflow



Input string contains **byte representation of executable code**
 Overwrite return address with address of exploit code
 When `bar()` executes `ret`, will jump to exploit code

Exploits Based on Buffer Overflows

Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines

Internet worm

Early versions of the finger server (fingerd) used `gets()` to read the argument sent by the client:

- **finger droh@cs.cmu.edu**

Morris worm, 1988

Worm attacked fingerd server by sending phony argument:

- **finger "exploit-code padding new-return-address"**

- exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.

how to stop this:

- want to make it impossible to execute code on the stack.
- want to make it hard to figure out where exploit code starts.
- want to protect return address.

Avoiding Overflow Vulnerability

use a function that checks how much you read.

```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    fgets(buf, 4, stdin);
    puts(buf);
}
```

Use library routines that limit string lengths

fgets instead of **gets**

strncpy instead of **strcpy**

Don't use **scanf** with **%s** conversion specification

- Use **fgets** to read the string
- Or use **%ns** where **n** is a suitable integer

System-Level Protections

Randomized stack offsets

At start of program, allocate random amount of space on stack
Makes it difficult for hacker to predict beginning of inserted code

Nonexecutable code segments

- In traditional x86, can mark region of memory as either “read-only” or “writeable”
- Can execute anything readable
- X86-64 added explicit “execute” permission

Stack Canaries

Idea

Place special value (“canary”) on stack just beyond buffer

Check for corruption before exiting function

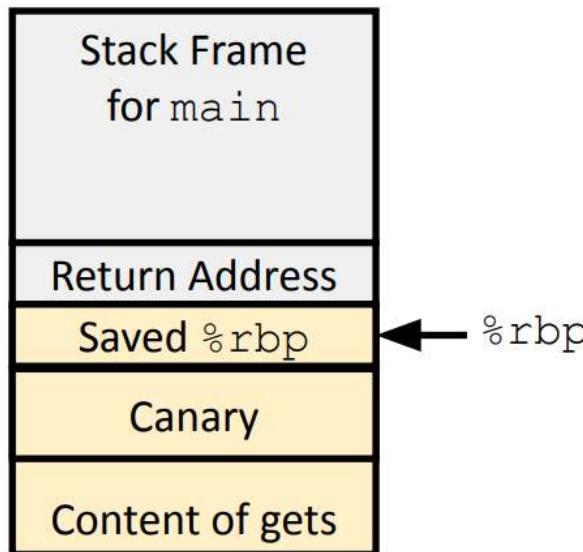
GCC Implementation

-fstack-protector

-fstack-protector-all (default)

Setting Up Canary

Before call to gets



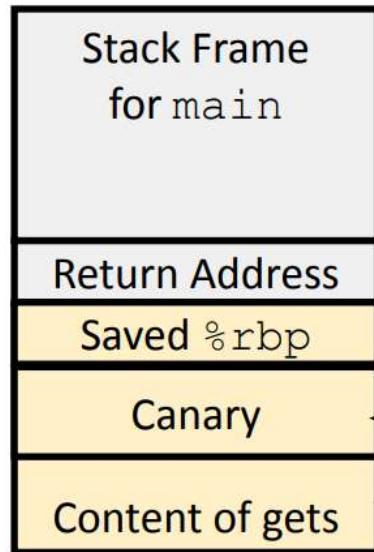
```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

<pre>echo: ... mov %fs:20, %rax mov %rax, -8(%rbp) xor %eax, %eax ... </pre>	random part of memory into rax put canary just below base pointer
	# Get canary # Put on stack # Erase canary

must overwrite canary to
overwrite return address

Checking Canary

Before call to gets



```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
echo:
    . . .
    mov -8(%rbp), %rax      # Retrieve from
stack
    xor %fs:20, %rax        # Compare with
Canary
    je .L24                  # Same: skip ahead
    call __stack_chk_fail    # ERROR
.L24:
    . . .
```

Canary Example

```

361
362 00000000004005d6 <echo>:
363 4005d6:    55          push   %rbp
364 4005d7:    48 89 e5    mov    %rsp,%rbp
365 4005da:    48 83 ec 10  sub    $0x10,%rsp
366 4005de:    64 48 8b 04 25 28 00  mov    %fs:0x28,%rax
367 4005e5:    00 00
368 4005e7:    48 89 45 f8  mov    %rax,-0x8(%rbp)
369 4005eb:    31 c0        xor    %eax,%eax
370 4005ed:    48 8d 45 f0  lea    -0x10(%rbp),%rax
371 4005f1:    48 89 c7    mov    %rax,%rdi
372 4005f4:    b8 00 00 00 00  mov    $0x0,%eax
373 4005f9:    e8 c2 fe ff ff  callq 4004c0 <gets@plt>
374 4005fe:    48 8d 45 f0  lea    -0x10(%rbp),%rax
375 400602:    48 89 c7    mov    %rax,%rdi
376 400605:    e8 86 fe ff ff  callq 400490 <puts@plt>
377 40060a:    90          nop
378 40060b:    48 8b 45 f8  mov    -0x8(%rbp),%rax
379 40060f:    64 48 33 04 25 28 00  xor    %fs:0x28,%rax
380 400616:    00 00
381 400618:    74 05        je     40061f <echo+0x49>
382 40061a:    e8 81 fe ff ff  callq 4004a0 <__stack_chk_fail@plt>
383 40061f:    c9          leaveq
384 400620:    c3          retq
385
386 0000000000400621 <main>:
387 400621:    55          push   %rbp
388 400622:    48 89 e5    mov    %rsp,%rbp
389 400625:    b8 00 00 00 00  mov    $0x0,%eax
390 40062a:    e8 a7 ff ff ff  callq 4005d6 <echo>
391 40062f:    b8 00 00 00 00  mov    $0x0,%eax
392 400634:    5d          pop    %rbp
393 400635:    c3          retq
394 400636:    66 2e 0f 1f 84 00 00  nopw   %cs:0x0(%rax,%rax,1)
395 40063d:    00 00 00
396

```

Canary Example

```
12345678901234567890123
% ./echo
12345678
12345678
% ./echo
123456789
123456789
*** stack smashing detected ***: ./echo terminated
[1] 23643 abort (core dumped) ./echo
```

when compiled w/
protection

GCC protections

Compile-time

```
[1] 17102 abort (core dumped) ./echo
cc      echo.c    -o echo
echo.c: In function 'echo':
echo.c:6:2: warning: implicit declaration of function 'gets' [-Wimplicit-function-declaration]
  gets(buf);
^
/tmp/cchUozYT.o: In function `echo':
echo.c:(.text+0x24): warning: the `gets' function is dangerous and should not be used.
[1]
```

Run-time

```
[1] 17102 abort (core dumped) ./echo
% ./echo
123456789
123456789
*** stack smashing detected ***: ./echo terminated
[1] 17139 abort (core dumped) ./echo
```

pretty clear that you
shouldn't use **gets**.